


CSE113: Parallel Programming

- **Topic:** Memory Hierarchy and C++ threads
 - Caches
 - Cache lines
 - Coherence
 - C++ threads
 - False Sharing



Announcements

- Homework due on Oct 8
 - Three free late days
 - Plenty of office hours remaining to get help
 - Work on your design doc before asking for help

Announcements

- Reminder on quiz: Not heavily graded, but low effort responses are liable to lose points
- (Hopefully) Last lecture of Module 1
 - Moving into Module 2: mutual exclusion next time!
- Should be able to do part 1 and part 2 of homework
 - Hopefully, part 3 by today, maybe next time though.

Quiz

The following statement in a language like C or Java would be compiled to how many instructions in low-level code?

`z = x + x + x + x;`

☐ 0

☐ 1

☐ 2

☐ 4

One instruction if compiler is smart. $z = 3 * x$

Quiz

How many levels of caches does a typical x86 system have?

☐ 1

☐ 2

☐ 3

☐ 4

x86: L1, L2, LLC, they have 3.

Apple only has L1 and a shared L2. Used in phone, more shallow cache hierarchy

Quiz

Write a few reasons why it may be difficult to reason about program performance when using a high-level language like Python

Quiz

Using your best guess, how much faster do you think a program written in C/Java is than a program written in Python? Give a few reasons explaining your guess. Feel free to run a simple experiment and see what happens!

It is said that python is 100 times slower than C.

Quiz

Modern-day compilers and runtimes will automatically make your code parallel. Because of this, most programmers do not need to think about parallelism when writing programs.

Review

Using Loop Unrolling to Exploit ILP

- Simple loop unrolling:

```
for (int i = 0; i < SIZE; i+=2) {  
    SEQ(i, 1);  
    SEQ(i, 2);  
    ...  
    SEQ(i, N); // end iteration for i  
    SEQ(i+1, 1);  
    SEQ(i+1, 2);  
    ...  
    SEQ(i+1, N); // end iteration for i + 1  
}
```

Let $SEQ(i, j)$ be the j th instruction of $SEQ(i)$.

Let each instruction chain have N instructions

Using Loop Unrolling to Exploit ILP

- Simple loop unrolling:

```
for (int i = 0; i < SIZE; i+=2) {  
    SEQ(i, 1);  
    SEQ(i+1, 1);  
    SEQ(i, 2);  
    SEQ(i+1, 2);  
    ...  
    SEQ(i, N);  
    SEQ(i+1, N);  
}
```

They can be interleaved

two instructions can be pipelined, or executed
on a superscalar processor


Loop Unrolling for Reduction Loops

- chunk array in equal sized partitions and do local reductions
- Consider size 2:

1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---

Loop Unrolling for Reduction Loops

- chunk array in equal sized partitions and do local reductions
- Consider size 2:



Main memory

```
int increment(int *a) {  
    a[0]++;  
}
```


```
%5 = load i32, i32* %4    200 cycles  
%6 = add nsw i32 %5, 1     1 cycles  
store i32 %6, i32* %4     200 cycles
```

200 cycles


1 cycles

200 cycles

401 cycles




Caches



Caches

```
int increment(int *a) {  
    a[0]++;  
}
```

```
%5 = load i32, i32* %4  
%6 = add nsw i32 %5, 1  
store i32 %6, i32* %4
```




Caches

```
int increment(int *a) {  
    a[0]++;  
}
```

%5 = load i32, i32* %4	4 cycles
%6 = add nsw i32 %5, 1	1 cycles
store i32 %6, i32* %4	4 cycles

9 cycles!



Quick overview of C/++ pointers/memory

Passing arrays in C++

```
int increment(int *a) {  
    a[0]++;  
}
```

```
int increment_alt1(int a[1]) {  
    a[0]++;  
}
```

```
int increment_alt2(int a[]) {  
    a[0]++;  
}
```

Not checked at compile time! but hints can help with compiler optimizations. Also good self documenting code.

Passing pointers

```
int foo0(int *a) {  
    increment(a)  
}
```

pass pointer directly through

```
int foo1(int *a) {  
    increment(&(a[8]))  
}
```

pass an offset of 8

```
int foo2(int *a) {  
    increment(a + 8)  
}
```

another way to pass an offset of 8

Memory Allocation

```
int allocate_int_array0() {  
    int ar[16];  
}
```

Stack allocation

```
int allocate_int_array1() {  
    int *ar = new int[16];  
    delete[] ar;  
}
```

Heap allocation, C++ style

```
int allocate_int_array2() {  
    int *ar = (int*)malloc(sizeof(int)*16);  
    free(ar);  
}
```

Heap allocation, C style

Cache lines


- Cache line size for x86: 64 bytes:
 - 64 chars
 - 32 shorts
 - 16 float or int
 - 8 double or long

Caches

```
int increment(int *a) {  
    a[0]++;  
}
```

```
%5 = load i32, i32* %4  
%6 = add nsw i32 %5, 1  
store i32 %6, i32* %4
```

Assume a[0] is not in the cache




Caches

```
int increment(int *a) {  
    a[0]++;  
}
```

```
%5 = load i32, i32* %4  
%6 = add nsw i32 %5, 1  
store i32 %6, i32* %4
```


$a[0] - a[15]$

Assume $a[0]$ is not in the cache




Caches

```
int increment_several(int *a) {  
    a[0]++;  
    a[15]++;  
    a[16]++;  
}
```



Caches


```
int increment_several(int *a) {  
    a[0]++;  
    a[15]++;  
    a[16]++;  
}
```



Caches

```
int increment_several(int *a) {  
    a[0]++;  
    a[15]++;  
    a[16]++;  
}
```

will be a hit because we've loaded a[0] cache line




Caches

```
int increment_several(int *a) {  
    a[0]++;  
    a[15]++;  
    a[16]++;  
}
```

Miss

Assume $a[0]$ is not in the cache




Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```


Assume a[0] is not in the cache



Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}  
  
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

Assume $a[0]$ is not in the cache




Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

This loads a[8]

Assume a[0] is not in the cache



Cache alignment


```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

This loads a[8]

This loads a[23], a miss!

Assume a[0] is not in the cache



Cache alignment

- Malloc typically returns a pointer with “good” alignment.
 - System specific, but will be aligned at least to a cache line, more likely a page
- For very low-level programming you can use special aligned malloc functions
- Prefetchers will also help for many applications (e.g. streaming)

Cache alignment

- Malloc typically returns a pointer with “good” alignment.
 - System specific, but will be aligned at least to a cache line, more likely a page
- For very low-level programming you can use special aligned malloc functions
- Prefetchers will also help for many applications (e.g. streaming)

```
for (int i = 0; i < 100; i++) {  
    a[i] += b[i];  
}
```


prefetcher will start collecting consecutive data in the cache if it detects patterns like this.

Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```


Assume a[0] is not in the cache



Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}  
  
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

Assume $a[0]$ is not in the cache




Cache alignment

```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

This loads a[8]

Assume a[0] is not in the cache



Cache alignment


```
int increment_several(int *b) {  
    b[0]++;  
    b[15]++;  
}
```

```
int foo(int *a) {  
    increment_several(&(a[8]))  
}
```

This loads a[8]

This loads a[23], a miss!

Assume a[0] is not in the cache




Cache coherence

Cache coherence


How to manage multiple values for the same address in the system?

Simplified view for illustration:
L1 cache and LLC


Consider 3 cores accessing the same memory location




Cache coherence




Cache coherence




Cache coherence




Cache coherence



Cache coherence




Cache coherence




Cache coherence

in parallel




Cache coherence



Cache coherence


Incoherent view of values!



Cache coherence


- MESI protocol
- Cache line can be in 1 of 4 states:
 - **Modified** - the cache contains a modified value, and it must be written back to the lower-level cache
 - **Exclusive** - only 1 cache has a copy of the value
 - **Shared** - more than 1 cache contains the value, they must all agree on the value
 - **Invalid** - the data is stale and a new value must be fetched from a lower-level cache

Cache coherence




Cache coherence

`load(a0)`




Cache coherence

*Exclusive states
are clean: they match
main memory*




Cache coherence


load(a0)




Cache coherence



Cache coherence




Cache coherence




Cache coherence

*Modified states
are dirty: they don't
match main memory*




Cache coherence


*Invalid states
are considered unused*




Cache coherence



Cache coherence



Cache coherence




Cache coherence

Takeaways:

Caches must agree on values across cores.

Caches are functionally invisible! Cannot tell with raw input and output

But performance measurements can expose caches, especially if they share the same cache line




C++ Threads

- Introduction
 - Learn as needed throughout class
- Multi-threading officially introduced in C++11
 - only widely available after ~2014
 - official specification
 - cross-platform
- Before C++ threads
 - pthreads

C++ Threads

- Main idea:
 - run functions concurrently




main

launch foo(a,b,c)


C++ Threads

- Main idea:
 - run functions concurrently



C++ Threads

- Main idea:
 - run functions concurrently




```
#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}
```

main waits for foo.
called **join()**

join() returns in main



```
#include <thread>
using namespace std;
```


header and namespace

```
void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}
```

main waits for foo.
called **join()**

join() returns in main



foo finishes

```

#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}

```


Launches a concurrent thread that executes foo

Stores a handle in thread_handle (don't lose the handle!)

constructor takes in the function, and all arguments

main waits for foo.
called **join()**

join() returns in main



```

#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo, 1, 2, 3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}


```

Requires C++14

clang++ -std=c++14 main.cpp

main waits for foo.
called **join()**

join() returns in main



```

#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}


int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}

```

calling join() on the thread handle will cause main to wait for the thread launched with thread_handle to finish.

main waits for foo.
called join()

join() returns in main



```

#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}


int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}

```

After foo finishes,
main starts executing again

main waits for foo.
called **join()**

join() returns in main



```
#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}
```

What happens if you don't join your threads?

```
#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}
```

What happens if you don't join your threads?

```
and, threads? /usr/...
libc++abi.dylib: terminating
Abort trap: 6
```

JOIN YOUR THREADS!!!

```
#include <thread>
using namespace std;

void foo(int a, int b, int c) {
    // some foo code
}

int main() {
    // some main code
    thread thread_handle (foo,1,2,3);
    // code here runs concurrently with foo
    thread_handle.join();
    return 0;
}
```

return value?

Doesn't have to be void,
but it is ignored

how to get values back
from threads?

```
#include <thread>
#include <iostream>
using namespace std;

void foo(int a, int b, int *c) {
    // return a + b;
    *c = a + b;
}

int main() {
    // some main code
    int ret = 0;
    thread thread_handle (foo, 1, 2, &ret);
    // code here runs concurrently with foo
    thread_handle.join();
    cout << ret << endl;
    return 0;
}
```

Pass by address (C++ or C)

```
#include <thread>
#include <iostream>
using namespace std;

int c;
void foo(int a, int b) {
    // return a + b;
    c = a + b;
}

int main() {
    // some main code
    int ret = 0;
    thread thread_handle (foo, 1, 2);
    // code here runs concurrently with foo
    thread_handle.join();
    cout << c << endl;
    return 0;
}
```

Options

global variable
(don't do this very often!)

```
#include <thread>
#include <iostream>
using namespace std;

void foo(int a, int b, int *c) {
    // return a + b;
    *c = a + b;
}

int main() {
    // some main code
    int ret = 0;
    thread thread_handle (foo, 1, 2, &ret);
    // code here runs concurrently with foo
    cout << ret << endl;
    thread_handle.join();
    return 0;
}
```

What if....

```
#include <thread>
#include <iostream>
using namespace std;

void foo(int a, int b, int *c) {
    // return a + b;
    *c = a + b;
}

int main() {
    // some main code
    int ret = 0;
    thread thread_handle (foo, 1, 2, &ret);
    // code here runs concurrently with foo
    cout << ret << endl;
    thread_handle.join();
    return 0;
}
```

What if....

Undefined behavior!
Cannot access the same
values concurrently
without protection!

Next module we will talk
protection (locks)

SPMD programming model

- Same program, multiple data
- Main idea: many threads execute the same function, but they operate on different data.
- How do they get different data?
 - each thread can access their own thread id, a contiguous integer starting at 0 up to the number of threads

SPMD programming model

```
void increment_array(int *a, int a_size) {  
    for (int i = 0; i < a_size; i++) {  
        a[i]++;  
    }  
}
```

lets do this in parallel!
each thread increments different
elements in the array

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = 0; i < a_size; i++) {  
        a[i]++;  
    }  
}
```

The function gets a thread id and the number of threads

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = 0; i < a_size; i++) {  
        a[i]++;  
    }  
}
```

*A few options on how to split up the work
lets do round robin*

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```



array a

Assume 2 threads
lets step through thread 0
i.e.
tid = 0
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```

iteration 1 computes index 0



array a

Assume 2 threads
lets step through thread 0
i.e.
tid = 0
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i += num_threads) {  
        a[i]++;  
    }  
}
```

iteration 2 computes index 2



array a

Assume 2 threads
lets step through thread 0
i.e.
tid = 0
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i += num_threads) {  
        a[i]++;  
    }  
}
```

iteration 3 computes index 4



array a

Assume 2 threads
lets step through thread 0
i.e.
tid = 0
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i += num_threads) {  
        a[i]++;  
    }  
}
```



array a

switch to thread 1

Assume 2 threads
lets step through thread 1
i.e.
tid = 1
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```

iteration 1 computes index 1



array a


switch to thread 1

Assume 2 threads
lets step through thread 1
i.e.
tid = 1
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```

iteration 2 computes index 3



array a


switch to thread 1

Assume 2 threads
lets step through thread 1
i.e.
tid = 1
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads) {  
    for (int i = tid; i < a_size; i+=num_threads) {  
        a[i]++;  
    }  
}
```

iteration 3 computes index 5



array a

switch to thread 1

Assume 2 threads
lets step through thread 1
i.e.
tid = 1
num_threads = 2

SPMD programming model

```
void increment_array(int *a, int a_size, int tid, int num_threads);
```

```
#define THREADS 8
#define A_SIZE 1024
int main() {
    int *a = new int[A_SIZE];
    // initialize a
    thread thread_ar[THREADS];
    for (int i = 0; i < THREADS; i++) {
        thread_ar[i] = thread(increment_array, a, A_SIZE, i, THREADS);
    }
    for (int i = 0; i < THREADS; i++) {
        thread_ar[i].join();
    }
    delete[] a;
    return 0;
}
```


Extra if time

Concurrency vs. Parallelism


- Abstract tasks:
 - In the abstract: a sequence of computation
 - *Given an input, produces an output*

Concurrency vs. Parallelism


- Abstract tasks:
 - In the abstract: a sequence of computation
 - *Given an input, produces an output*
- Concrete tasks:
 - Application (e.g. Spotify and Chrome)
 - Function
 - Loop iterations
 - Individual instructions
 - Circuit level?



Concurrency vs. Parallelism




Concurrency vs. Parallelism



Concurrency vs. Parallelism

Sequential execution


Not concurrent or parallel



Concurrency vs. Parallelism



The OS can preempt a thread
(remove it from the hardware resource)




time

Concurrency vs. Parallelism



The OS can preempt a thread
(remove it from the hardware resource)




Concurrency vs. Parallelism



The OS can preempt a thread
(remove it from the hardware resource)

tasks are interleaved on the same processor




Concurrency vs. Parallelism



- Definition:
 - 2 tasks are **concurrent** if there is a point in the execution where both tasks have started and neither has ended.

The OS can preempt a thread
(remove it from the hardware resource)




Concurrency vs. Parallelism




- Definition:

- 2 tasks are **concurrent** if there is a point in the execution where both tasks have started and neither has ended.


The OS can preempt a thread
(remove it from the hardware resource)




Concurrency vs. Parallelism



Concurrency vs. Parallelism




Concurrency vs. Parallelism




Concurrency vs. Parallelism

- 2 tasks are **concurrent** if there is a point in the execution where both tasks have started and neither has ended.




Concurrency vs. Parallelism




Concurrency vs. Parallelism

- Definition:
 - An execution is **parallel** if there is a point in the execution where computation is happening simultaneously




Concurrency vs. Parallelism

- Examples:
 - Neither concurrent or parallel (sequential)




Concurrency vs. Parallelism

- Examples:
 - Concurrent but not parallel




Concurrency vs. Parallelism

- Examples:
 - Parallel and Concurrent




Concurrency vs. Parallelism

- Examples:
 - Parallel but not concurrent?




Concurrency vs. Parallelism

- Examples:
 - Parallel but not concurrent?



Concurrency vs. Parallelism

- Examples:
 - Parallel execution but task 0 and task 1 are not concurrent?



Concurrency vs. Parallelism

- In practice:
 - Terms are often used interchangeably.
 - *Parallel programming* is often used by high-performance engineers when discussing using parallelism to accelerate things
 - *Concurrent programming* is used more by interactive applications, e.g. event driven interfaces.

Code

- Folder thread_introduction
 - race
 - spmd
- Folder false_sharing