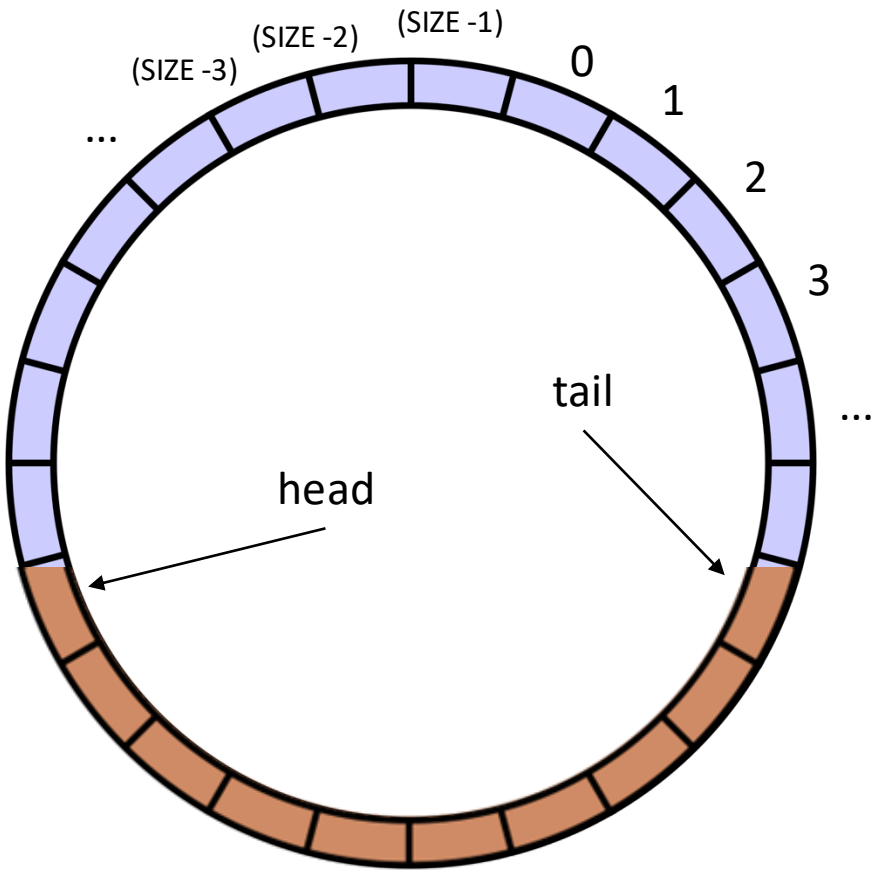
CSE113: Parallel Programming

- **Topics**:
  - Linearizability
  - Input/output queues
  - Producer/consumer queues



Announcements

- HW 2 deadline stays Oct 22.

- Midterm on Tuesday.

Announcements

- HW 3 released.

- You will have what you need for Part 1 by the end of today's lecture

Demo

Plain CAS vs. Relaxed peeking

- Folder: RMW_mutexes

Build with -O2

- At -O0 the plain CAS loop is slow, so its waiters barely disturb the holder.
- Peeking still pays its extra trip per handoff: see the flag free, then CAS.
- At -O2 the plain CAS contention is high, and peeking's trip is cheap.

Watch the cache line

- The lock flag sat on the same line as the counter it protects.
- Every increment in the critical section knocked the flag out of every spinner's cache.
- Peeking is only local if nobody writes the line. Pad the lock: alignas(128).

Previous quiz + Review

Quiz 9.

Lock-free data structures are technically undefined because they contain data conflicts

Select one

A. True

B. False

Scan to answer

forms.gle/YtF1gSqYHRqyMh1X9

Previous quiz + Review

It is impossible to use objects that are not thread-safe in a concurrent program.

- ◯ True
- ◯ False

global variables:

```
bank_account tylers_account;
mutex m;
```

Tyler's coffee addiction:

```
for (int i = 0; i < HOURS; i++) {
   m.lock();
   tylers_account.buy_coffee();
   m.unlock();
}
```

Tyler's employer

```
for (int j = 0; j < HOURS; j++) {
   m.lock();
   tylers_account.get_paid();
   m.unlock();
}
```

what if you have multiple objects?

First solution: The client (user of the object) can use locks.

We might decide to wrap my bank account in an object

```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      balance -= 1;
    }

    void get_paid() {
      balance += 1;
    }

  private:
    int balance;
};
```

The object is not "thread safe"

Previous quiz + Review

Non-locking objects do not use mutexes in their implementation. This is beneficial because:

- ◯ it is potentially faster
- ◯ it is easier to reason about
- ◯ it is easier to extend

Bank account example

global variables:

```
bank_account tylers_account;
```

Tyler's coffee addiction:

```
for (int i = 0; i < HOURS; i++) {
   tylers_account.buy_coffee();
}
```

Tyler's employer

```
for (int j = 0; j < HOURS; j++) {
   tylers_account.get_paid();
}
```

```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      atomic_fetch_add(&balance, -1);
    }

    void get_paid() {
      atomic_fetch_add(&balance, 1);
    }

  private:
    atomic_int balance;
};
```

Demo

Folder: concurrent_data_structures

Observation

- Descheduling Mutex beats lock-free beats spinlock

Why: count cache line transfers per update

- Spinlock: both holder and waiting threads run. Lock, unlock, and the data write each move the line, two to three transfers per update.
- Lock-free: both threads run. One atomic per update, so the line moves once per update.
- Mutex: the waiter is descheduled until the unlocker wakes it. The holder keeps the line and runs long batches. Almost no transfers.

Conclusion

- Lock-free guarantees progress, not speed.
- Lock-free beats the spinlock with one atomic instead of two.
- Shared counters do not scale. Reduce sharing.

Previous quiz + Review

Write a few sentences about the pros and cons of using a concurrent data structure vs. using mutexes to protect data structures that are not thread-safe.

Previous quiz + Review

Write a few sentences about the pros and cons of using a concurrent data structure vs. using mutexes to protect data structures that are not thread-safe.

Pros: Easier to use
Cons: Composability

Previous quiz + Review

When multiple threads access a concurrent object, only 1 possible execution is allowed. We reason about that execution by sequentializing object method calls and it is called sequential consistency

- ◯ True
- ◯ False

Global variable:

```
CQueue<int> q;
```

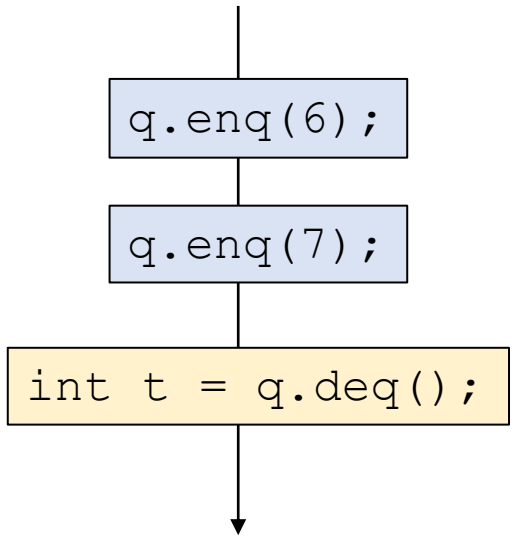
Thread 0:

```
q.enq(6);
q.enq(7);
```

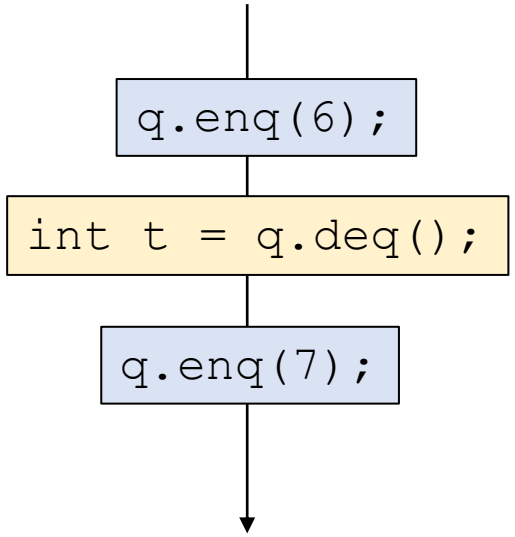
Thread 1:

```
int t = q.deq();
```

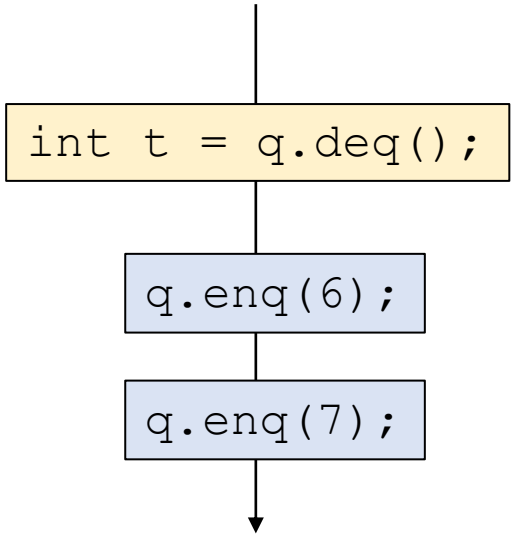
Construct a sequential timeline of API calls
Any sequence is valid:



t is 6



t is 6



t is None

Can t ever be 7?

Previous quiz + Review

What is the relationship between linearizable (L) and sequentially consistent (SC)?

- ◯ Objects can be one or the other, but not both
- ◯ Objects that are L are also SC, but not the other way around
- ◯ Objects that are SC are also L, but not the other way around
- ◯ SC and L are the different definitions for the same concept

Review

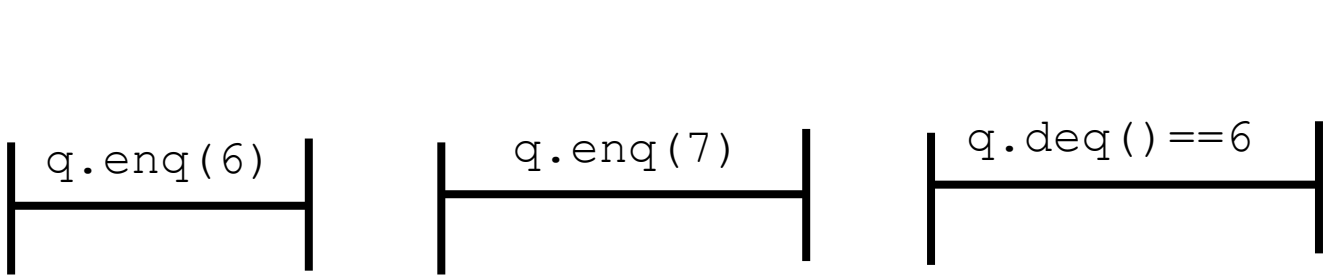
Sequential consistency and real time

- Add in real time:

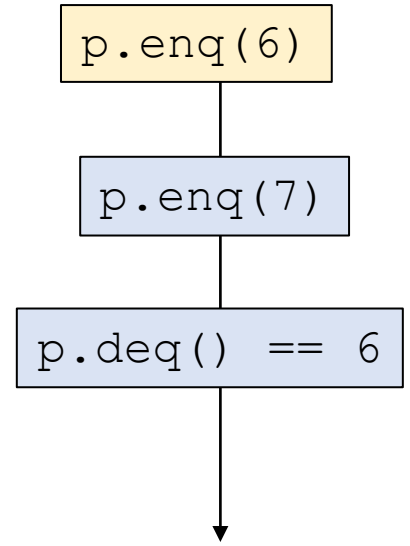
This execution is allowed in sequential consistency!

SC doesn't care about real time, only if it can construct its virtual sequential timeline

Thread 0



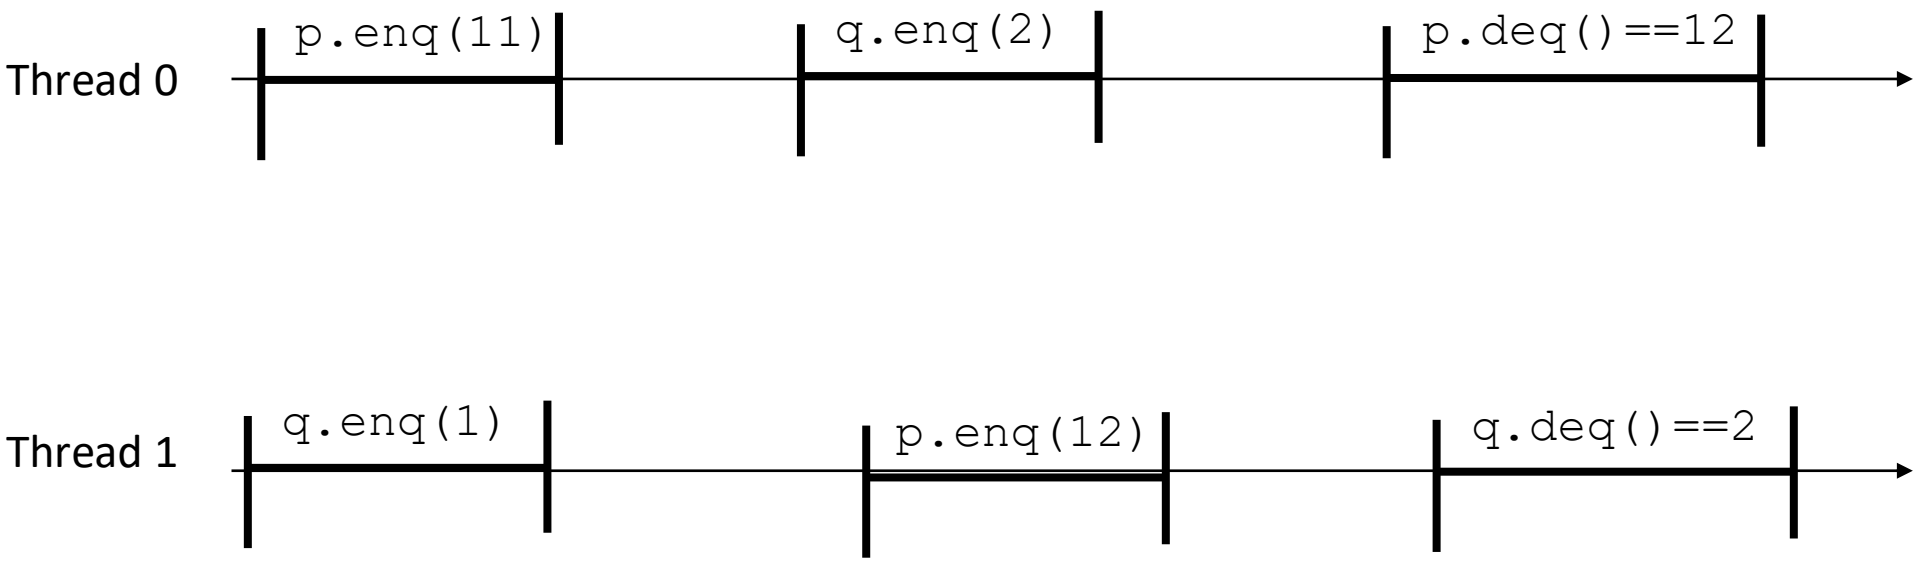
Thread 1



Sequential consistency and real time

- Add in real time:

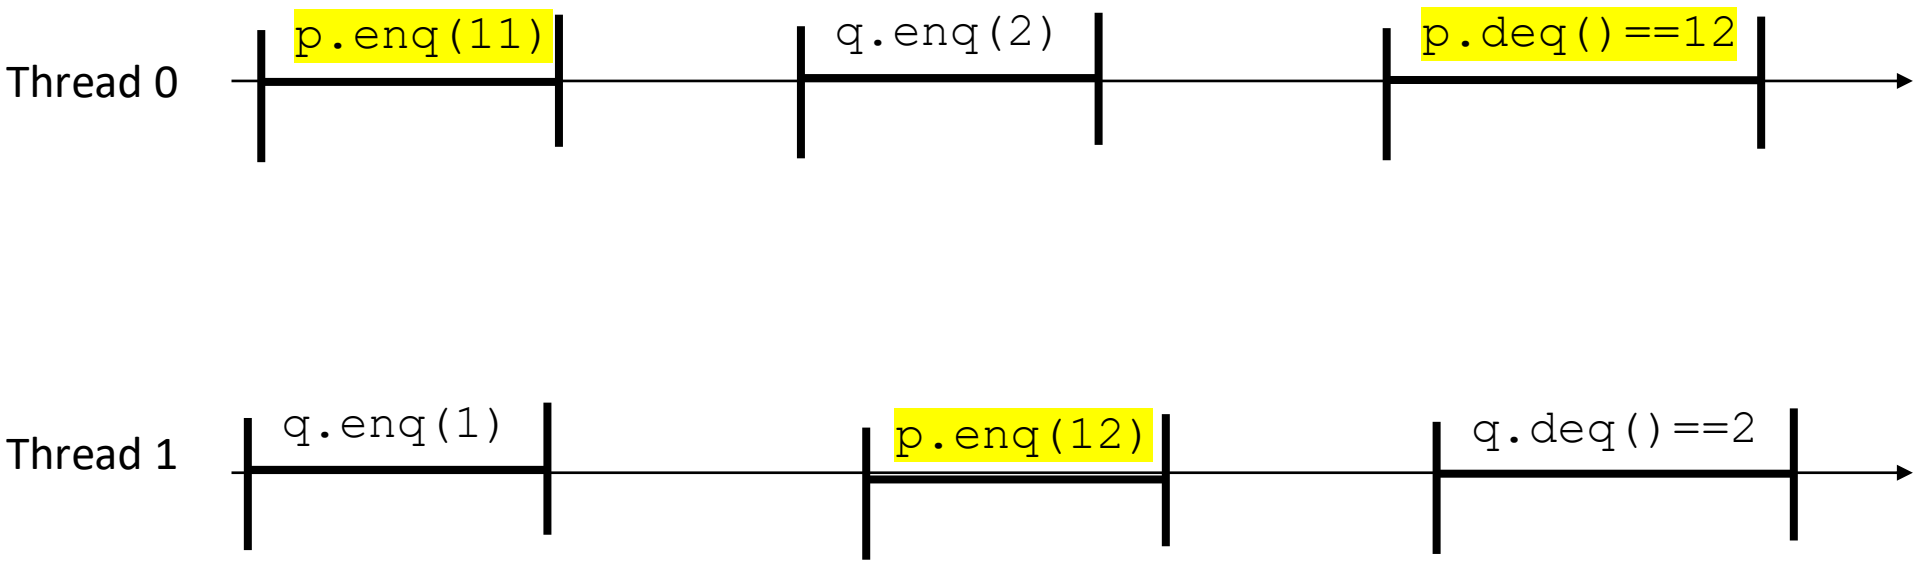
2 objects now: p and q



Sequential consistency and real time

- Add in real time:

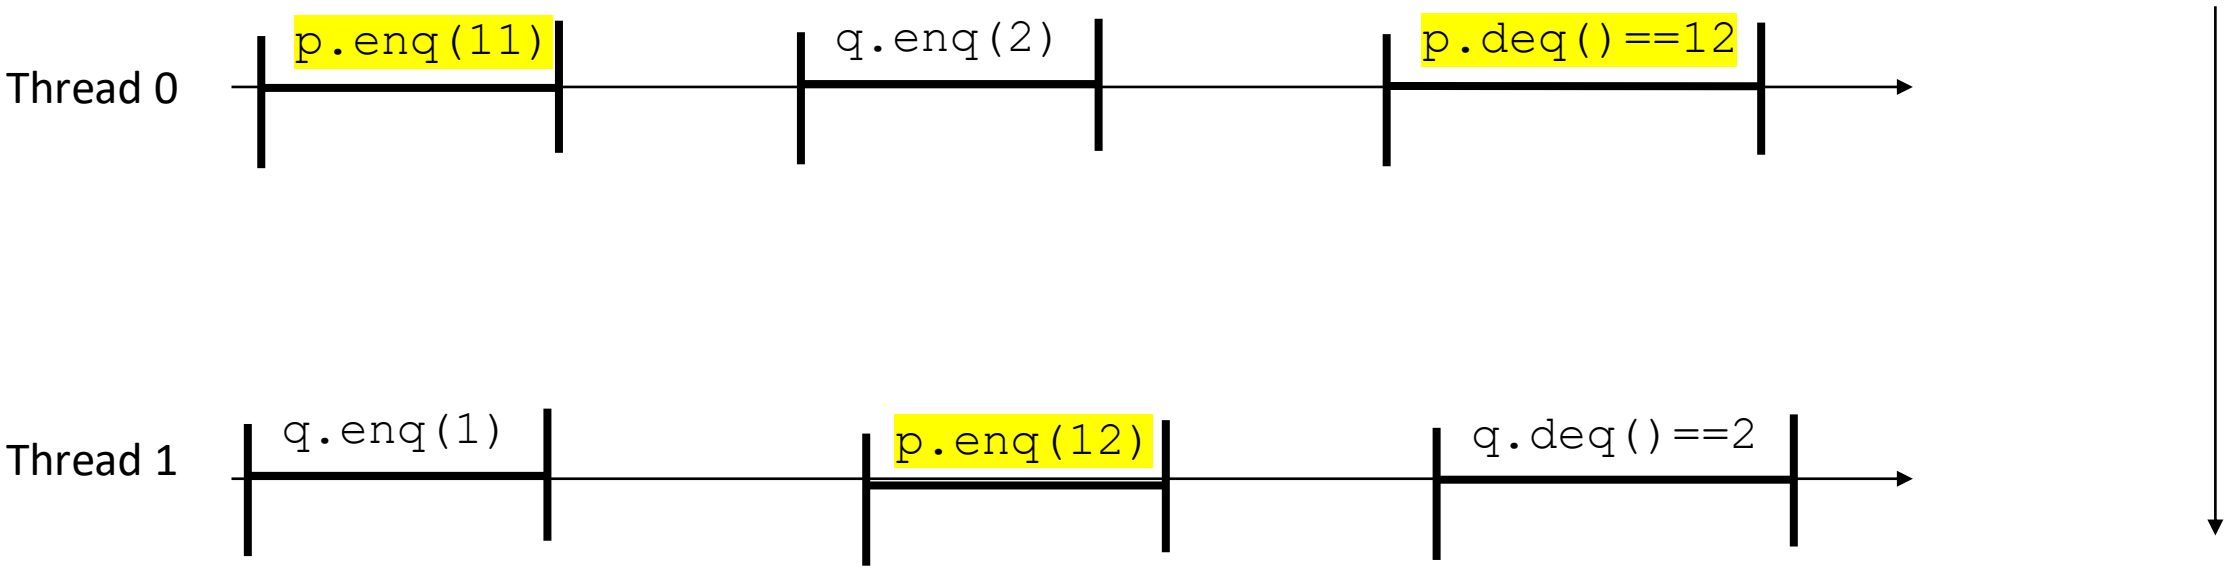
2 objects now: p and q
Consider each object in isolation



Sequential consistency and real time

- Add in real time:

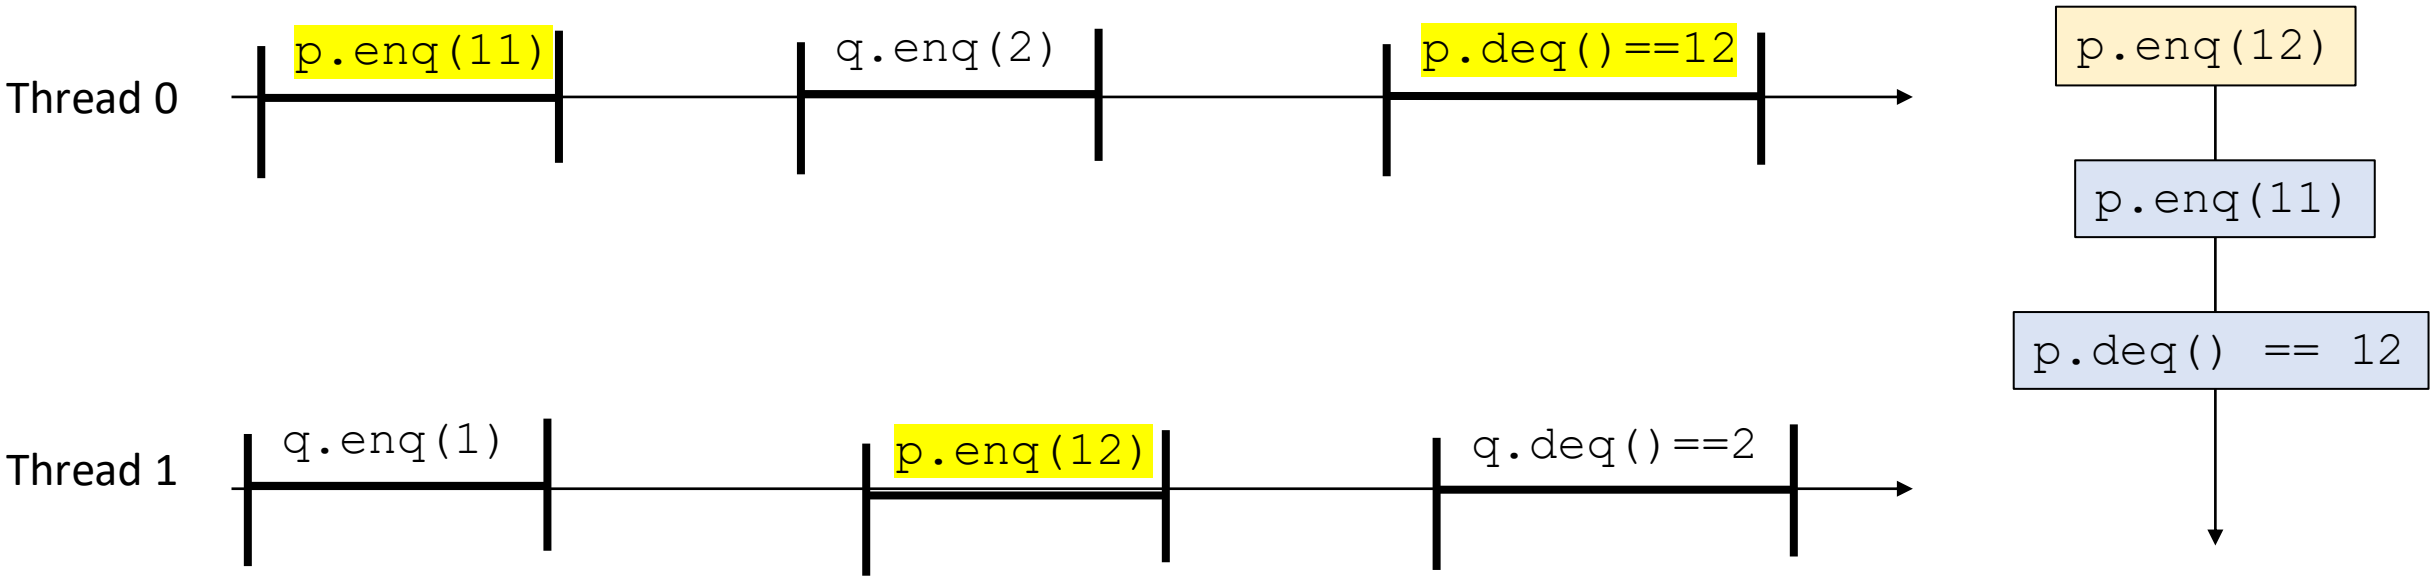
2 objects now: p and q
Consider each object in isolation



Sequential consistency and real time

- Add in real time:

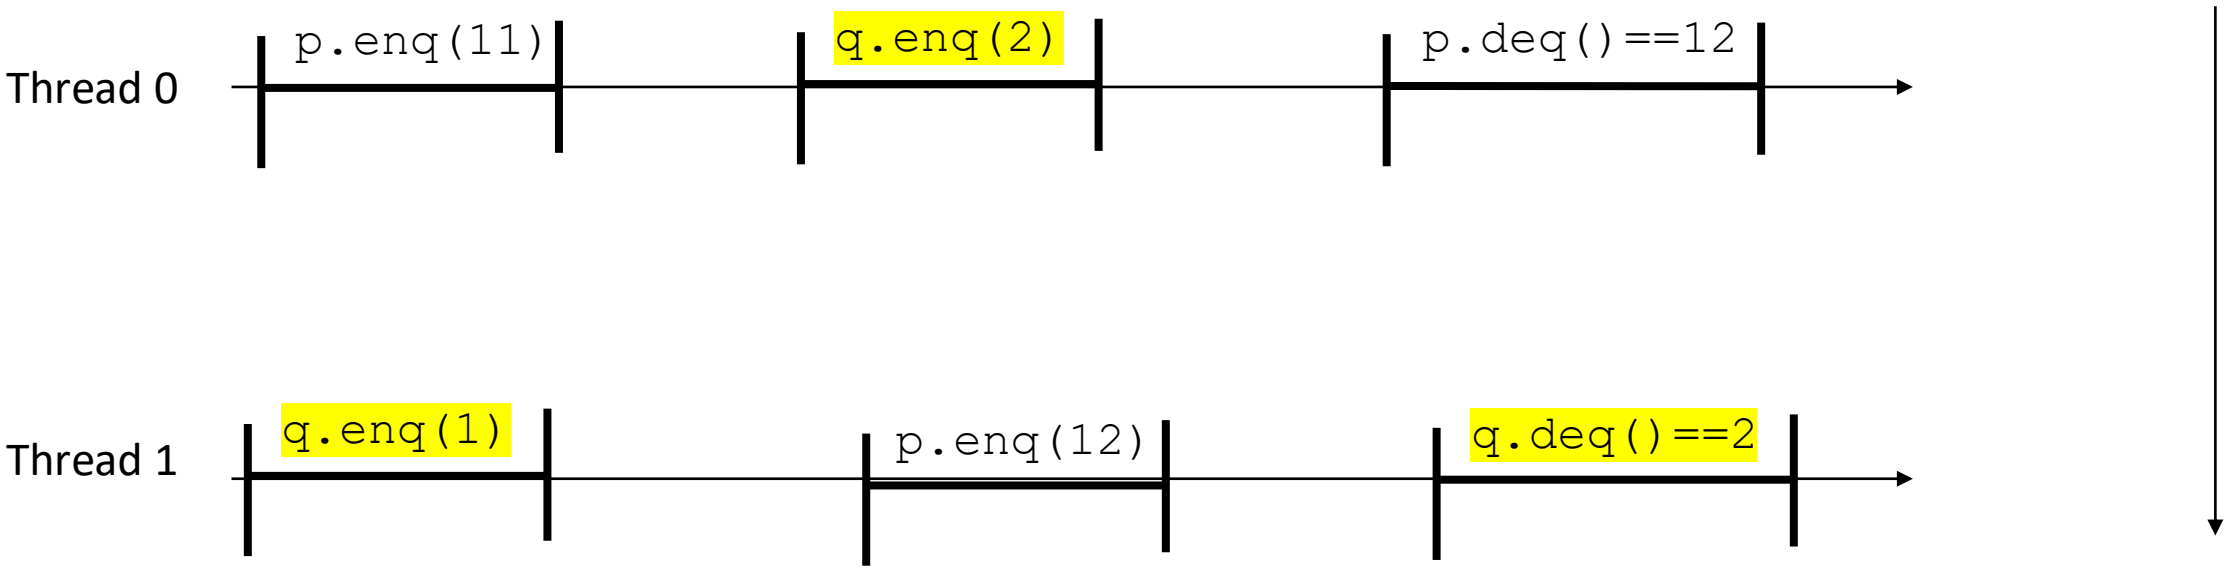
2 objects now: p and q
Consider each object in isolation



Sequential consistency and real time

- Add in real time:

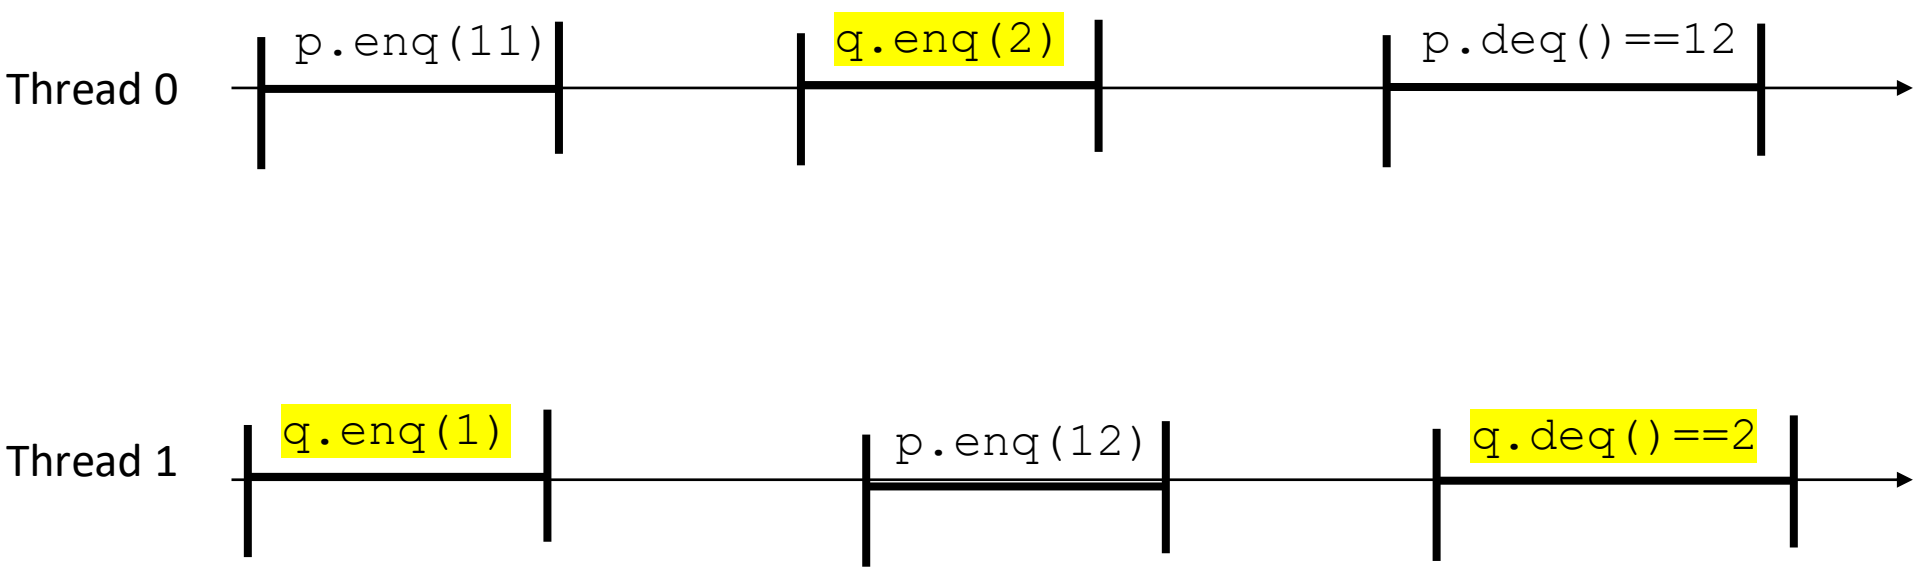
2 objects now: p and q
Consider each object in isolation

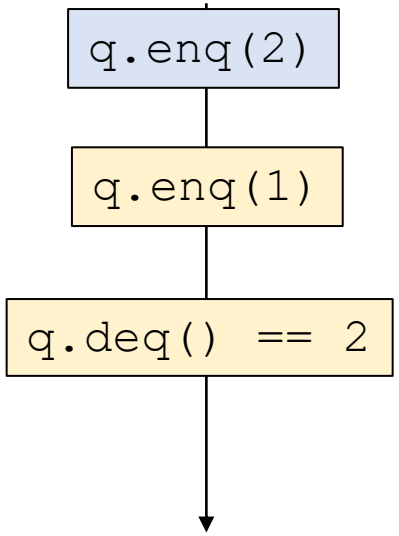


Sequential consistency and real time

- Add in real time:

2 objects now: p and q
Consider each object in isolation

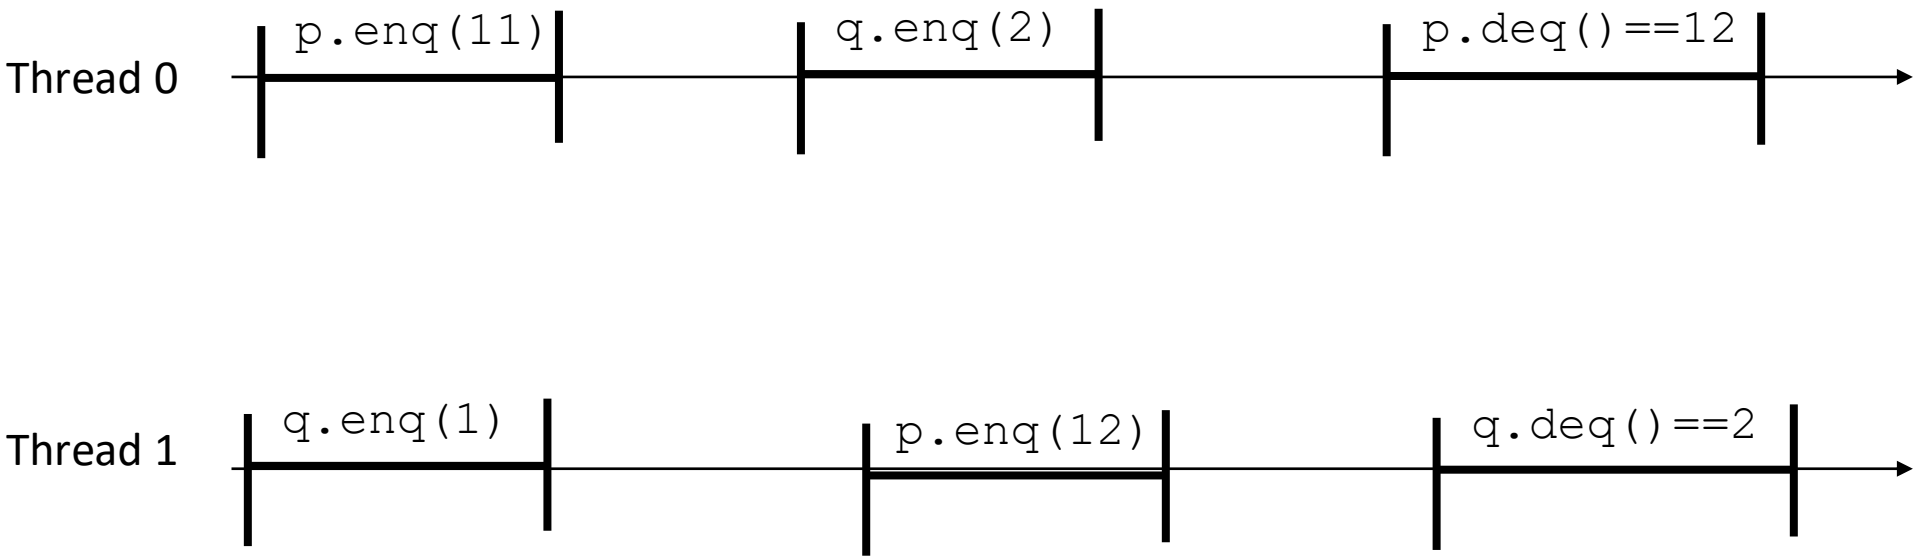




Sequential consistency and real time

- Add in real time:

Now consider them all together



Global variable:

```
CQueue<int> p,q;
```

Thread 0:

```
p.enq(11)
q.enq(2)
p.deq()==12
```

Order 1

```
p.enq(12)
p.enq(11)
p.deq() == 12
```

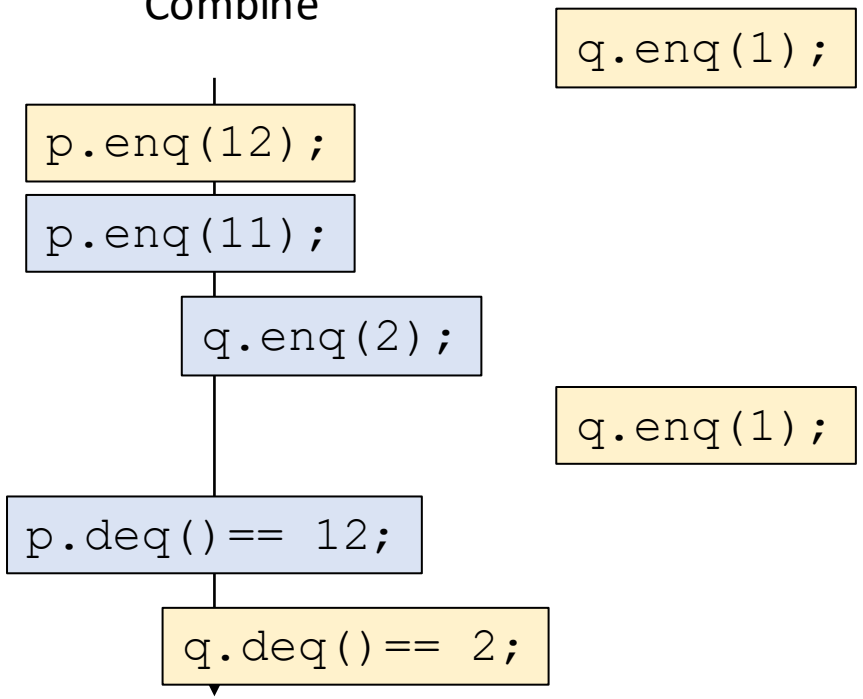
Order 2

```
q.enq(2)
q.enq(1)
q.deq() == 2
```

Thread 1:

```
q.enq(1)
p.enq(12)
q.deq()==2
```

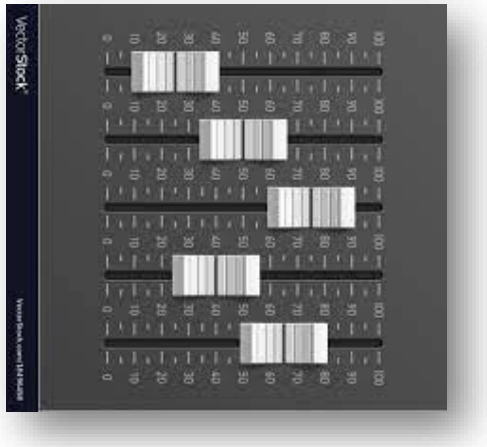
Combine



Linearizability

- Linearizability
  - Defined in term of real-time histories
  - SC + real-time preservation

- Slightly different game:
  - Sequential consistency is a game about stacking lego bricks
  - Linearizability is about sliders



Linearization point

Each operation has a linearization point (similar to the location of a call in sequential consistency)

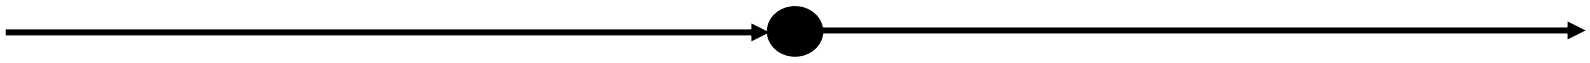
- does not overlap with other with other linearizability points
- indivisible computation (critical section, atomic RMW, atomic load, atomic store)
- object update (or read) occurs exactly at this point



Linearizability

each operation has a linearizability point

- does not overlap with other with other linearizability points
- indivisible computation (critical section, atomic RMW, atomic load, atomic store)
- object update (or read) occurs exactly at this point



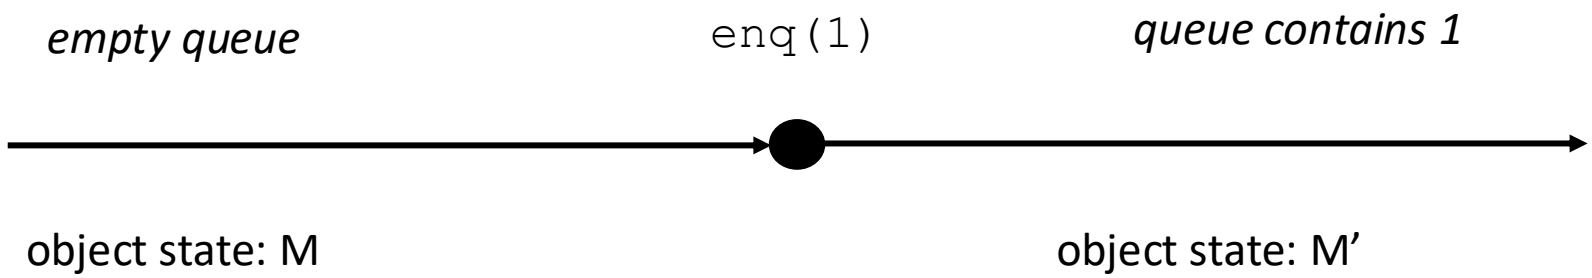
object state: M

object state: M’

Linearizability

each operation has a linearizability point

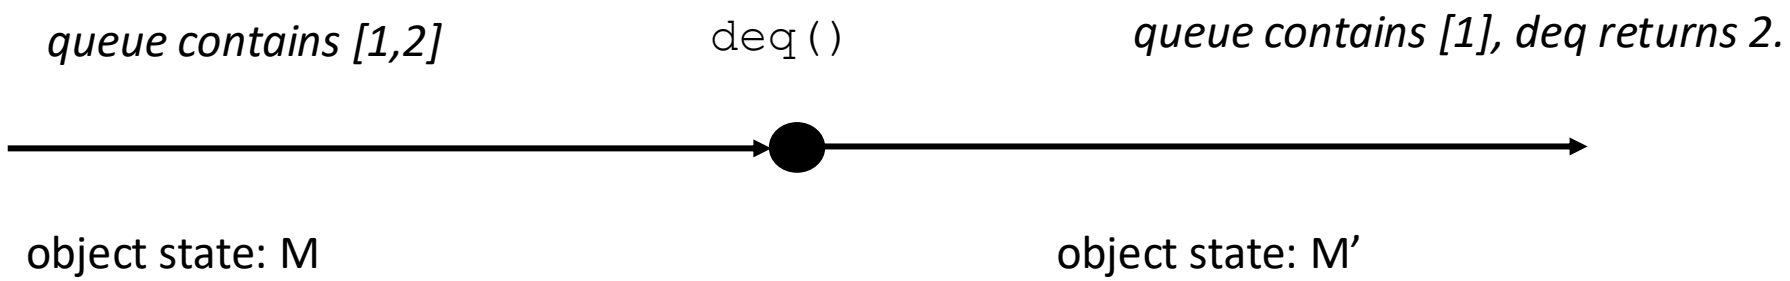
- does not overlap with other with other linearizability points
- indivisible computation (critical section, atomic RMW, atomic load, atomic store)
- object update (or read) occurs exactly at this point



Linearizability

each operation has a linearizability point

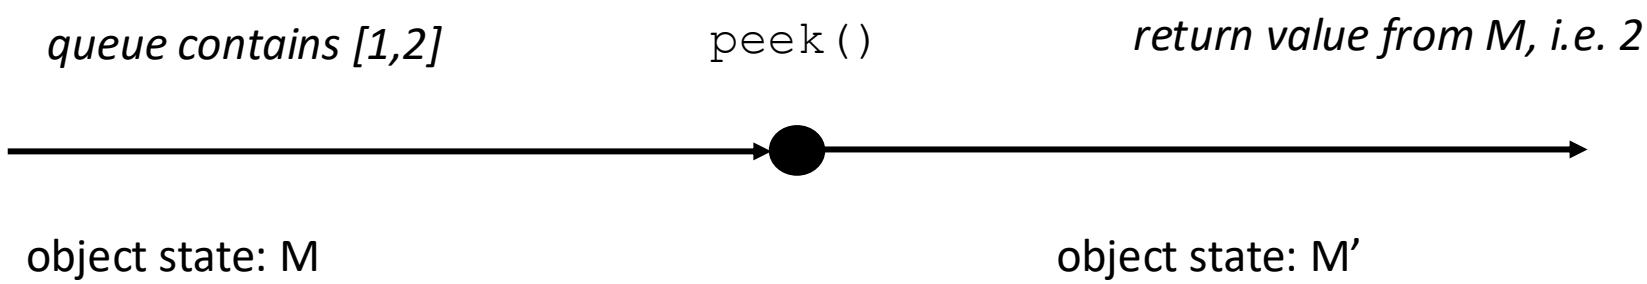
- does not overlap with other with other linearizability points
- indivisible computation (critical section, atomic RMW, atomic load, atomic store)
- object update (or read) occurs exactly at this point



Linearizability

each operation has a linearizability point

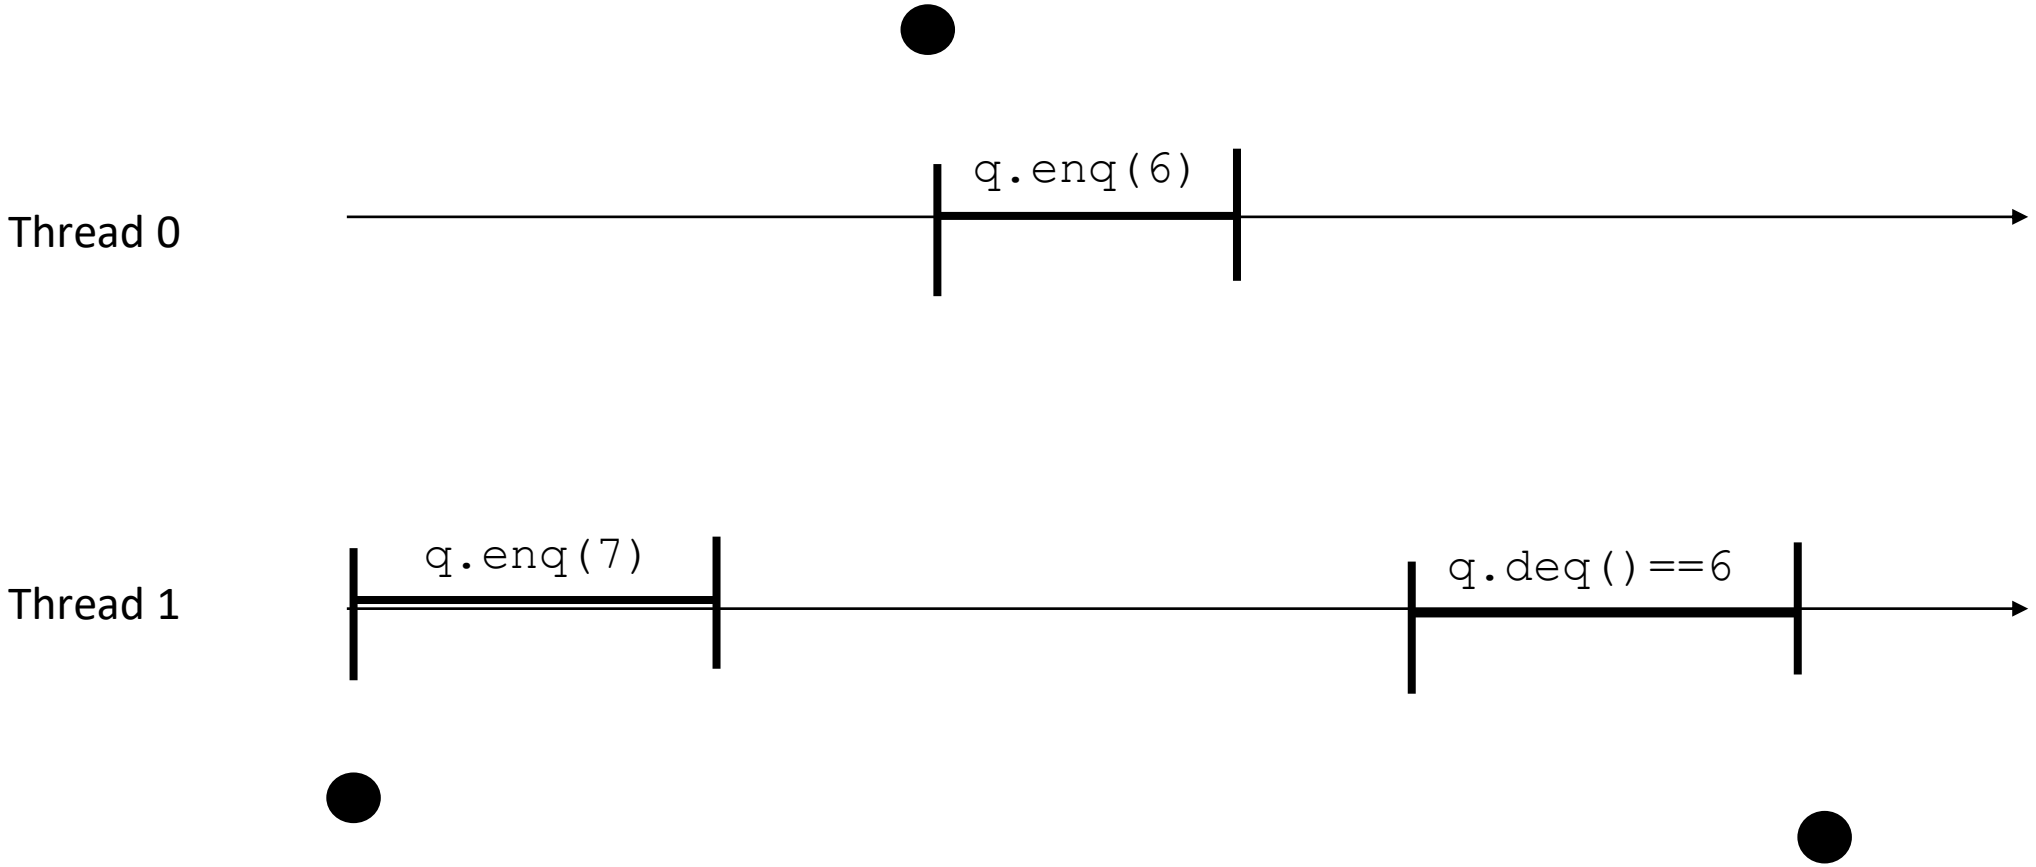
- does not overlap with other with other linearizability points
- indivisible computation (critical section, atomic RMW, atomic load, atomic store)
- object update (or read) occurs exactly at this point



Linearizability

each command gets a linearization point.

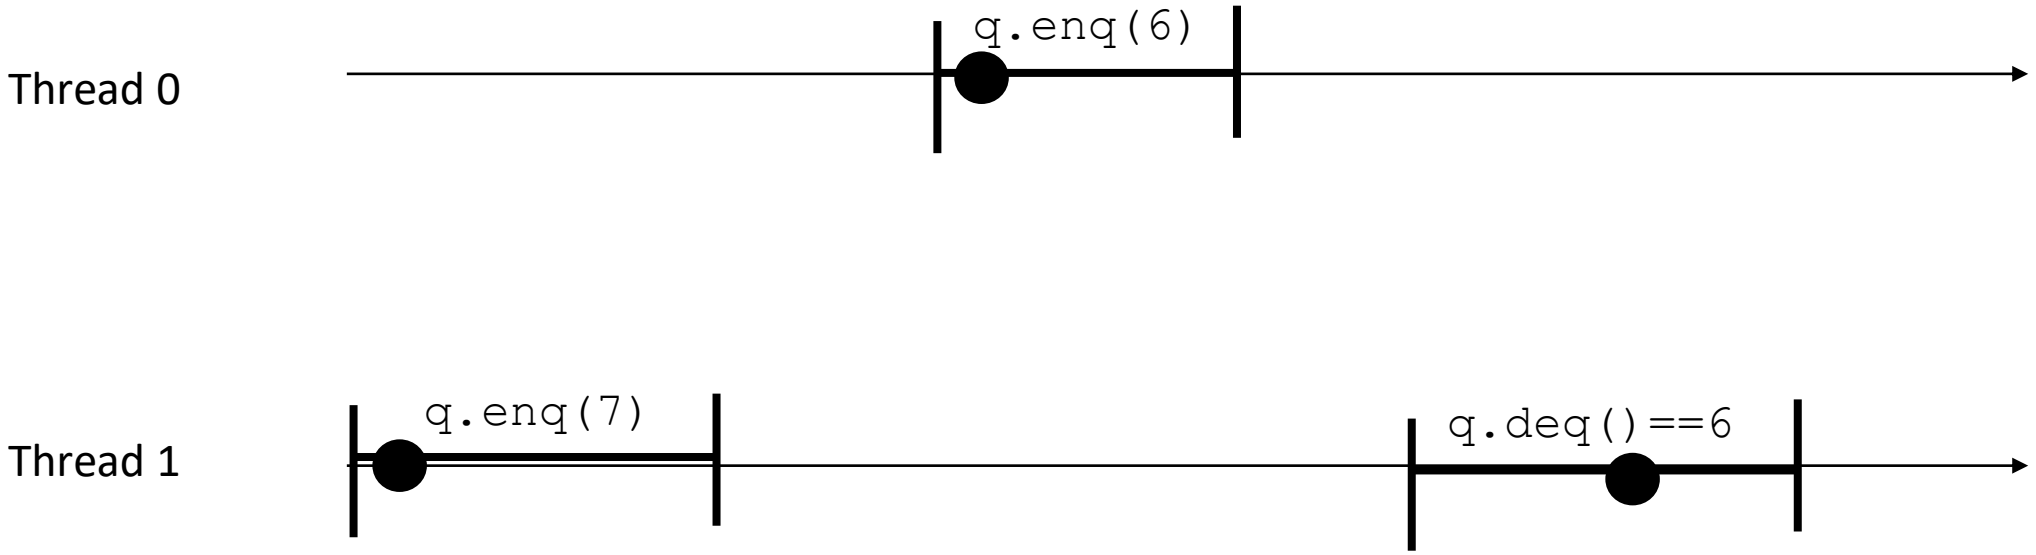
You can place the point anywhere **between its innovation and response! (respecting the real-time order)**



Linearizability

each command gets a linearization point.

You can place the point anywhere between its innovation and response!

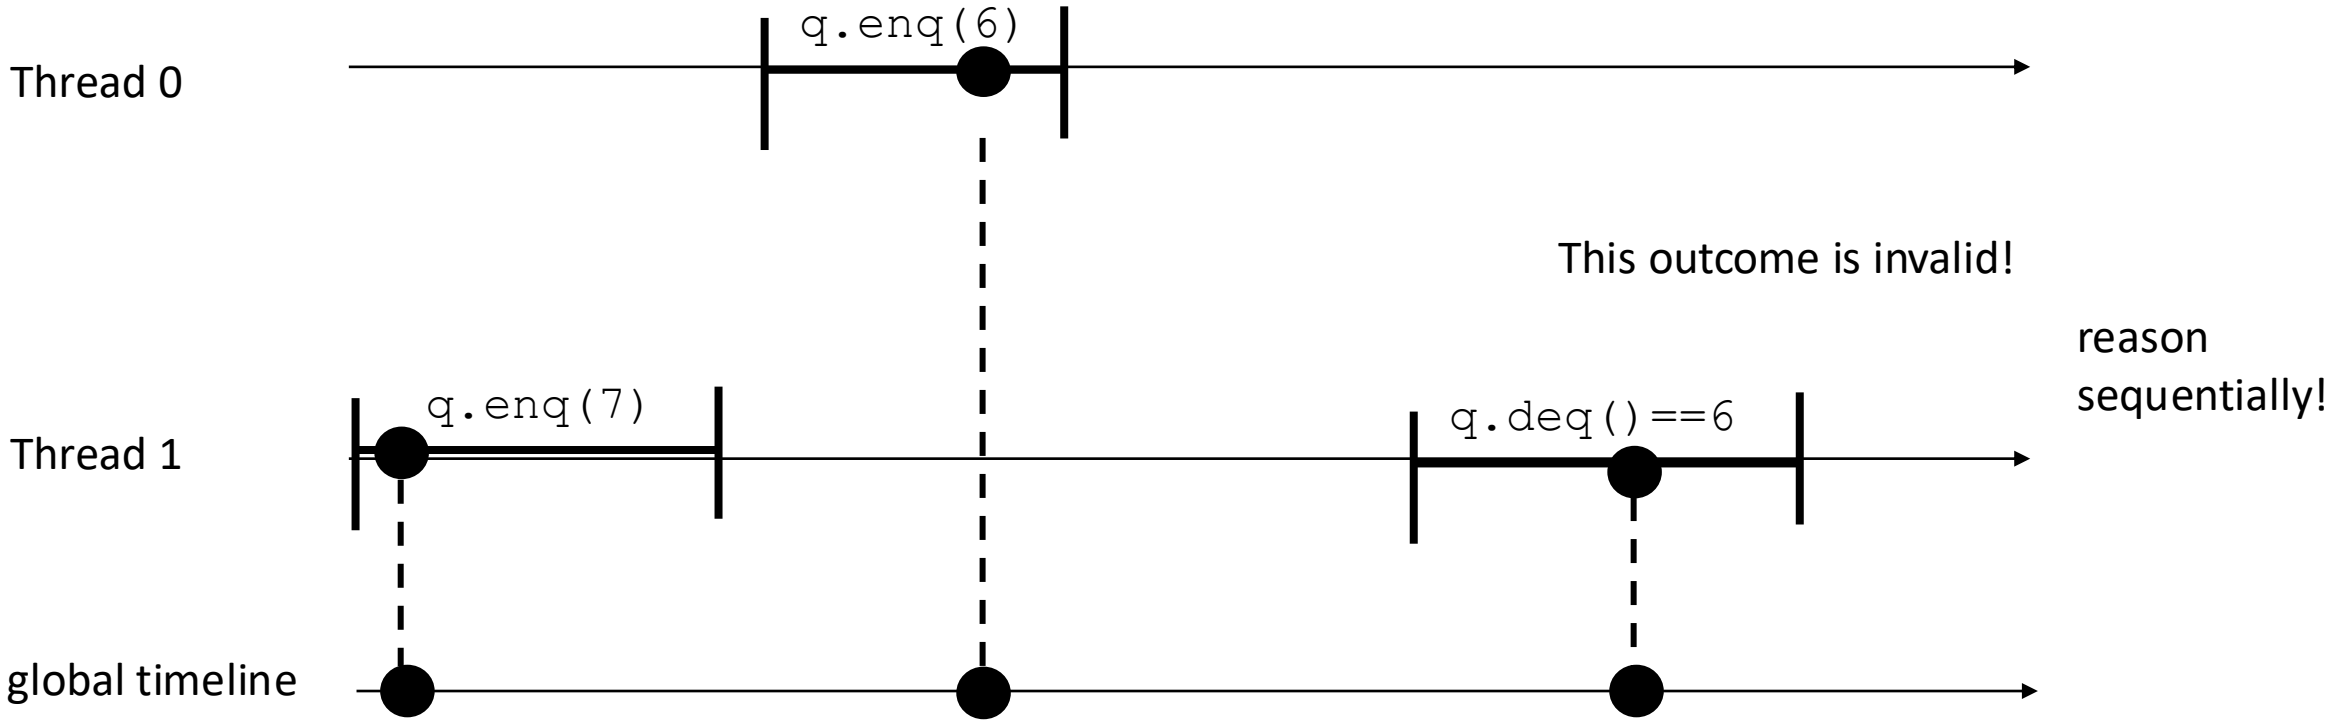


Linearizability

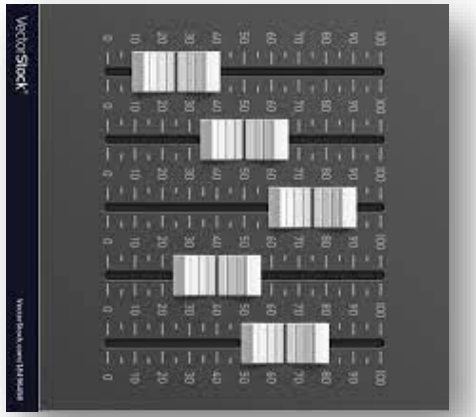
each command gets a linearization point.

You can place the point anywhere between its innovation and response (so long as they don't overlap)!

Project the linearization points to a global timeline



Linearizability

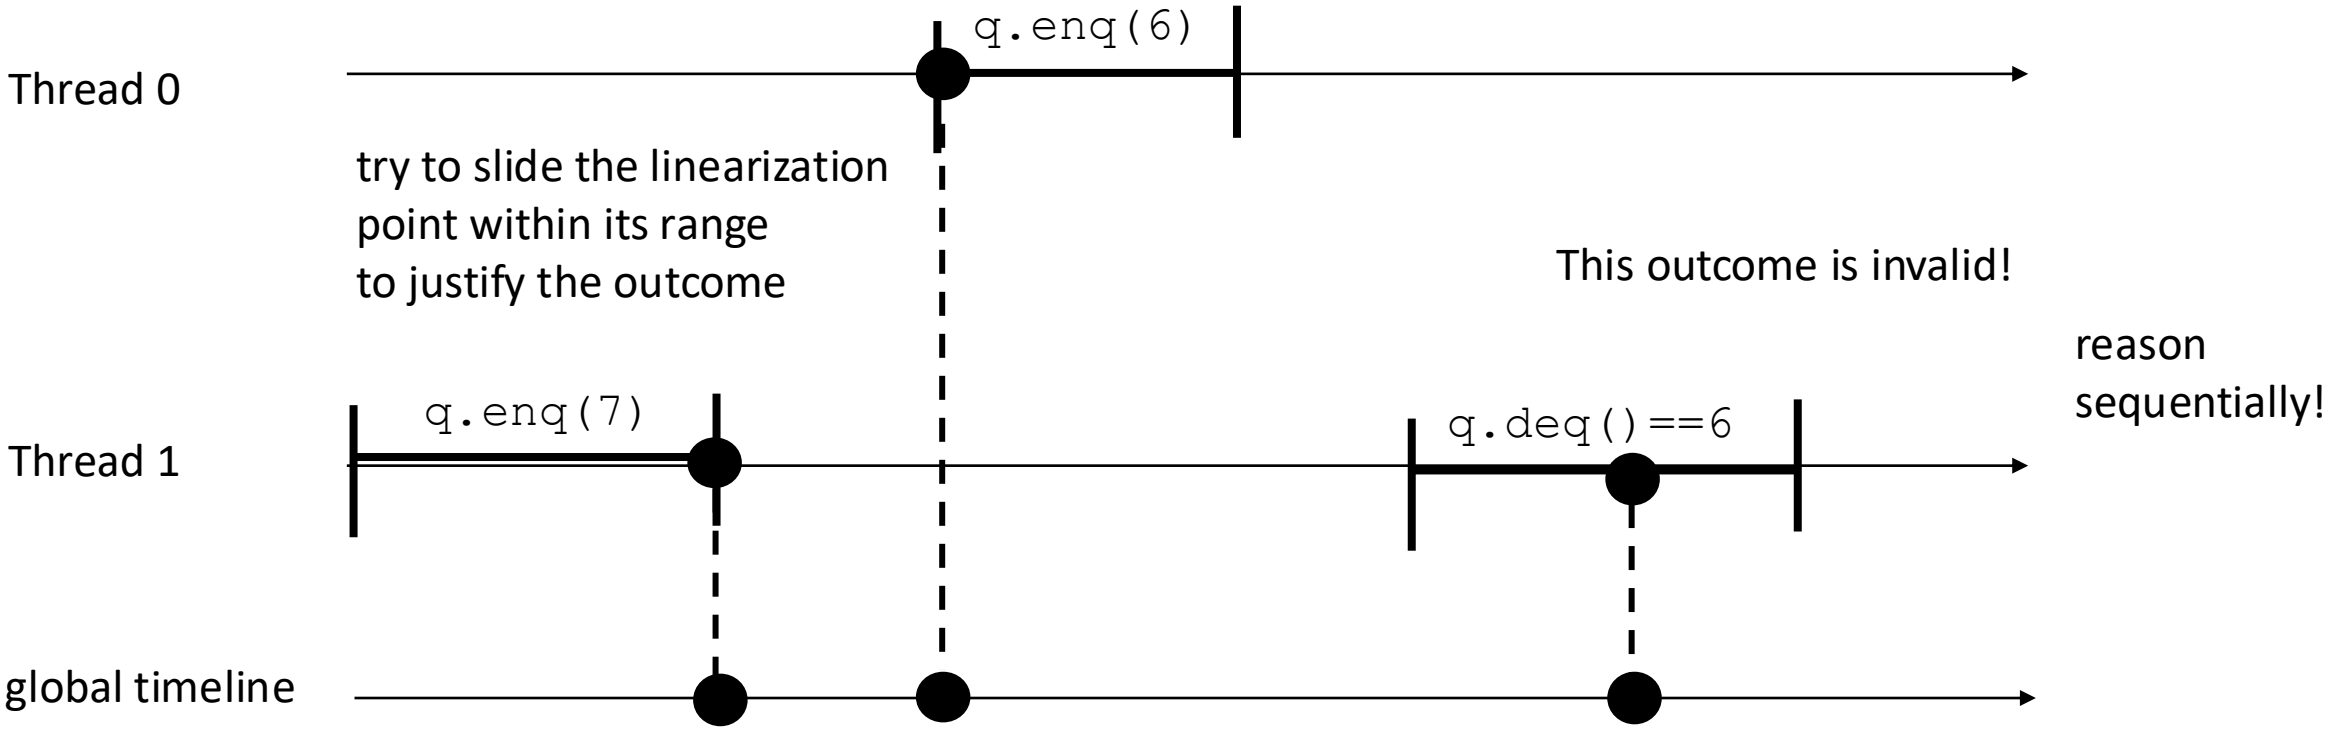


each command gets a linearization point.

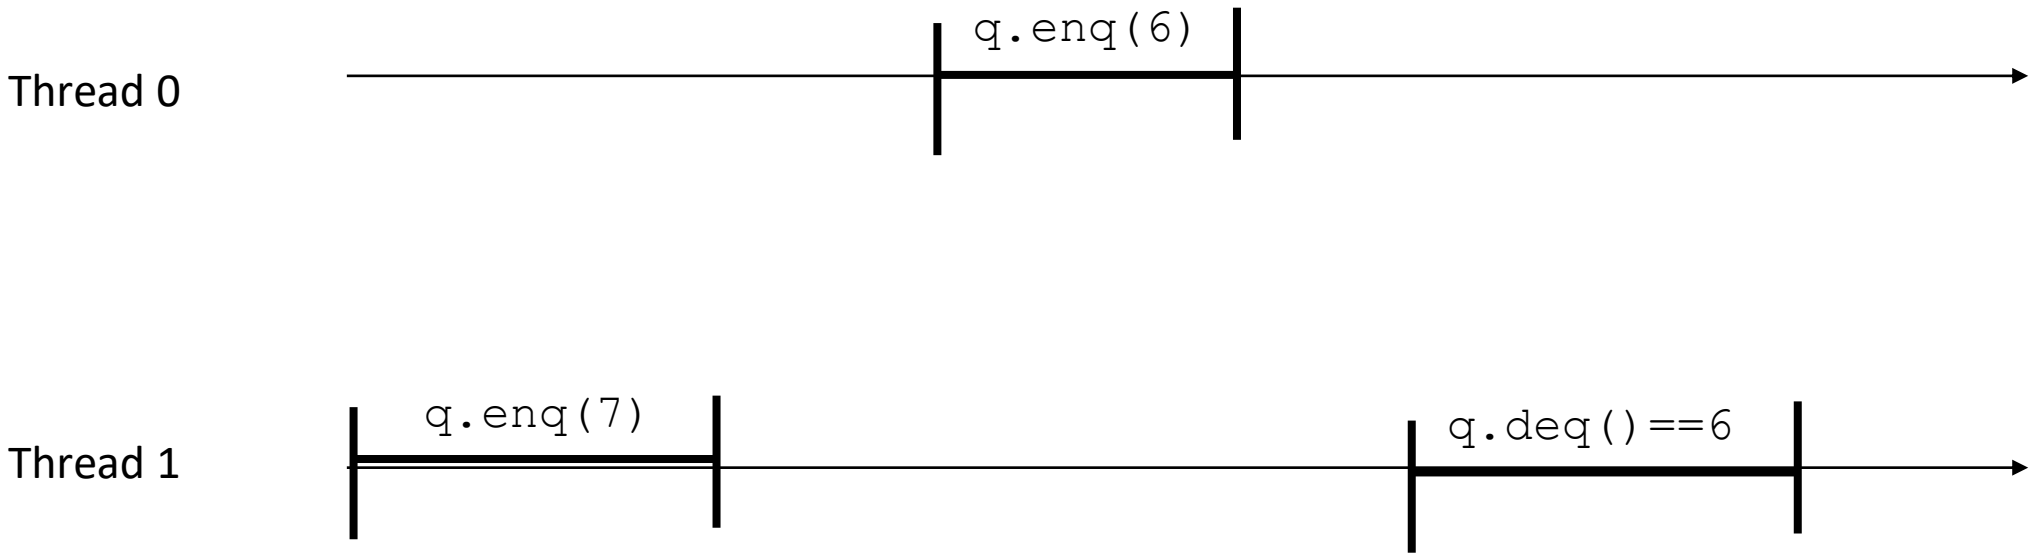
You can place the point anywhere between its innovation and response!

Project the linearization points to a global timeline

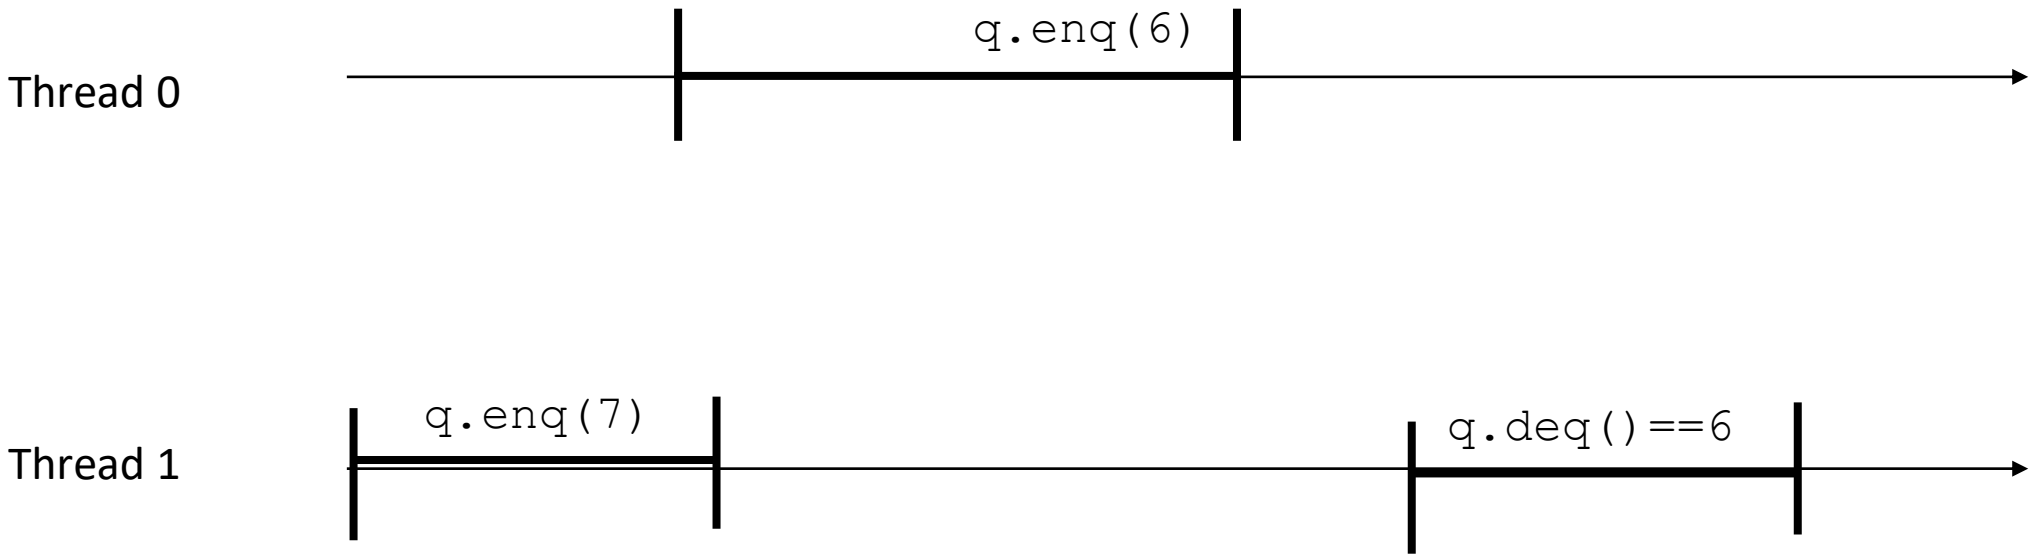
slider game!



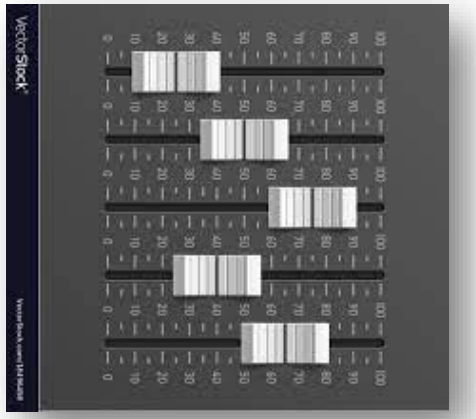
Linearizability

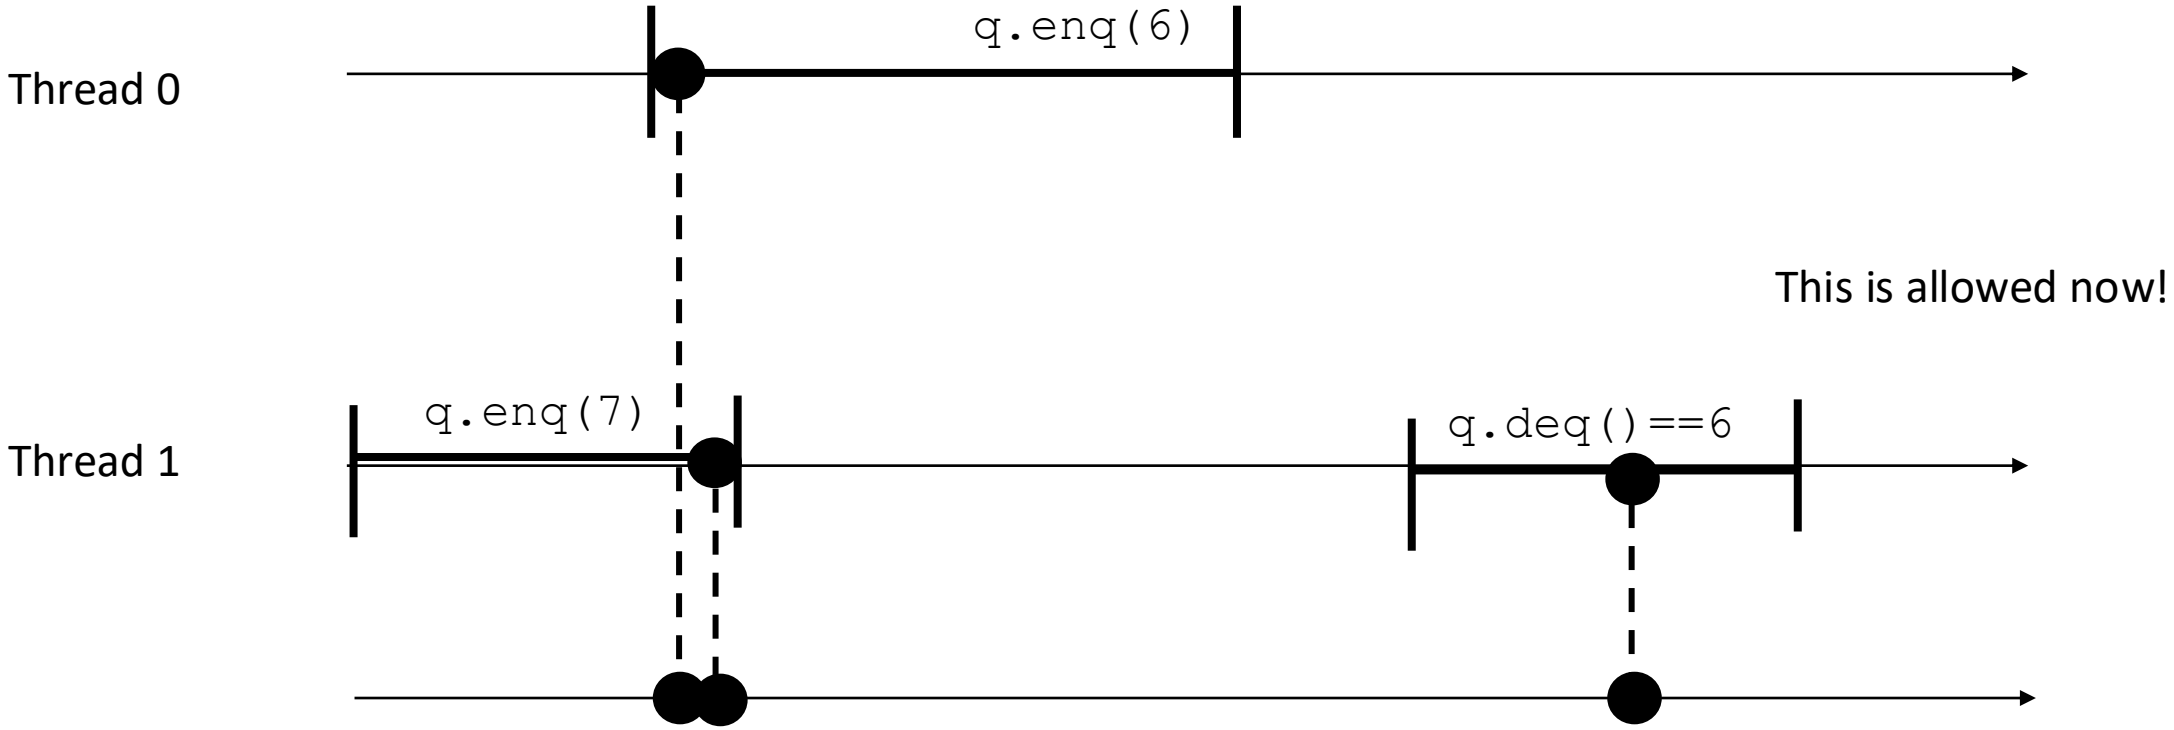


Linearizability

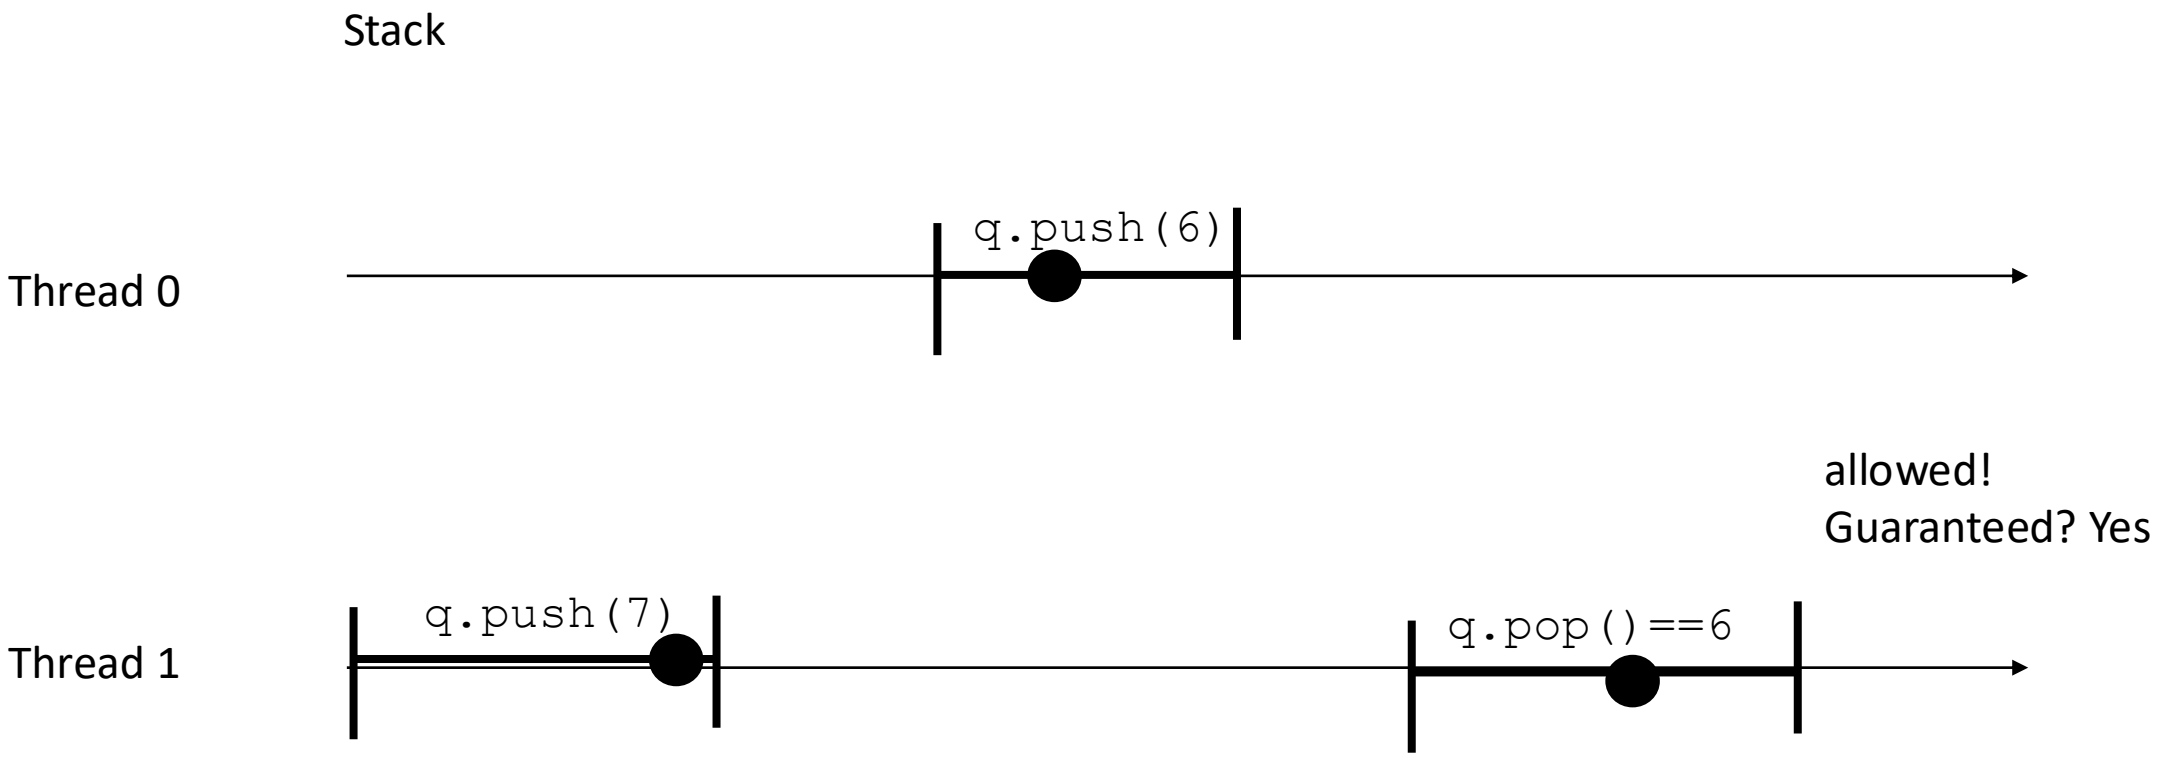


Linearizability

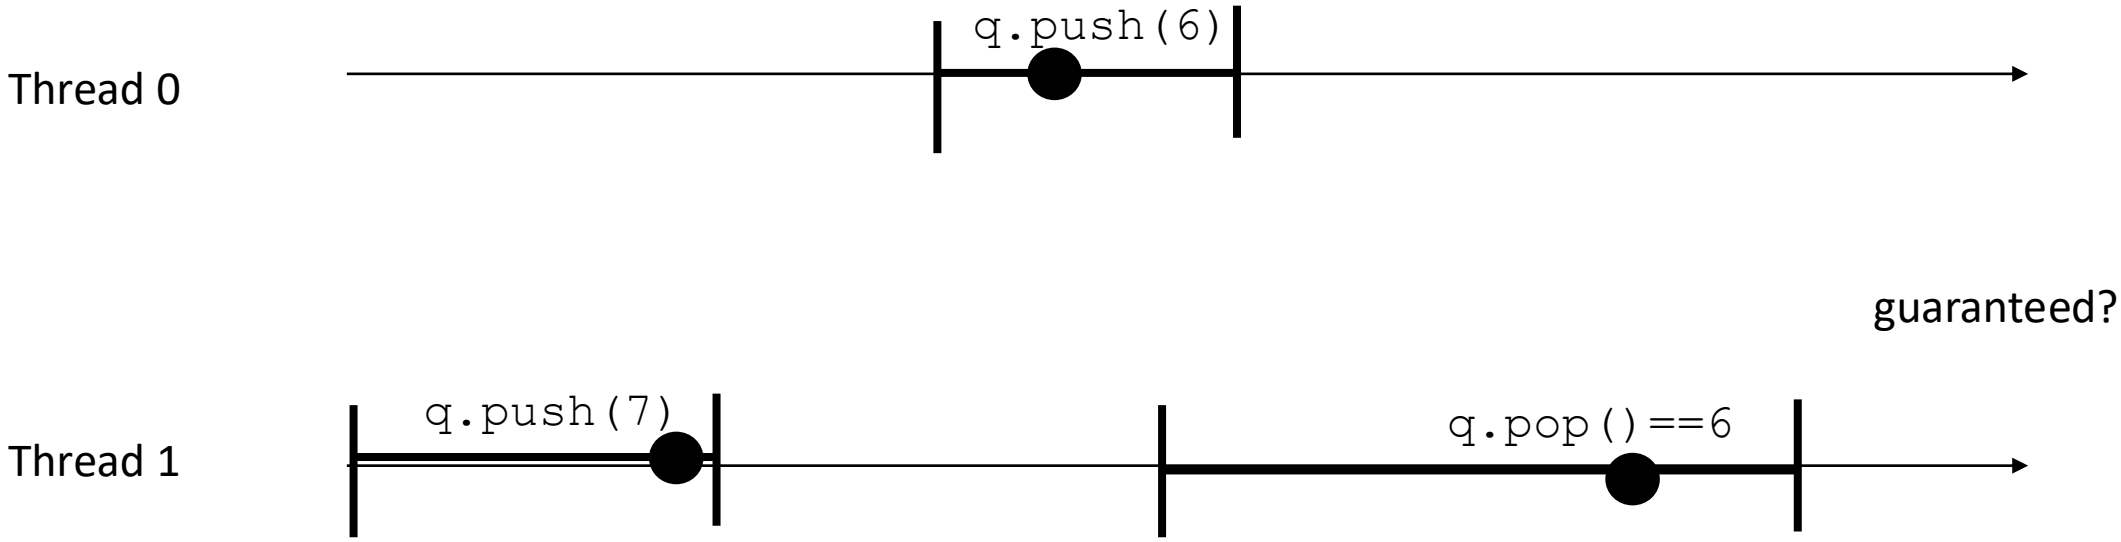




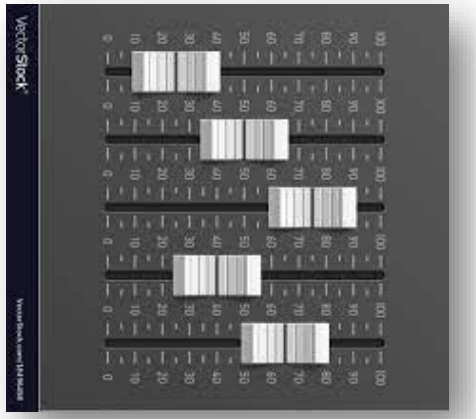
Linearizability

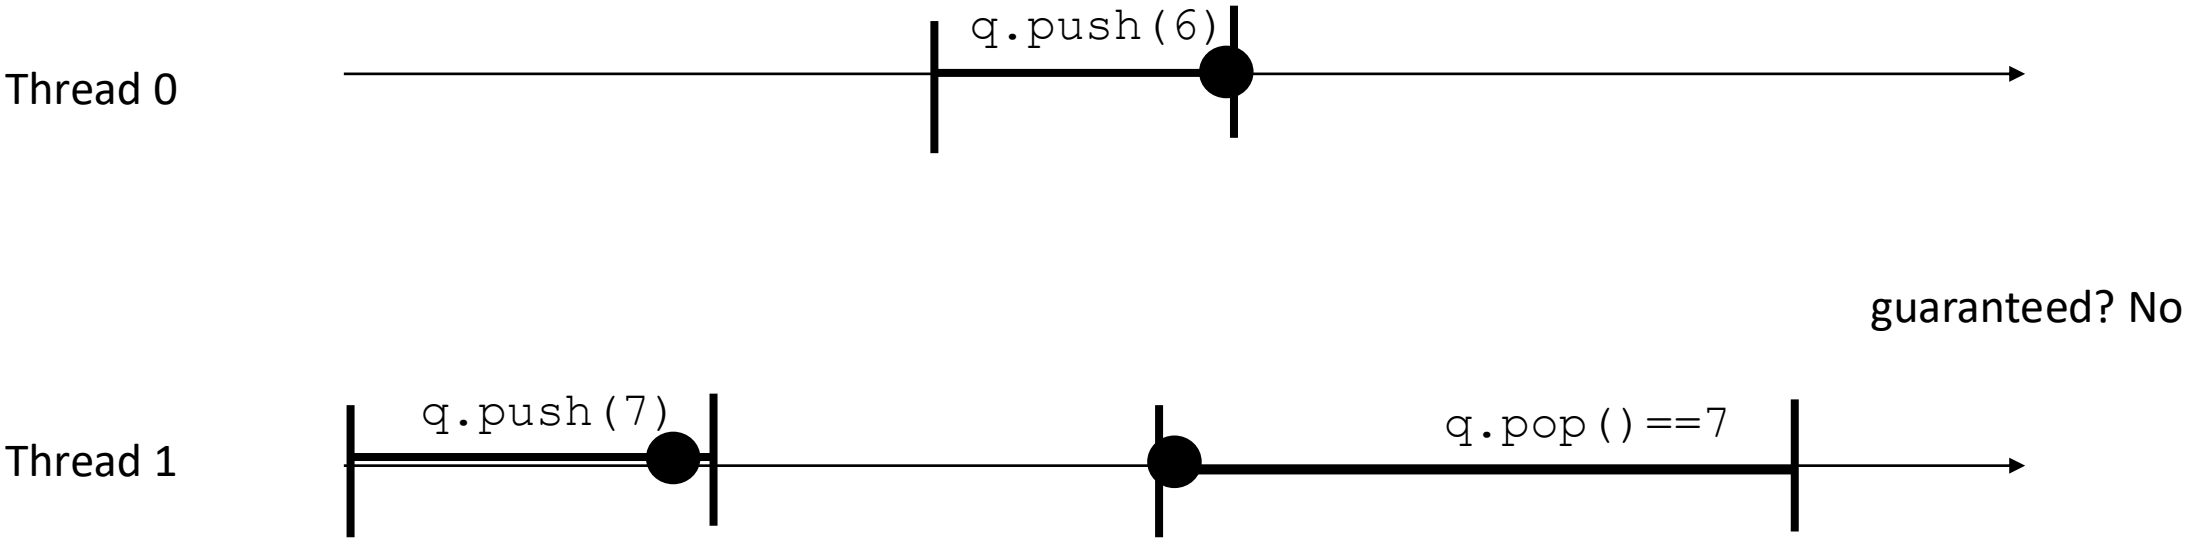


Linearizability



Linearizability





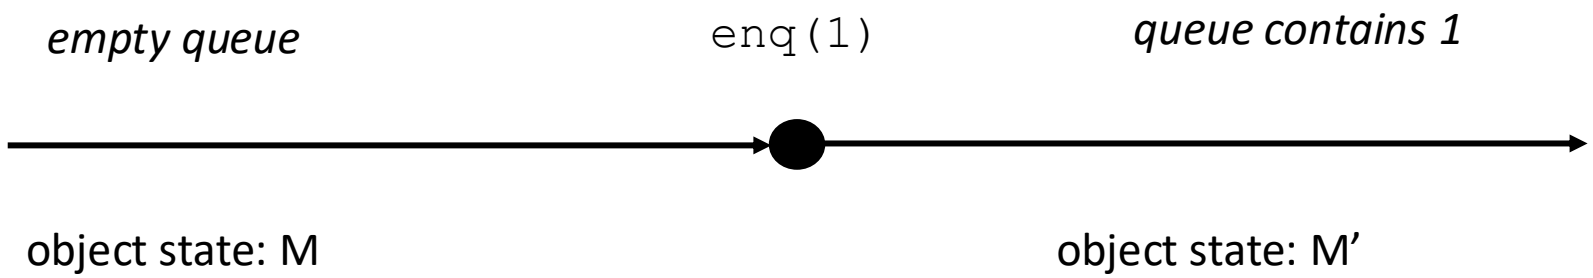
guaranteed? No

Linearizability

- We spent a bunch of time on SC... did we waste our time?
  - No!
  - Linearizability is strictly stronger than SC. Every linearizable execution is SC, but not the other way around.
  - If a behavior is disallowed under SC, it is also disallowed under linearizability.

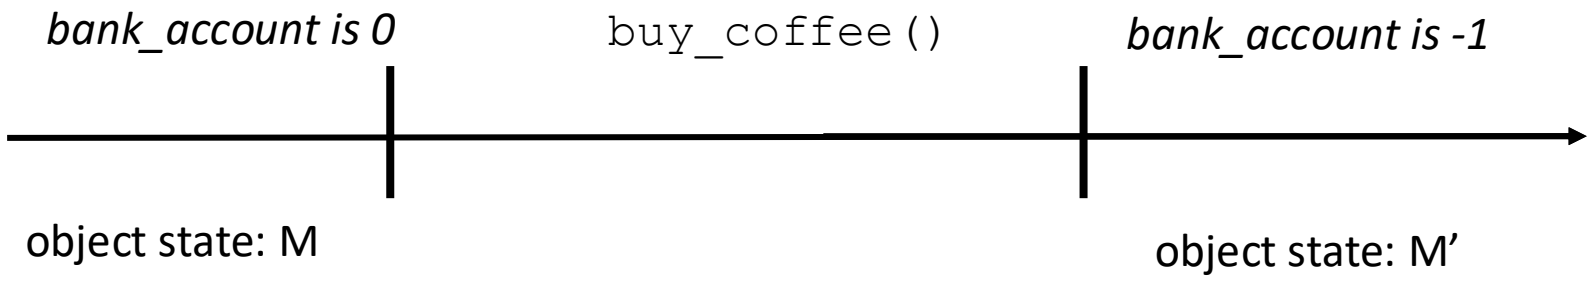
Providing Linearizability

- How do we write our programs to be linearizable?
  - Identify the linearizability point
  - One indivisible region (e.g. an atomic store, atomic load, atomic RMW, or critical section) where the method call takes effect. Modeled as a point.



Providing Linearizability

- Locked data structures are linearizable.




```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      m.lock();
      balance -= 1;
      m.unlock();
    }

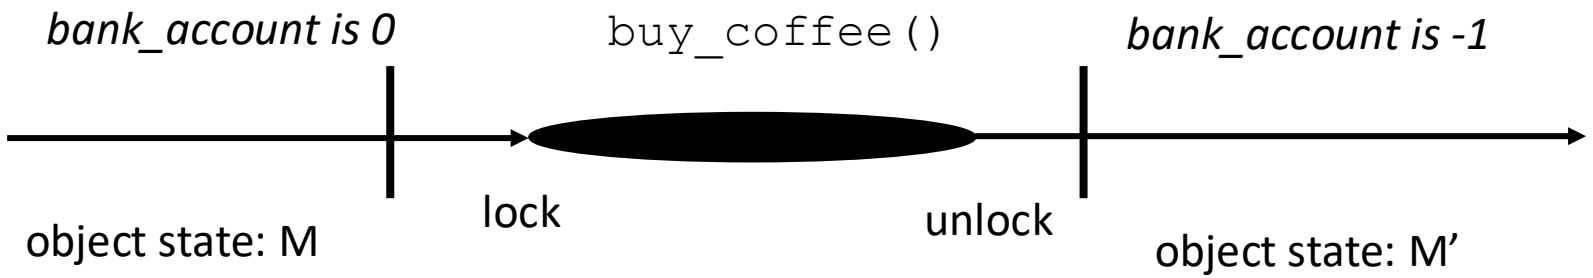
    void get_paid() {
      m.lock();
      balance += 1;
      m.unlock();
    }

  private:
    int balance;
    mutex m;
};
```

Linearizability

- Locked data structures are linearizable.

typically modeled as the point the lock is acquired or released




```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      m.lock();
      balance -= 1;
      m.unlock();
    }

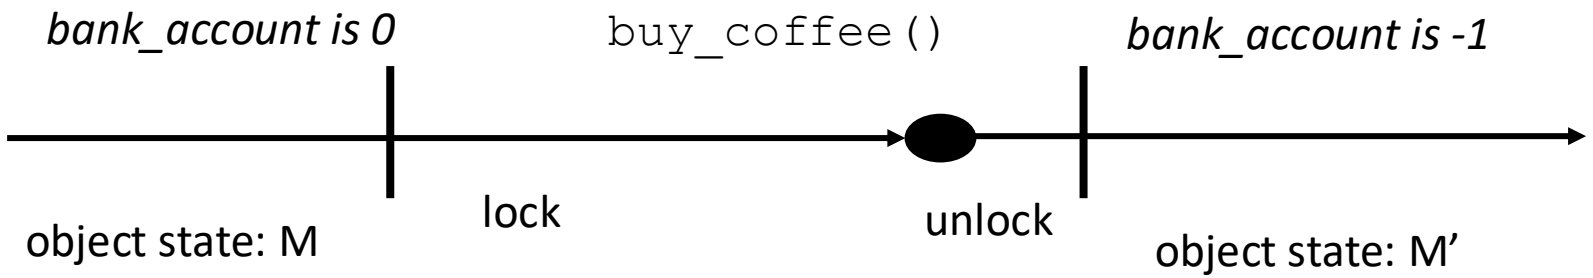
    void get_paid() {
      m.lock();
      balance += 1;
      m.unlock();
    }

  private:
    int balance;
    mutex m;
};
```

Linearizability

- Locked data structures are linearizable.

typically modeled as the point the lock is acquired or released
let's say released.




```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      m.lock();
      balance -= 1;
      m.unlock();
    }

    void get_paid() {
      m.lock();
      balance += 1;
      m.unlock();
    }

  private:
    int balance;
    mutex m;
};
```

Linearizability

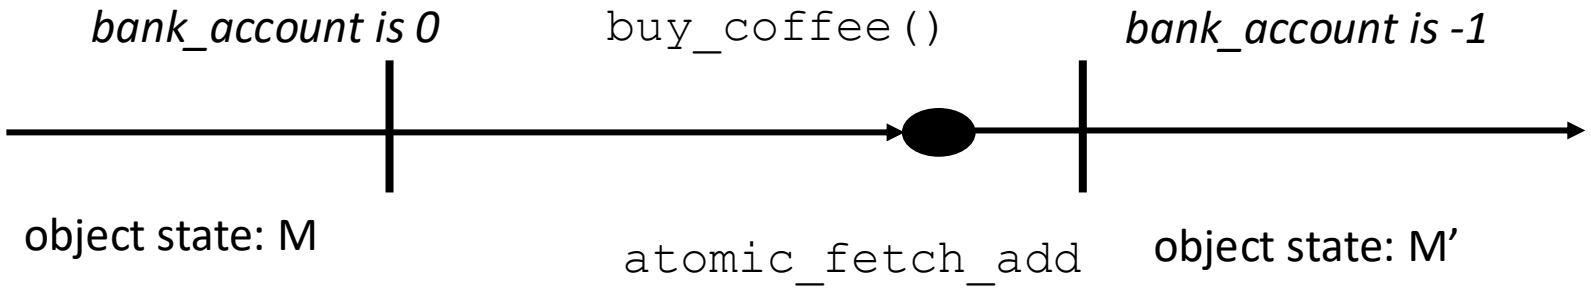
- Our lock-free bank account is linearizable:
  - The atomic operation is the linearizable point

```
class bank_account {
  public:
    bank_account() {
      balance = 0;
    }

    void buy_coffee() {
      atomic_fetch_add(&balance, -1);
    }

    void get_paid() {
      atomic_fetch_add(&balance, 1);
    }

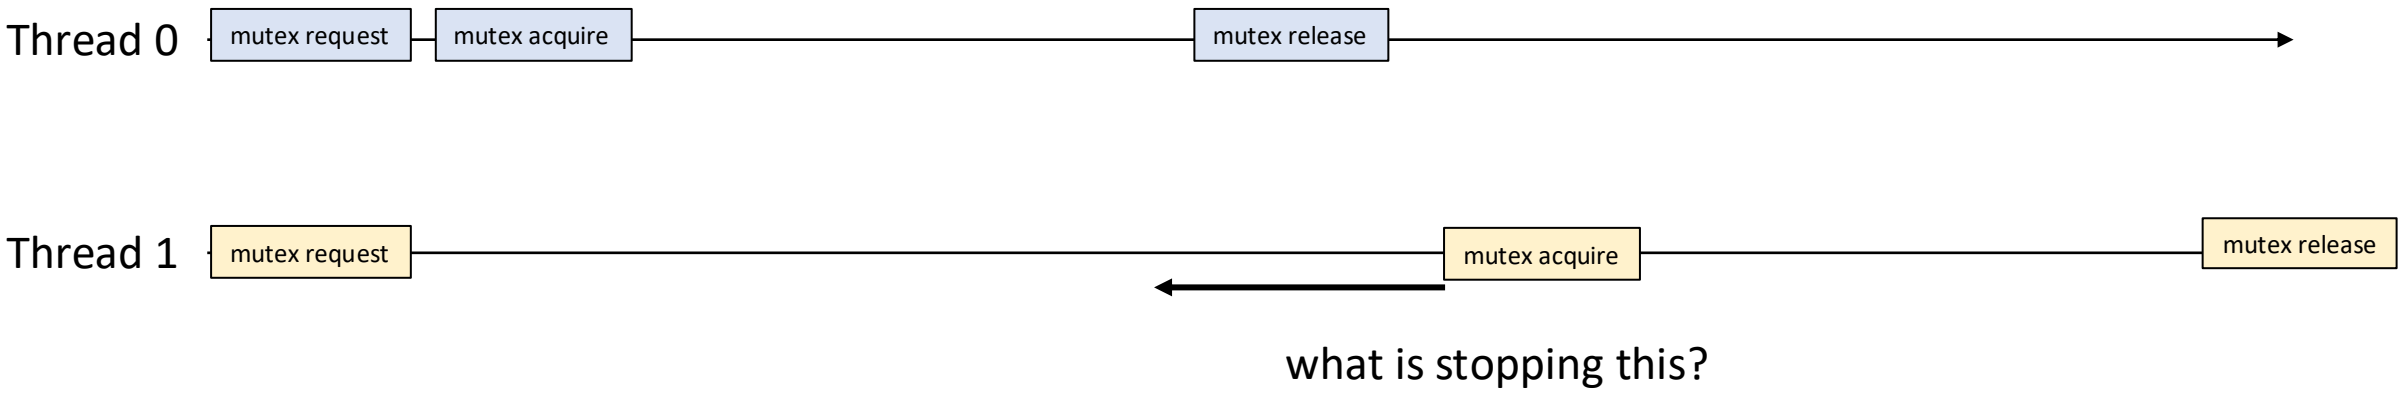
  private:
    atomic_int balance;
};
```



Progress properties

- Going back to specifications:

Recall the mutex

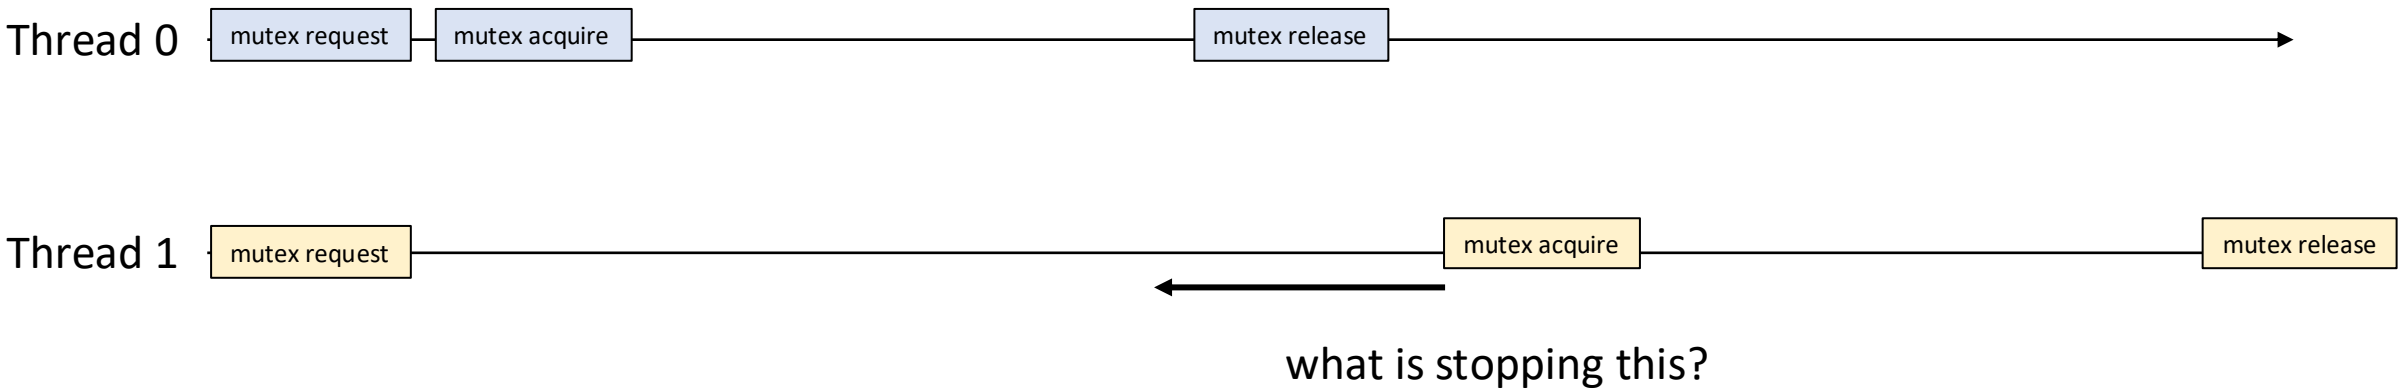


Progress properties

- Going back to specifications:

Recall the mutex

Thread 0 is stopping Thread 1 from making progress.
If delays in one thread can cause delays in other threads, we say that it is blocking

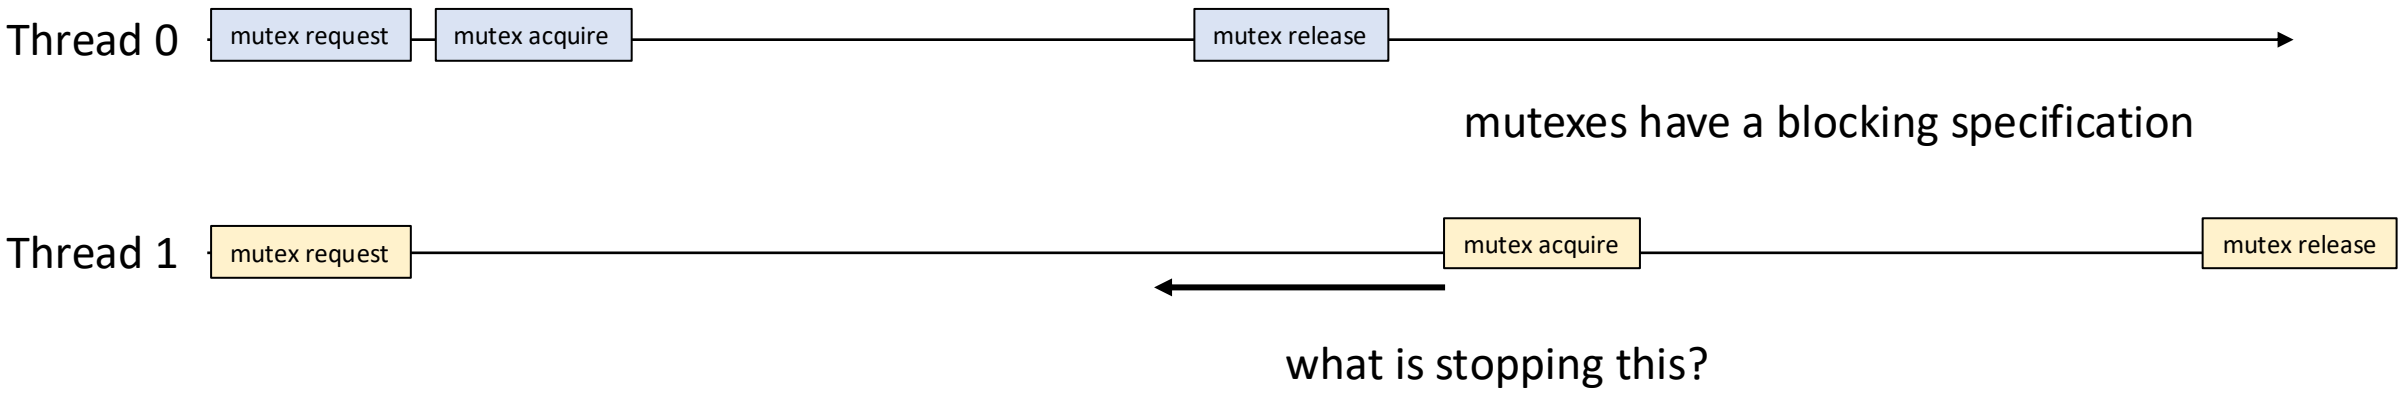


Progress properties

- Going back to specifications:

Thread 0 is stopping Thread 1 from making progress.
If delays in one thread can cause delays in other threads, we say that it is blocking

Recall the mutex

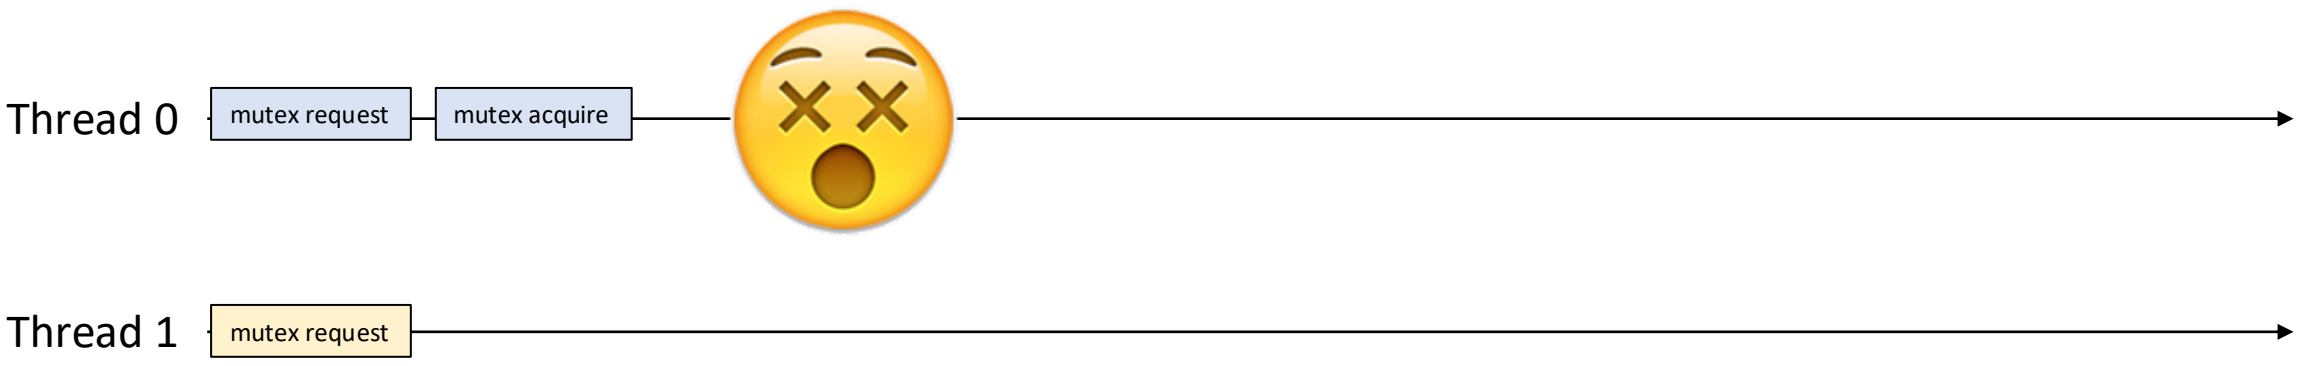


Progress properties

- Going back to specifications:

Thread 0 is stopping Thread 1 from making progress.
If delays in one thread can cause delays in other threads, we say that it is blocking

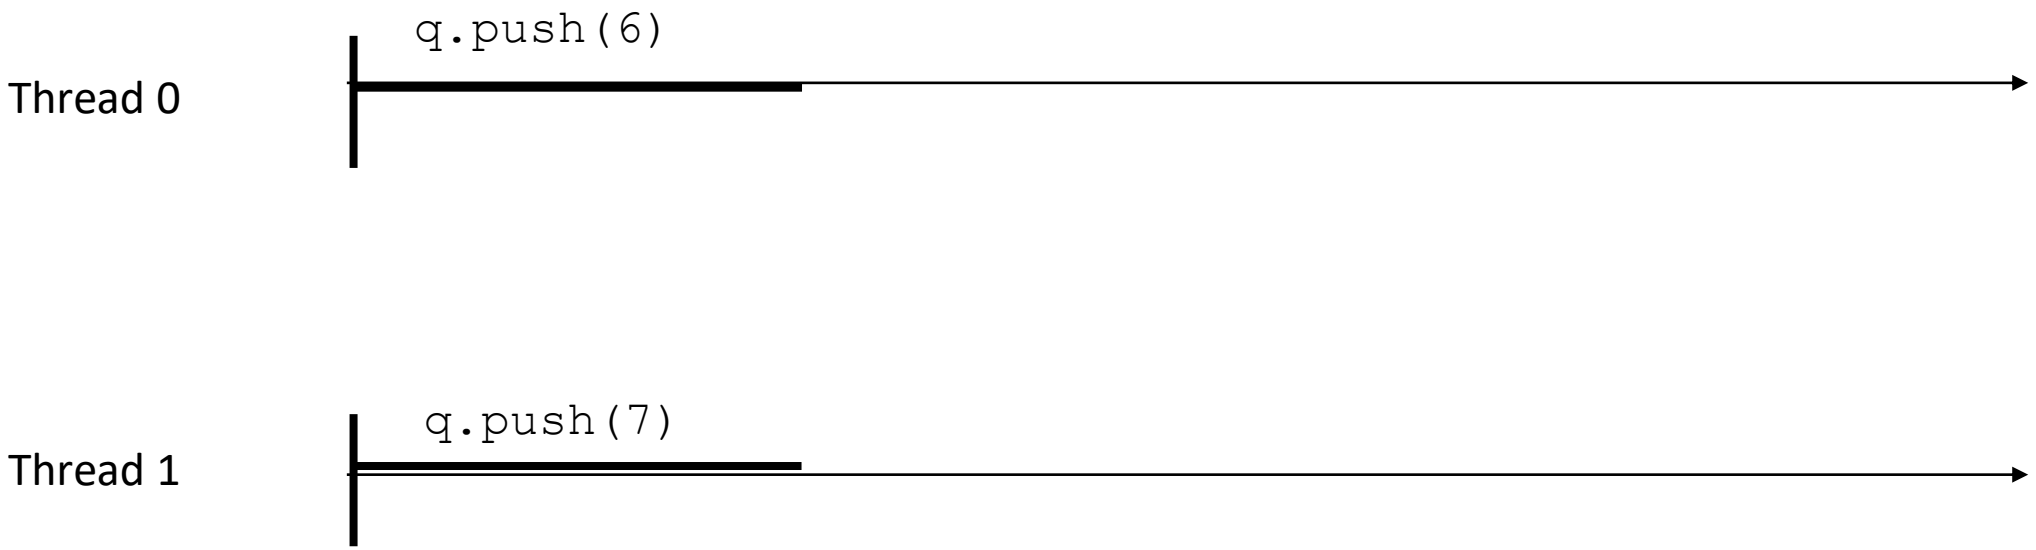
Recall the mutex



What now?!

Linearizability

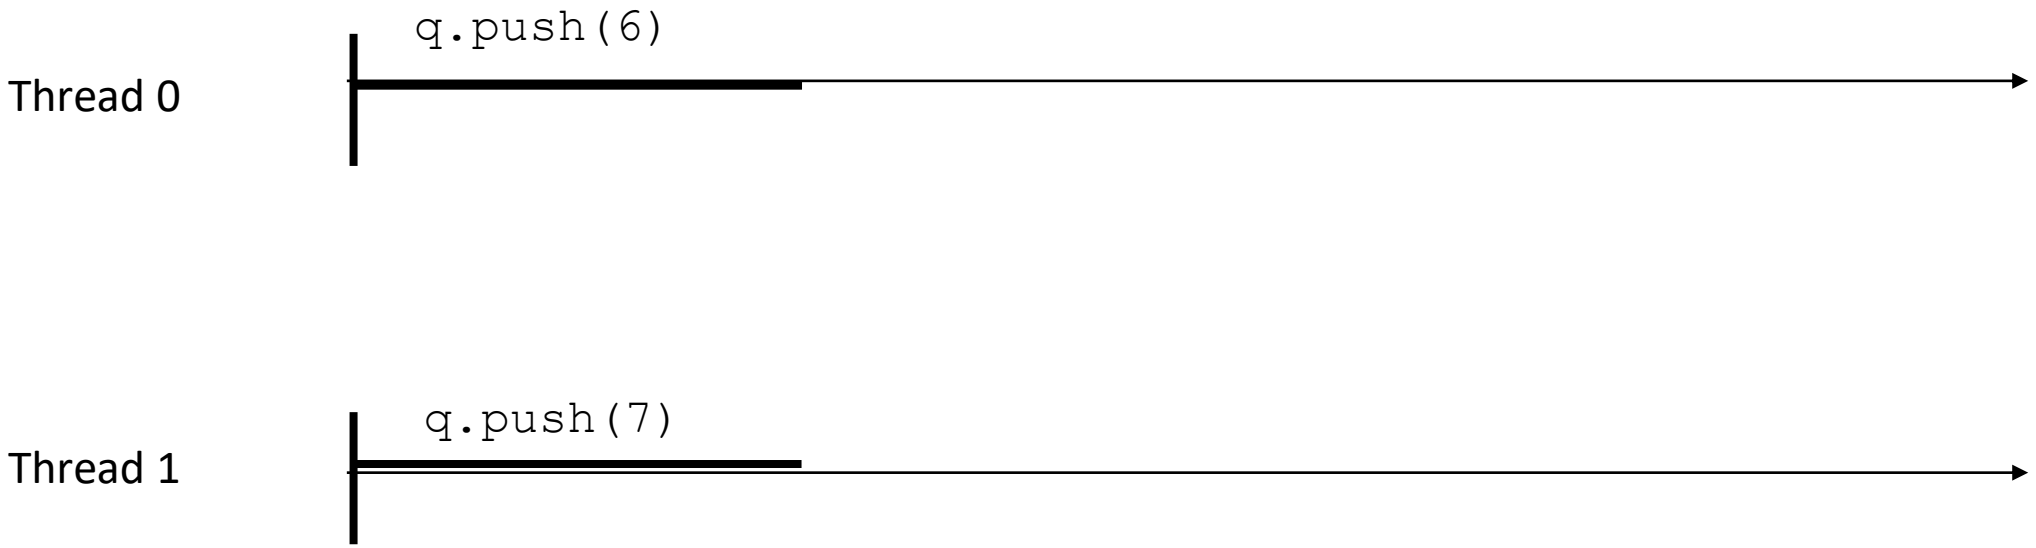
Two unfinished commands.



Linearizability

Two unfinished commands.

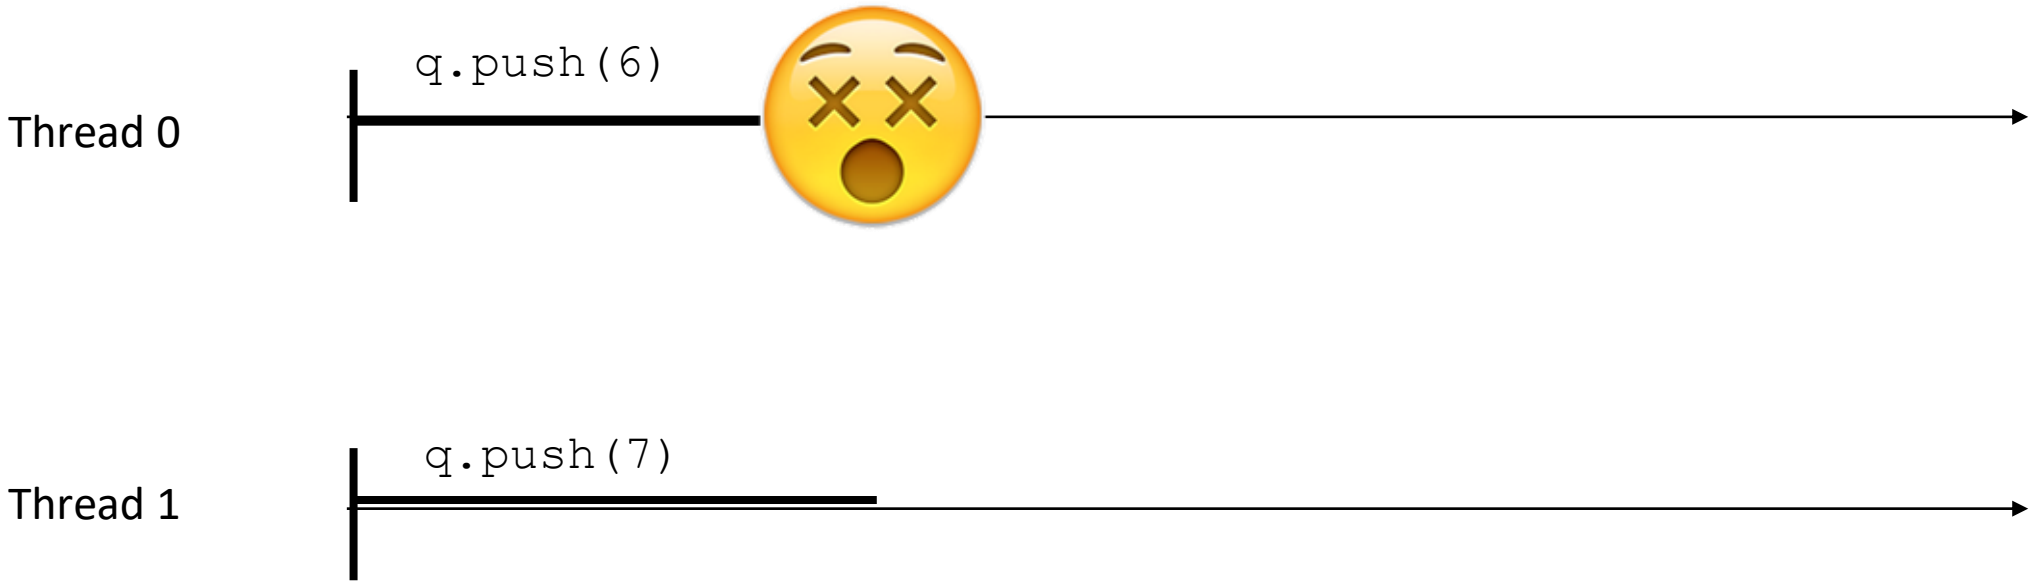
Linearizability does not dictate that one needs to wait for another



Linearizability

Two unfinished commands.

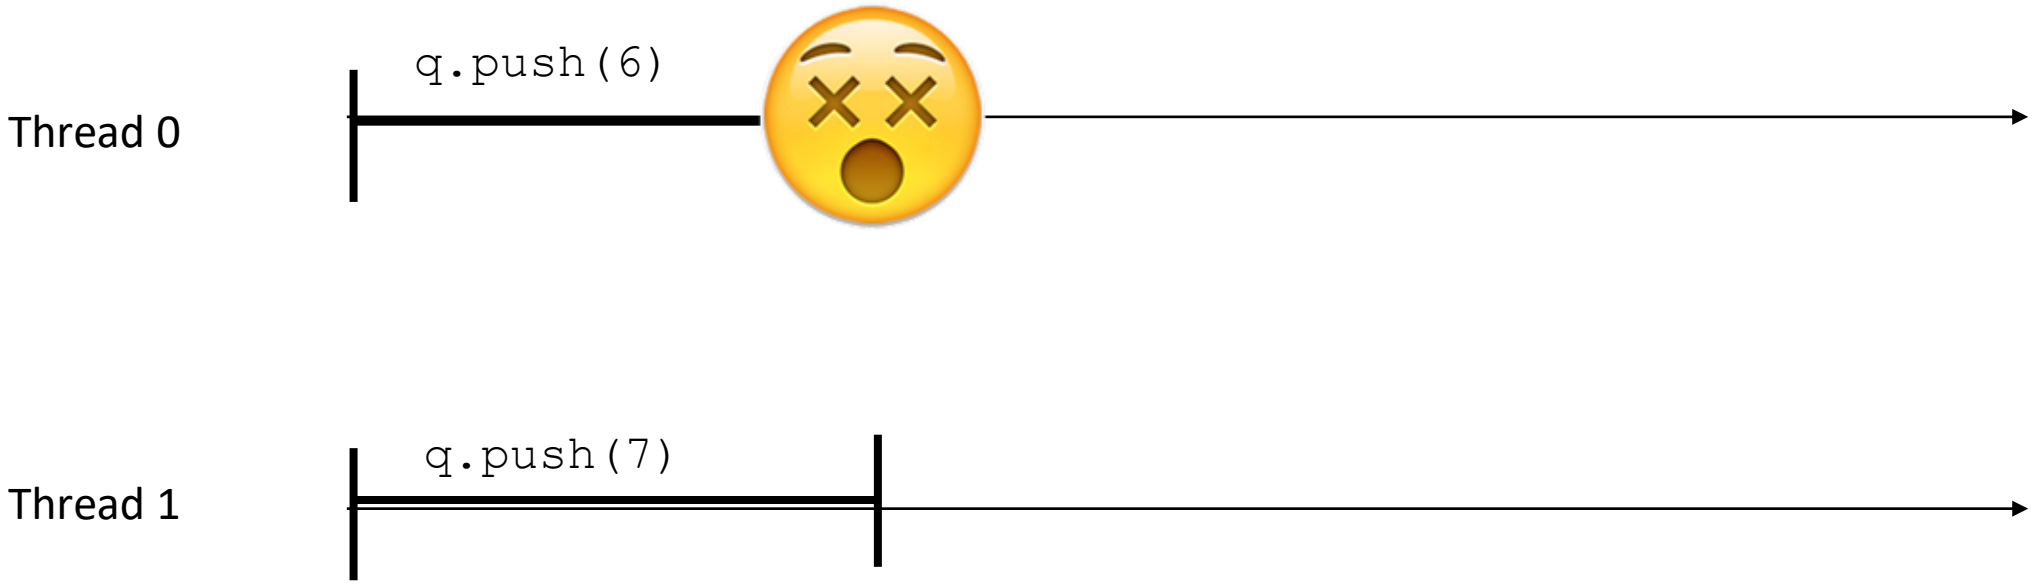
Linearizability does not dictate that one needs to wait for another



for mutexes, the specification required that the system hang.
no such specification here.

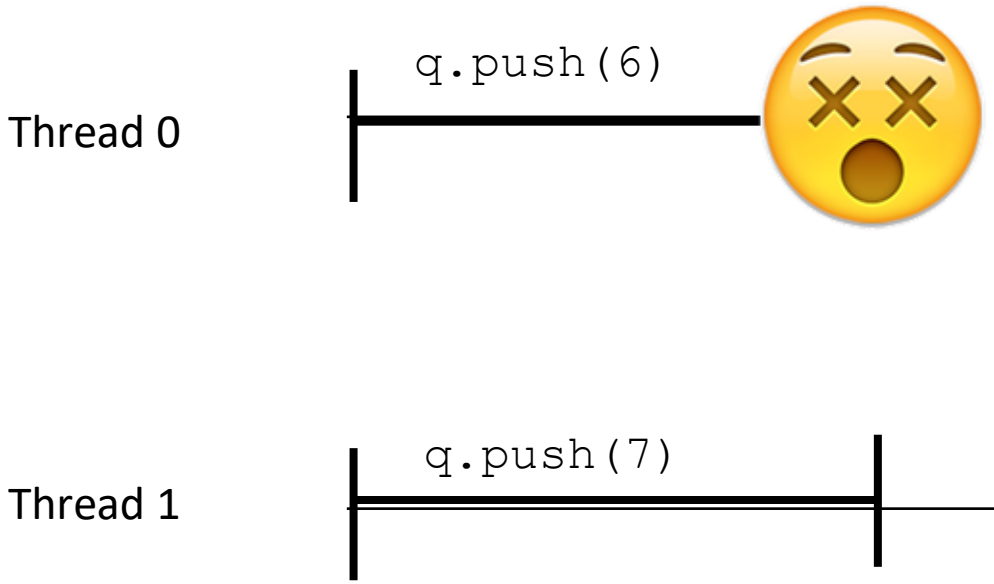
Linearizability

Non-blocking specification:
Every thread is allowed to continue executing
REGARDLESS of the behavior of other threads



Linearizability

Non-blocking specification:
Every thread is allowed to continue executing
REGARDLESS of the behavior of other threads



This is a specification property, not an implementation property! You can implement your concurrent objects with locks and have a “blocking implementation”.

But that is because of implementation choice, not because of specification requirements.

Terminology overview

- Thread-safe implementation:
- Sequential consistency:
- Linearizability:
- Linearizability point:

Terminology overview

- Lock-free implementation:
- (Non-)blocking specification:
- (non-)blocking implementation:

Concurrent Queues

- List of items, accessed in a first-in first-out (FIFO) way
- *duplicates allowed*
- Methods
  - **`enq(x)`** put **`x`** in the list at the end
  - **`deq()`** remove the item at the front of the queue and return it.
  - **`size()`** returns how many items are in the queue

Concurrent Queues

- General implementation given in Chapter 10 of the book.
- Similar types of reasoning as the linked list
  - Lots of reasoning about node insertion, node deletion
  - Using atomic RMWs (CAS) in clever ways

- We will think about specialized queues
  - Implementations can be simplified!

Quiz 9.

Write a few sentences about the pros and cons of using a specialized concurrent queue (e.g. an IO queue) and a fully general concurrent queue.

Free response: a few sentences on the form

Scan to answer

forms.gle/fBRaHAzLLFsHZqJq8

Three Variants of Concurrent Queues

- Input/Output Queues
  - Multiple threads enqueue, and multiple threads dequeue but not both
- Producer/Consumer Queues
  - 1 thread enqueues, 1 thread dequeues
  - Two variants:
    - Synchronous
    - Asynchronous

Input/Output Queues

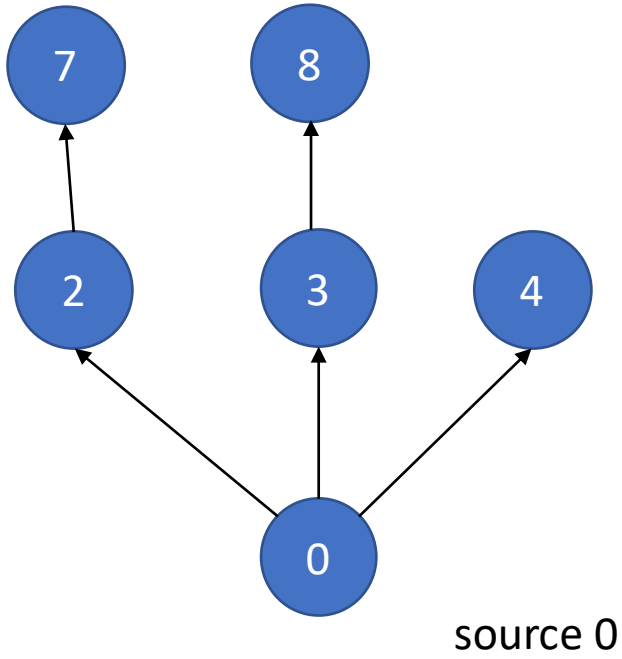
- Queue in which multiple threads read (deq), or write (enq), but not both.
- Why would we want a thing?
- Computation done in phases:
  - First phase prepares the queue (by writing into it)
  - All threads join
  - Second phase reads values from the queue.

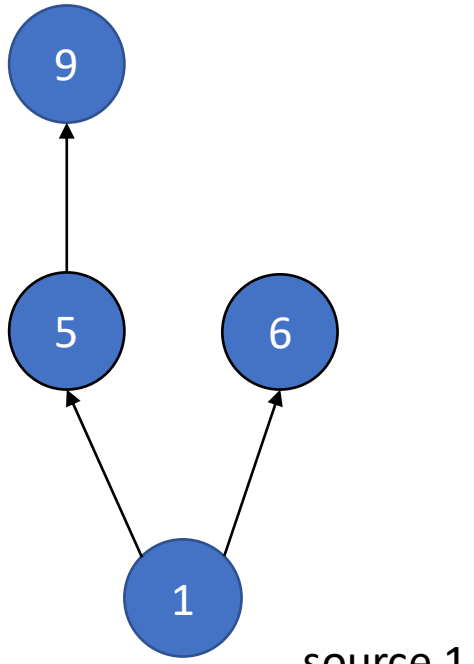
Input/Output Queues

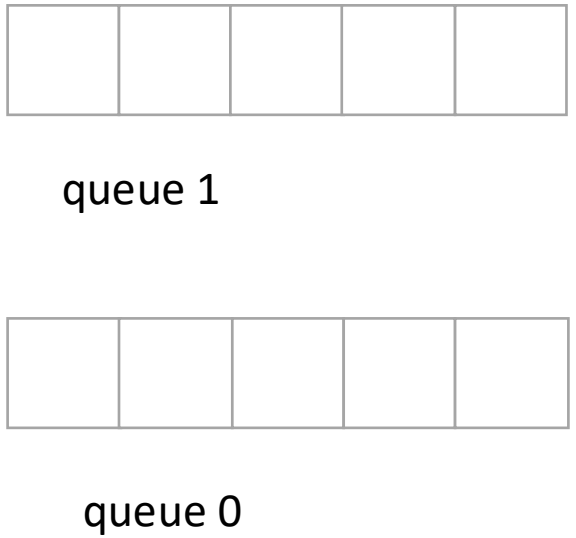
- Example: Information flow in graph applications:

Input/Output Queues

- Example: Information flow in graph applications:

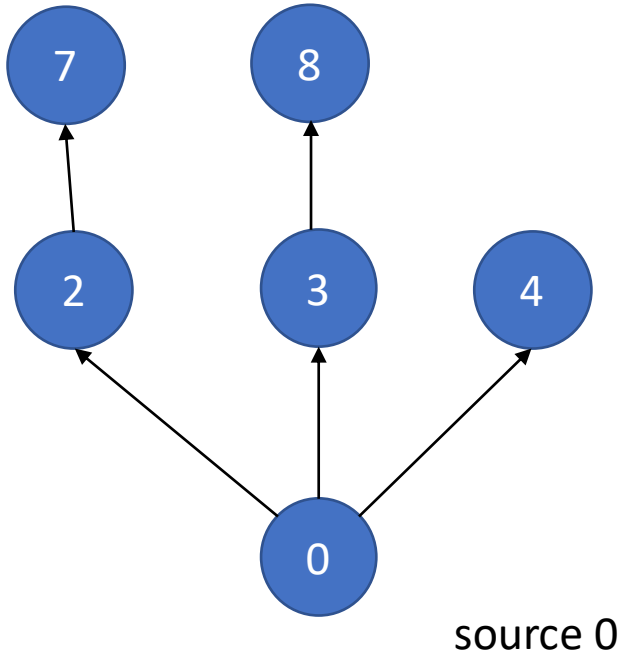


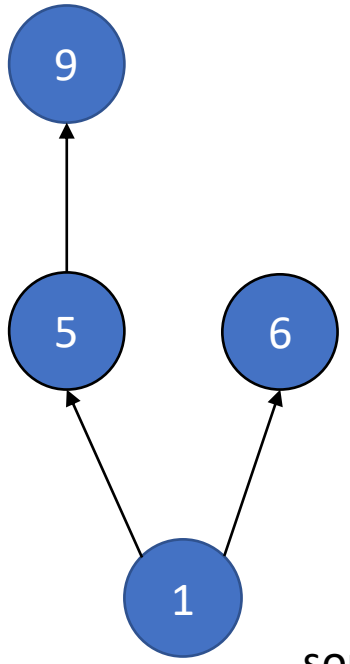


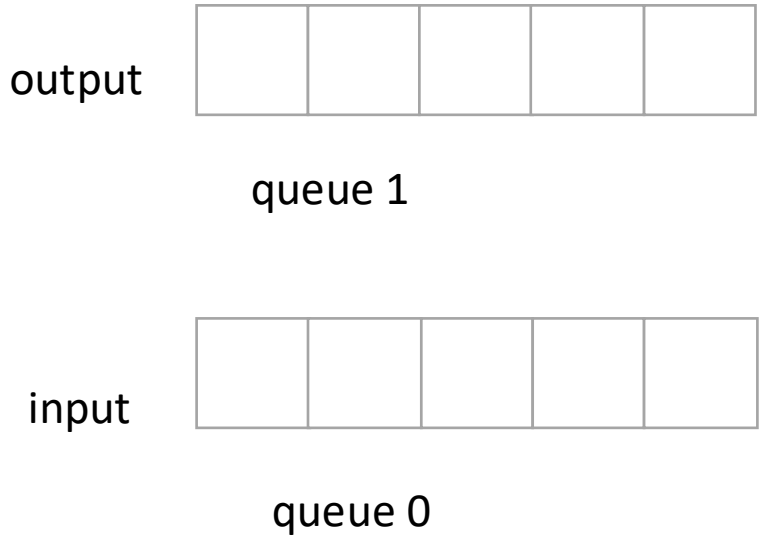


Input/Output Queues

- Example: Information flow in graph applications:

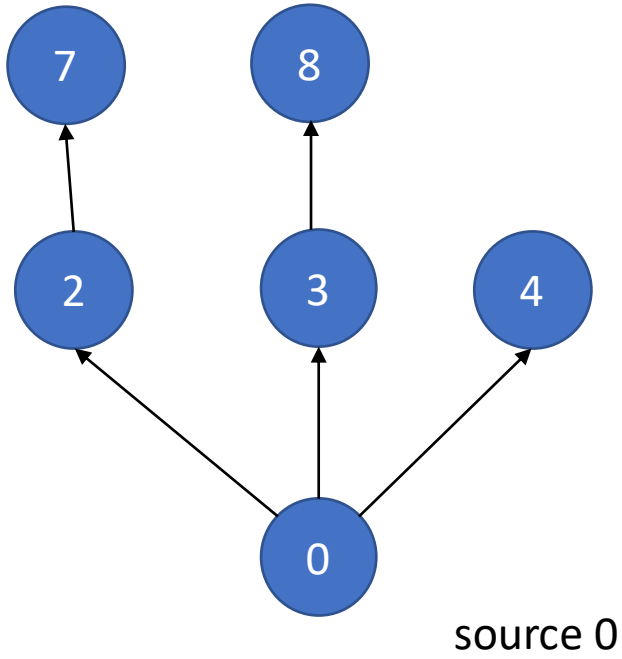


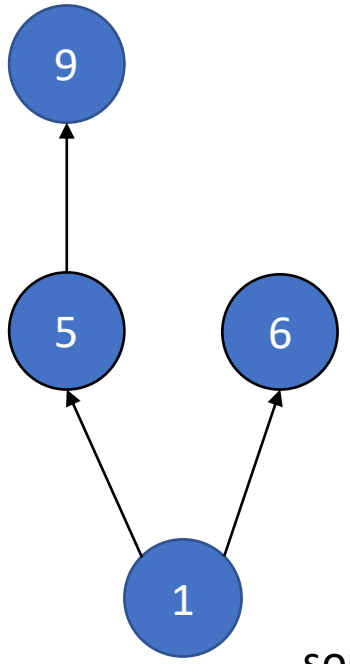


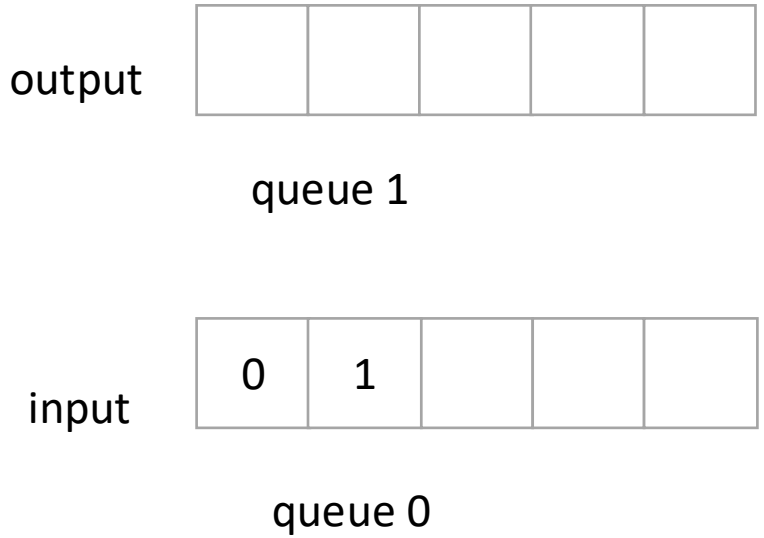


Input/Output Queues

- Example: Information flow in graph applications:

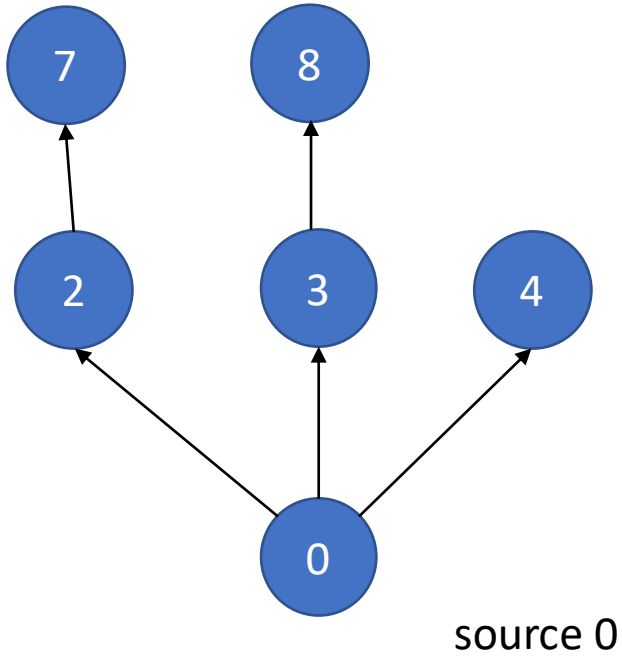


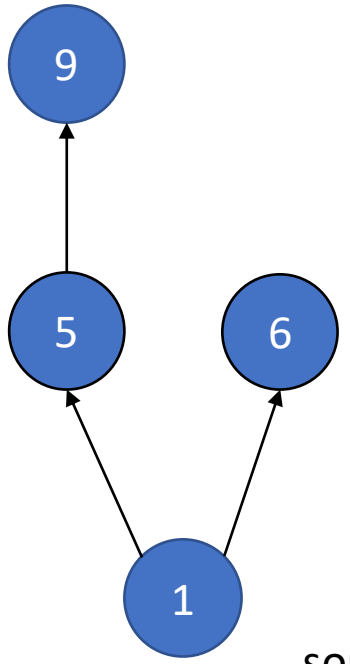


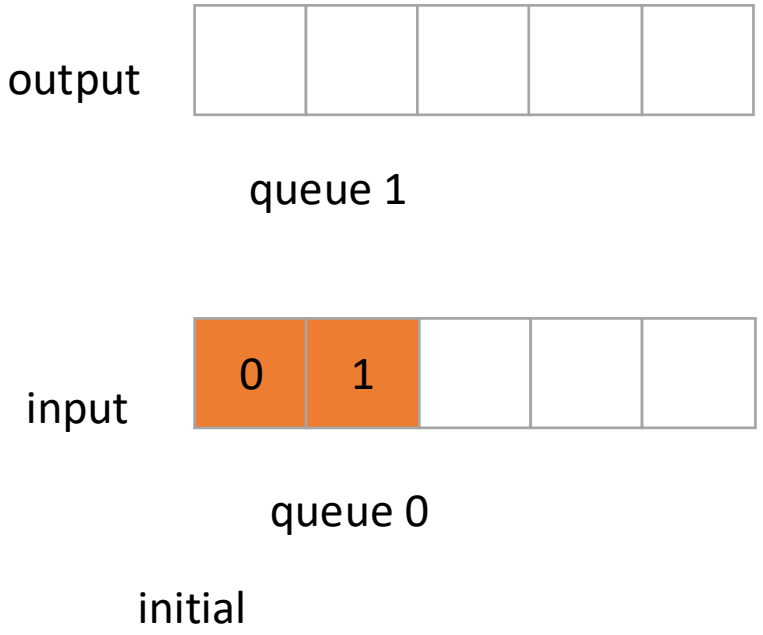


Input/Output Queues

- Example: Information flow in graph applications:

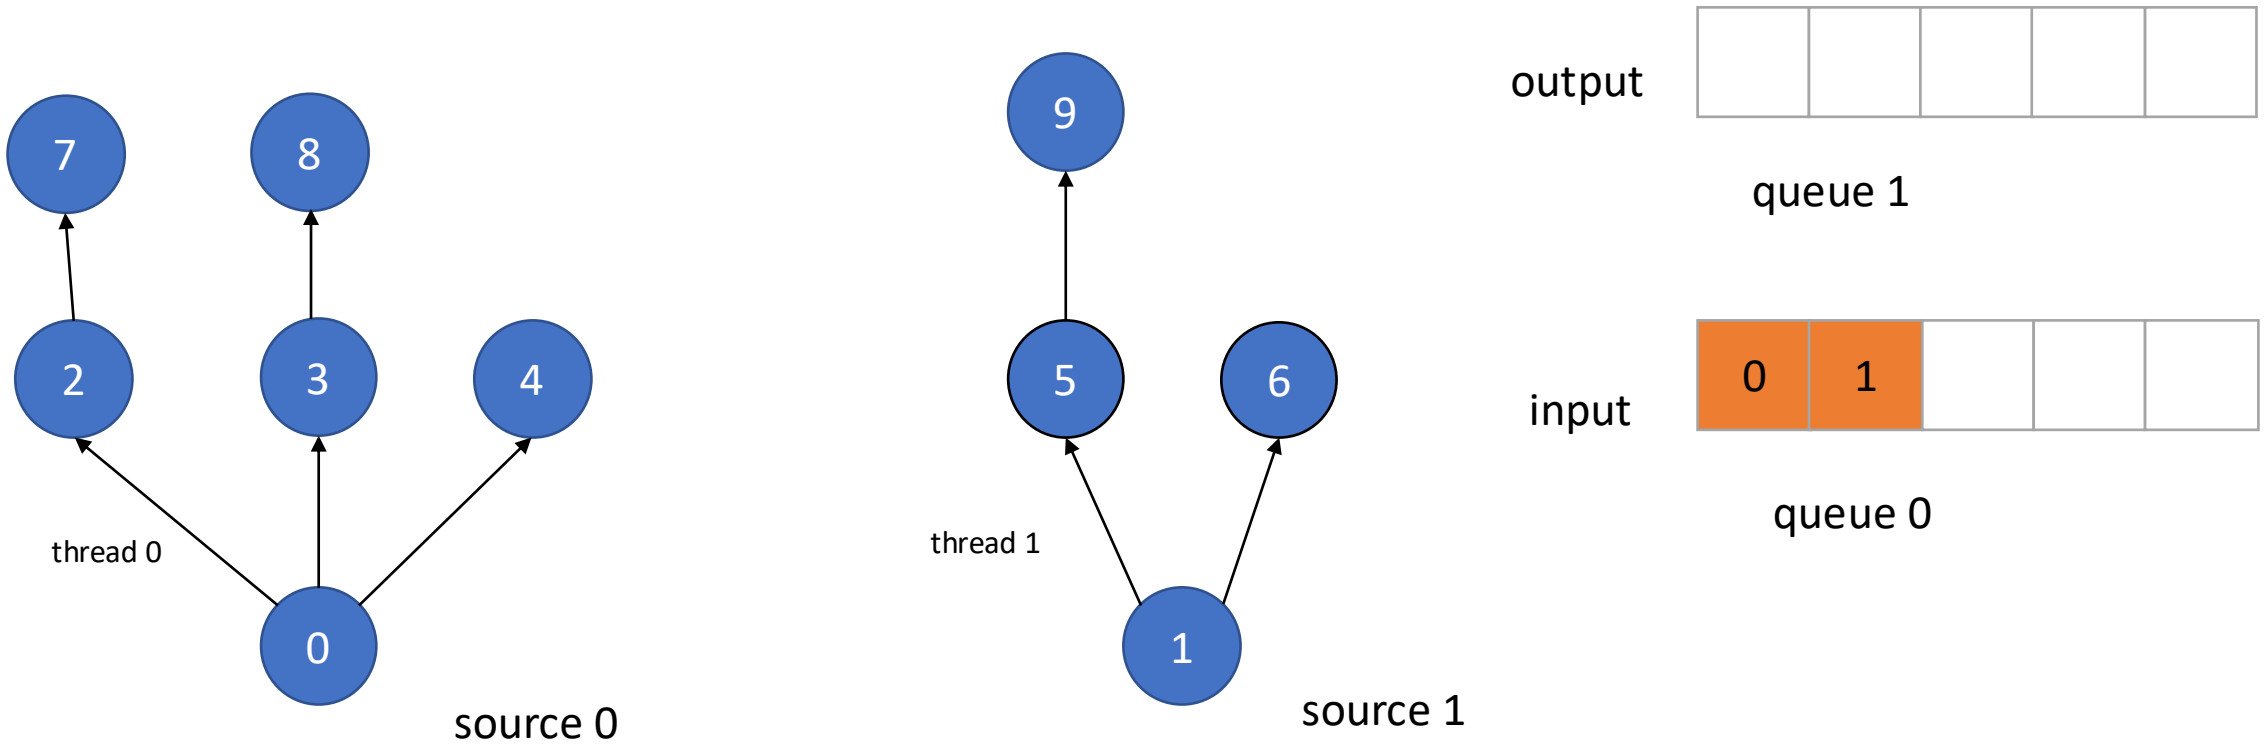






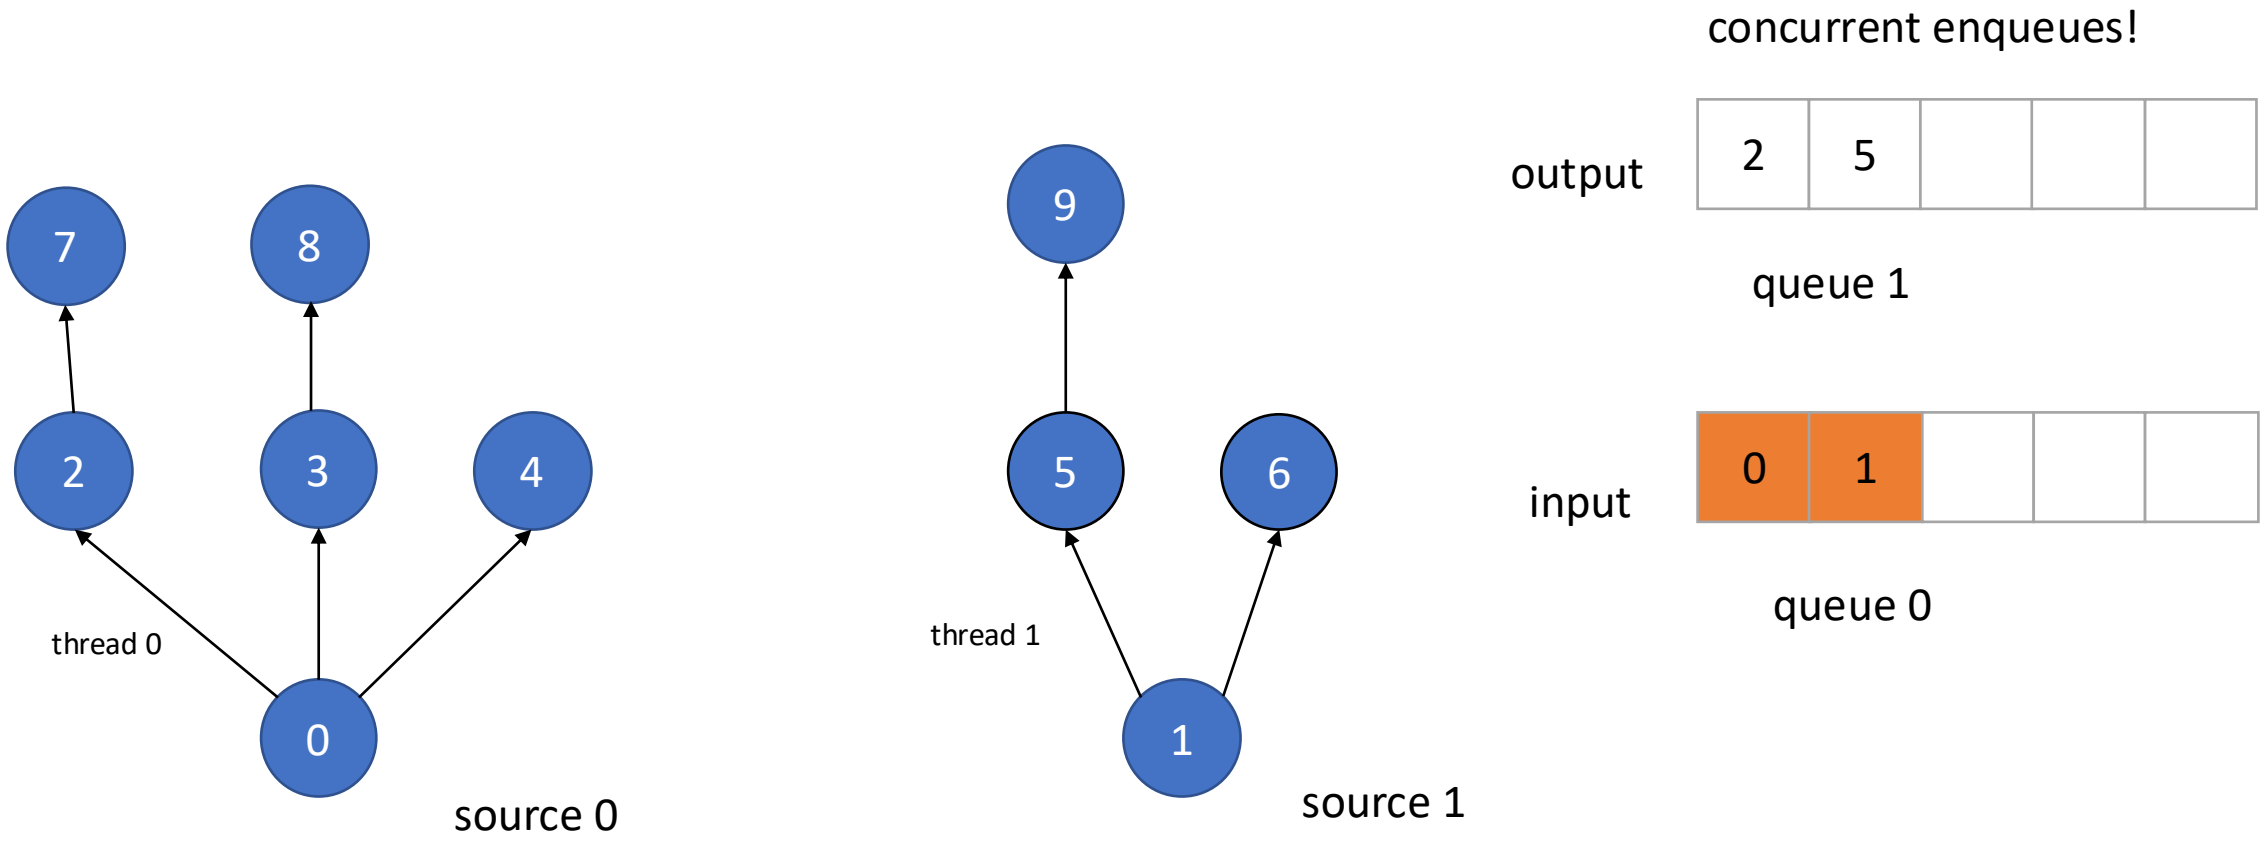
Input/Output Queues

- Example: Information flow in graph applications:



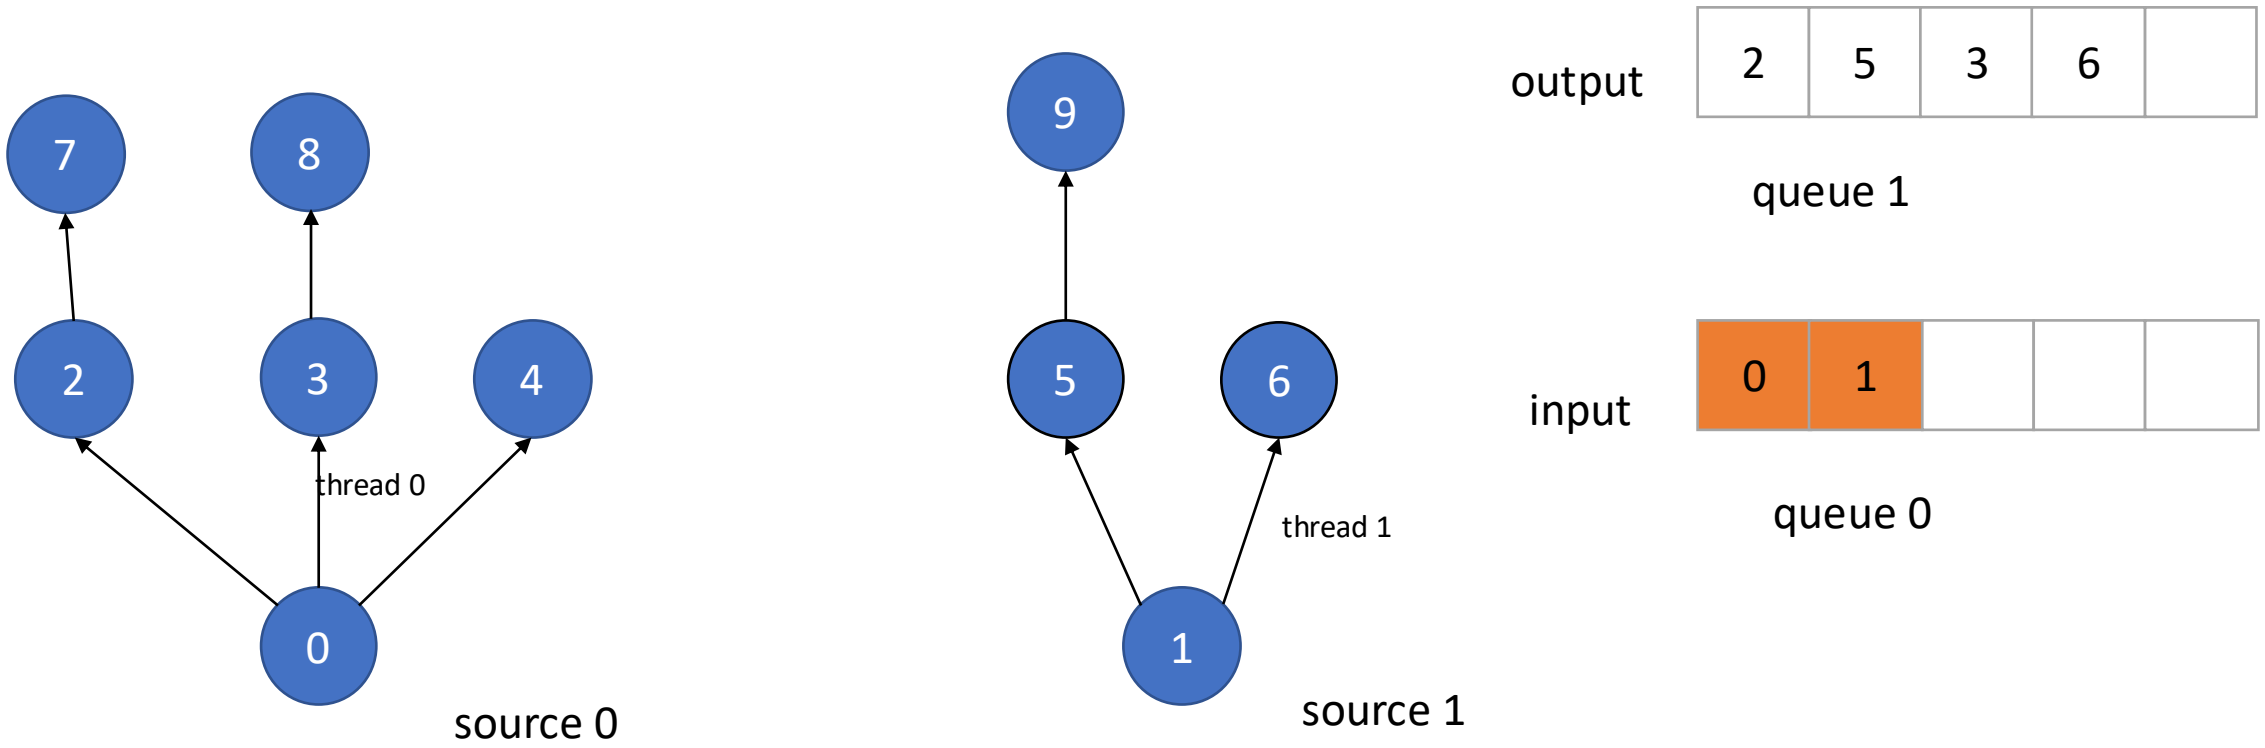
Input/Output Queues

- Example: Information flow in graph applications:



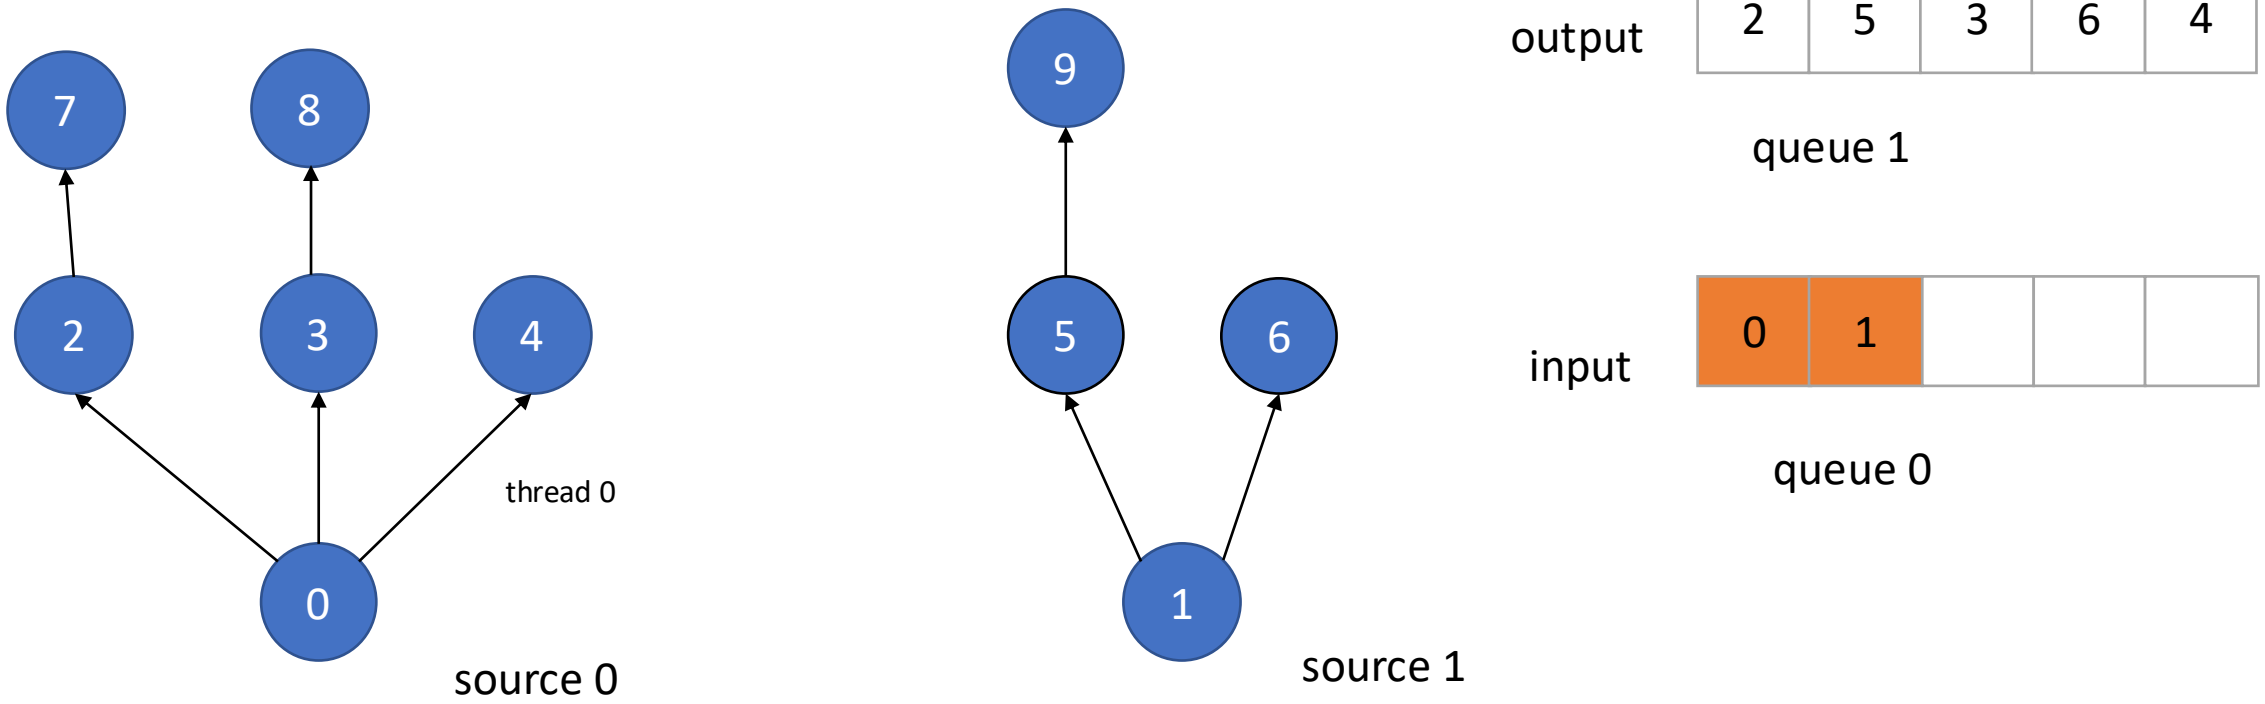
Input/Output Queues

- Example: Information flow in graph applications:



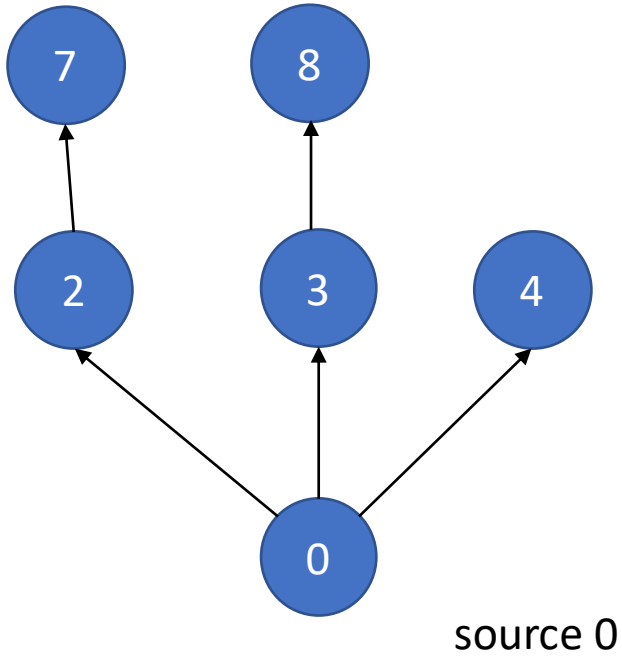
Input/Output Queues

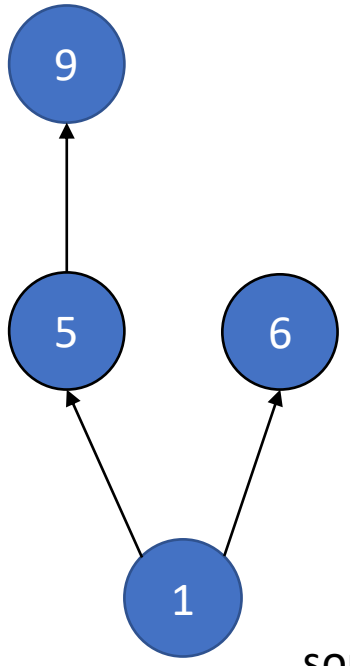
- Example: Information flow in graph applications:

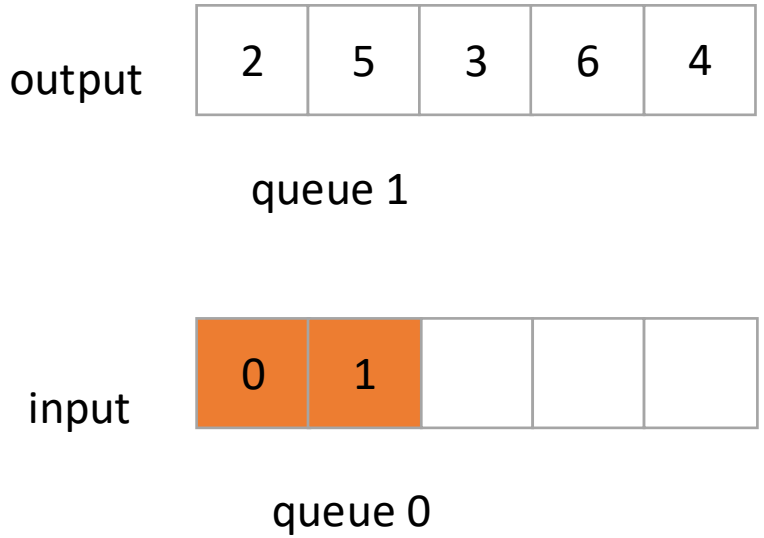


Input/Output Queues

- Example: Information flow in graph applications:

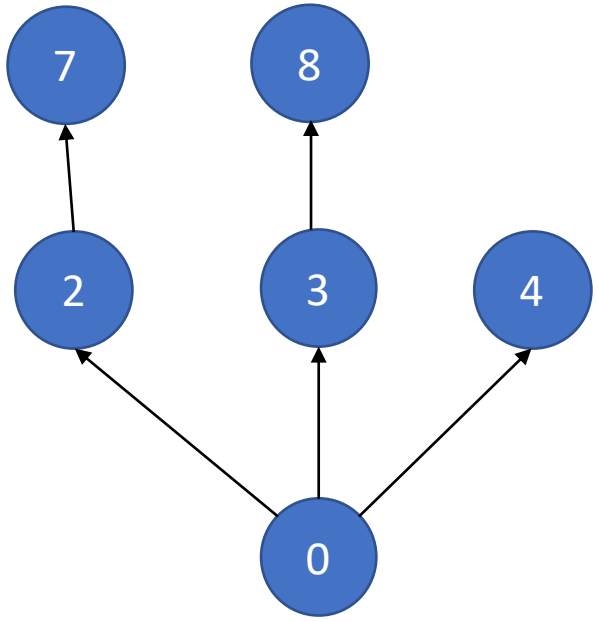




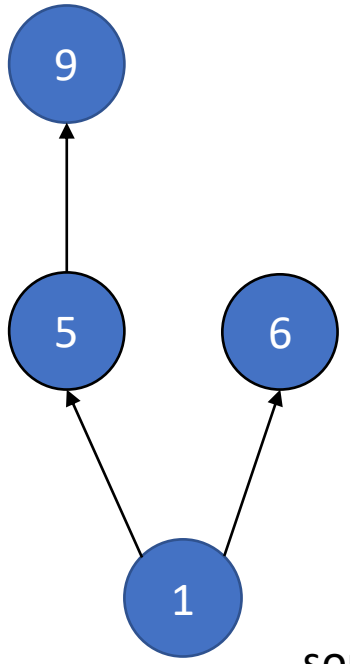


Input/Output Queues

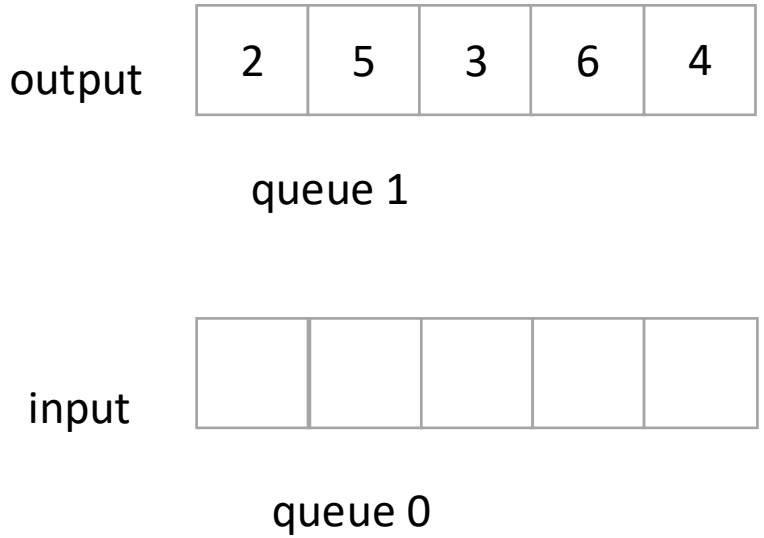
- Example: Information flow in graph applications:



source 0



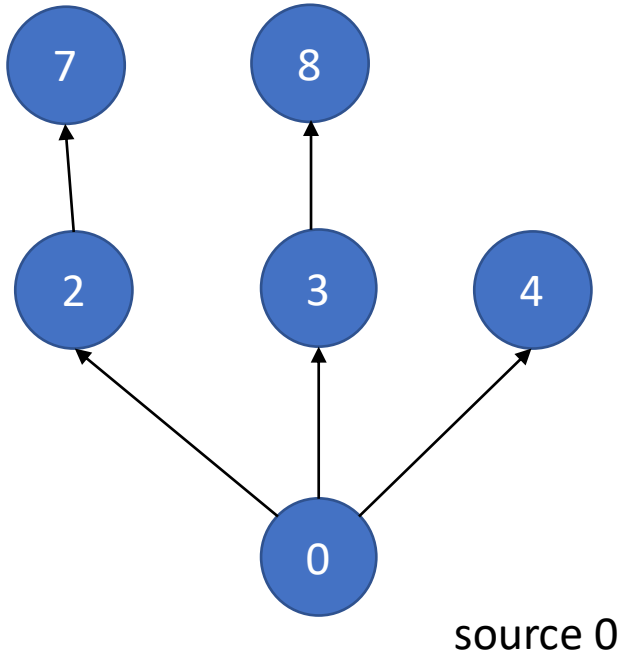
source 1

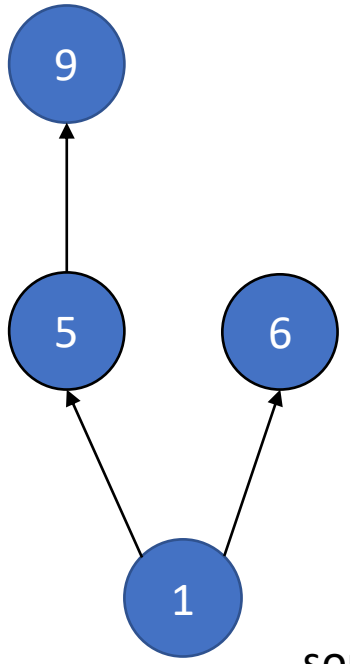


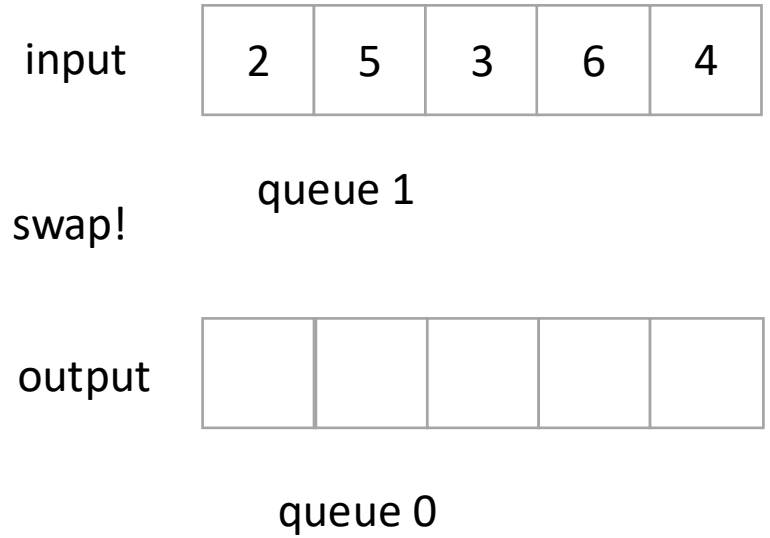
join all threads and clear input

Input/Output Queues

- Example: Information flow in graph applications:

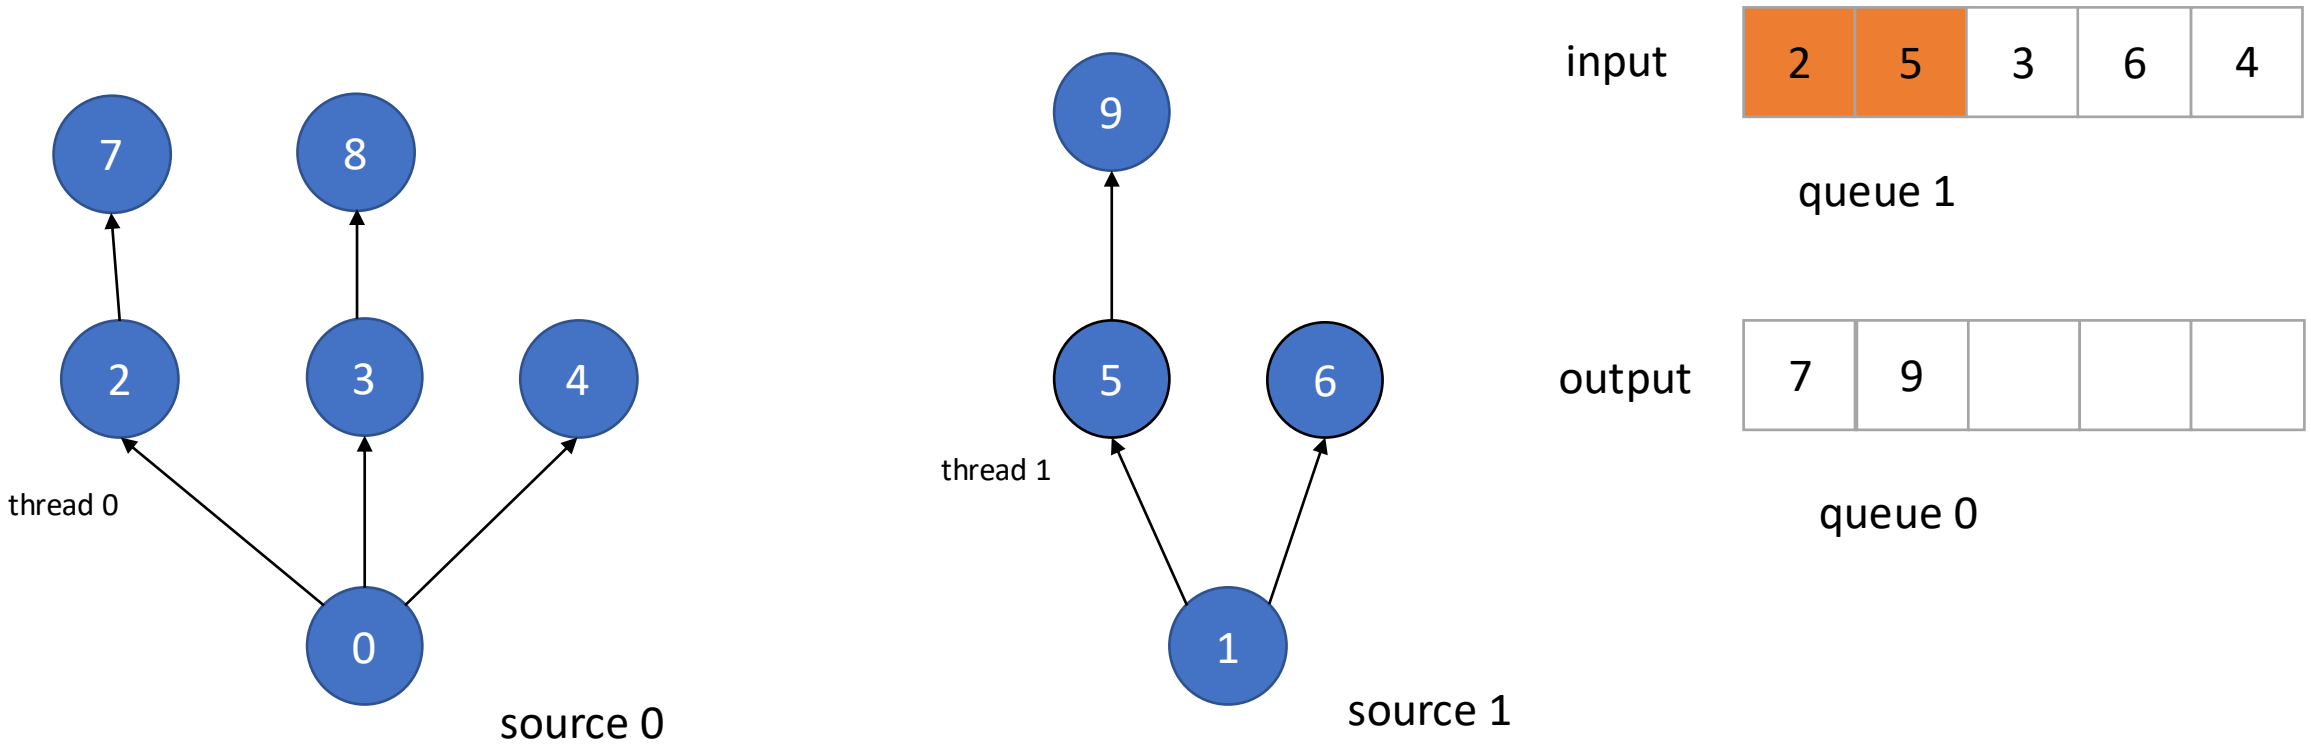






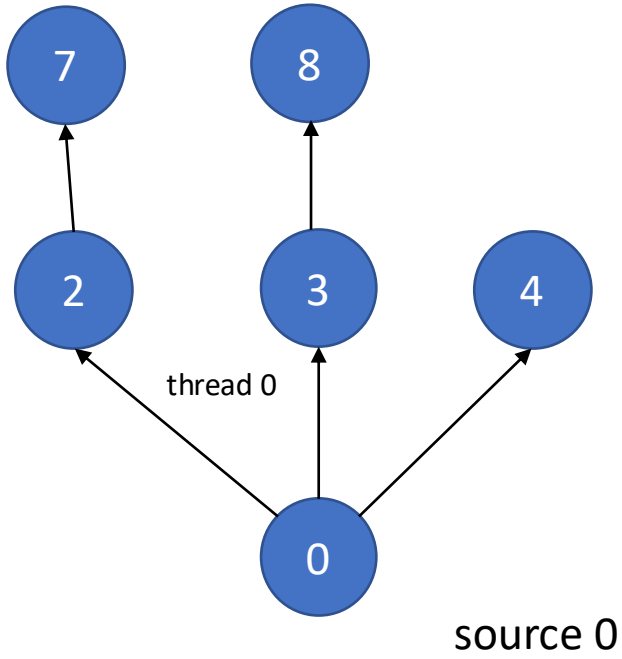
Input/Output Queues

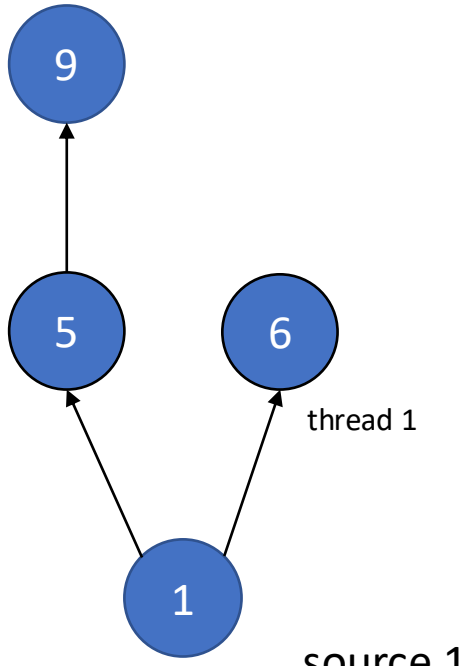
- Example: Information flow in graph applications:

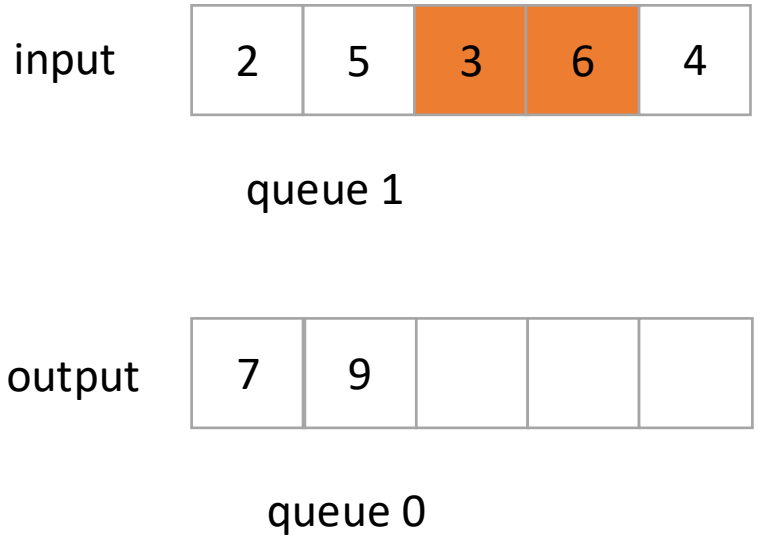


Input/Output Queues

- Example: Information flow in graph applications:

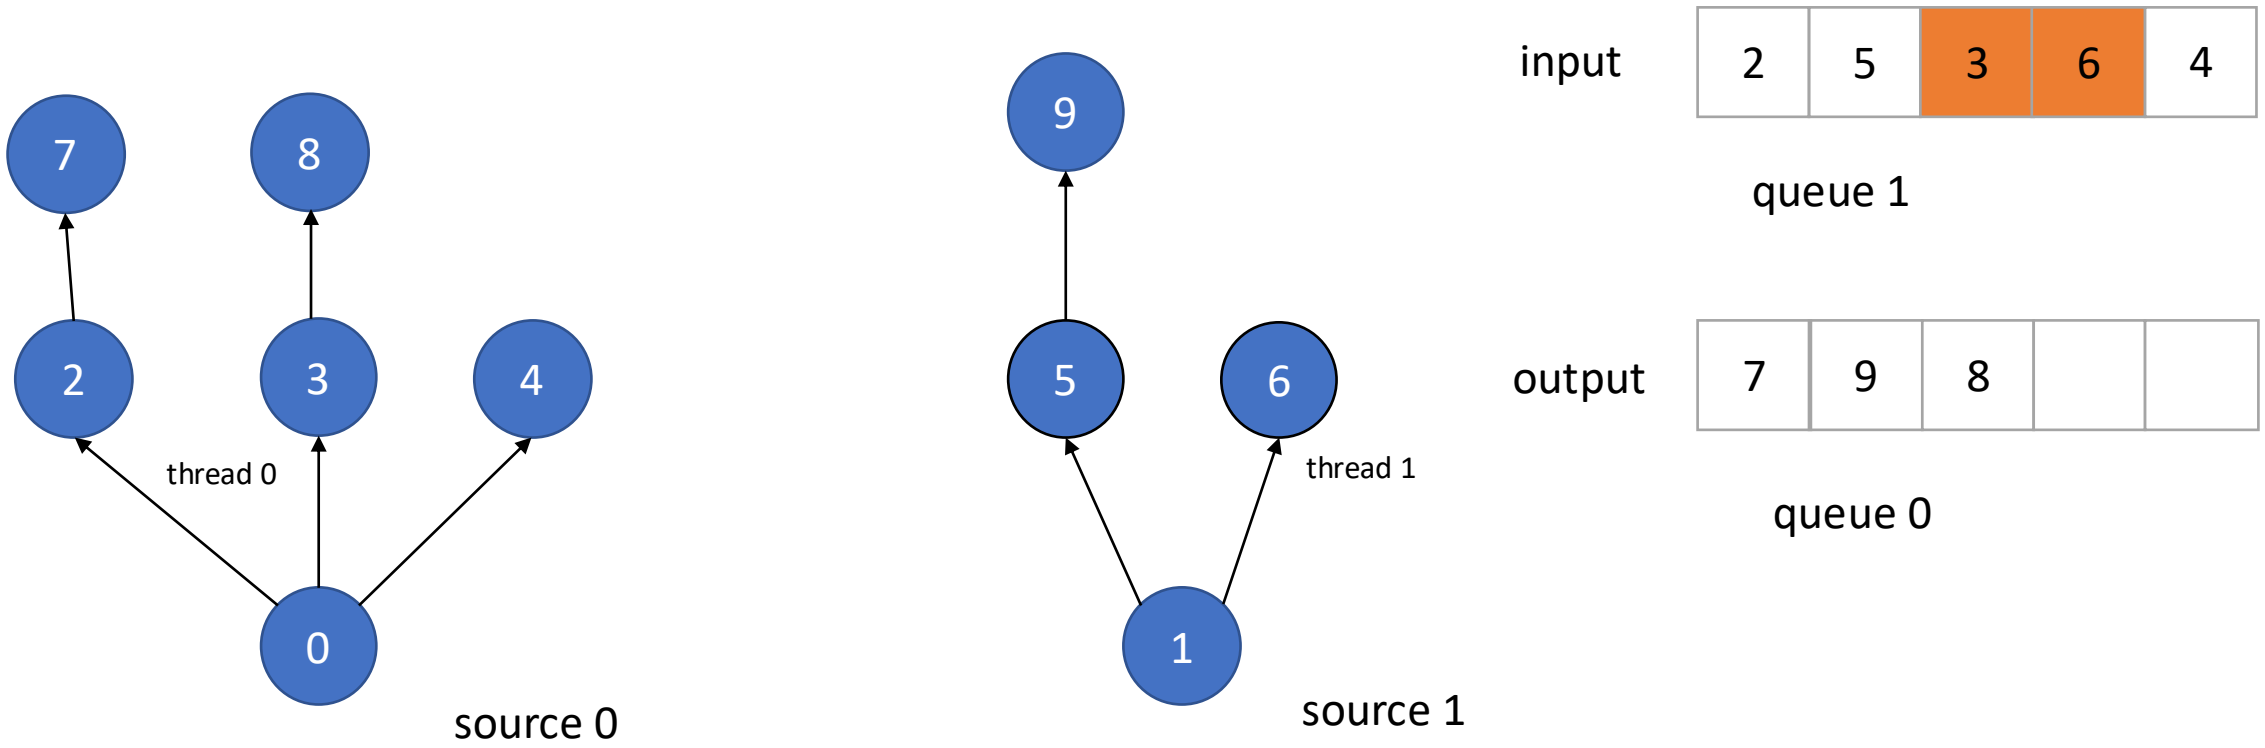






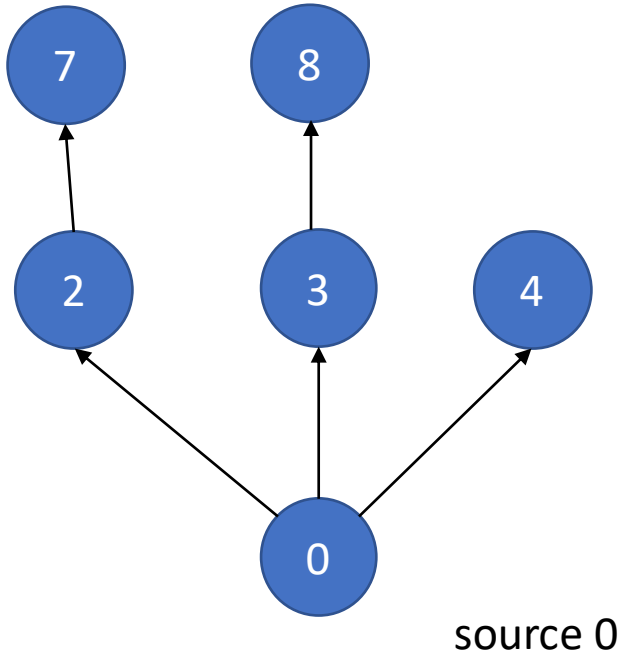
Input/Output Queues

- Example: Information flow in graph applications:

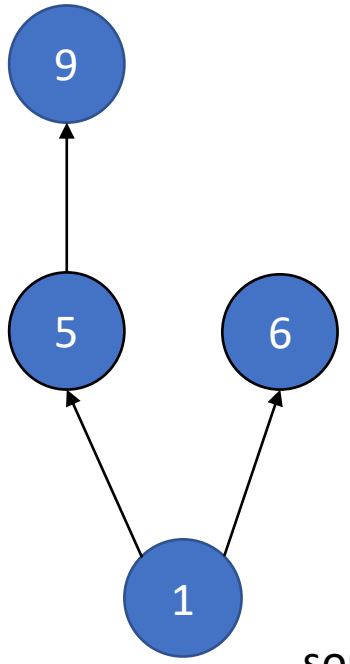


Input/Output Queues

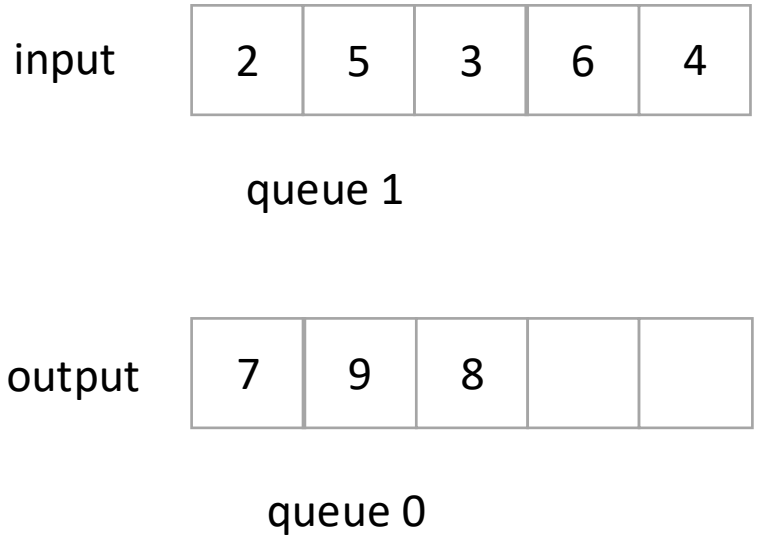
- Example: Information flow in graph applications:



source 0

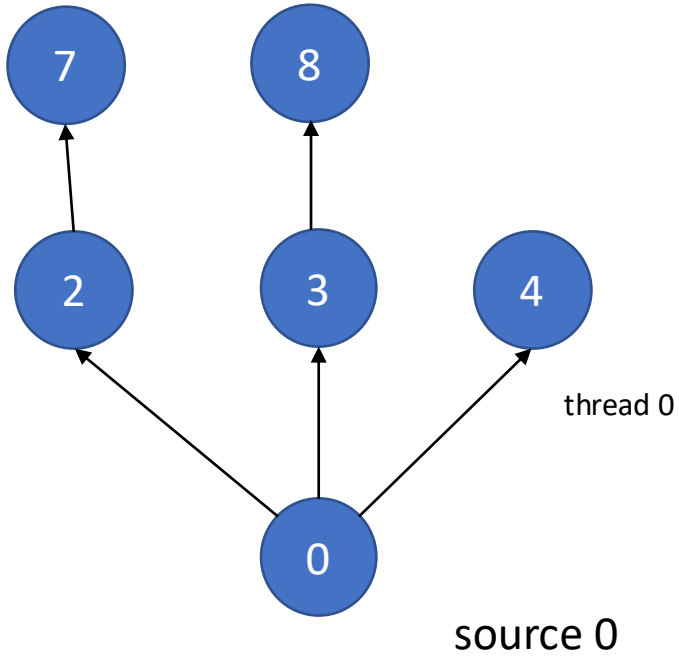


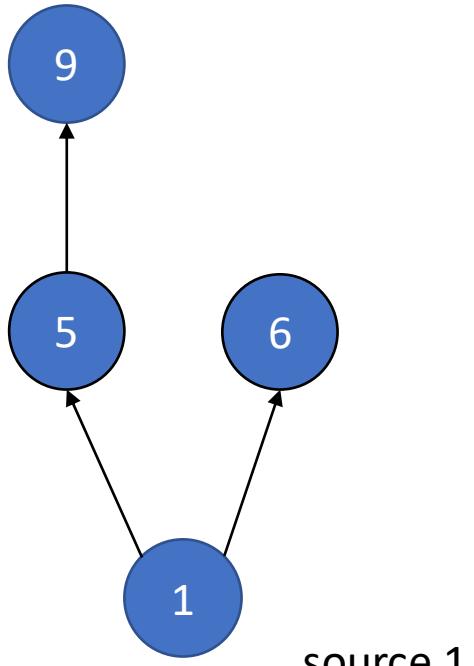
source 1

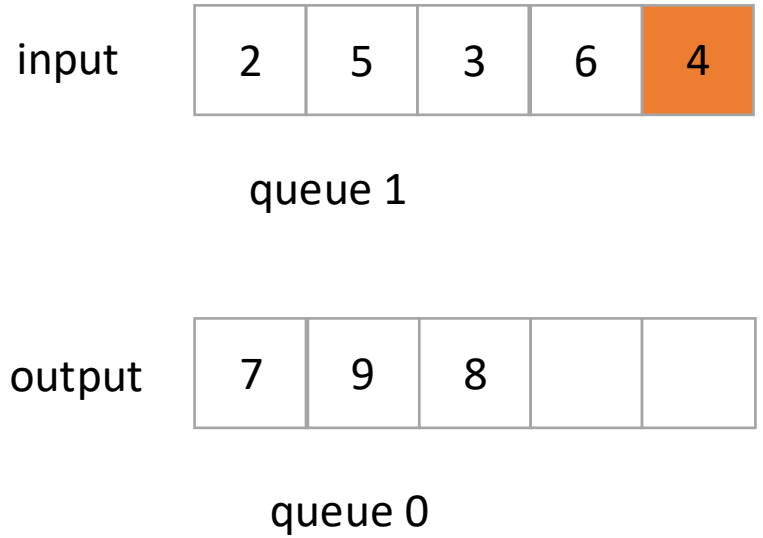


Input/Output Queues

- Example: Information flow in graph applications:

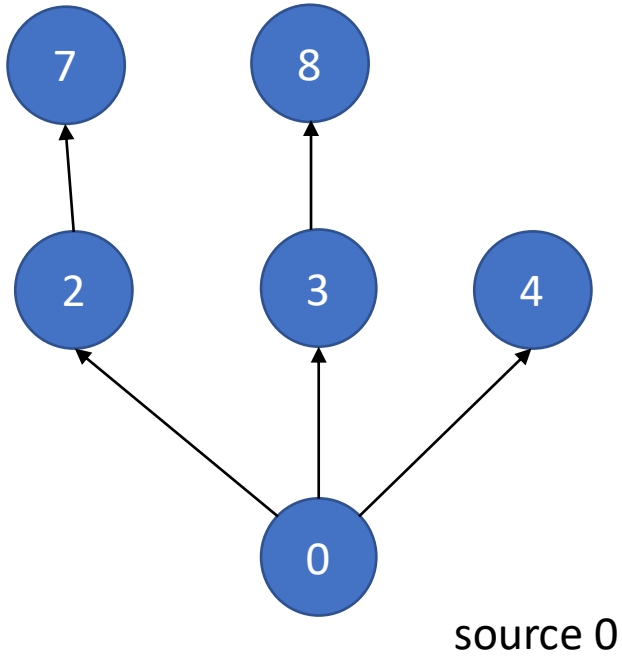




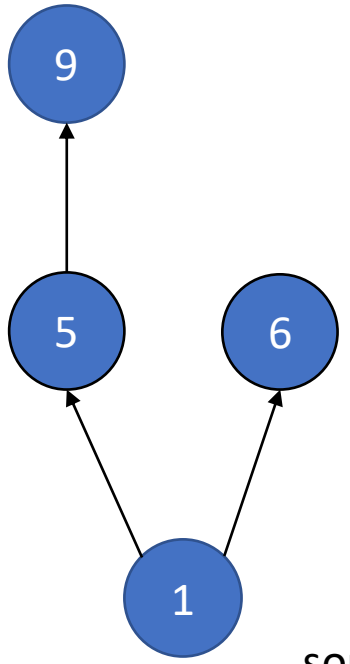


Input/Output Queues

- Example: Information flow in graph applications:



source 0



source 1

