

CSE113: Parallel Programming

- **Topics:**

- Finish up GPUs
- Homework 4
- Start on memory models



Announcements

Grading

- Grades for HW 3 should be out by the end of the week
- HW 4 is released.

Previous quiz + Review

Previous quiz + Review

What is a warp?

-
- ☐ Group of 32 threads in a GPU
 - ☐ Group of 16 threads in a GPU
 - ☐ Group of 2 threads in a GPU
 - ☐ Group of some threads in a GPU
-



Previous quiz + Review

What is a warp?

-
- > ☐ Group of 32 threads in a GPU
-
- ☐ Group of 16 threads in a GPU
-
- ☐ Group of 2 threads in a GPU
-
- > ☐ Group of some threads in a GPU

Previous quiz + Review

Like the CPU cache, the Load/Store Unit reads in memory in chunks. Is this affirmation true or false?

☐ True

☐ False


4 cores are accessing memory. What can happen

Read contiguous values

Like the CPU cache, the Load/Store Unit reads in memory in chunks. 16 bytes

Can easily distribute the values to the threads

1 request to GPU memory



Previous quiz + Review

Like the CPU cache, the Load/Store Unit reads in memory in chunks. Is this affirmation true or false?

-> ☐ True

☐ False

Previous quiz + Review

What could we observe with the demonstrations made in class about memory accesses on GPUs?

Chunked Pattern

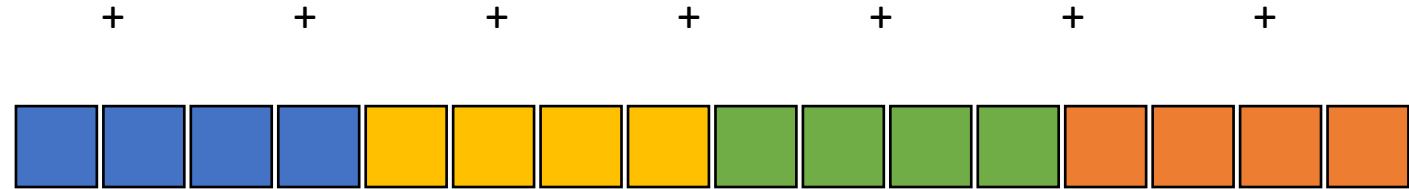
the first element accessed
by the 4 threads sharing a
load store unit. What
sort of access is this?

array a



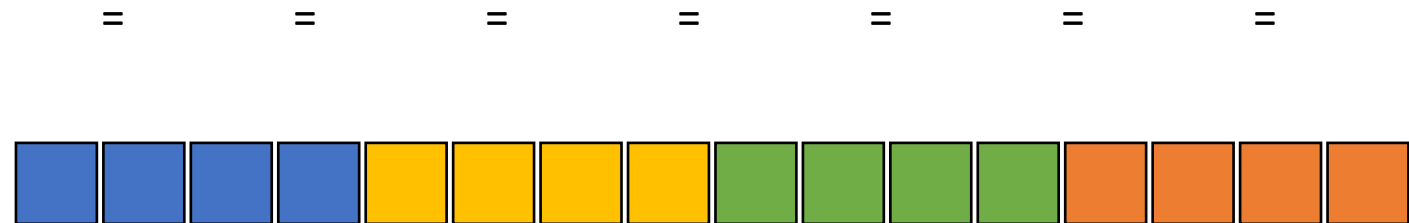
Computation
can easily be
divided into
threads

array b



Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

array c



How can we fix this

Stride Pattern

array a



+ + + + + + +

array b



= = = = = = =

array c



Computation
can easily be
divided into
threads

Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

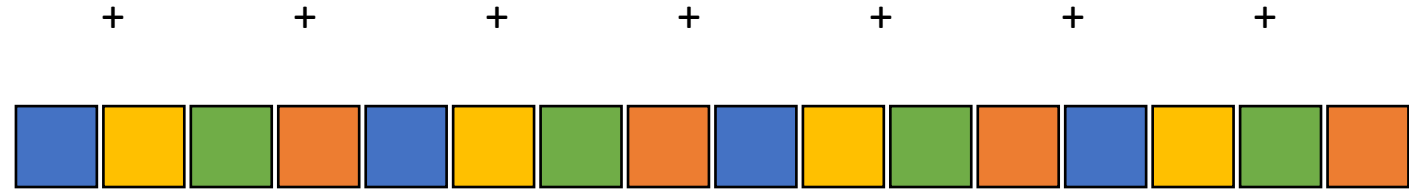
Stride Pattern

What sort of pattern is this?


array a



array b



array c



Computation
can easily be
divided into
threads

Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange


4 cores are accessing memory. What can happen

Read contiguous values

Like the CPU cache, the Load/Store Unit reads in memory in chunks. 16 bytes

Can easily distribute the values to the threads

1 request to GPU memory



Previous quiz + Review

What could we observe with the demonstrations made in class about memory accesses on GPUs?

Memory coalescing accelerates execution.

Previous quiz + Review

Why do we need to calculate $\text{int } i = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$?

- ☐ To get the index in the matrix we want to compute in that thread
- ☐ To get the whole matrix we want to compute in that thread
- ☐ To get the index in the matrix we want to compute in that warp

Previous quiz + Review

Why do we need to calculate $\text{int } i = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$?

-> ☐ To get the index in the matrix we want to compute in that thread

☐ To get the whole matrix we want to compute in that thread

☐ To get the index in the matrix we want to compute in that warp

Homework 4 - first look

- Prerequisites
 - Google Chrome
 - Should be stable on Linux, Windows and Mac.
- No docker container for this assignment.

Homework 4 - first look

- Javascript shared array buffer:
 - How javascript threads can actually share memory
 - Similar to memory in C++

```
var buffer = new SharedArrayBuffer(Float32Array.BYTES_PER_ELEMENT * NUM);  
var array = new Float32Array(buffer);
```

Shared Array Buffer

- Like Malloc, allocates a “pointer” to a contiguous array of bytes
- Can pass the “pointer” to different threads
- Need to instantiate a typed array to access the values

Web Workers

- How to do multi-threading in javascript
- Async
 - Concurrent (executes on the same thread)
 - Good for I/O and user interactions
- Web Workers will execute on multiple cores
 - Better for compute intensive applications
 - Better performance

How to use?

- Create a new worker with a file
 - Doesn't do anything yet
- File contains a function: “on message”
- Main file calls “post message” along with arguments to start the thread
- Worker sends a message back to the main file, it can catch the data

WebGPU

- The language is wgsi (WebGPU Shading Language)
 - It is new, there are not many examples (and the specification changes!)
 - Official specification is here: <https://www.w3.org/TR/WGSL/>

WebGPU

- wgsi is NOT javascript
- Javascript is interpreted: not possible on GPUs
- wgsi is compiled
 - into Vulkan on Linux
 - into Metal on Apple
 - into HLSL on Windows
- No printing (so GPU code can be difficult to debug)

WebGPU

- variables (optional types):

var <name> = <value>;

var cluster_dist = 3.0;

var <name> : <type> = <value>;

var cluster_dist : f32 = 3.0;

WebGPU

- types:
 - i32
 - u32
 - f32
 - vec2<f32>
 - array<type>

```
struct Particle {  
    pos : vec2<f32>;  
};
```

```
struct Particles {  
    particles : array<Particle>;  
};
```

```
var index_pos : vec2<f32> = particlesA.particles[index].pos;
```

- structures

```
var index : u32 = GlobalInvocationID.x;
```

- Built-ins (global id) *you have one thread for each particle!*

WebGPU

- Built in functions:
 - `arrayLength`
 - `sqrt`
 - `pow`
 - `distance`

WebGPU

For loops:

```
for (var i : u32 = 0u; i < arrayLength(&particlesA.particles); i = i + 1u) {  
    ...  
}
```

WebGPU

- Types can be frustrating
- But compiler errors will help you, and you can do casts.

Moving on to memory consistency!

- One of my favorite topics!

Moving on to memory consistency!

- One of my favorite topics!
- What do other people think?

*Look, memory ordering pretty much is
the rocket science of CS,*

Linus Torvalds

Memory Consistency

- We have been very strict about using atomic types in this class
 - and the methods (.load and .store)
 - why?
- Architectures do very strange things with memory loads and stores
- Compilers do too (but we won't talk too much about them today)
- C++ gives us sequential consistency if we use atomic types and operations.
- What do we remember sequential consistency from?

Sequential consistency for atomic memory

- Let's play our favorite game:

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

Is it possible for

`t0 == 0 and t1 == 1`



Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:


```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

Is it possible for

$t0 == 0$ and $t1 == 1$



Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:


```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

Is it possible for

`t0 == 0 and t1 == 1`



Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

```
x.store(1);
```

How about:

Is it possible for

`t0 == 1 and t1 == 0`

```
y.store(1);
```

```
int t0 = y.load();
```

```
int t1 = x.load();
```

Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

```
x.store(1);
```

no where for this one to go!

How about:

Is it possible for

$t0 == 1$ and $t1 == 0$

```
y.store(1);
```

```
int t0 = y.load();
```

```
int t1 = x.load();
```

Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
int t0 = y.load();
```

Dekker mutual exclusion:

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
y.store(1);  
int t1 = x.load();
```



Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
int t0 = y.load();
```

```
x.store(1);
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
y.store(1);  
int t1 = x.load();
```

```
int t0 = y.load();
```

```
y.store(1);
```

```
int t1 = x.load();
```


Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
int t0 = y.load();
```

```
x.store(1);
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
y.store(1);  
int t1 = x.load();
```

```
int t0 = y.load();
```

```
y.store(1);
```

```
int t1 = x.load();
```

no place for this one!

C++

- Plain atomic accesses are documented to be sequentially consistent (SC)
- Why wasn't SC very good for concurrent data structures?
 - Compossibility: two objects that are SC might not be SC when used together
 - Programs contain only 1 shared memory though; no reason to compose different main memories.

What about ISAs?

- Remember, it is important for us to understand how our code executes on the architecture to write high performing programs
- Lets think about x86
 - Instructions:
 - `MOV %t0 [x]` - loads the value at x to register t0
 - `MOV [y] 1` - stores the value 1 to memory location y

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

```
mov [x], 1
```

```
mov %t0, [y]
```

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```

```
mov [y], 1
```

```
mov %t1, [x]
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:


```
mov [x], 1  
mov %t0, [y]
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```



mov [y], 1

no place for this event!

ISA is not SC

- We'd like to be able to compile atomic instructions just to regular ISA loads and stores
- What if we actually run this code?

Schedule

- Memory consistency models:
 - **Total store order**
 - Relaxed memory consistency

Thread 0:

mov [x], 1

mov %t0, [y]

Core 0

Thread 1:

mov [y], 1

mov %t1, [x]

Core 1



Thread 0:

mov %t0, [y]

Core 0

mov [x], 1

execute first instruction
what happens to the stores?

Thread 1:

mov %t1, [x]



Core 1

mov [y], 1



Thread 0:



```
mov %t0, [y]
```



X86 cores contain a store buffer; holds stores before going to main memory

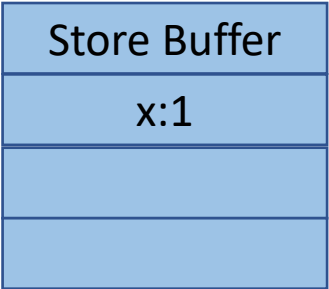

Thread 1:

```
mov %t1, [x]
```



Thread 0:



```
mov %t0, [y]
```




X86 cores contain a store buffer; holds stores before going to main memory

Thread 1:

```
mov %t1, [x]
```

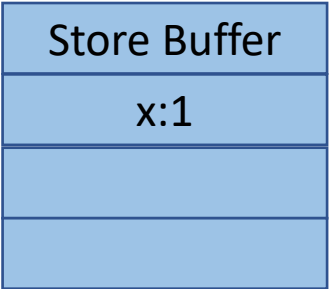



eventually they flush to main memory



Thread 0:



```
mov %t0, [y]
```




X86 cores contain a store buffer; holds stores before going to main memory

Thread 1:

```
mov %t1, [x]
```



eventually they flush to main memory



Thread 0:

mov [x], 1

mov %t0, [y]



rewind

Thread 1:

mov [y], 1

mov %t1, [x]

Core 0

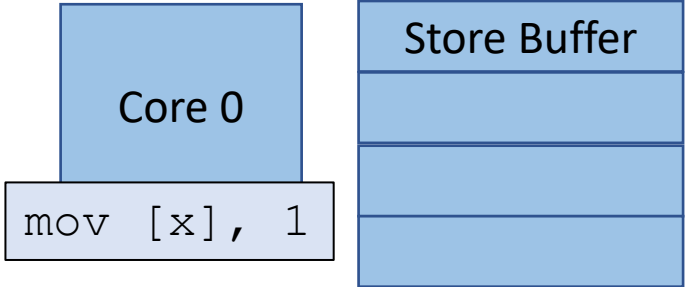


Core 1



Thread 0:

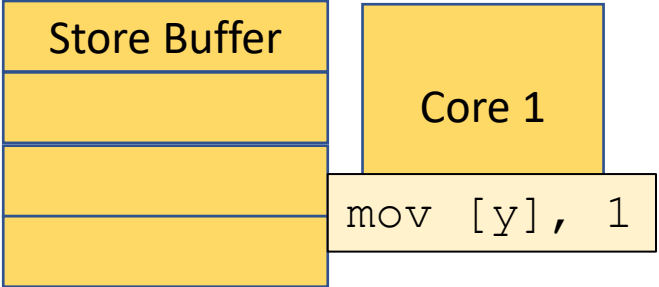
mov %t0, [y]



execute first instruction

Thread 1:

mov %t1, [x]



Thread 0:

mov %t0, [y]

Core 0

| Store Buffer |
|--------------|
| x:1 |
| |
| |

values get stored in SB

Thread 1:

mov %t1, [x]

| Store Buffer |
|--------------|
| y:1 |
| |
| |


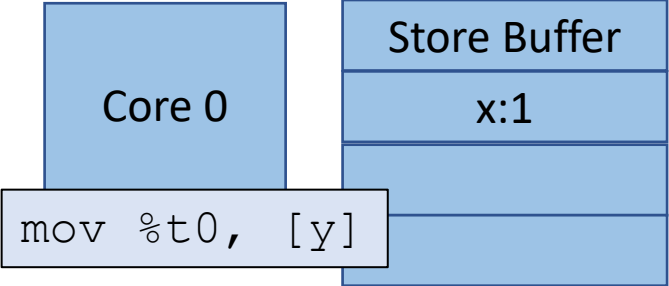
Core 1

| | |
|-----|-------------|
| x:0 | Main Memory |
| y:0 | |
| | |

Thread 0:

Thread 1:

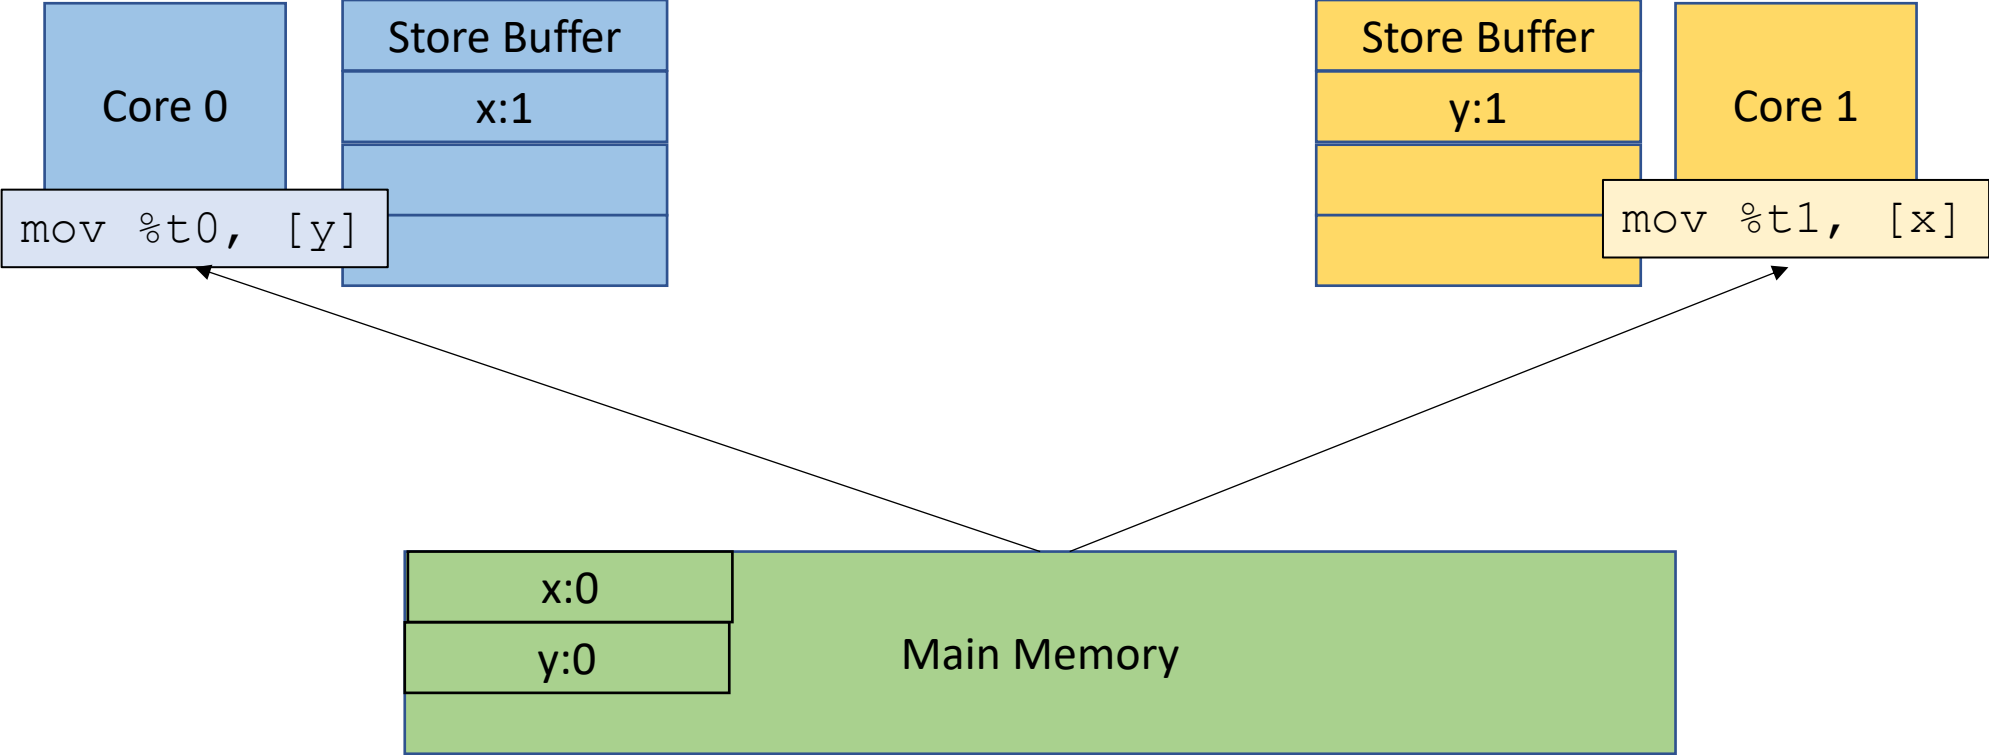
Execute next instruction



Thread 0:

Thread 1:


Values get loaded from memory



Thread 0:

Thread 1:

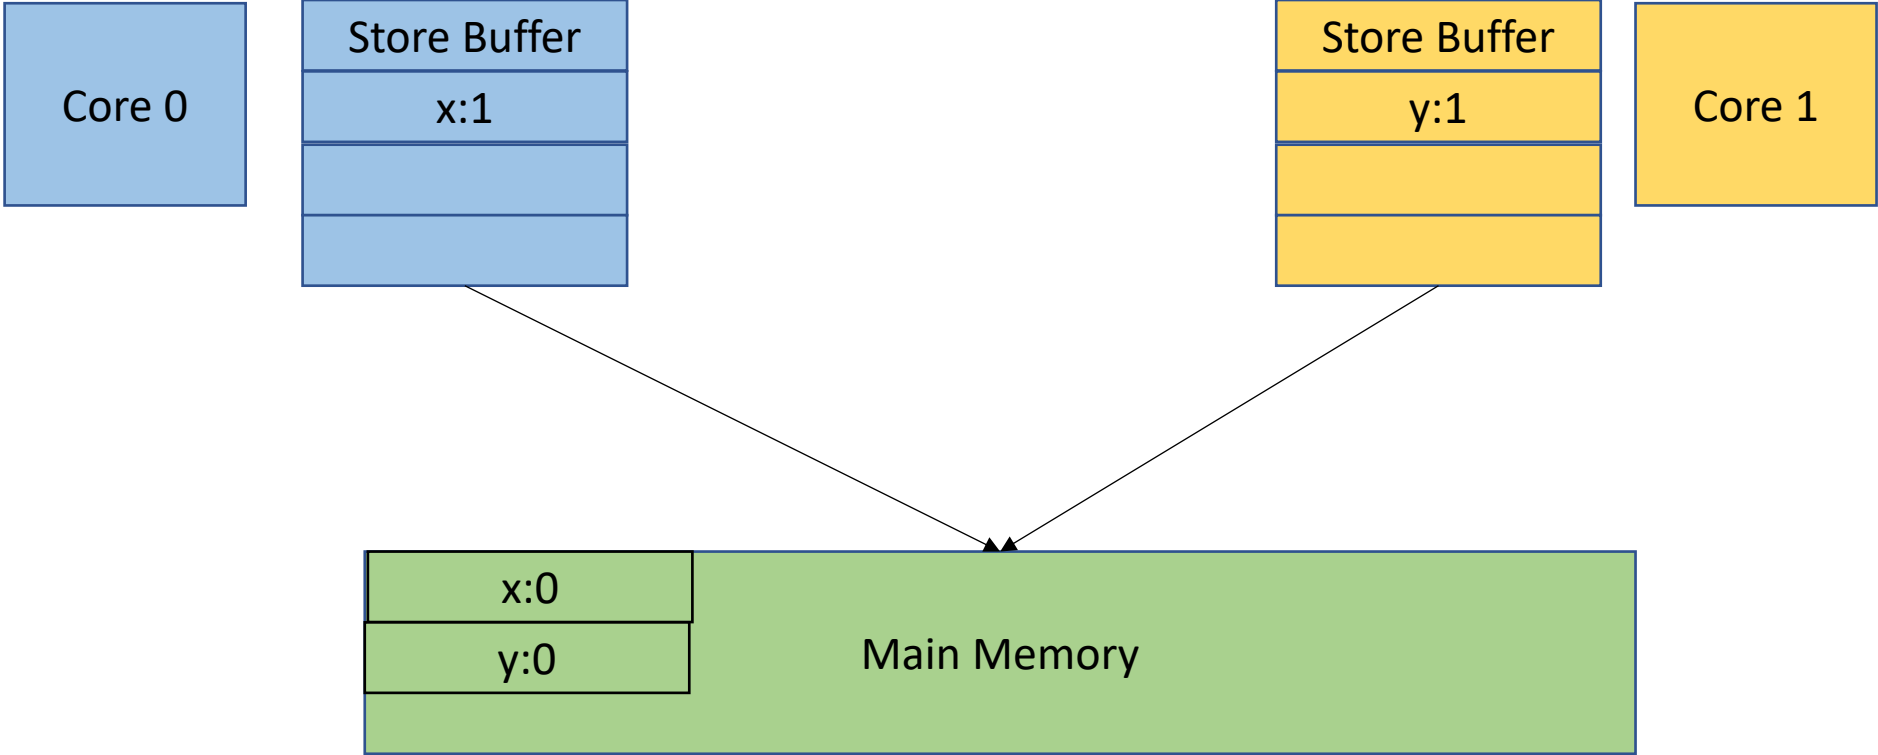
we see `t0 == t1 == 0!`



Thread 0:

Thread 1:

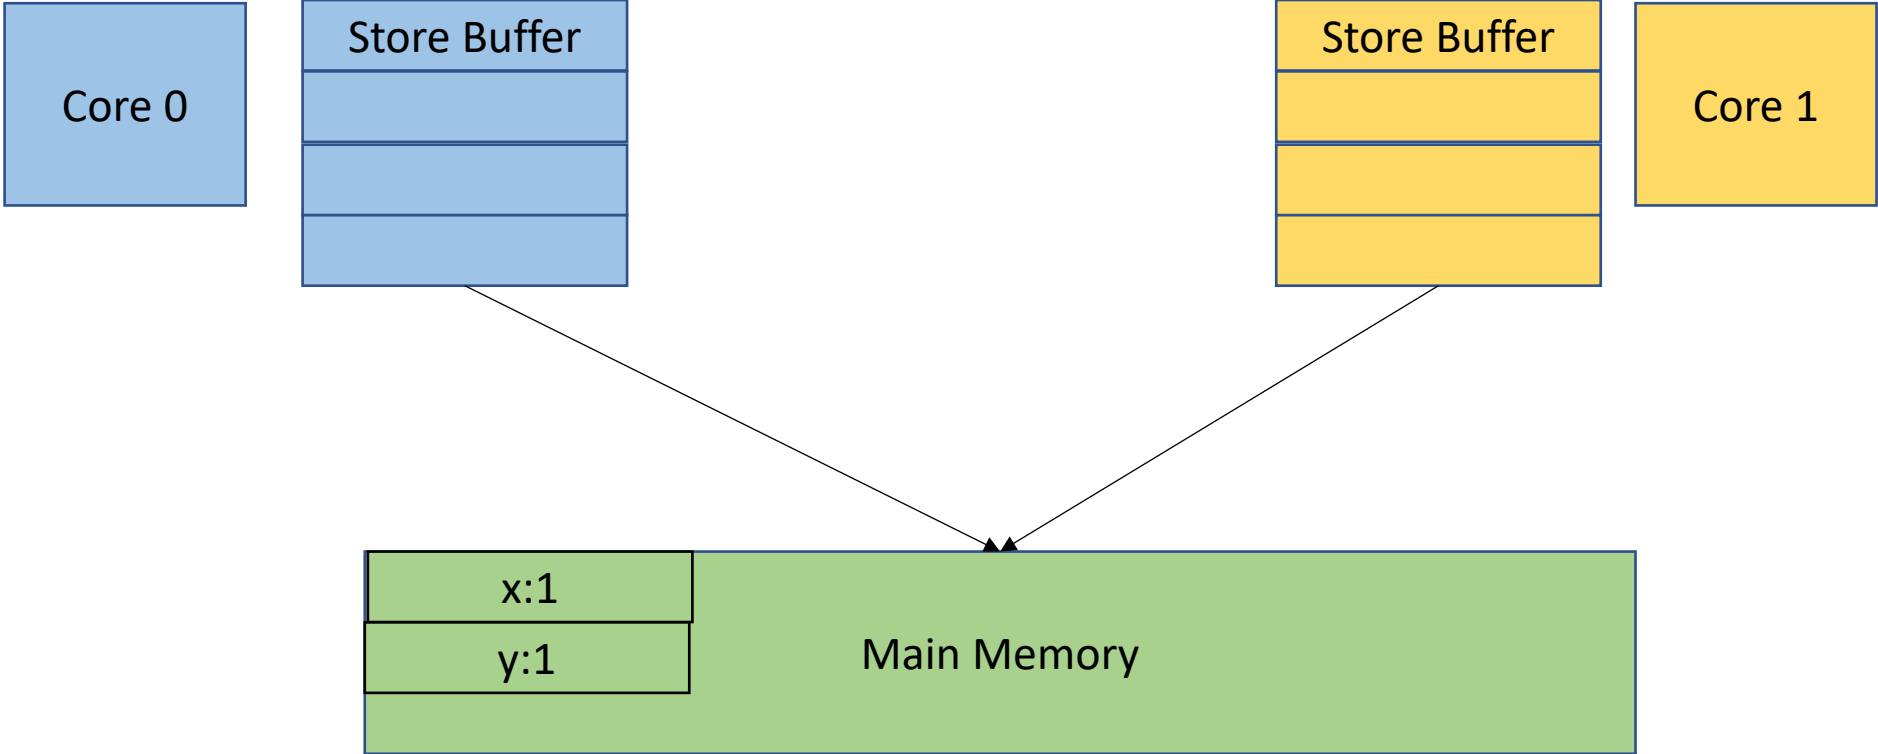
Store buffers are drained eventually



Thread 0:

Thread 1:

Store buffers are drained eventually
but we've already done our loads



Our first relaxed memory execution!

- **Relaxed memory** (also known as **Weak memory**) behaviors
- An execution that is **NOT allowed by sequential consistency**
- **Relaxed memory model**: a memory model that allows relaxed memory executions
 - **X86** has a relaxed memory model due to store buffering
 - If you restrict yourself to use only default atomic operations, **C++** does NOT have a weak memory model

Litmus tests

- Small concurrent programs that check for relaxed memory behaviors
- Vendors have a long history of under documented memory consistency models
- Academics have empirically explored the memory models
 - X86 behaviors were documented by researchers before Intel!
 - Many vendors have unofficially endorsed academic models

Litmus tests

This test is called “store buffering”

Thread 0:

```
mov  [x], 1  
mov  %t0, [y]
```

Thread 1:

```
mov  [y], 1  
mov  %t1, [x]
```

Can `t0 == t1 == 0`?

Fences: restoring sequential consistency



- It is typical that relaxed memory models provide special instructions which can be used to disallow weak behaviors.
- These instructions are called Fences
- The X86 fence that flushes the store buffer is called `mfence`.

Thread 0:

mov [x], 1

mfence

mov %t0, [y]





Thread 1:

mov [y], 1

mfence


mov %t1, [x]



Thread 0:

mfence

mov %t0, [y]




Execute first instruction

Thread 1:

mfence

mov %t1, [x]



Thread 0:

mfence

mov %t0, [y]

Core 0

| Store Buffer |
|--------------|
| x:1 |
| |
| |

Thread 1:

mfence

mov %t1, [x]

Core 1

| Store Buffer |
|--------------|
| y:1 |
| |
| |

Values go into the store buffer

| | |
|-----|-------------|
| x:0 | Main Memory |
| y:0 | |
| | |

Thread 0:

Thread 1:

Execute next instruction

```
mov %t0, [y]
```

```
mov %t1, [x]
```

Core 0

mfence

Store Buffer

x:1

Store Buffer

y:1

Core 1

mfence

x:0

y:0

Main Memory






Thread 0:

Thread 1:

store buffers are flushed

```
mov %t0, [y]
```

```
mov %t1, [x]
```







Thread 0:

Thread 1:

store buffers are flushed

```
mov %t0, [y]
```



```
mov %t1, [x]
```



Thread 0:

Thread 1:

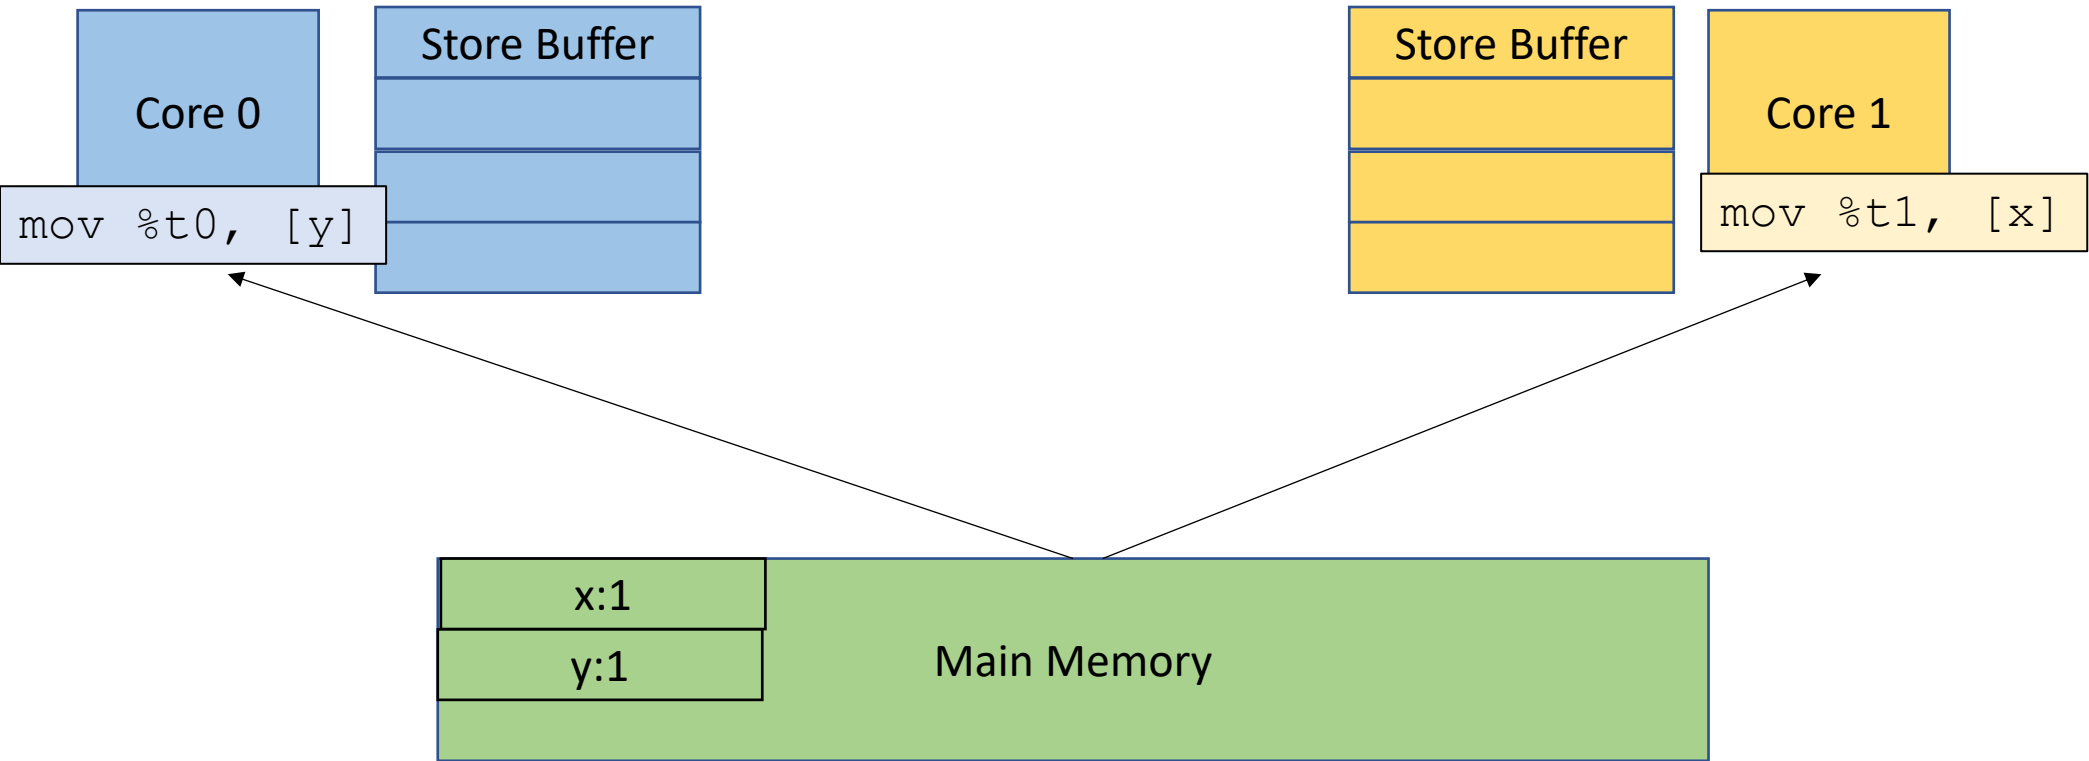
execute next instruction



Thread 0:

Thread 1:

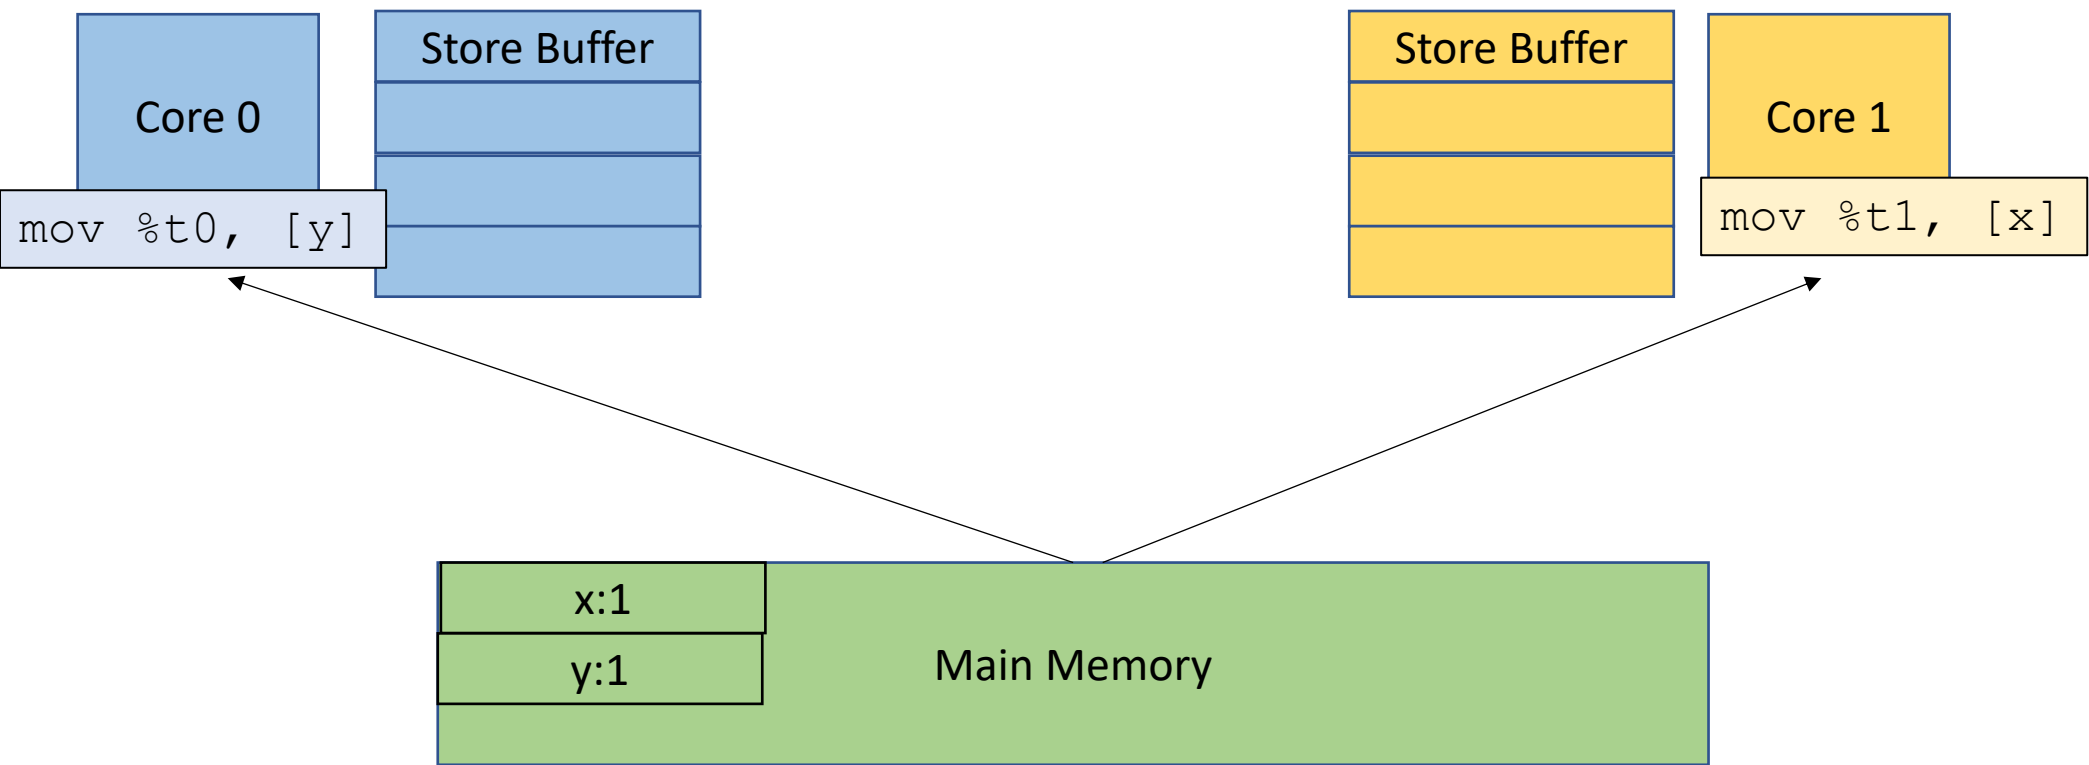
values are loaded from memory



Thread 0:

Thread 1:

We don't get the problematic behavior: `t0 == t1 == 0`



Thread 0:

```
mov [x], 1
```

```
mov %t0, [x]
```




single thread
same address

possible outcomes:

t0 = 1

t0 = 0

Which one do you expect?





Thread 0:

mov [x], 1

mov %t0, [x]


How does this execute?



Thread 0:

execute first instruction


mov %t0, [x]



Thread 0:


Store the value in the store buffer

```
mov %t0, [x]
```



Thread 0:

Next instruction





Thread 0:

Where to load??

Store buffer?

Main memory?





Thread 0:

Where to load??

Threads check store buffer before going to main memory

It is close and cheap to check.



Memory Consistency

- How to specify a relaxed memory model?
- We can do it operationally
 - by constructing a high-level machine and reasoning about operations through the machine.
 - or we can talk about instructions that are allowed to “break” program order.

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```



We will annotate instructions with S for store, and L for loads

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```



We will annotate instructions with S for store, and L for loads

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

```
S:mov [x], 1
```

```
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

```
L:mov %t1, [x]
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

L:mov %t1, [x]

S:mov [x], 1

L:mov %t0, [y]

S:mov [y], 1

Now we make a new rule:

S(tores) followed by a L(oad)
do not have to follow program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

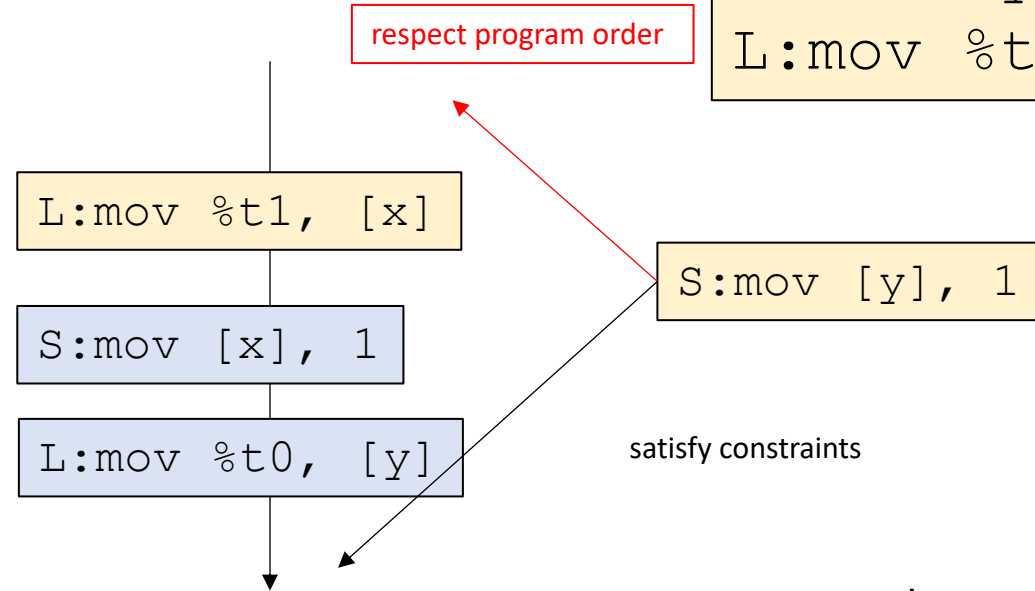
Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

we can ignore this condition!!



Now we make a new rule:

S(tores) followed by a L(oad)
do not have to follow program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

L:mov %t1, [x]

S:mov [x], 1

L:mov %t0, [y]

S:mov [y], 1

Now we can satisfy the condition!

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

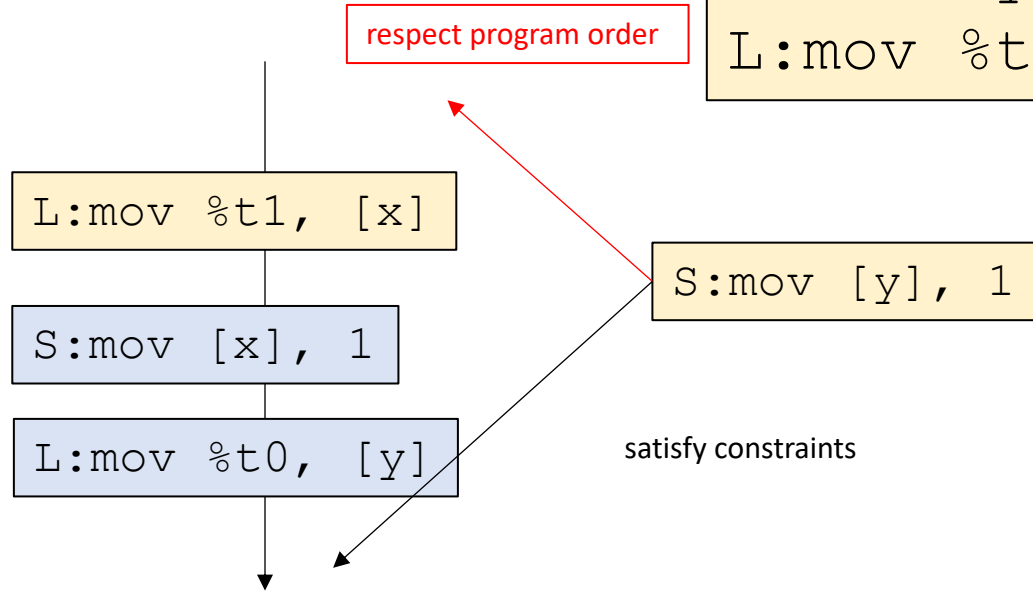
```
S:mov [x], 1  
L:mov %t0, [y]
```

Lets peak under the hood here

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

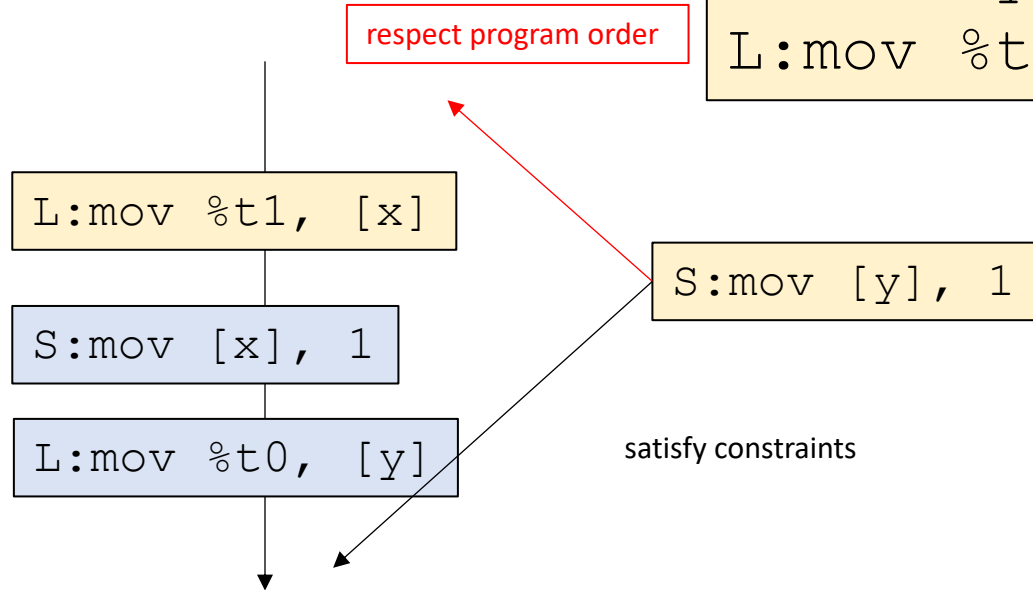
Lets peak under the hood here

Global timeline is when the
Store operation becomes visible
to other threads

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```


Lets peak under the hood here

Global timeline is when the
Store operation becomes visible
to other threads

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```


Lets peak under the hood here

Global timeline is when the
Store operation becomes visible
to other threads

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

Questions

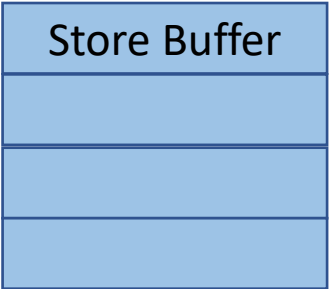
- Can stores be reordered with stores?

Thread 0:

mov [x], 1

mov [y], 1


Core 0



Thread 0:

mov [y], 1

execute the first instruction



Thread 0:

mov [y], 1

value goes into store buffer

Core 0

| Store Buffer |
|--------------|
| x:1 |
| |
| |

| | |
|-----|-------------|
| x:0 | Main Memory |
| y:0 | |
| | |

Thread 0:

mov [y], 1

execute next instruction

Core 0

| Store Buffer |
|--------------|
| x:1 |
| |
| |

| | |
|-----|-------------|
| x:0 | Main Memory |
| y:0 | |
| | |


Thread 0:

execute next instruction



Thread 0:

value goes into the store buffer



Thread 0:

Core 0


| Store Buffer |
|--------------|
| x:1 |
| y:1 |
| |

On x86, the store buffer trains in a FIFO way:
thus stores cannot be reordered


| | |
|-----|-------------|
| x:0 | Main Memory |
| y:0 | |
| | |

Thread 0:

Core 0




On x86, the store buffer trains in a FIFO way:
thus stores cannot be reordered




Thread 0:

Core 0



On x86, the store buffer trains in a FIFO way:
thus stores cannot be reordered



Questions

- Can stores be reordered with stores?
- How do we make rules about mfence?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
mfence  
L:mov %t0, [y]
```

```
S:mov [x], 1
```

```
mfence
```

```
L:mov %t0, [y]
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
S:mov [y], 1  
mfence  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

```
mfence
```

```
L:mov %t1, [x]
```

Rules: S(tores) followed by a L(oad)
do not have to follow program order.

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

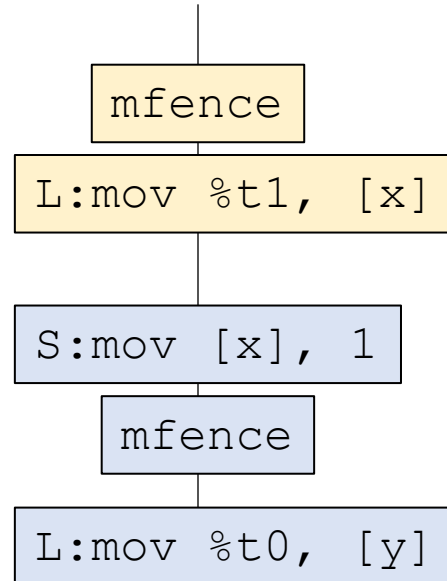
Thread 0:

```
S:mov [x], 1  
mfence  
L:mov %t0, [y]
```

*So we can't
reorder
this instruction
at all!*

Another test

Can `t0 == t1 == 0`?



Thread 1:

```
S:mov [y], 1  
mfence  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

Rules:

S(tores) followed by a L(oad)
do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order

Rules

- Are we done?

Rules:

S(tores) followed by a L(oad)

do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

```
S:mov [x], 1
```

```
L:mov %t0, [x]
```



Rules:

S(tores) followed by a L(oad)

do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

Another test
Can `t0 == 0`?

S:mov [x], 1

where to put this store?

L:mov %t0, [x]

Rules:
S(tores) followed by a L(oad)
do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order

Global variable:


```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

Another test

Can `t0 == 0`?



where to put this store?

Rules:

S(tores) followed by a L(oad)
do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order

S(tores) cannot be reordered past L(oads)
from the same address

TSO - Total Store Order

Rules:

S(tores) followed by a L(oad)
do not have to follow program order.

S(tores) cannot be reordered past a fence
in program order


S(tores) cannot be reordered past L(oads)
from the same address

Schedule

- Memory consistency models:
 - Total store order
 - **Relaxed memory consistency**

Other memory models?

- We can specify them in terms of what reorderings are allowed



If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Other memory models?

- We can specify them in terms of what reorderings are allowed

| | | memory access 0 | |
|-----------------|---|-----------------|----|
| | | L | S |
| memory access 1 | L | NO | NO |
| | S | NO | NO |

Sequential Consistency

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Other memory models?

- We can specify them in terms of what reorderings are allowed

| | | memory access 0 | |
|-----------------|---|-----------------|-------------------|
| | | L | S |
| memory access 1 | L | NO | Different address |
| | S | NO | NO |

TSO - total store order

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Other memory models?

- We can specify them in terms of what reorderings are allowed

| | | memory access 0 | |
|-----------------|---|-----------------|---|
| | | L | S |
| memory access 1 | L | ? | ? |
| | S | ? | ? |

Weaker models?

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Other memory models?

- We can specify them in terms of what reorderings are allowed

| | | memory access 0 | |
|-----------------|---|-----------------|-------------------|
| | | L | S |
| memory access 1 | L | NO | Different address |
| | S | NO | Different address |

PSO - partial store order

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Allows stores to drain from the store buffer in any order

Other memory models?

- We can specify them in terms of what reorderings are allowed

| | | memory access 0 | |
|-----------------|---|-------------------|-------------------|
| | | L | S |
| memory access 1 | L | YES | Different address |
| | S | Different address | Different address |

RMO - Relaxed Memory Order

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

Very relaxed model!

Other memory models?

- FENCE: can always restore order using fences. Accesses cannot be reordered past fences!

| | | memory access 0 | |
|-----------------|---|-----------------|----|
| | | L | S |
| memory access 1 | L | NO | NO |
| | S | NO | NO |

Any Memory Model

If memory access 0 appears before memory access 1 in program order, and there is a FENCE between the two accesses, can it bypass program order?

Schedule

- Memory consistency models:
 - Total store order
 - Relaxed memory consistency
- **Some Examples**