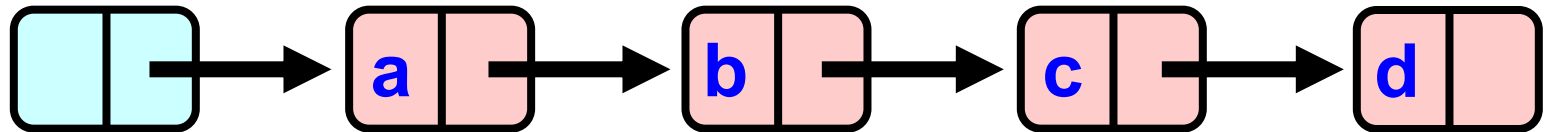
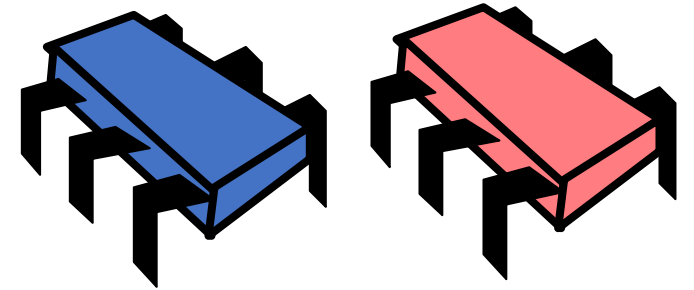


CSE113: Parallel Programming

- **Topics:**

- Concurrent general set



Announcements

- HW 4 submission by tomorrow.

Announcements

- HW 5 will be released on this week.

Announcements

SETs are out, please do them! It helps us out a lot.

Previous quiz + Review

Previous quiz + Review

The C++ relaxed memory order provides

- ☐ no orderings at all
- ☐ orderings only between accesses of the same address
- ☐ TSO memory behaviors when run on an x86 system
- ☐ an easy way to accidentally introduce horrible bugs into your program

Relaxed memory order

language
C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

language
C++11 (memory_order_relaxed)

	L	S
L	different address	different address
S	different address	different address

basically no orderings except for accesses to
the same address

Previous quiz + Review

The C++ relaxed memory order provides

- ☐ no orderings at all
- > ☐ orderings only between accesses of the same address
- ☐ TSO memory behaviors when run on an x86 system
- ☐ an easy way to accidentally introduce horrible bugs into your program

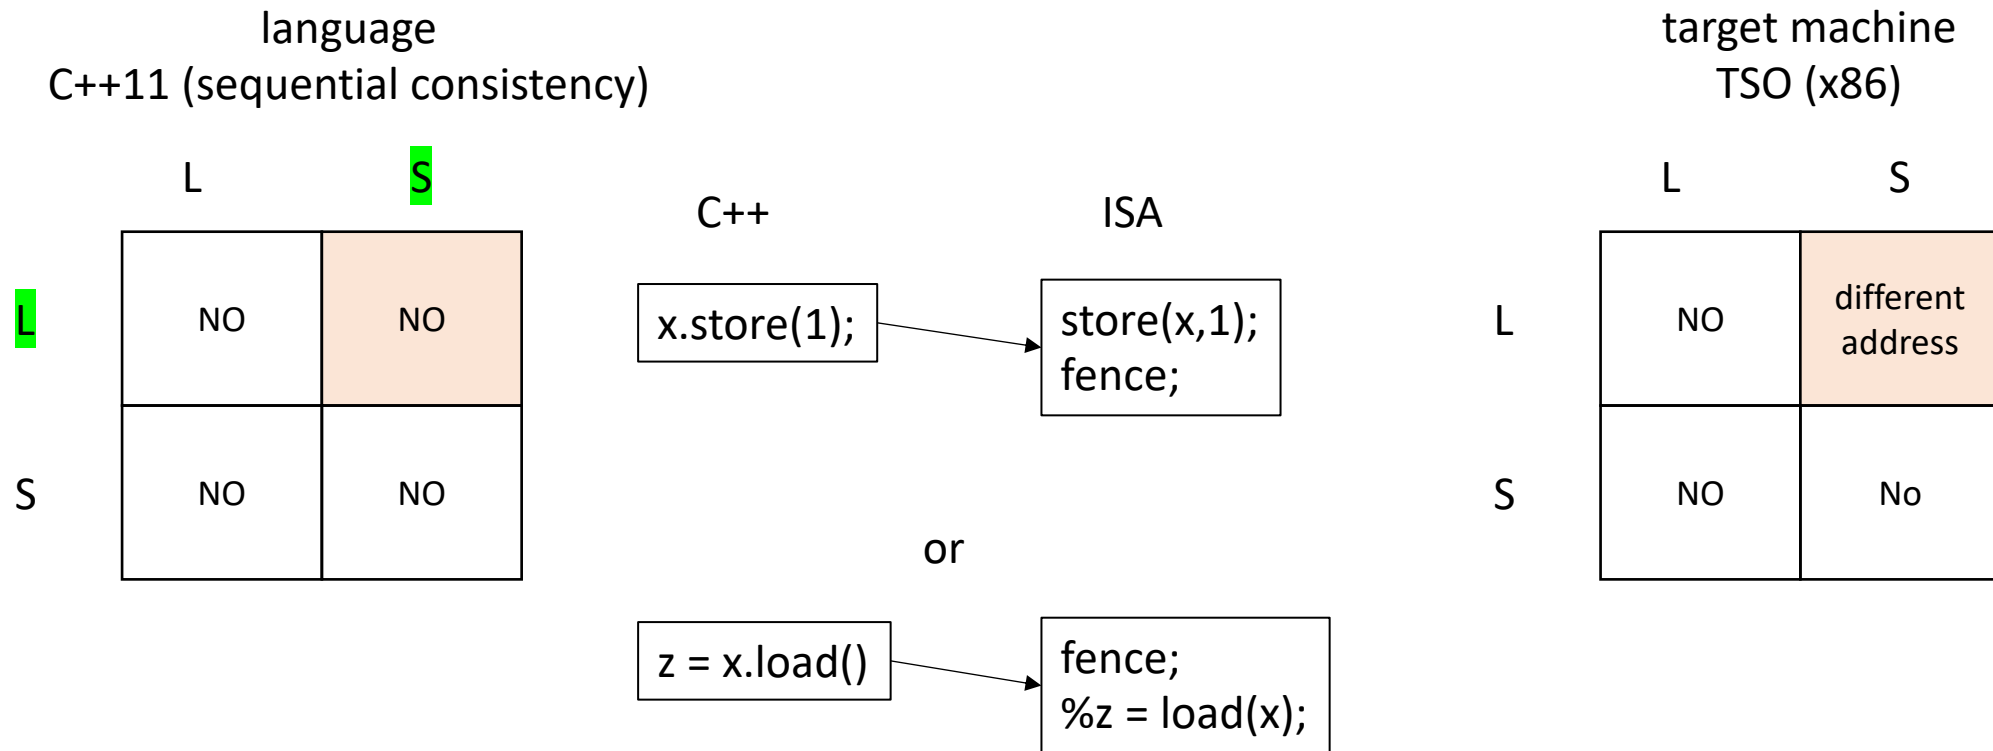
Previous quiz + Review

In terms of memory models, the compiler needs to ensure the following property:

- ☐ Any weak behavior allowed in the language is also allowed in the ISA
- ☐ Any weak behaviors that are disallowed in the language need to be disallowed in the ISA
- ☐ The compilation ensures that the program has sequentially consistent behavior at the ISA level
- ☐ The compiler does not need to reason about relaxed memory

C++11 atomic operation compilation

start with both both of the grids for the two different memory models



Compiling memory order relaxed

language
C++11 (memory_order_relaxed)

	L	S
L	different address	different address
S	different address	different address

lots of mismatches!

But language is more relaxed than machine

so no fences are needed

target machine
TSO (x86)

	L	S
L	NO	different address
S	NO	No

Previous quiz + Review

In terms of memory models, the compiler needs to ensure the following property:

- ☐ Any weak behavior allowed in the language is also allowed in the ISA
- > ☐ Any weak behaviors that are disallowed in the language need to be disallowed in the ISA
- ☐ The compilation ensures that the program has sequentially consistent behavior at the ISA level
- ☐ The compiler does not need to reason about relaxed memory

Previous quiz + Review

A program that uses mutexes and has no data conflicts does not have weak memory behaviors for which of the following reasons?

-
- ☐ Mutexes prevent memory accesses from happening close enough in time for weak behaviors to occur
-
- ☐ The OS has built in support for Mutexes that disable architecture features, such as the store buffer
-
- ☐ A correct mutex implementation uses fences in lock and unlock to disallow weak behaviors

Previous quiz + Review

A program that uses mutexes and has no data conflicts does not have weak memory behaviors for which of the following reasons?

- ☐ Mutexes prevent memory accesses from happening close enough in time for weak behaviors to occur
- ☐ The OS has built in support for Mutexes that disable architecture features, such as the store buffer
- > ☐ A correct mutex implementation uses fences in lock and unlock to disallow weak behaviors

Previous quiz + Review

Assuming you had a sequentially consistent processor, any C/++ program you ran on it would also be sequentially consistent, regardless of if there are data-conflicts or not.

☐ True

☐ False

Previous quiz + Review

Assuming you had a sequentially consistent processor, any C/++ program you ran on it would also be sequentially consistent, regardless of if there are data-conflicts or not.

☐ True

-> ☐ False

In addition to processor, the compiler can reorder instruction as well.

Previous quiz + Review

If you put a fence after every memory instruction, would that be sufficient to disallow all weak behaviors on a weak architecture? Please write a few sentences explaining your answer.

Previous quiz + Review

If you put a fence after every memory instruction, would that be sufficient to disallow all weak behaviors on a weak architecture? Please write a few sentences explaining your answer.

It is sufficient but makes execution slower.

General Concurrent Set

Set Interface

- Unordered collection of items
- No duplicates
- We will implement set as a sorted linked list.

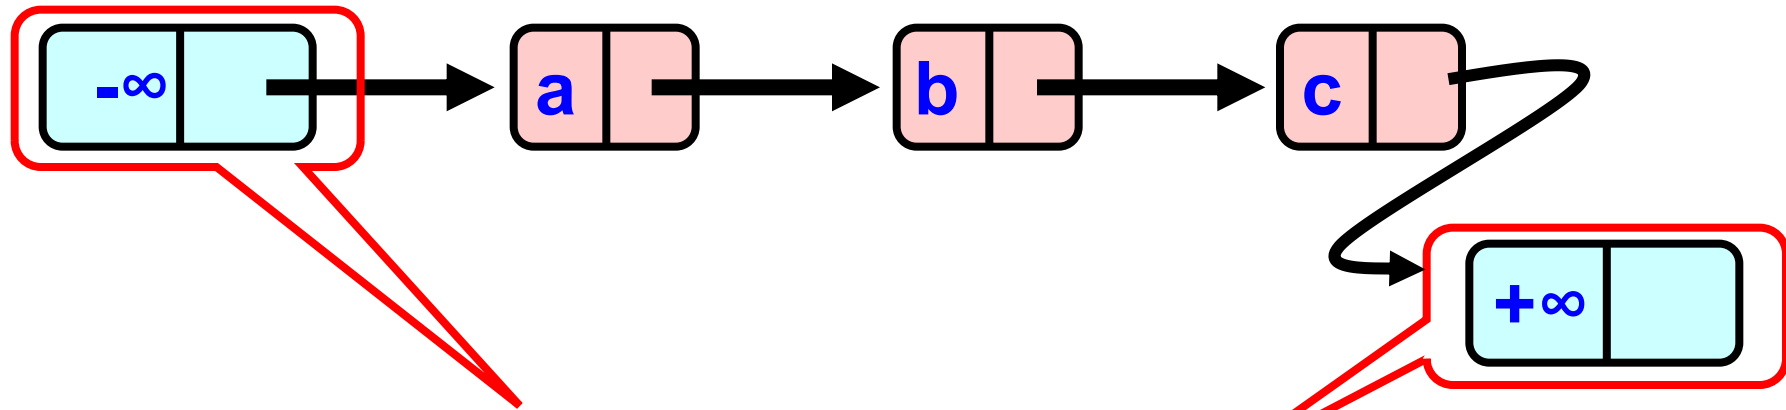
Set Interface

- Methods
 - **add(x)** put **x** in set
 - **remove(x)** take **x** out of set
 - **contains(x)** tests if **x** in set

List Node

```
class Node {  
    public:  
        Value v;  
        int key;  
        Node *next;  
}
```

The List-Based Set



Sorted with Sentinel nodes
(min & max possible keys)

Sequential List Based Set

add(b)

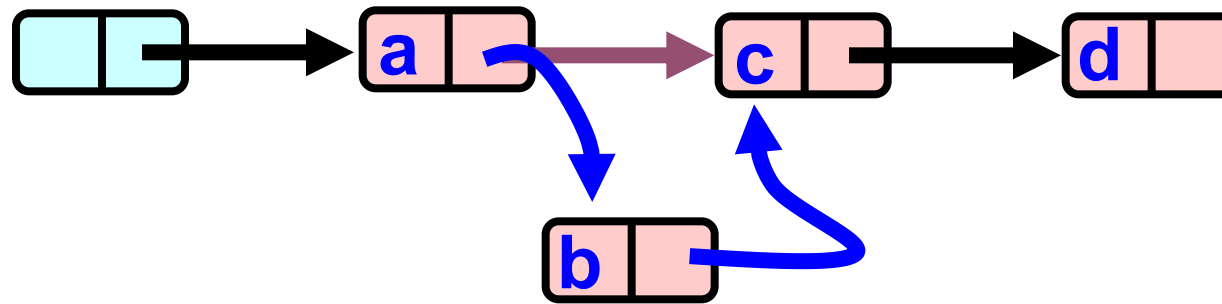


remove(b)

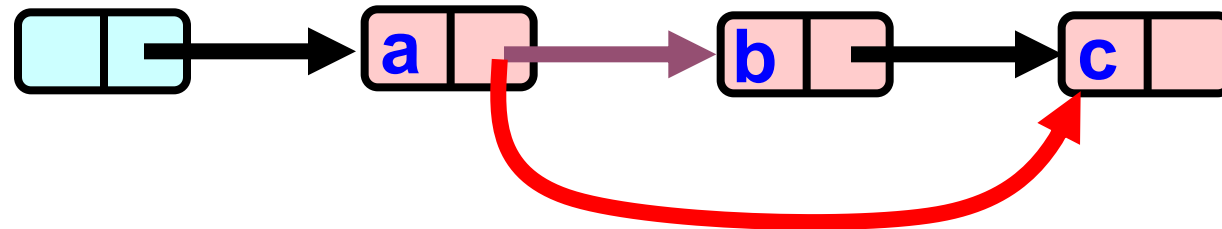


Sequential List Based Set

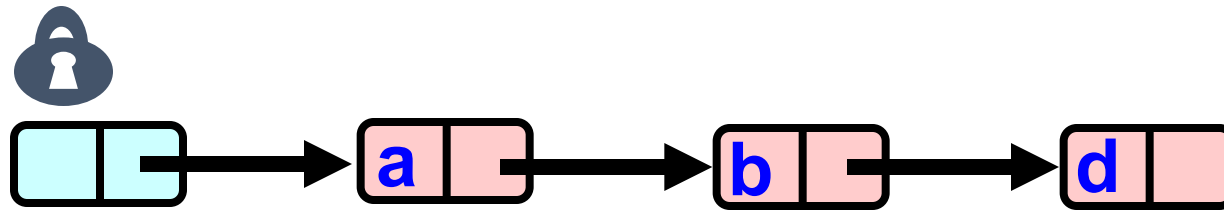
add(b)



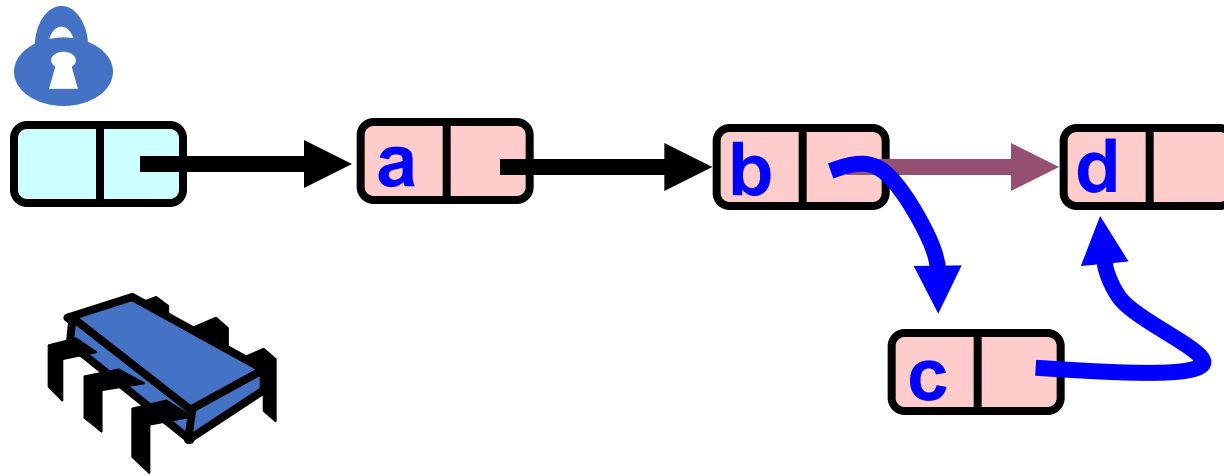
remove(b)



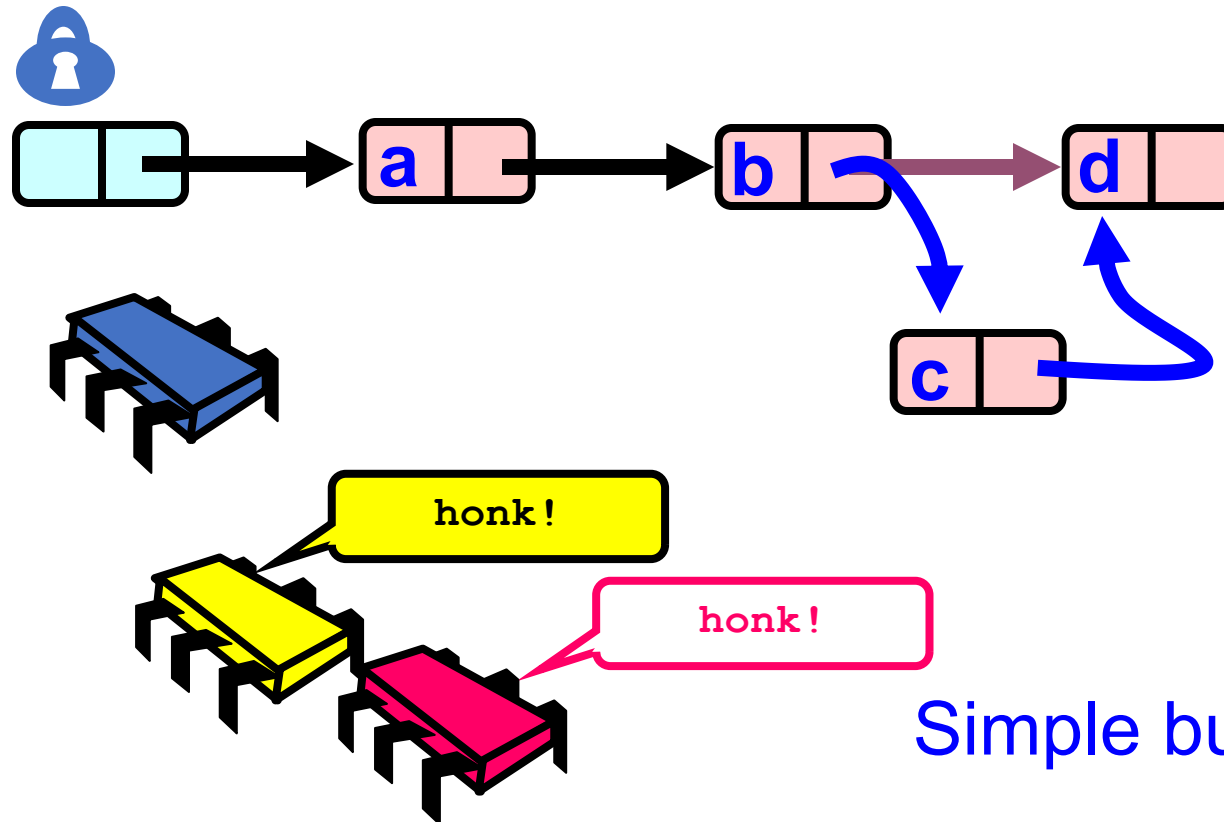
Coarse-Grained Locking



Coarse-Grained Locking



Coarse-Grained Locking



Simple but inefficient!

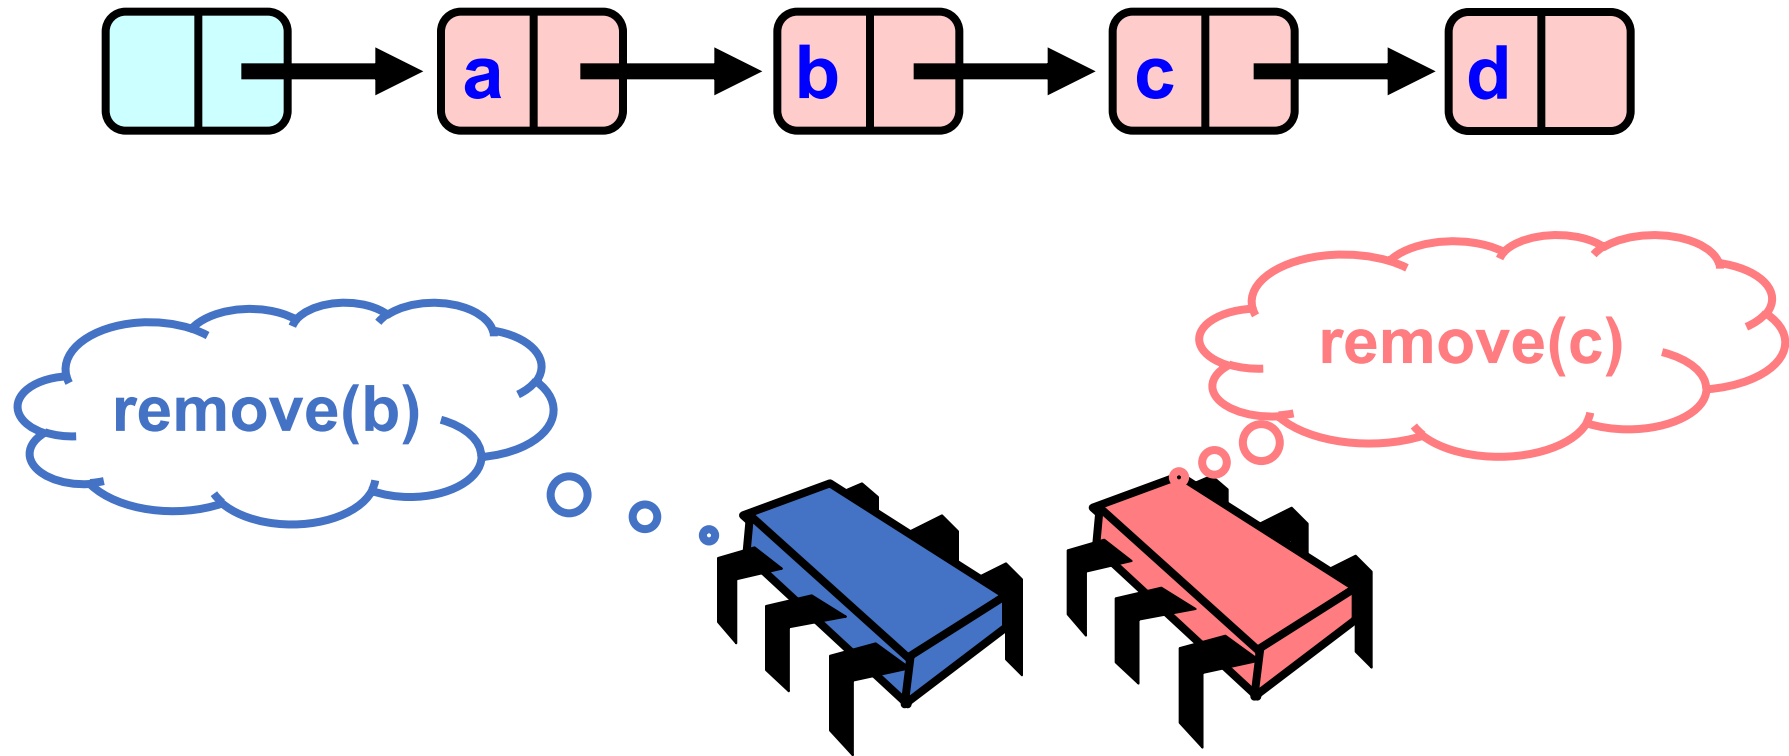
Schedule

- Concurrent set
 - Coarse-grained lock
 - **fine-grained lock**
 - optimistic locking

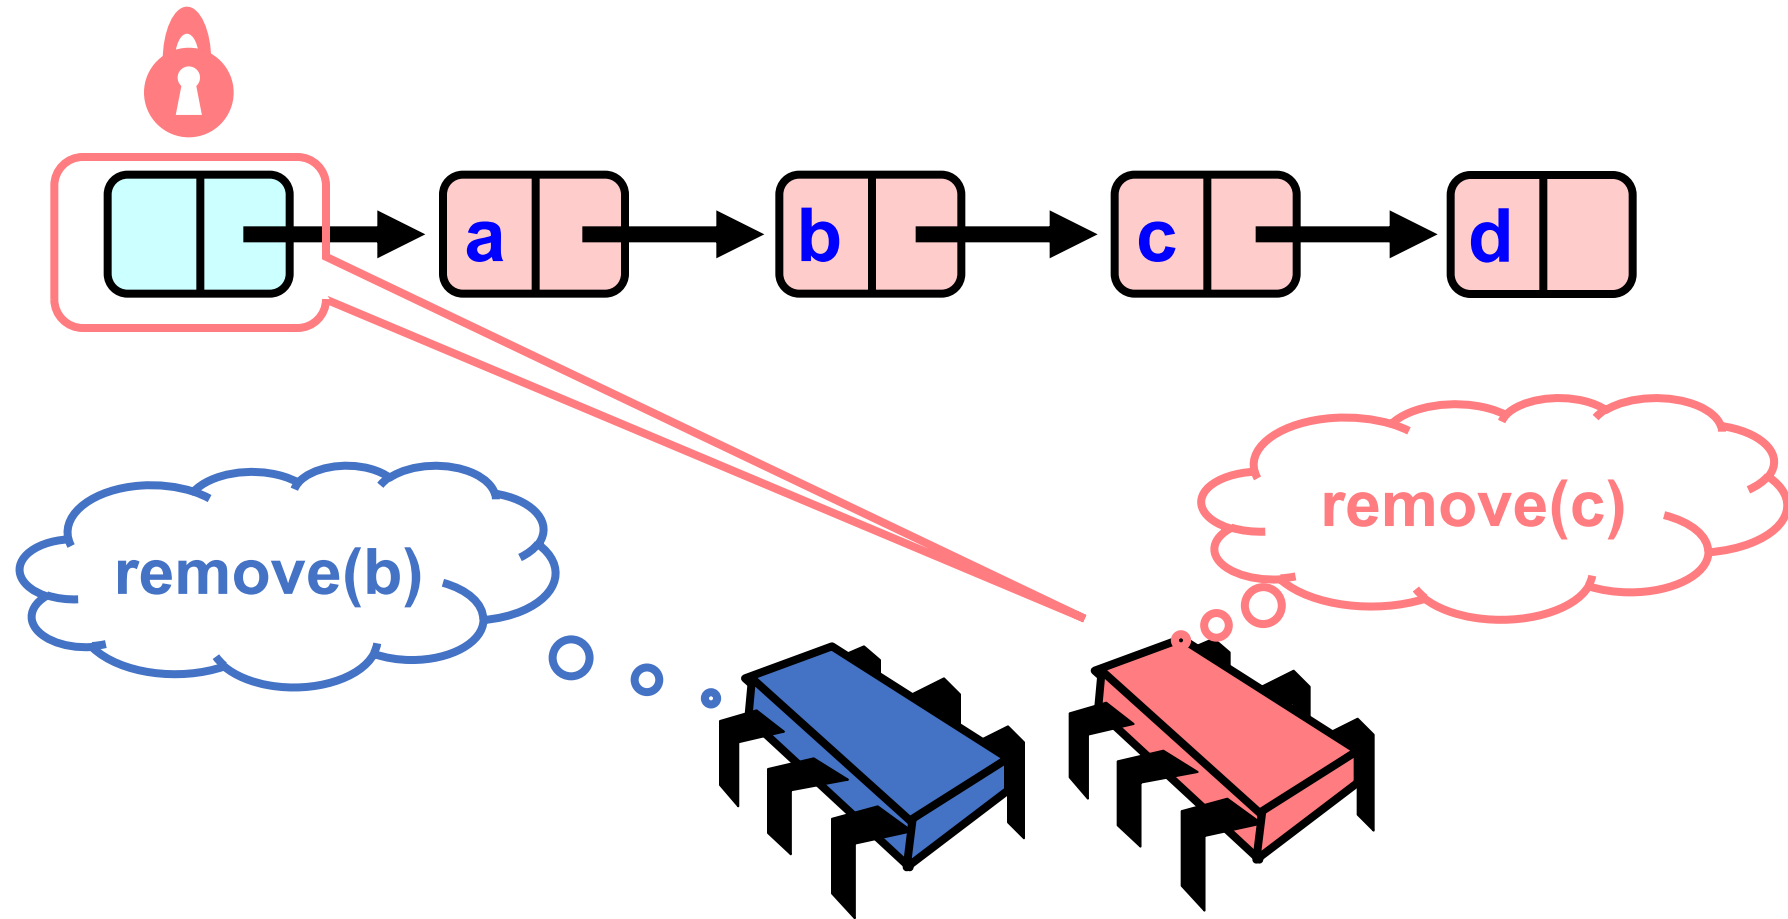
Fine-grained Locking

- Split object into pieces
 - Each piece has own lock
 - Methods that work on disjoint pieces need not exclude each other
- Requires **careful** thought
 - Acquire all required locks
 - Acquire them in the same order

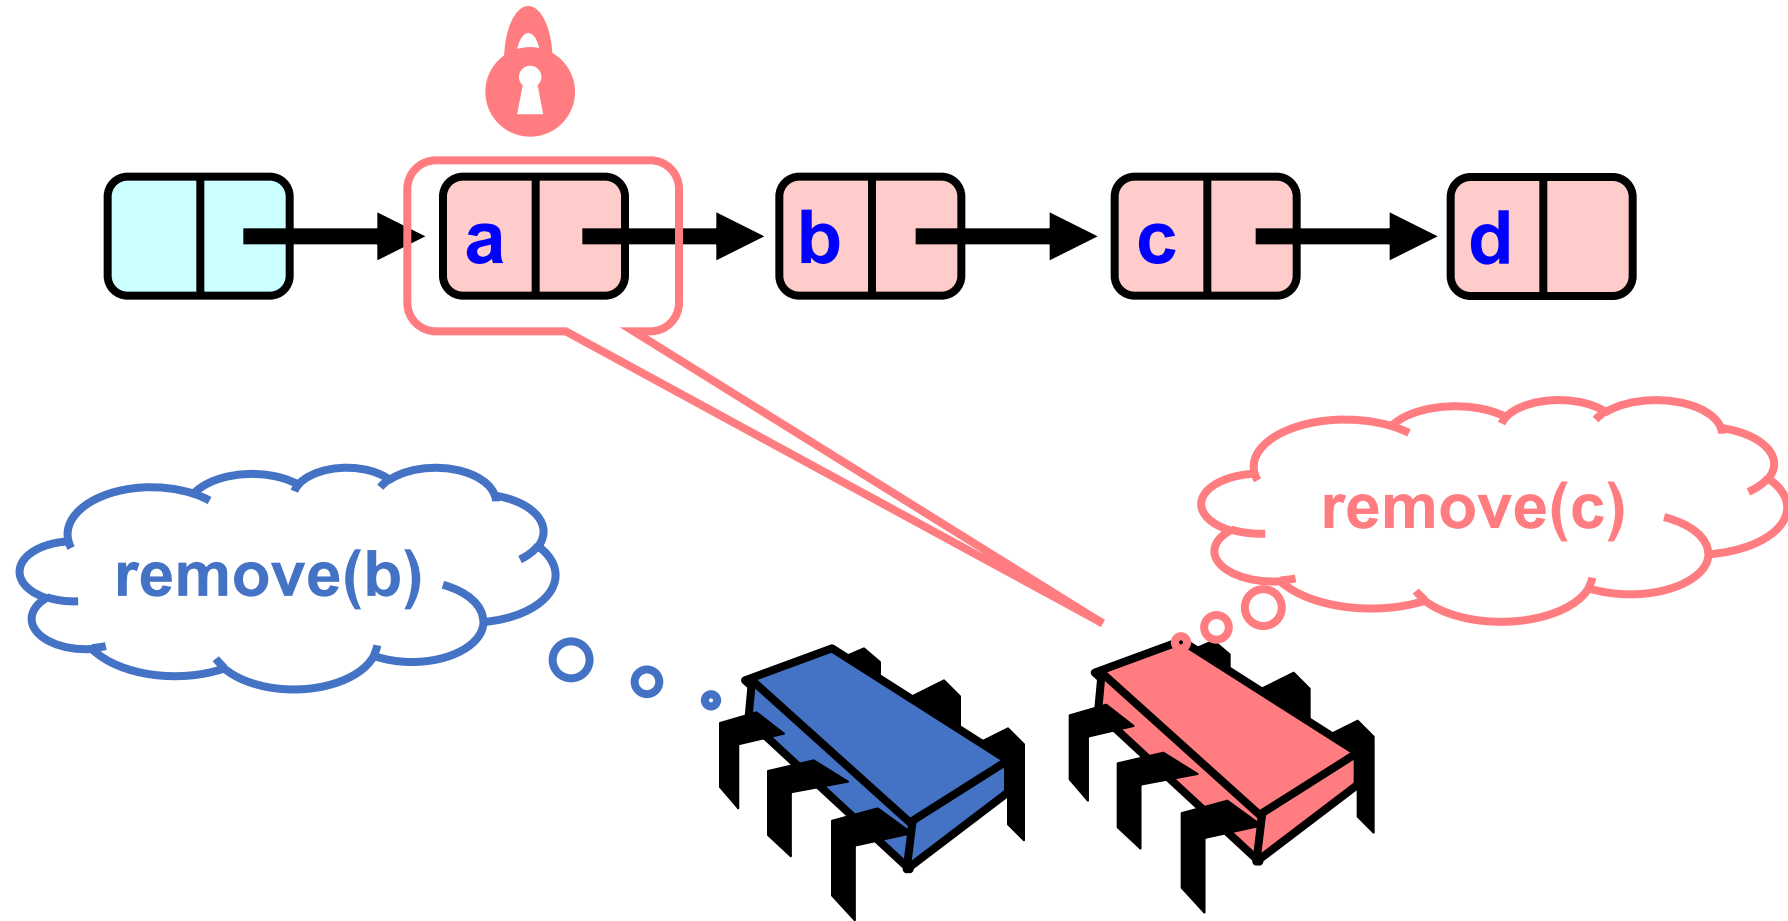
Concurrent Removes



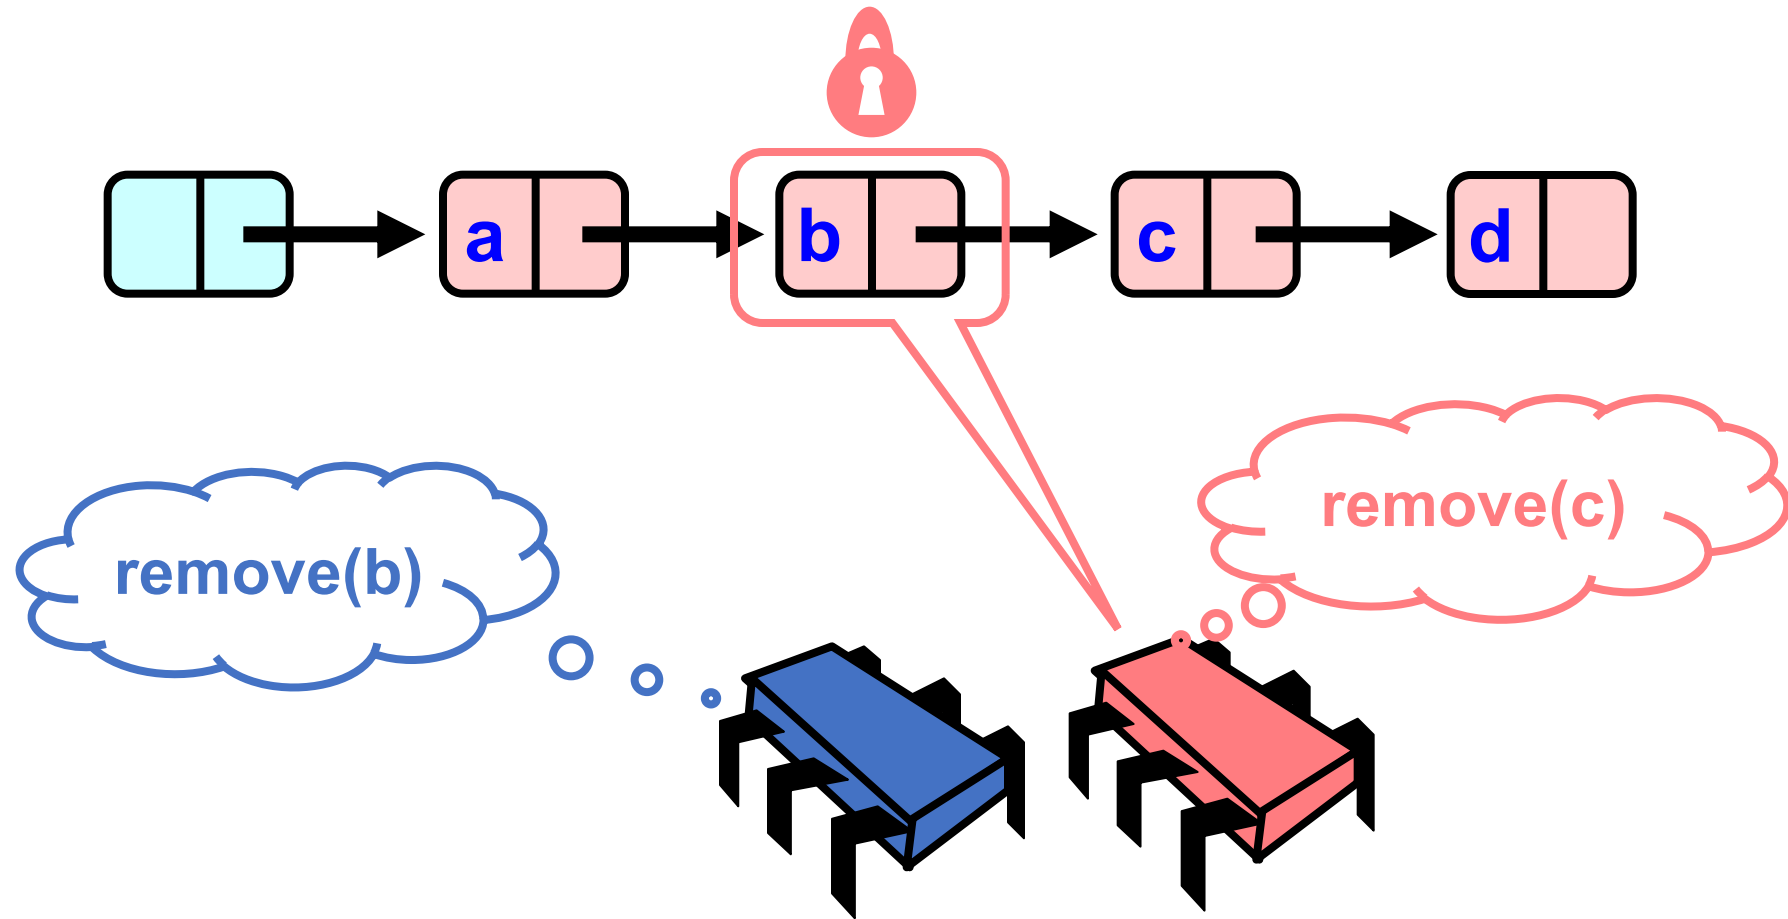
Concurrent Removes



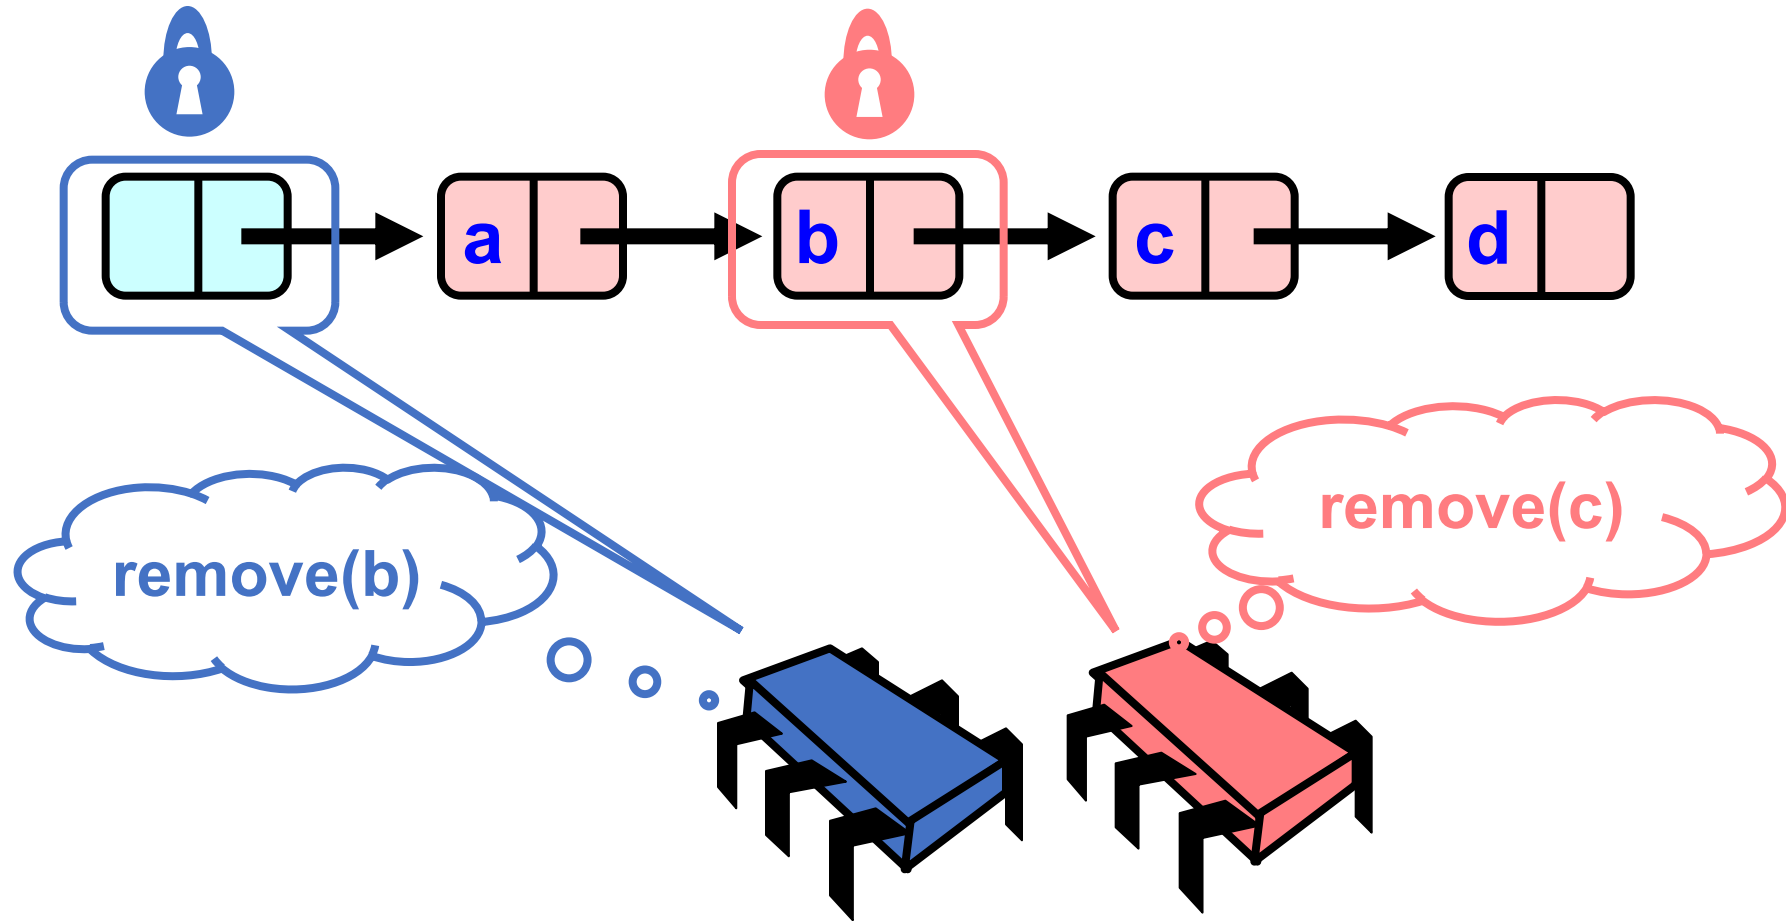
Concurrent Removes



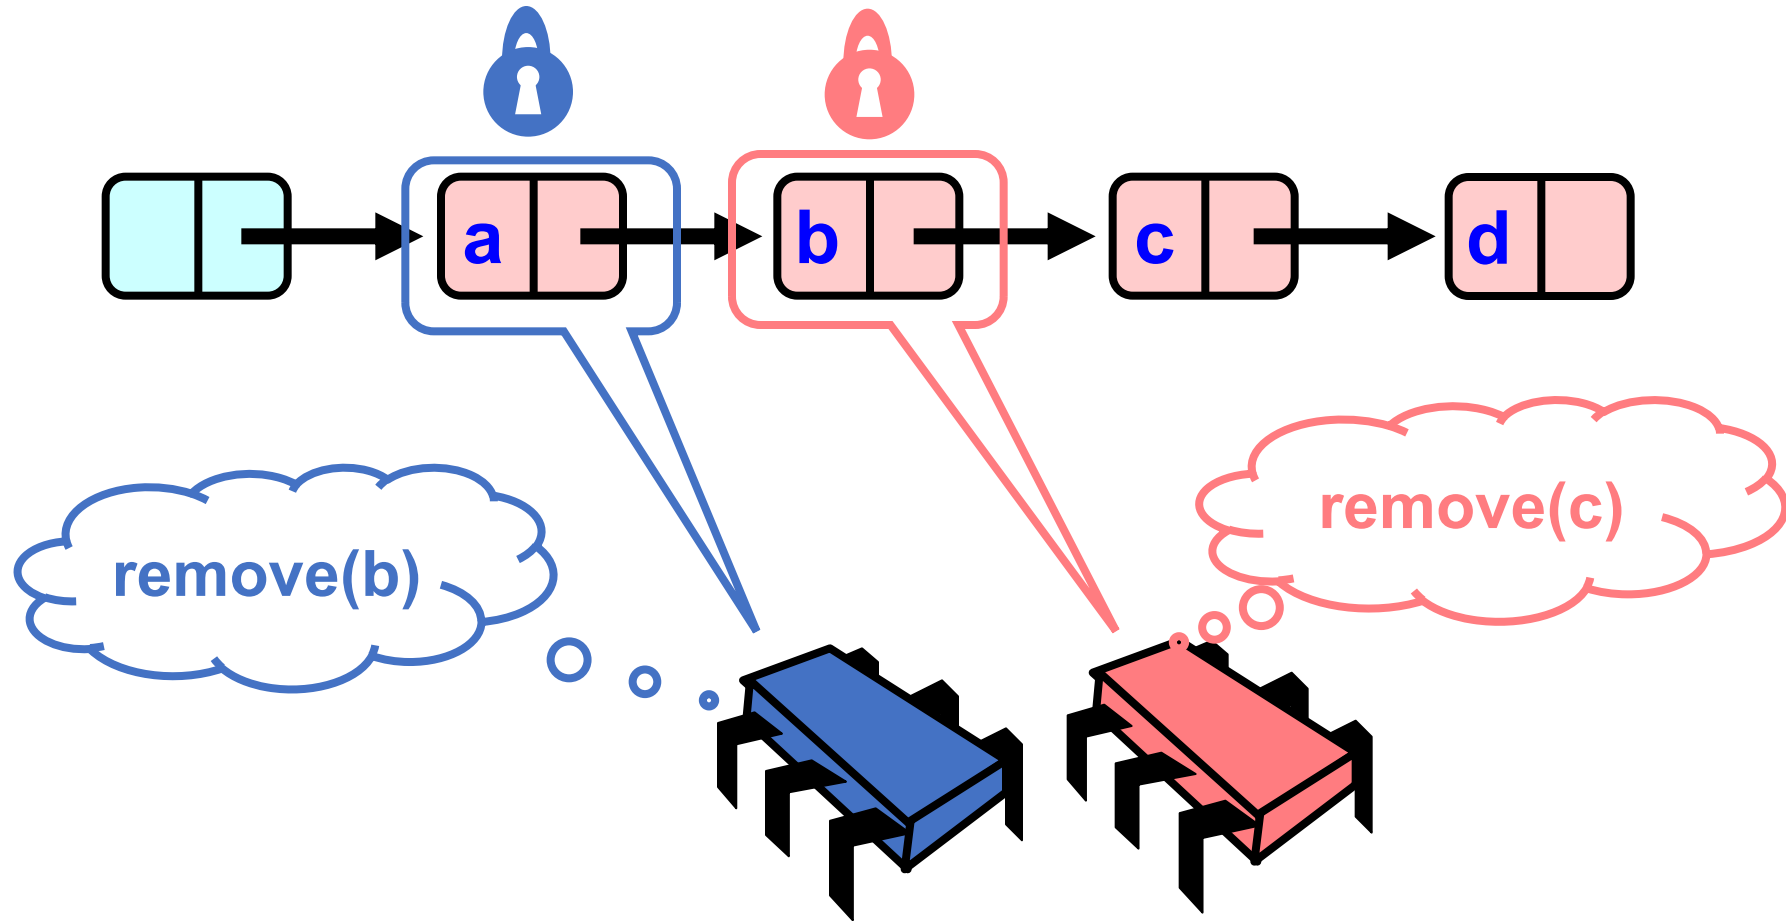
Concurrent Removes




Concurrent Removes



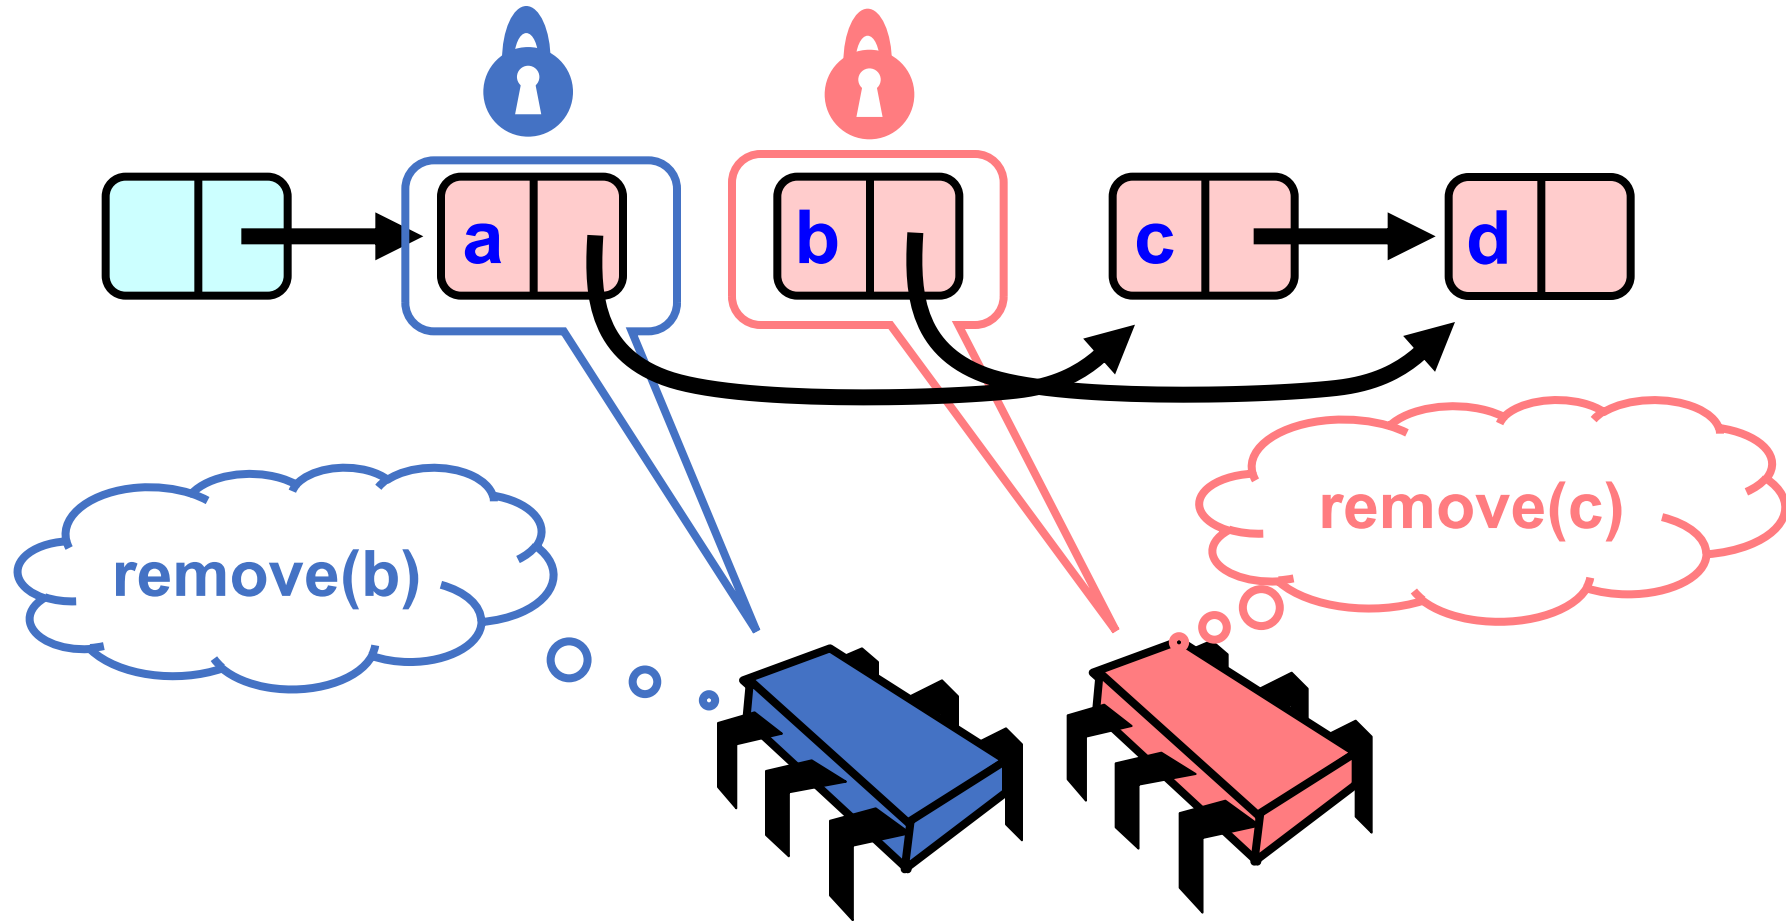
Concurrent Removes



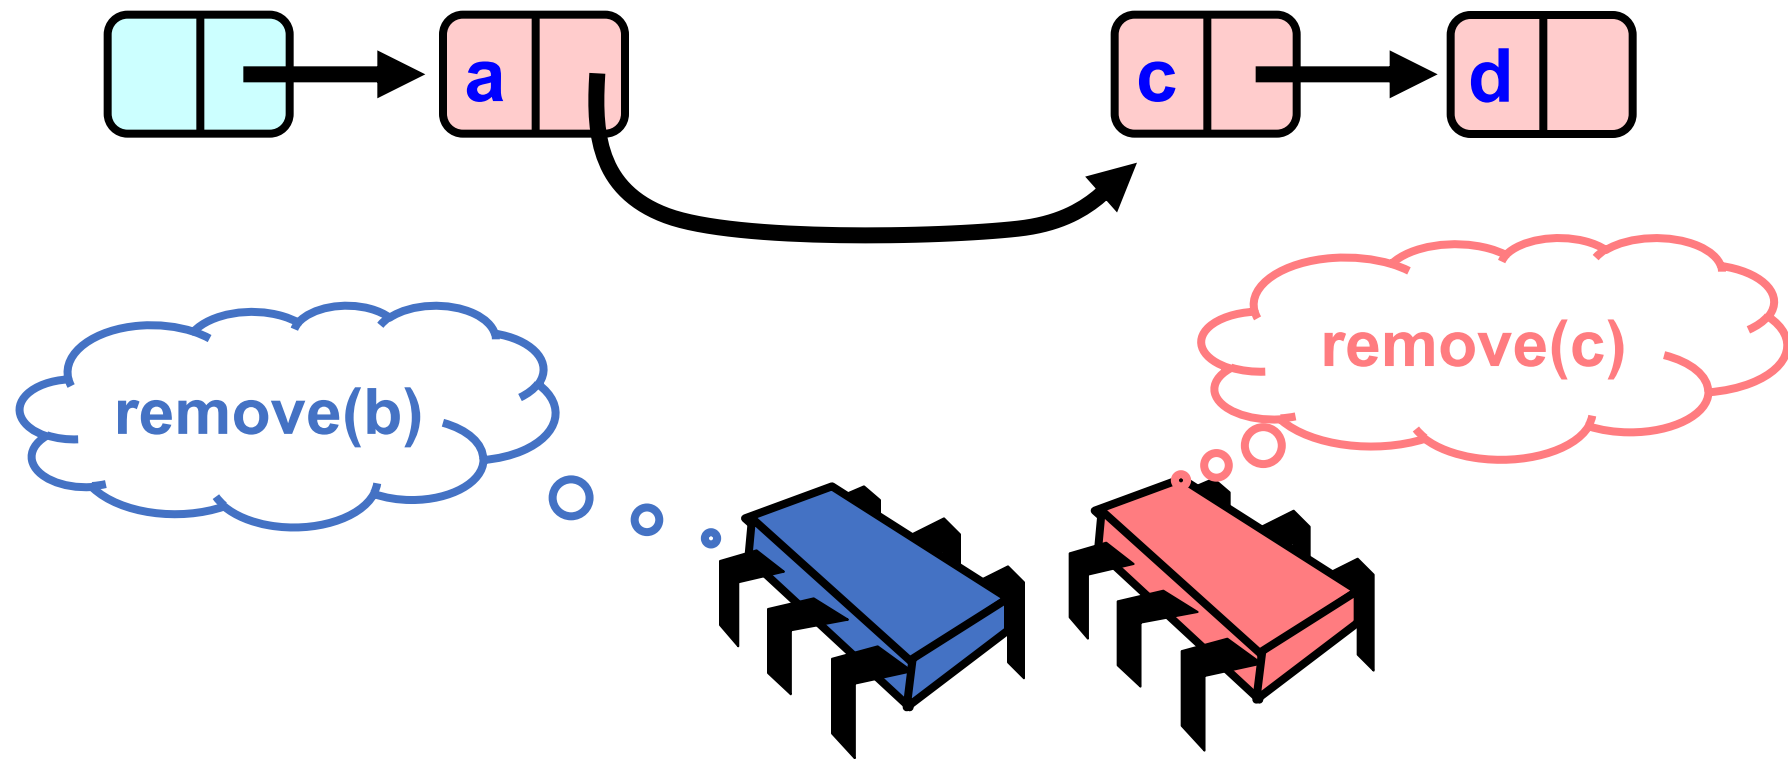
Concurrent Removes



Concurrent Removes




Uh, Oh



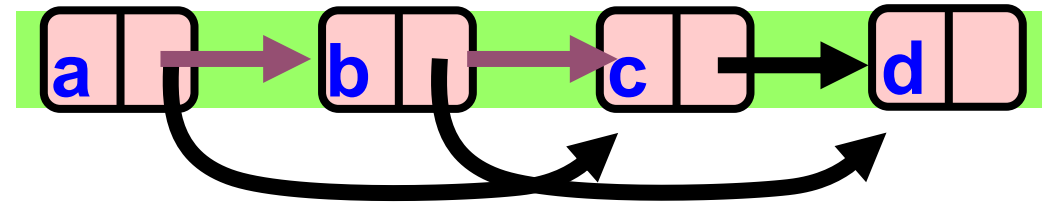
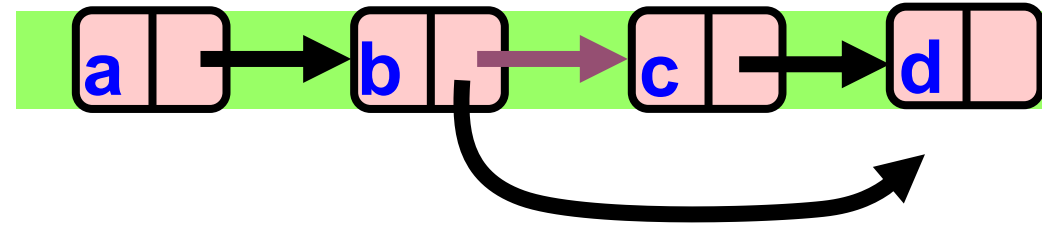
Uh, Oh

Bad news, c not removed

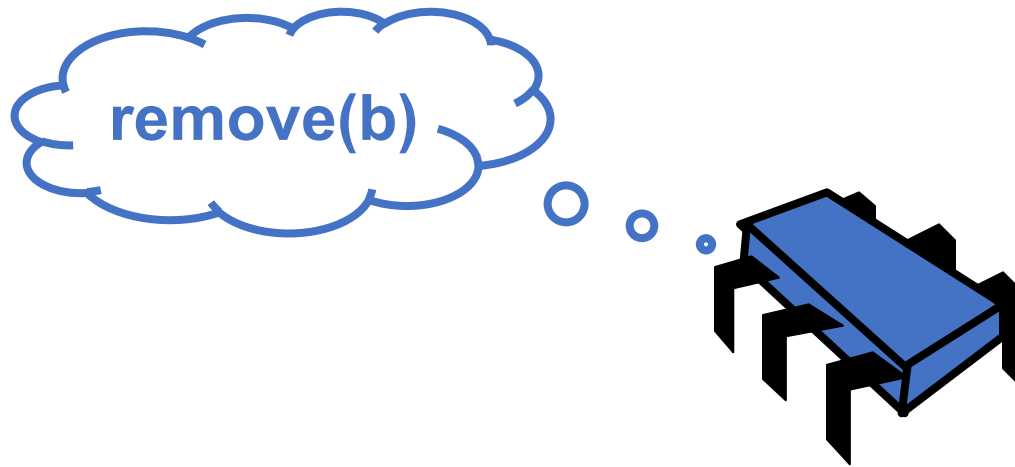
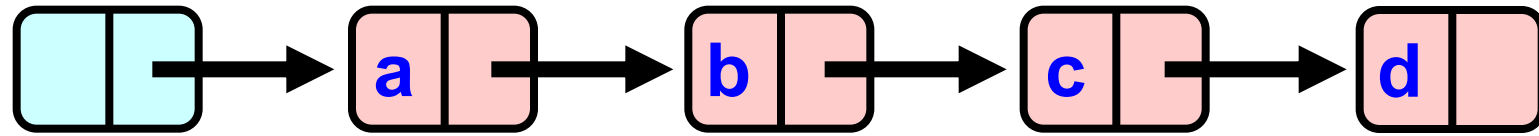


Problem

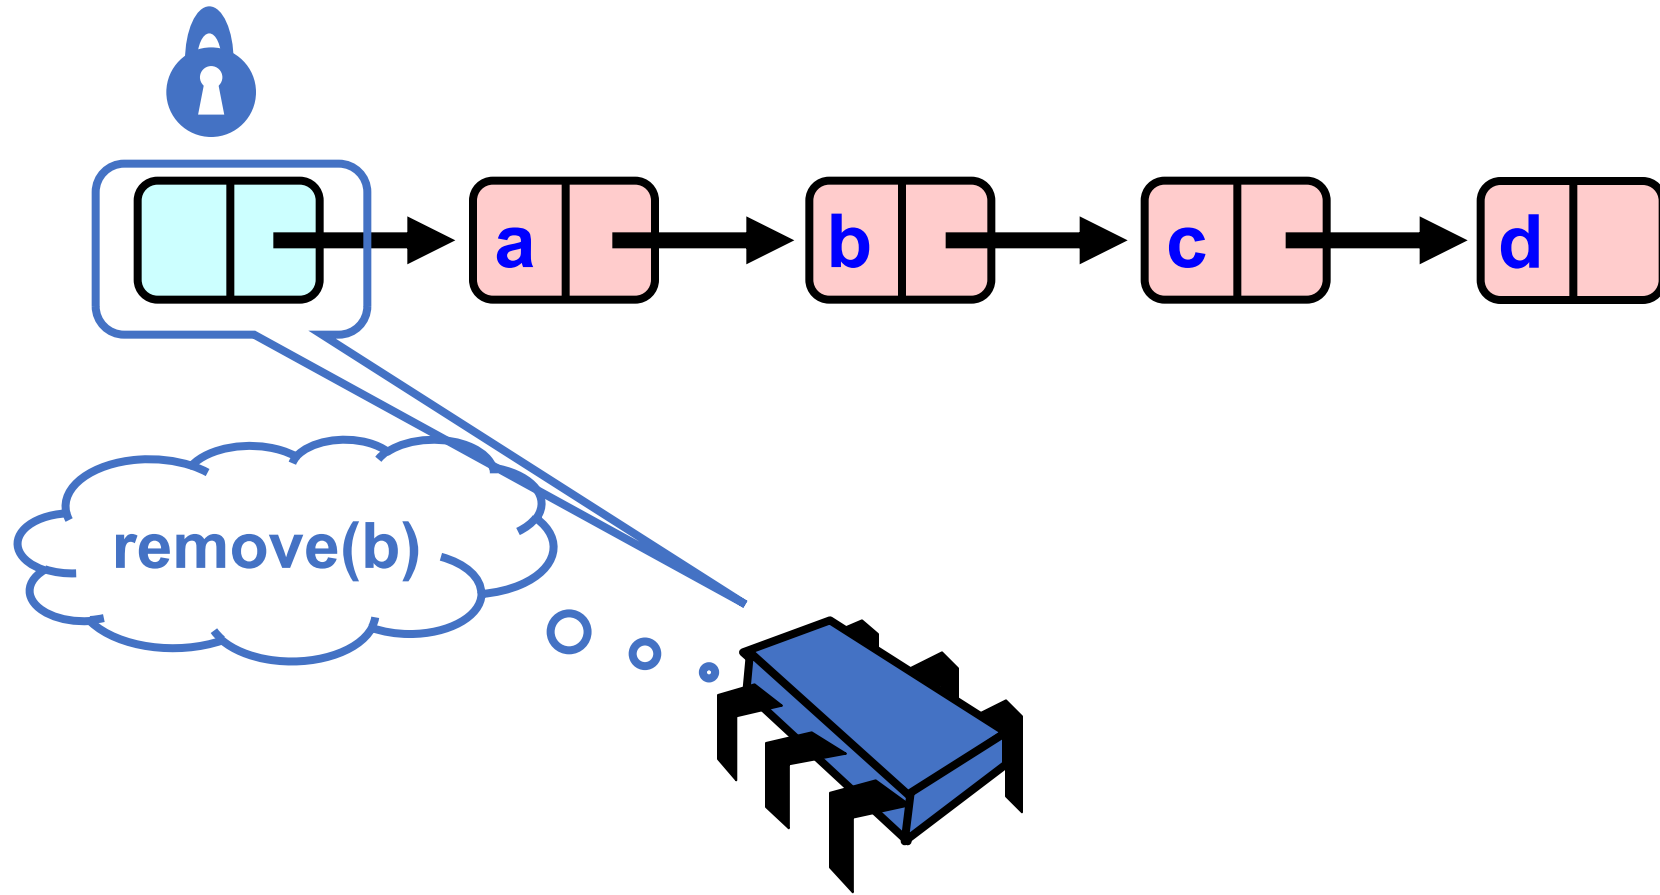
- To delete node c
 - Swing node b's next field to d
- Problem is,
 - Someone deleting b concurrently could direct a pointer to C



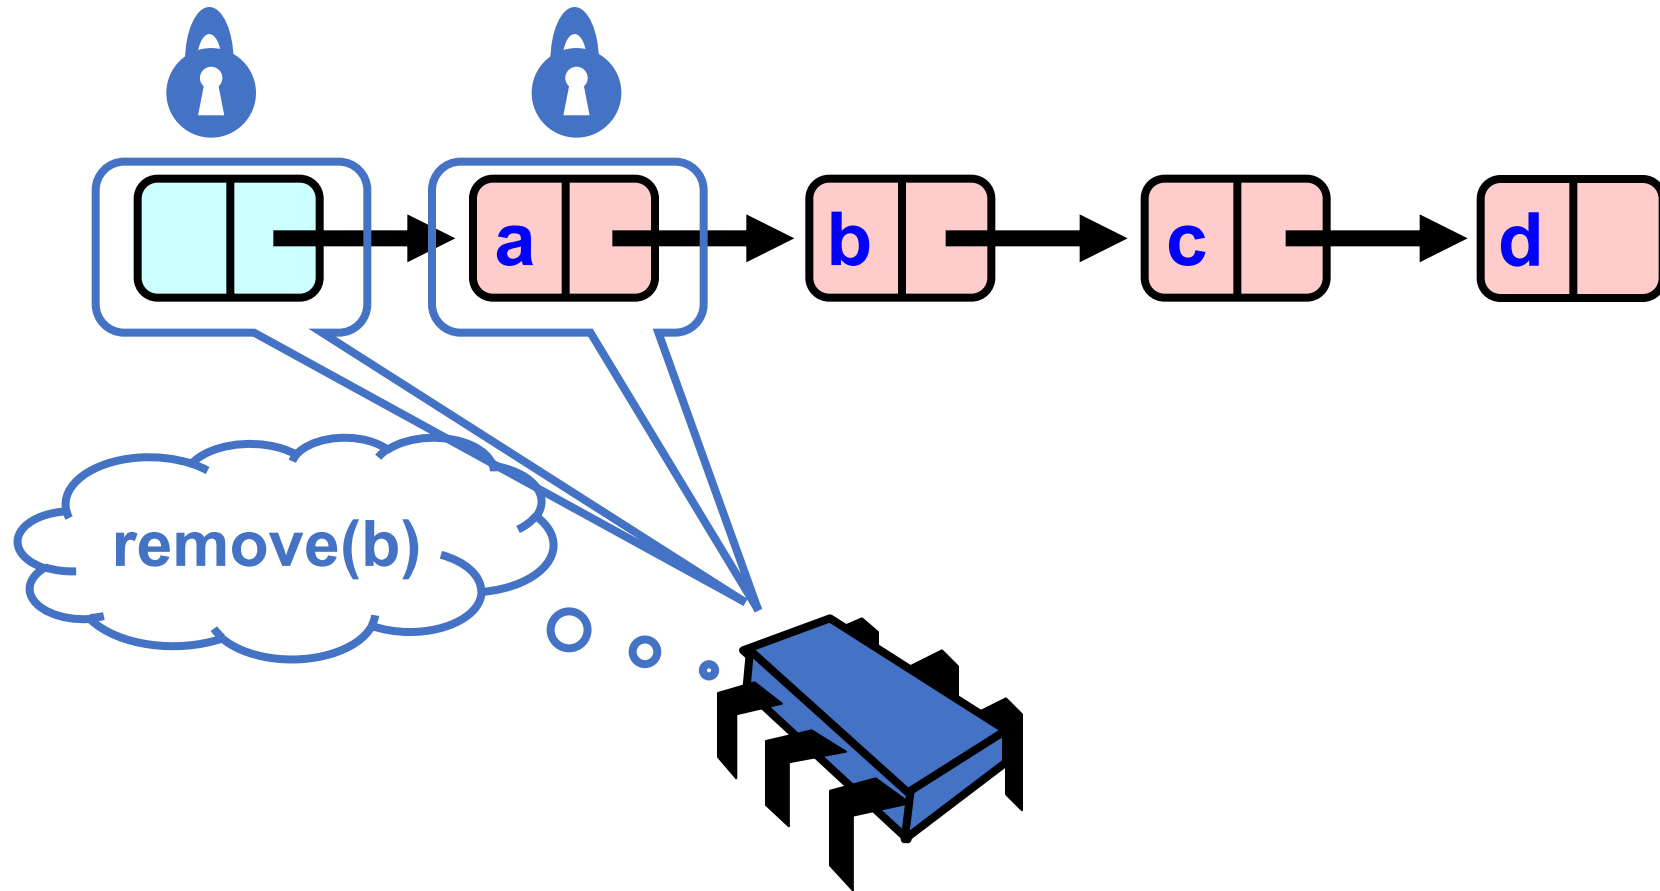
Hand-Over-Hand Again



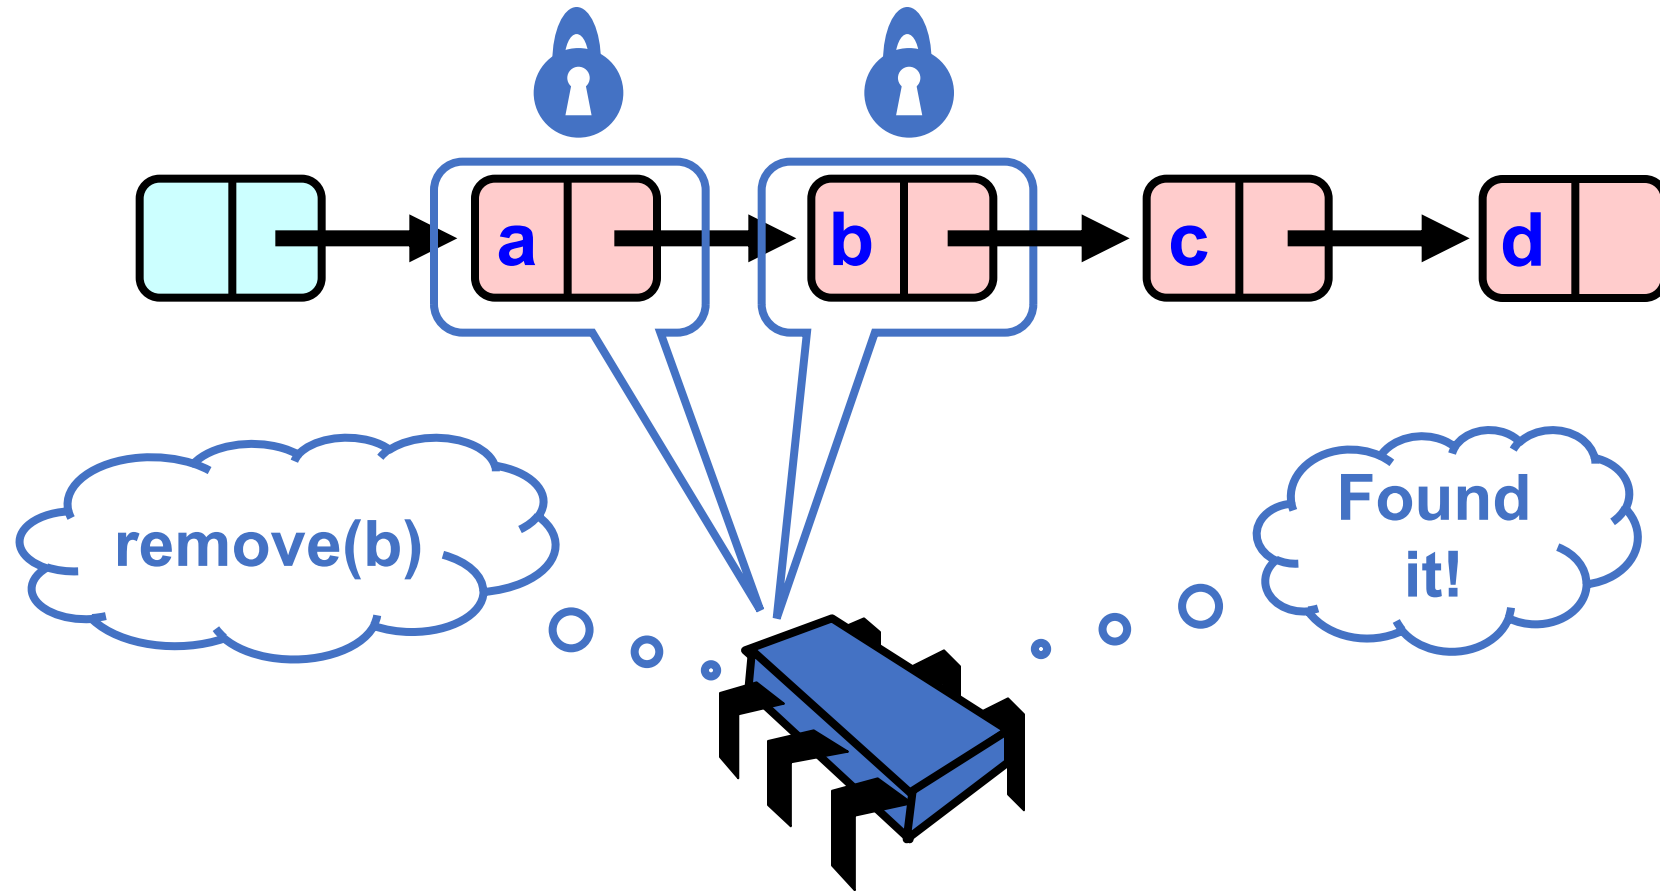
Hand-Over-Hand Again



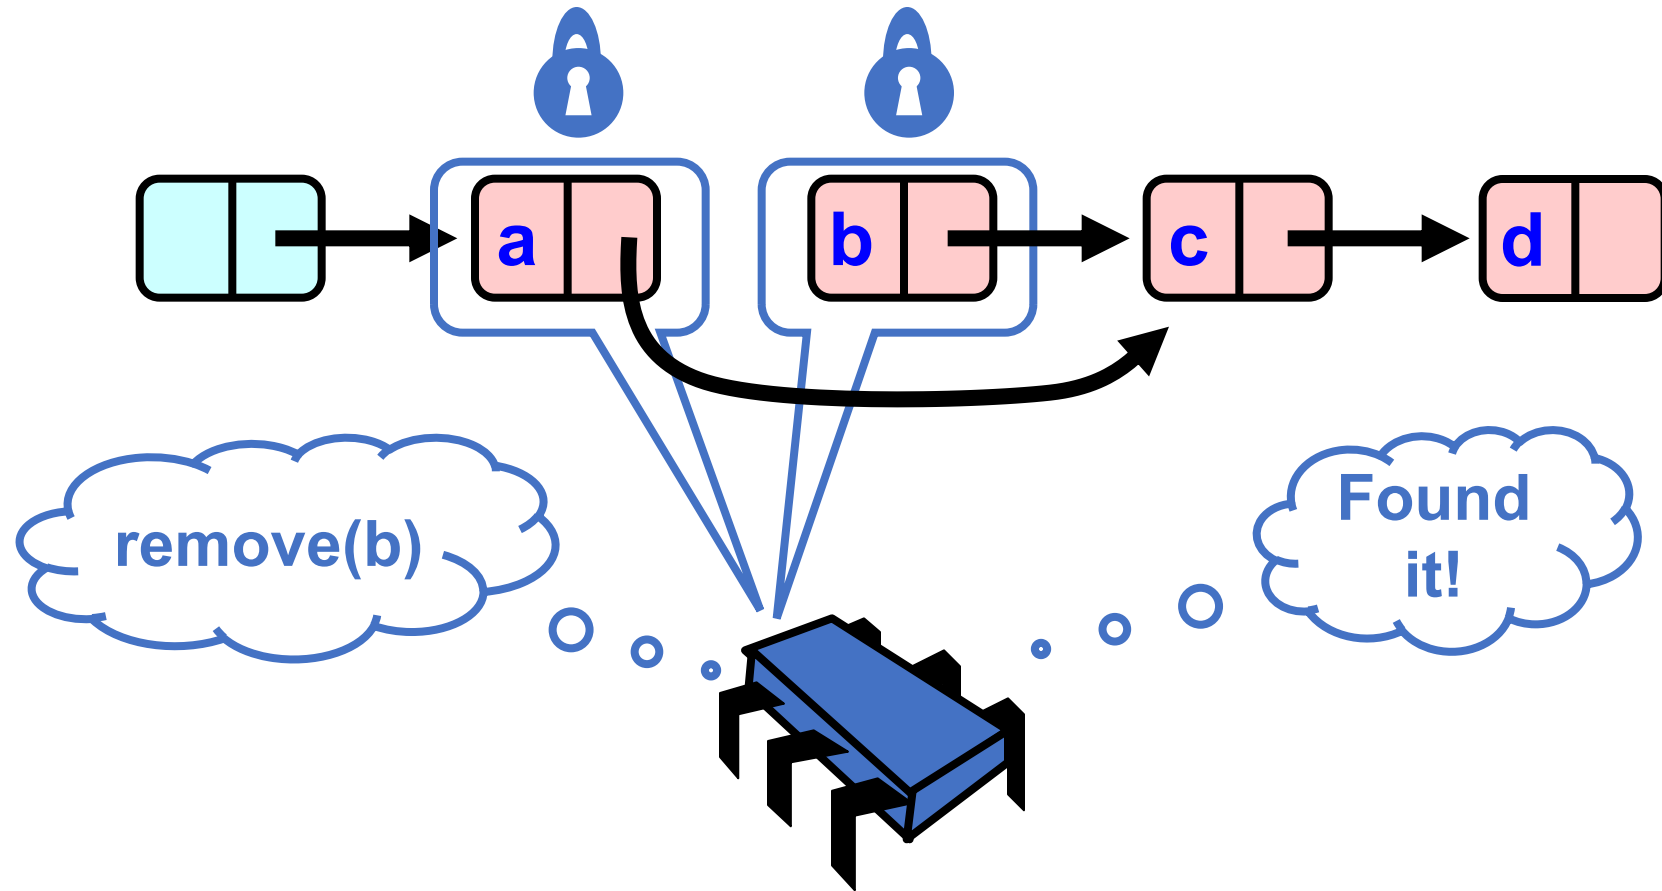
Hand-Over-Hand Again




Hand-Over-Hand Again



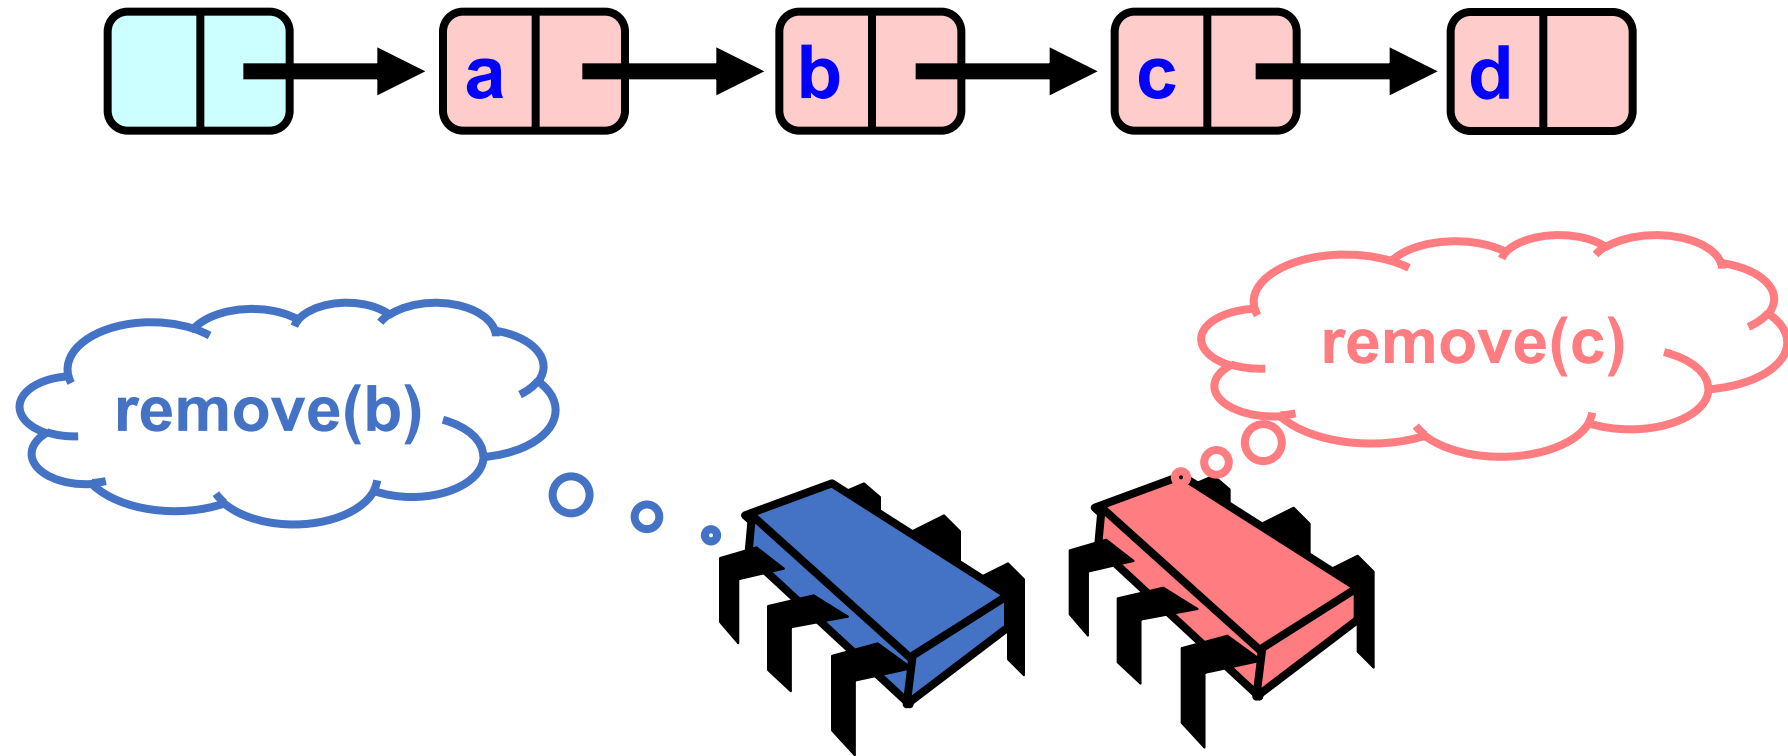
Hand-Over-Hand Again



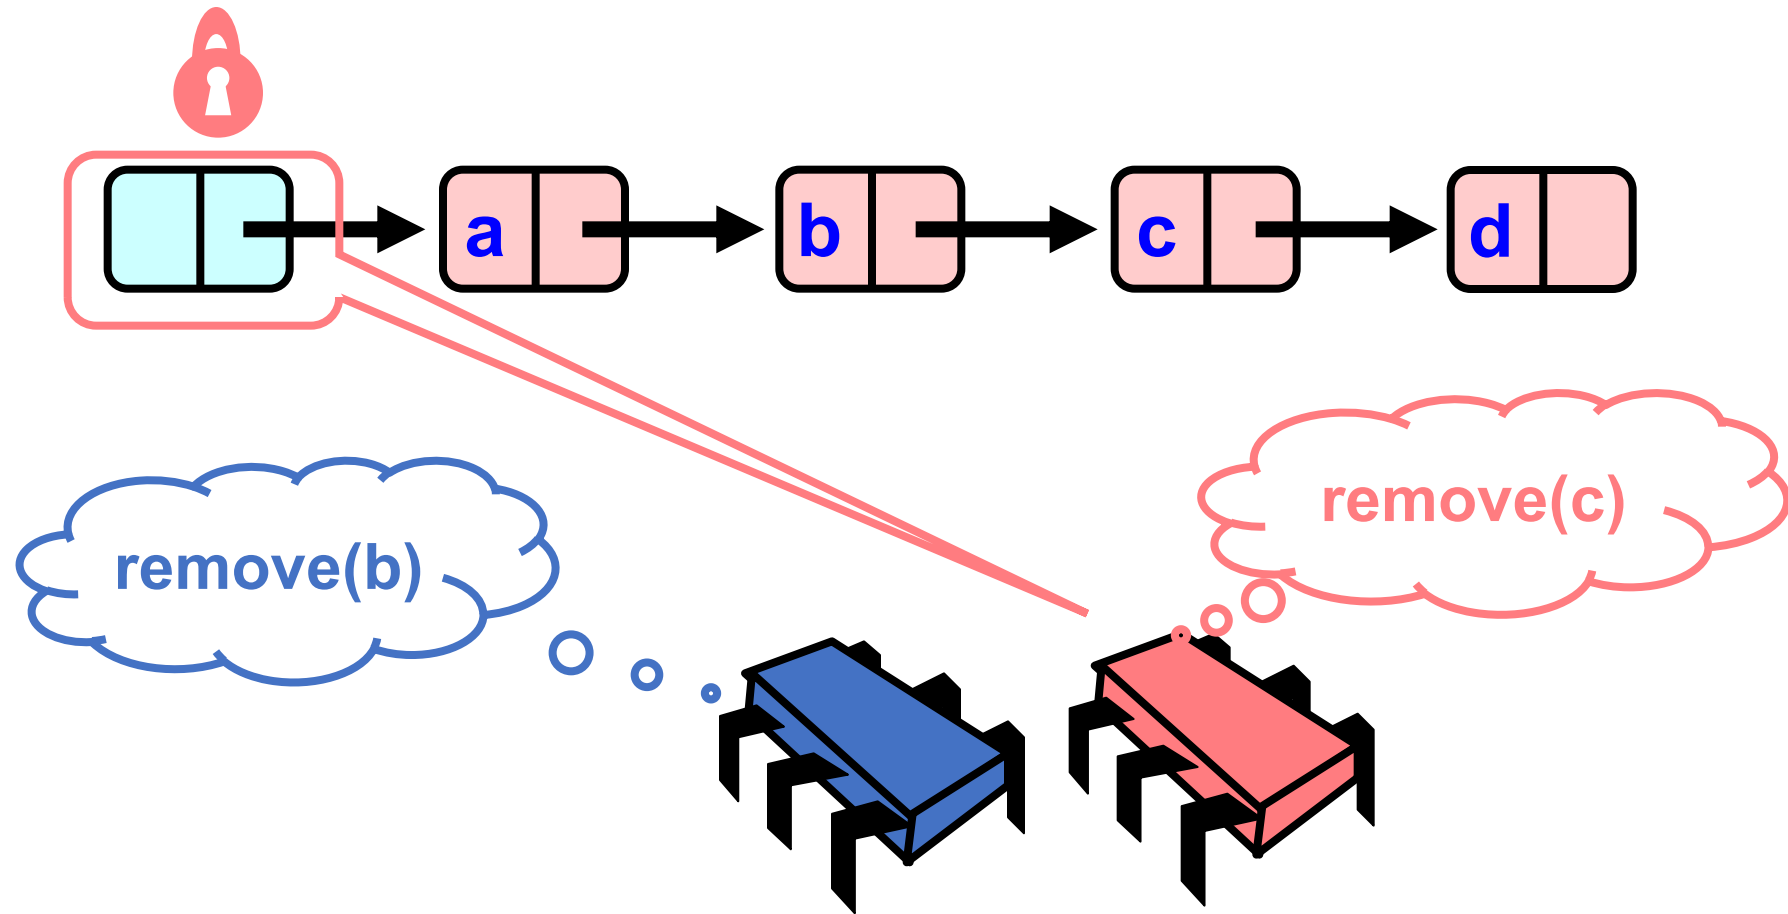
Hand-Over-Hand Again



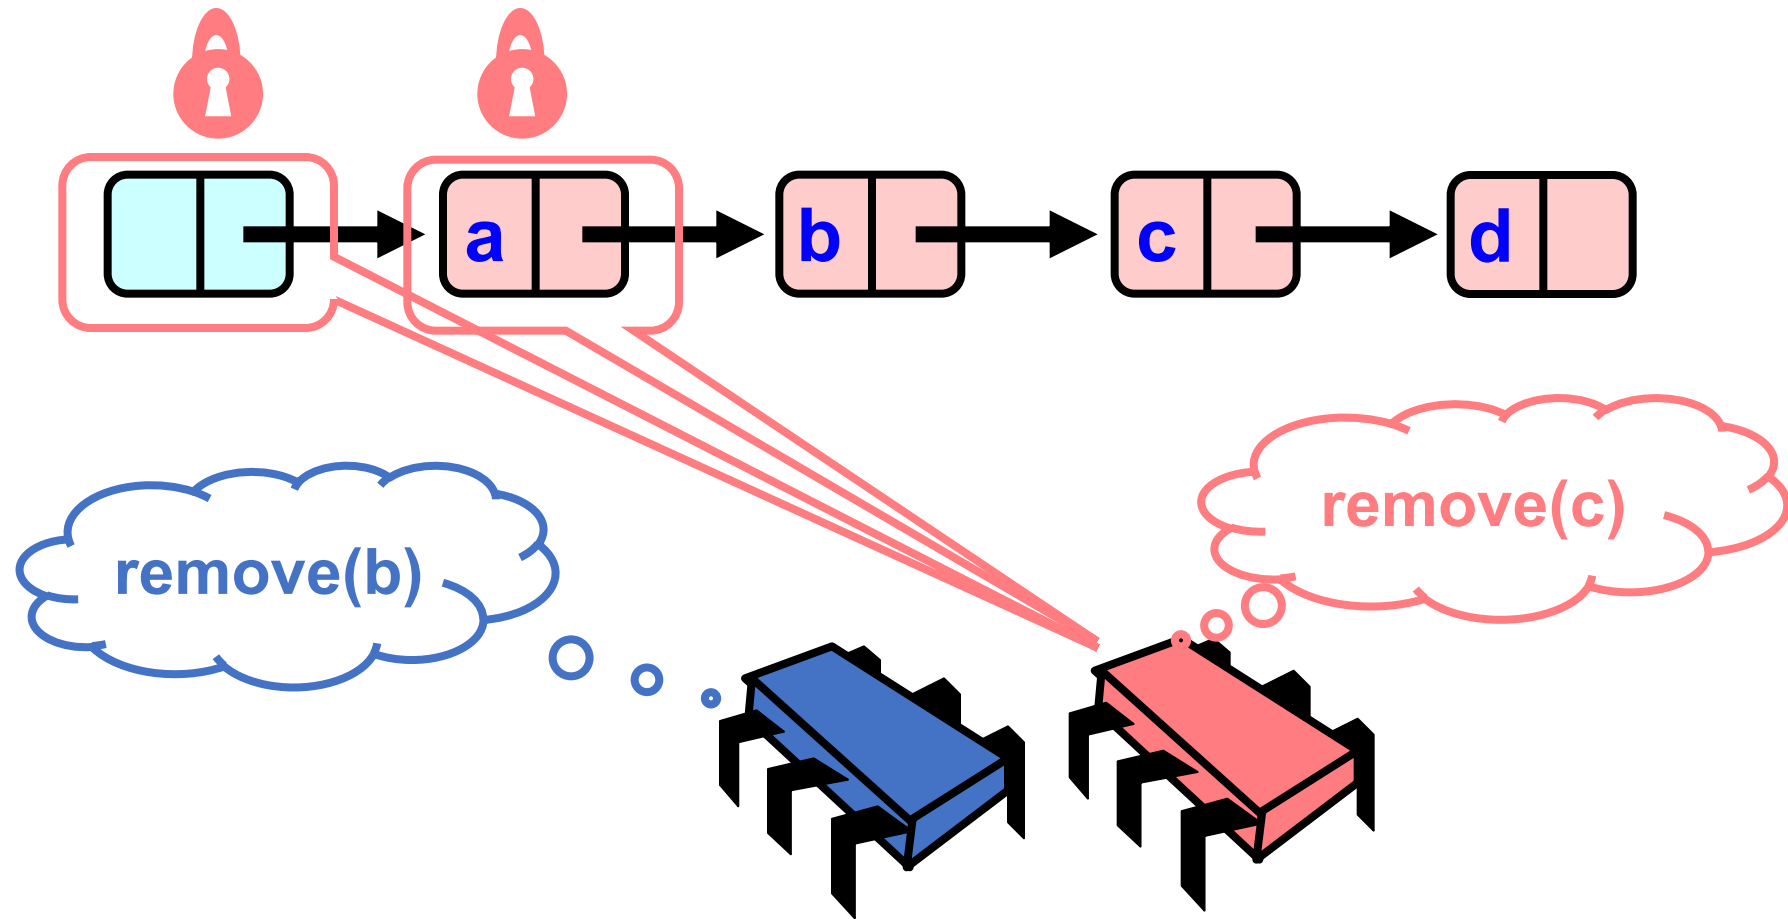
Removing a Node



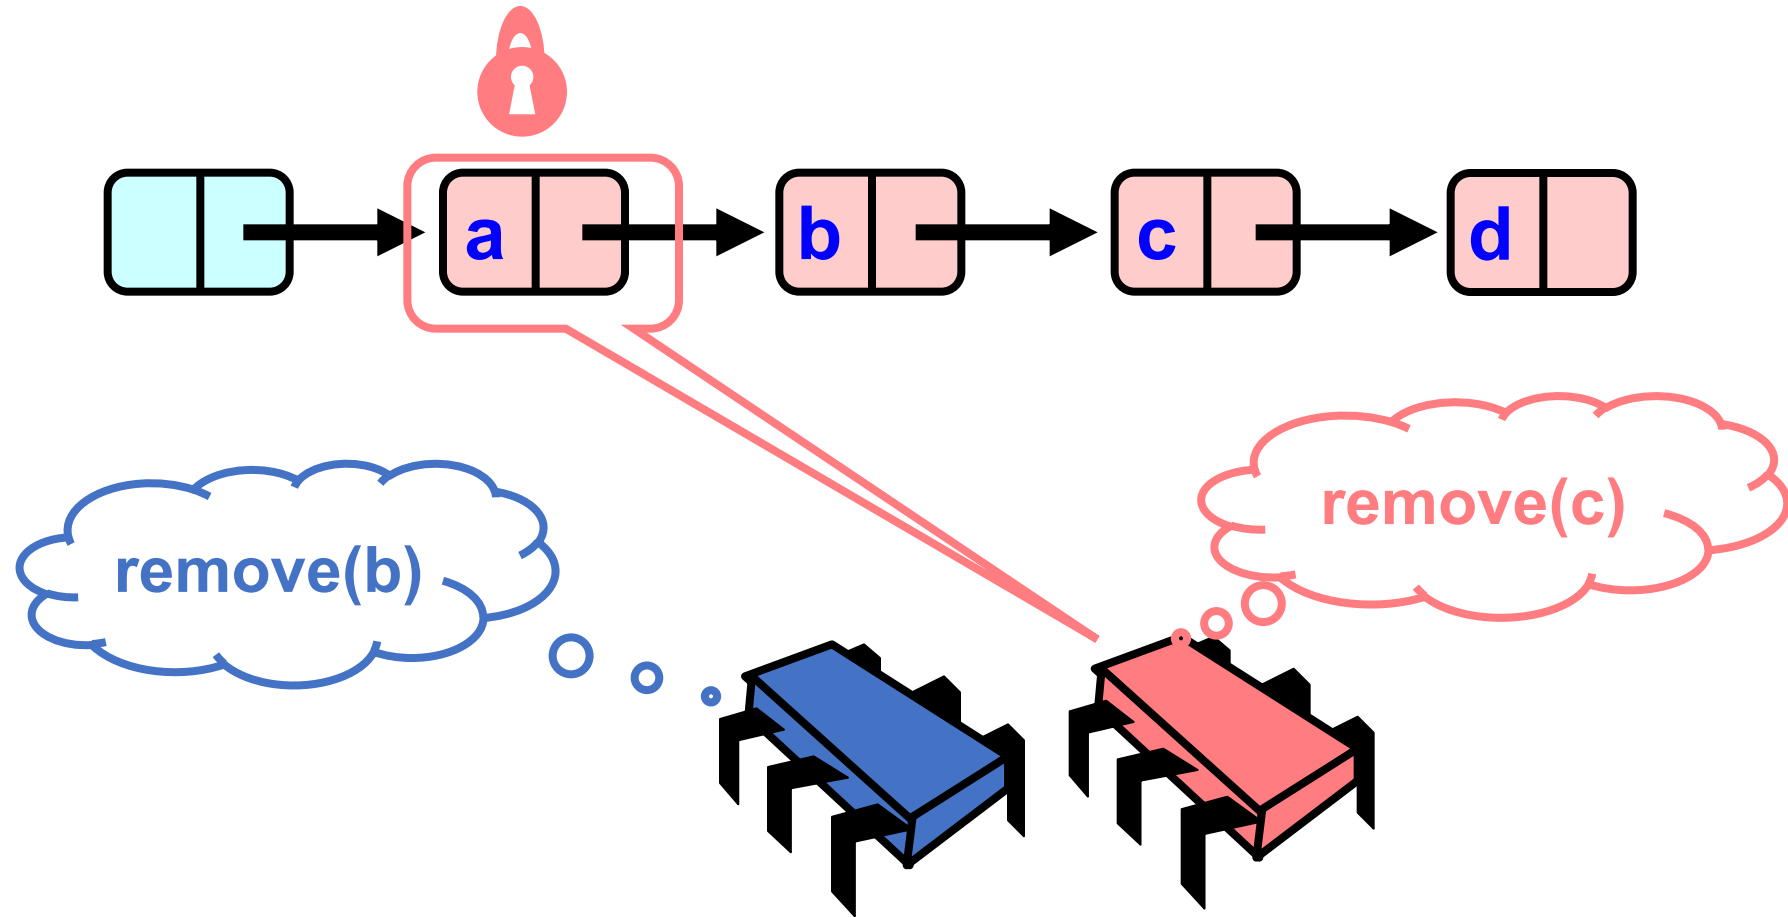
Removing a Node



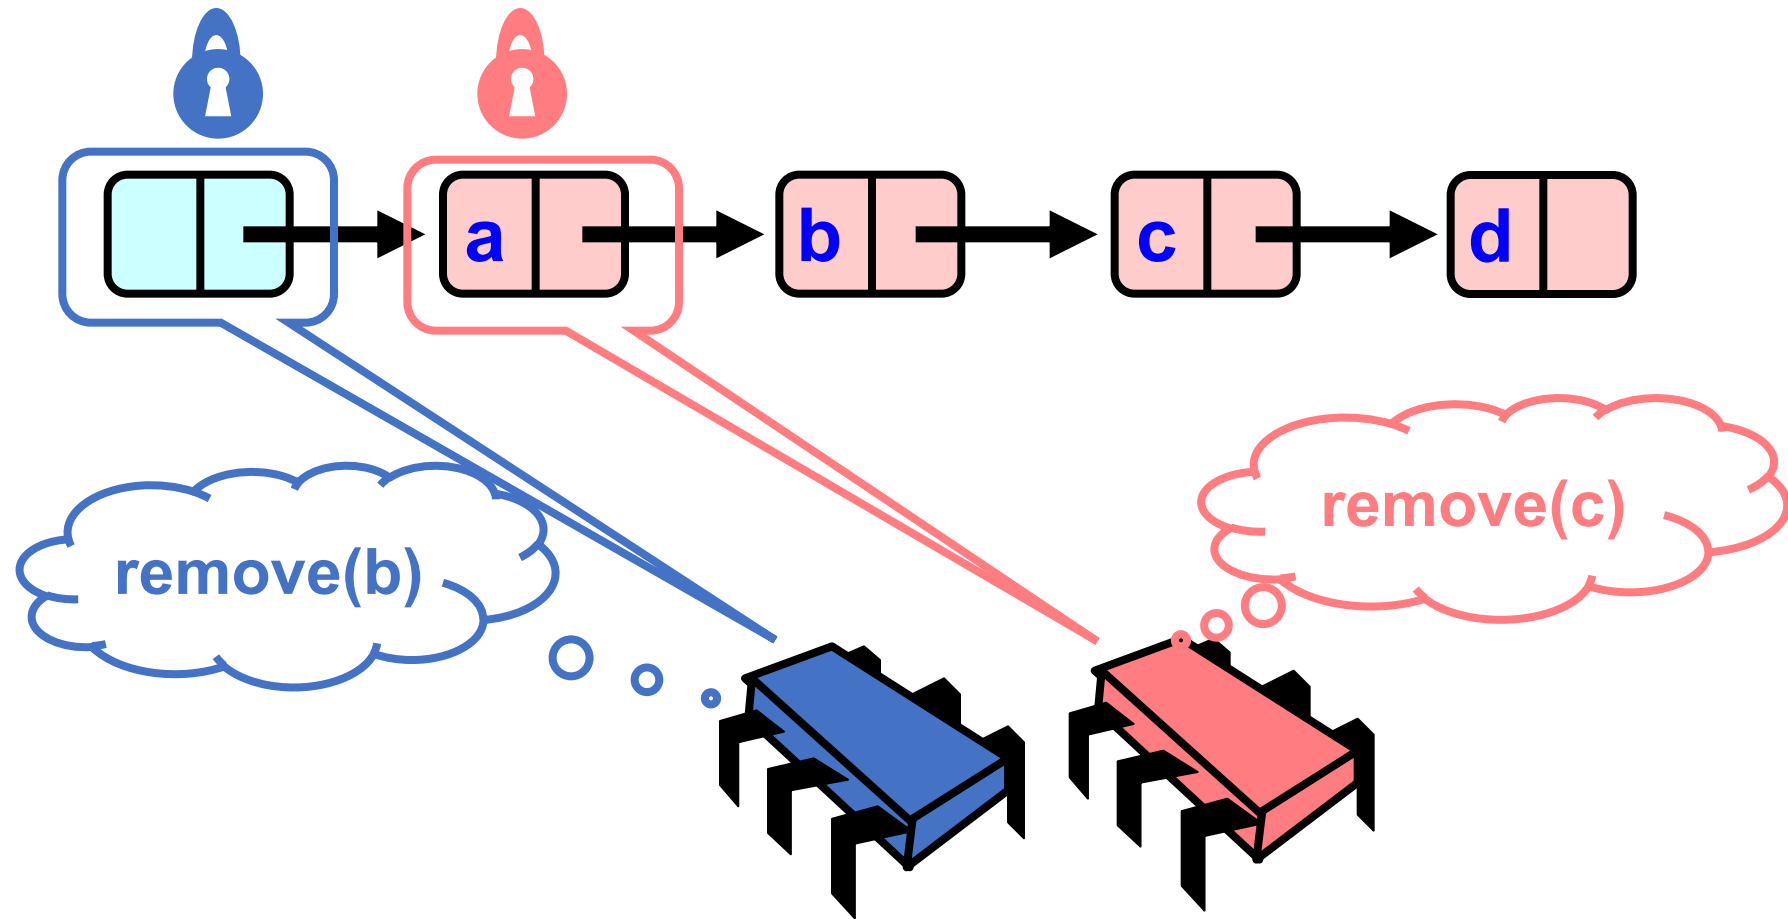
Removing a Node



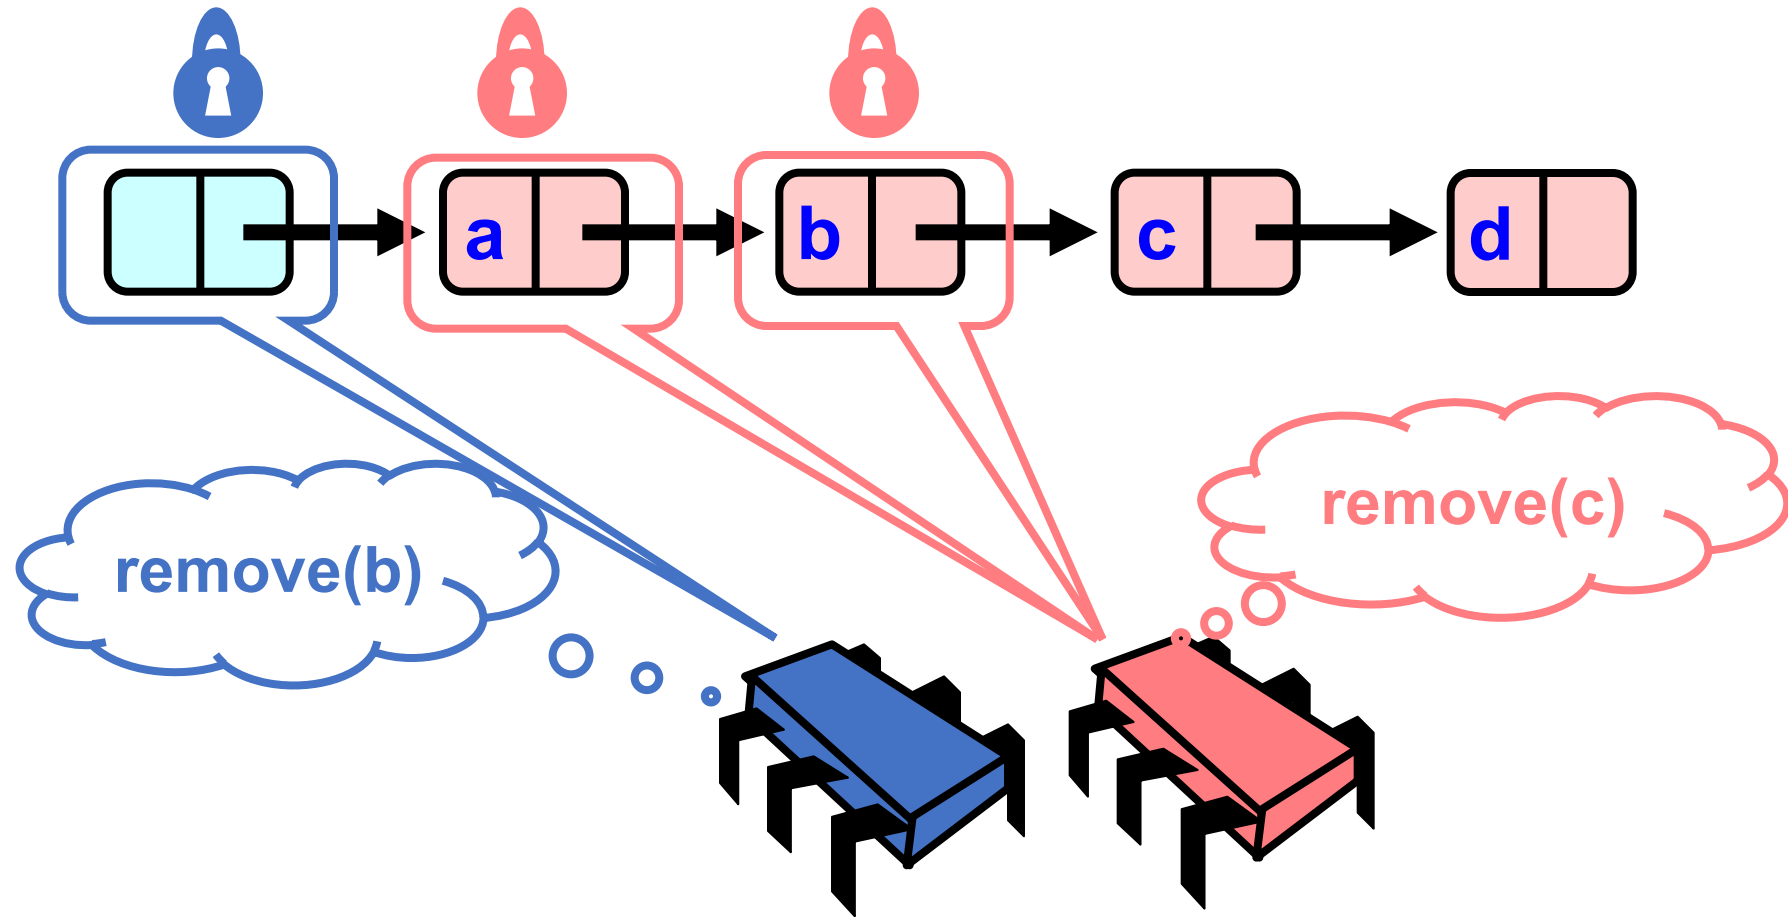
Removing a Node




Removing a Node



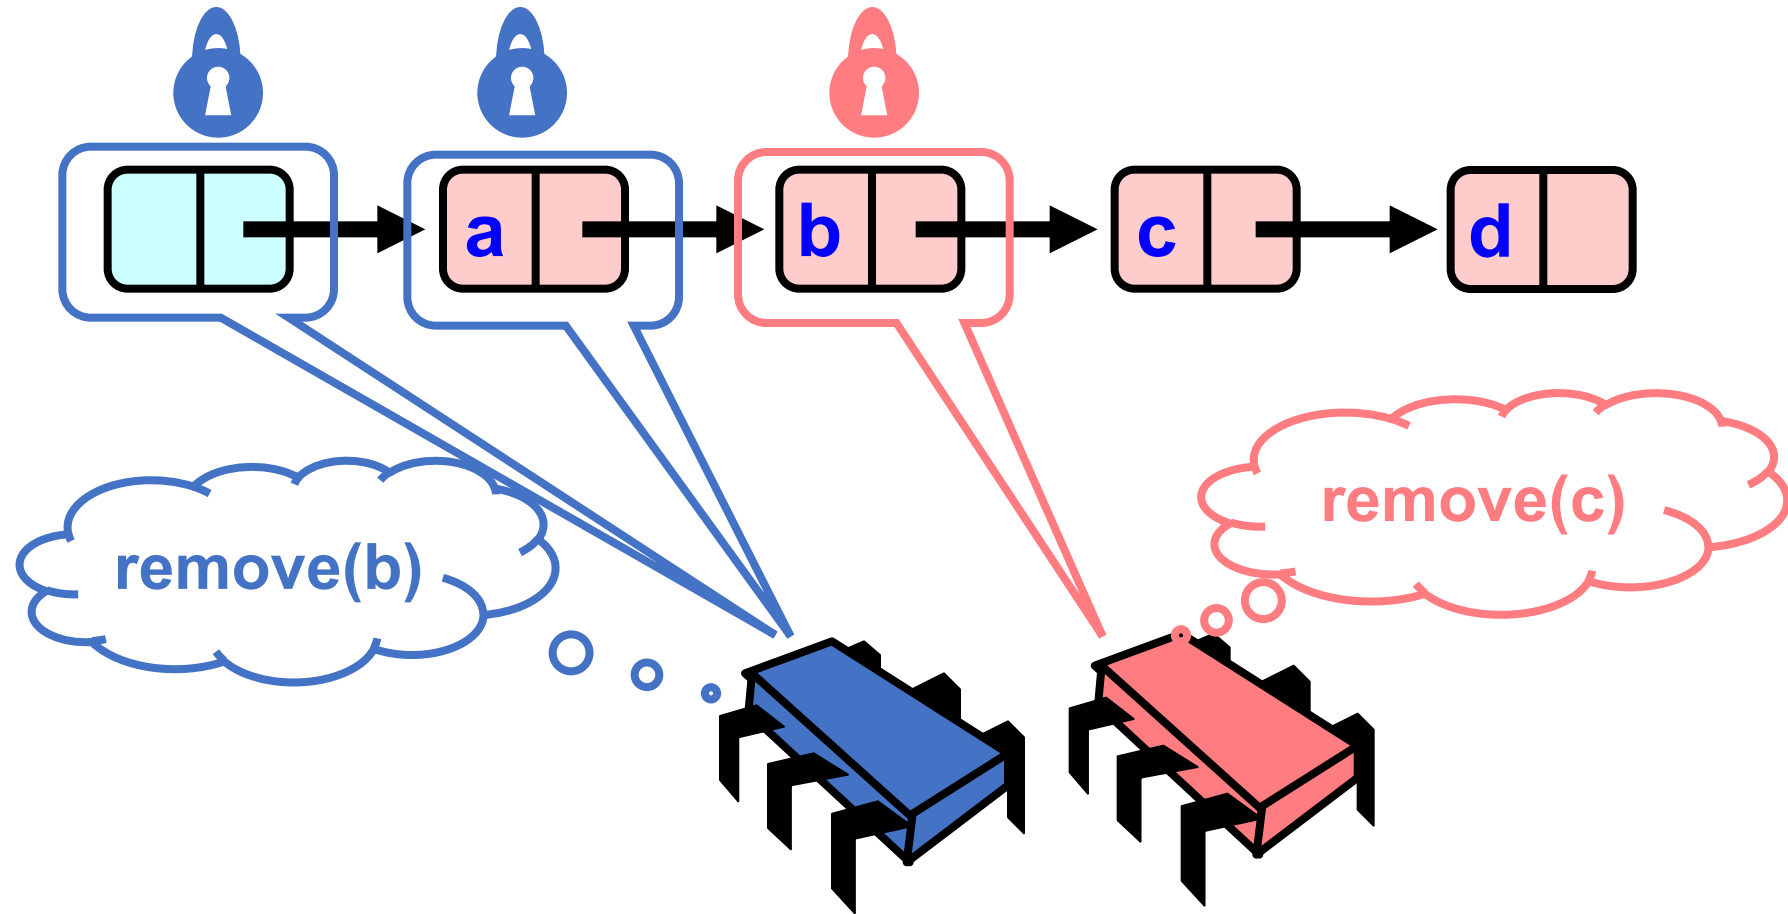
Removing a Node



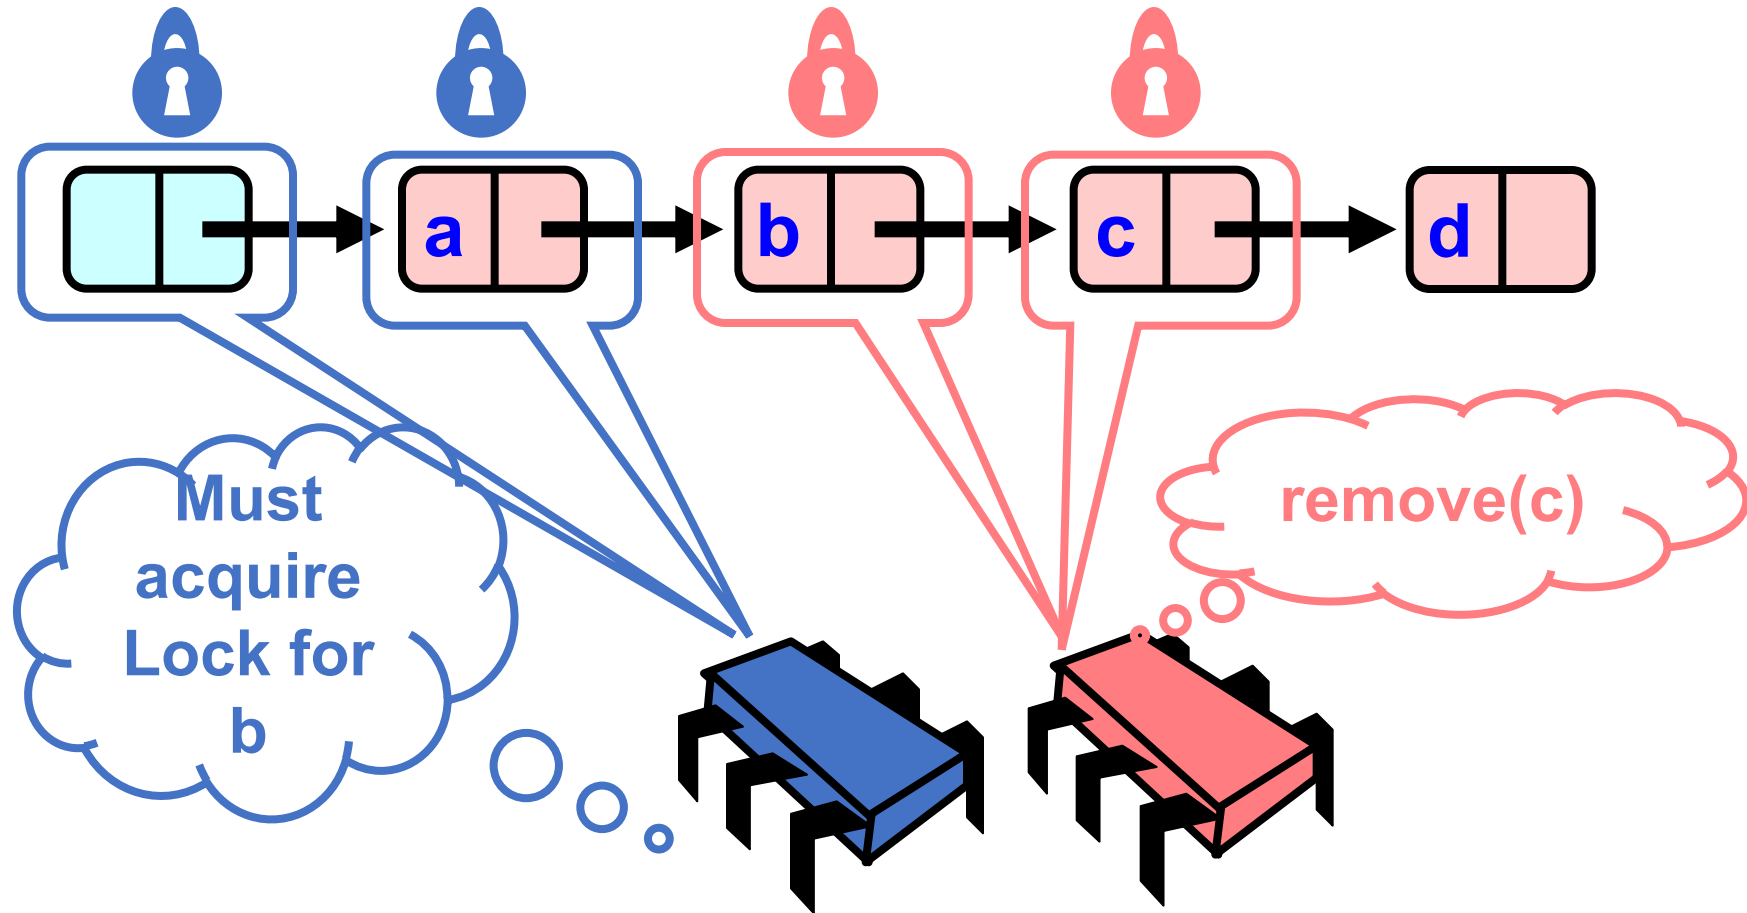
Removing a Node



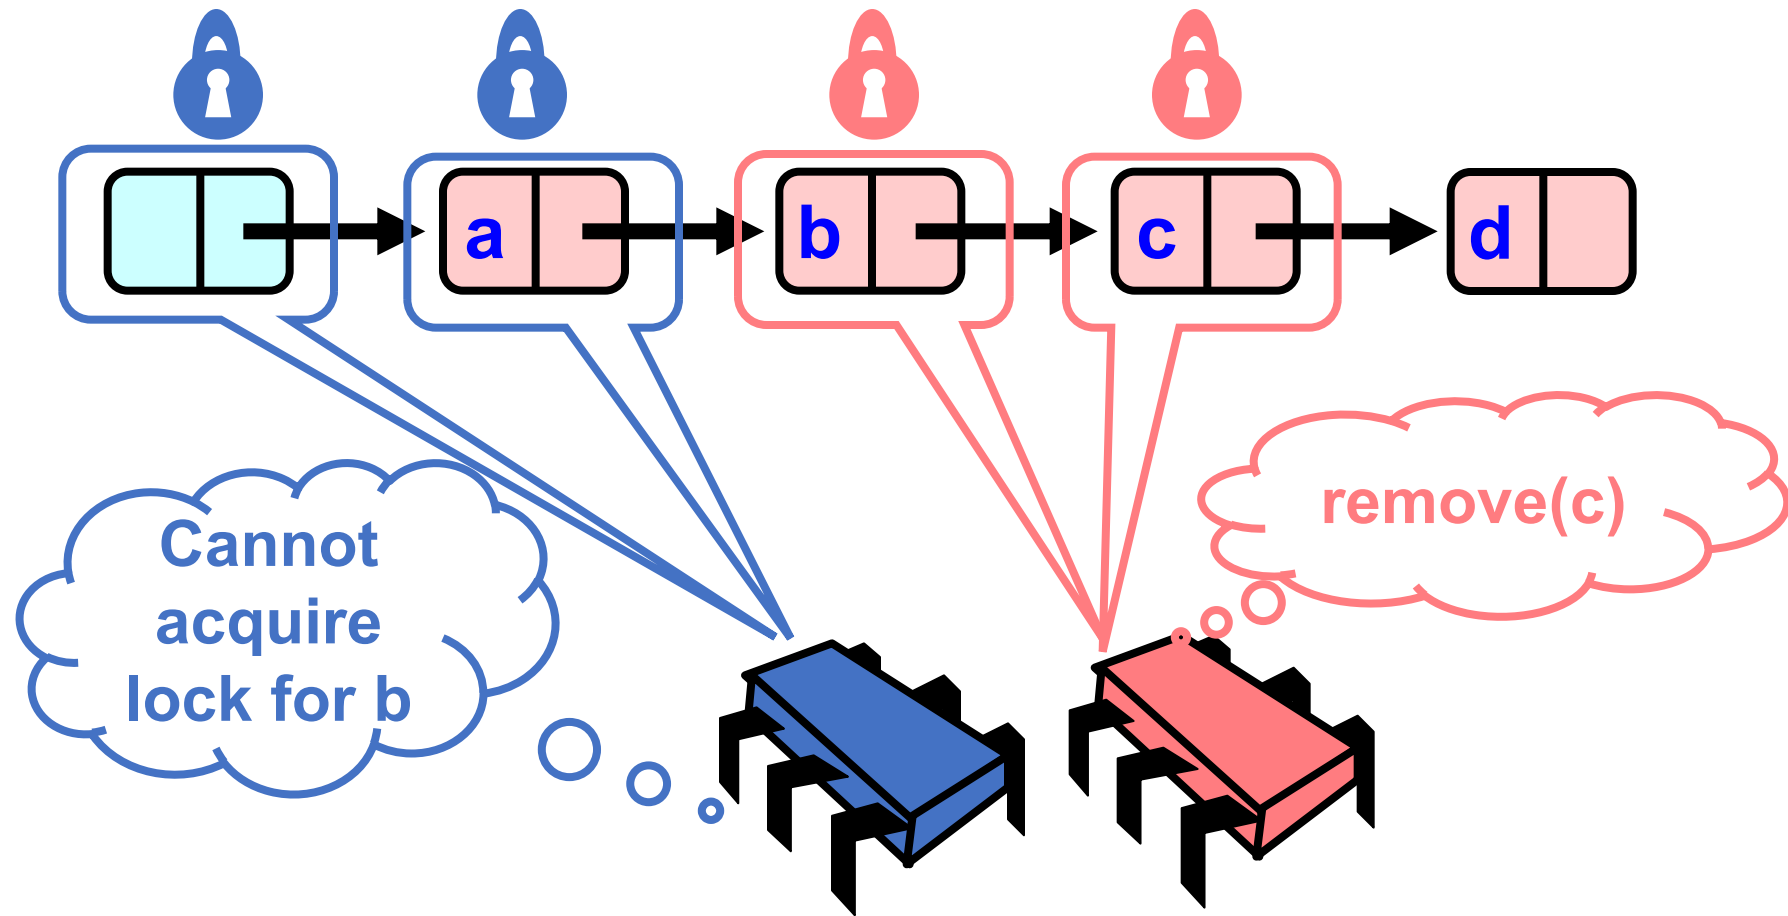
Removing a Node




Removing a Node



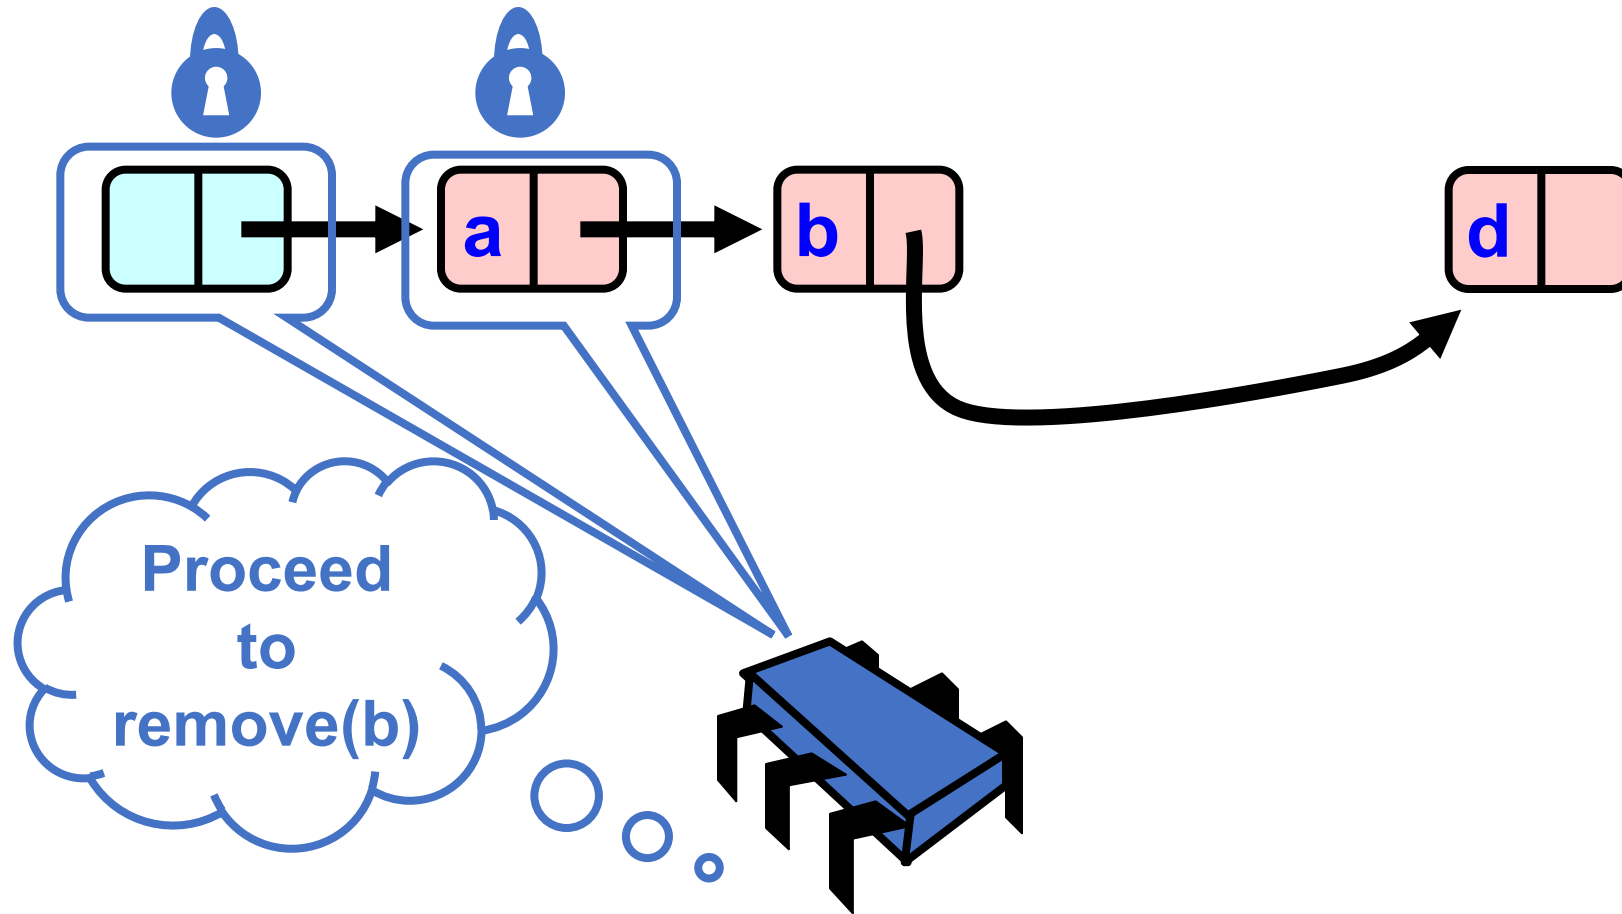
Removing a Node



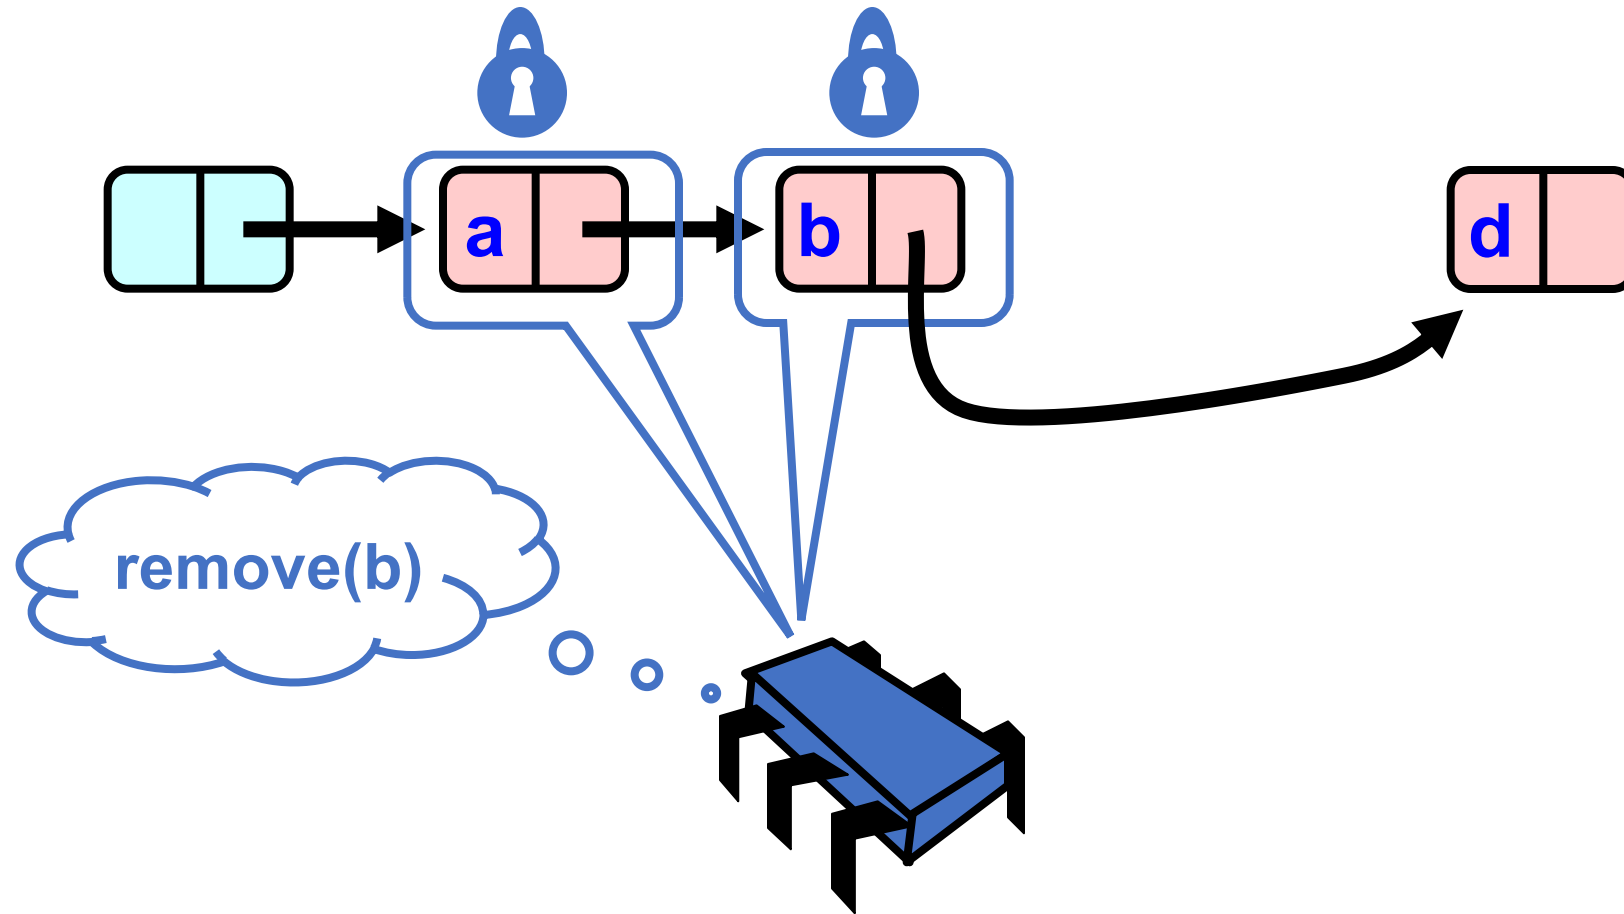
Removing a Node



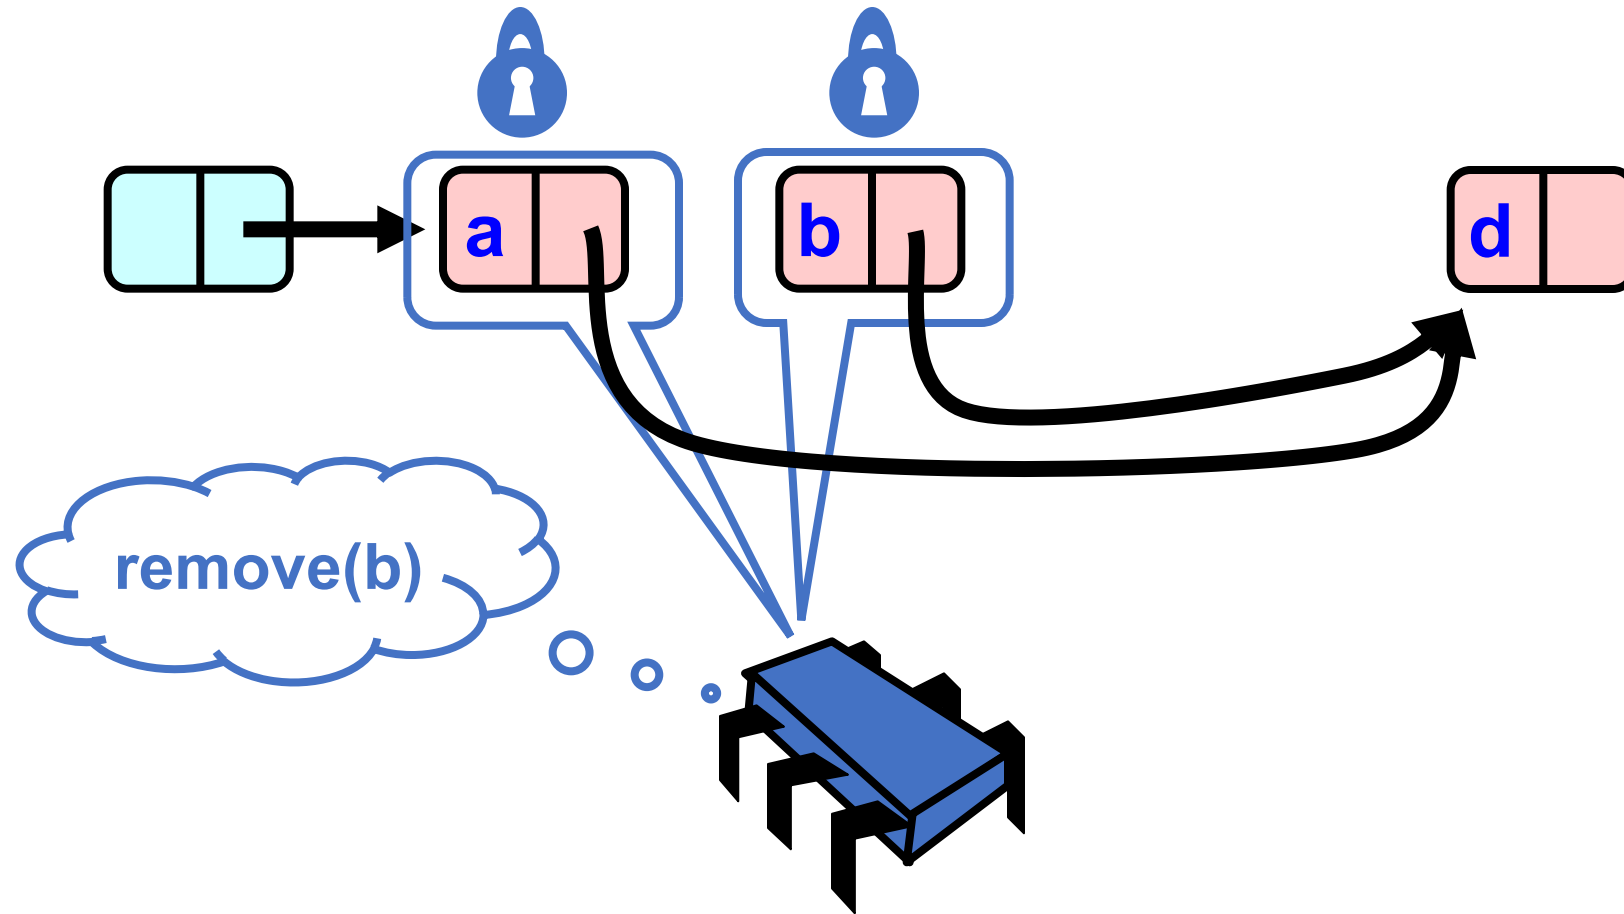
Removing a Node



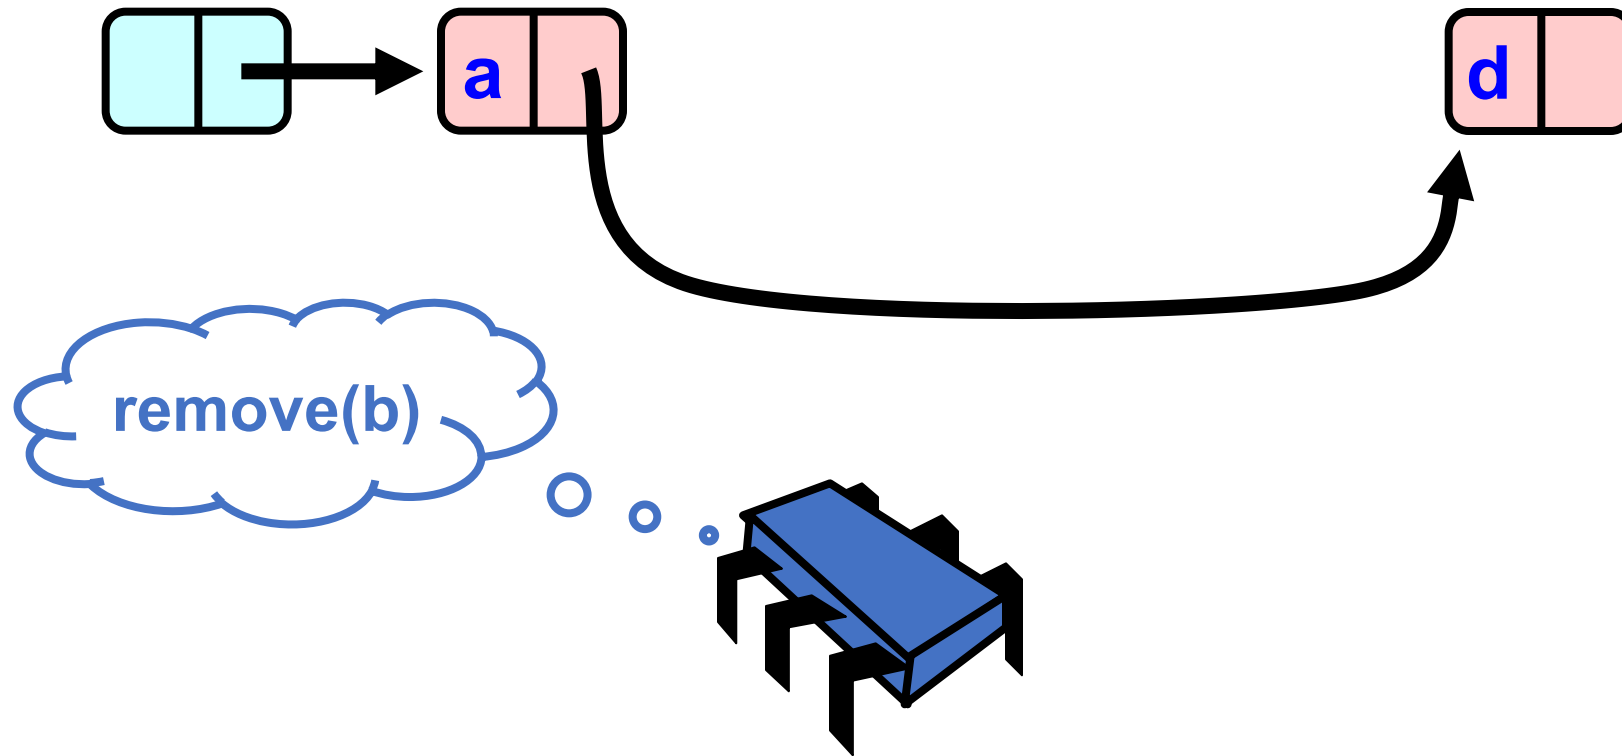
Removing a Node



Removing a Node



Removing a Node



Removing a Node



Adding Nodes

- Similar hand-over-hand locking
- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted

Drawbacks of fine-grained locking

- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Acquires and releases lock for every node traversed. Long chain of acquire/release.
 - Reduces concurrency (traffic jams). Inefficient.

Optimistic Synchronization

- Assume there will be no conflicts.
- Check before committing.
- If there was a conflict, try again.

Optimistic Synchronization

- Find nodes without locking. Then, lock the two nodes.

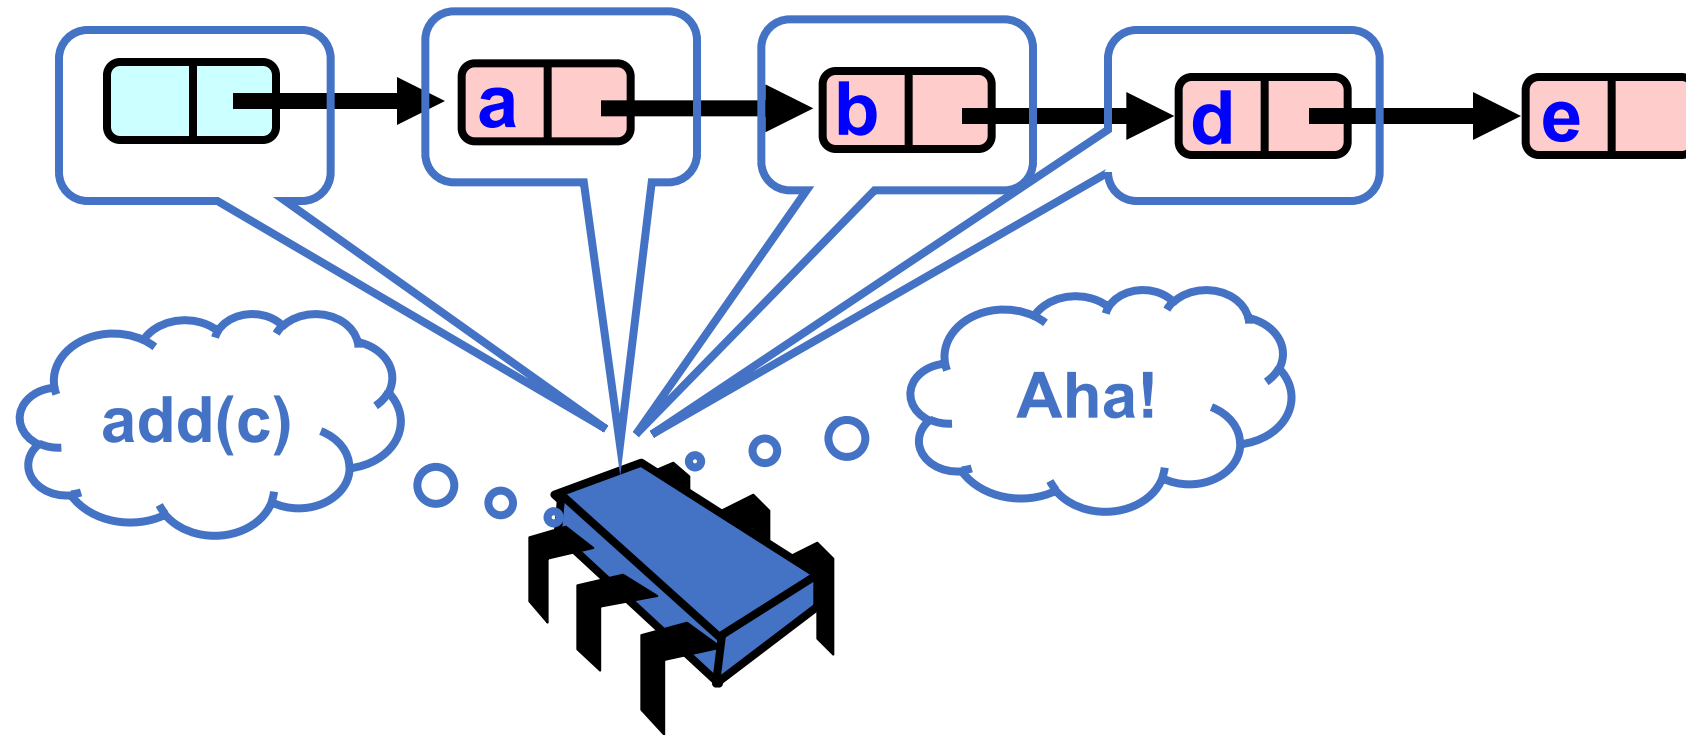
Optimistic Synchronization

- Find nodes without locking. Then, lock the two nodes.
- Check that the two nodes are still reachable.

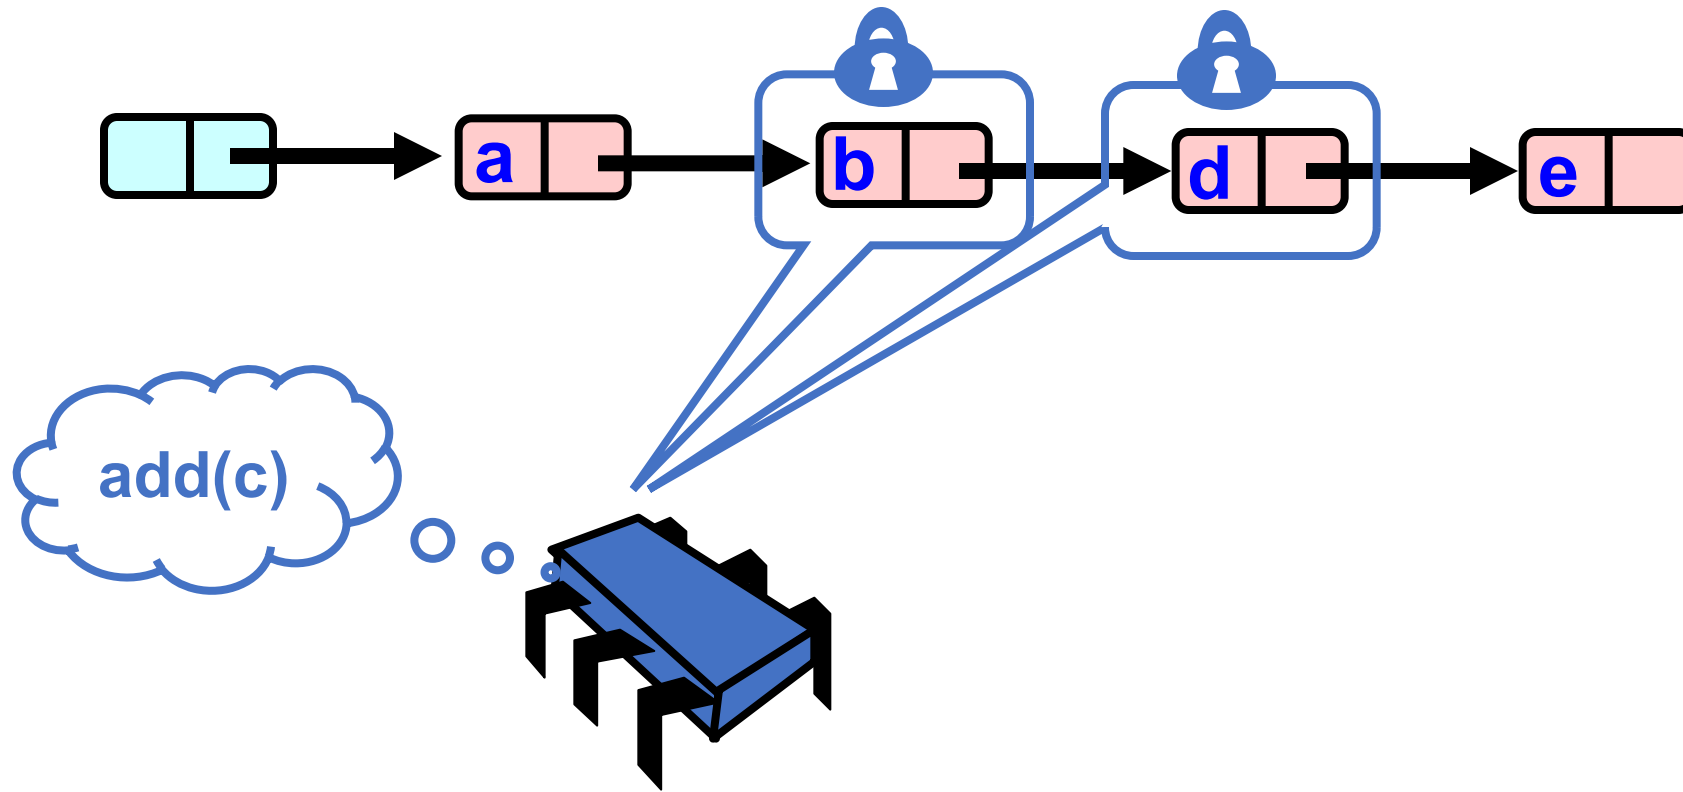
Optimistic Synchronization

- Find nodes without locking. Then, lock the two nodes.
- Check that the two nodes are still reachable.
- If they are not, search again from the beginning of the list.

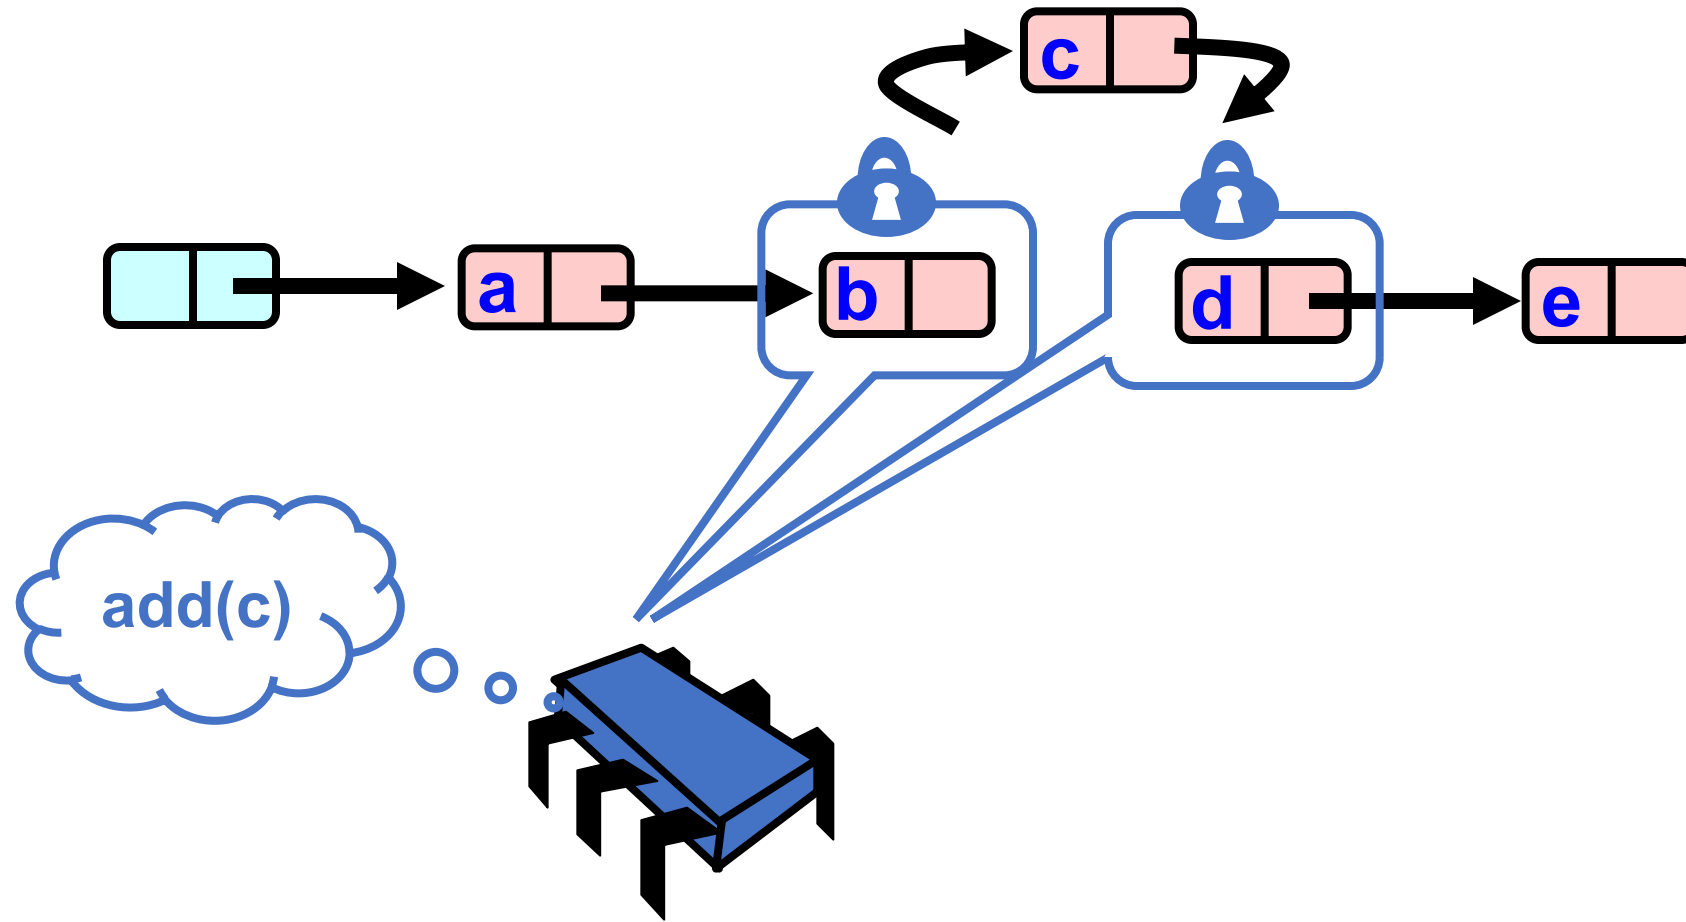
Optimistic: Traverse without Locking



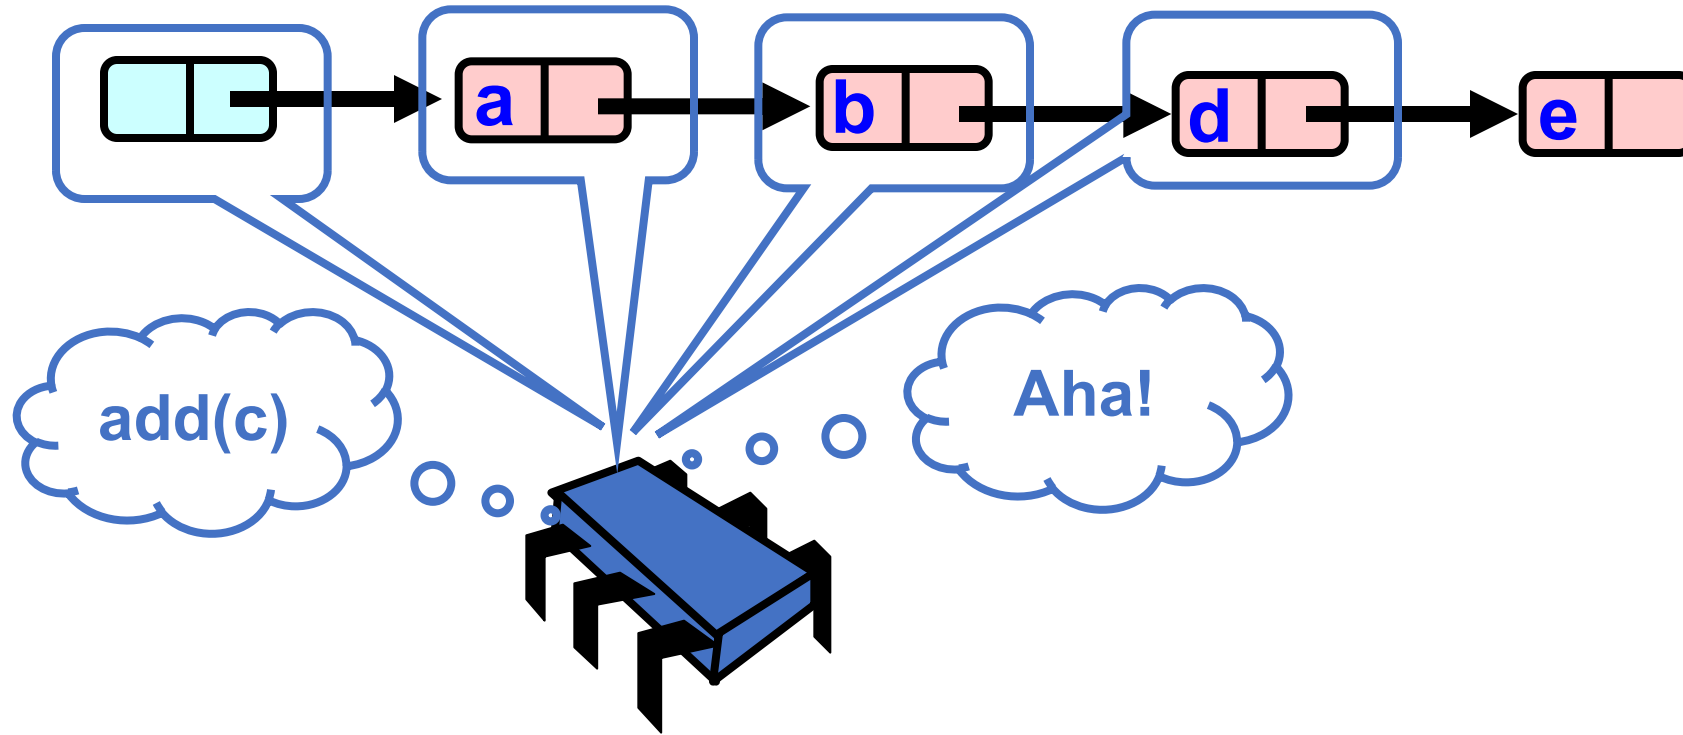
Optimistic: Lock



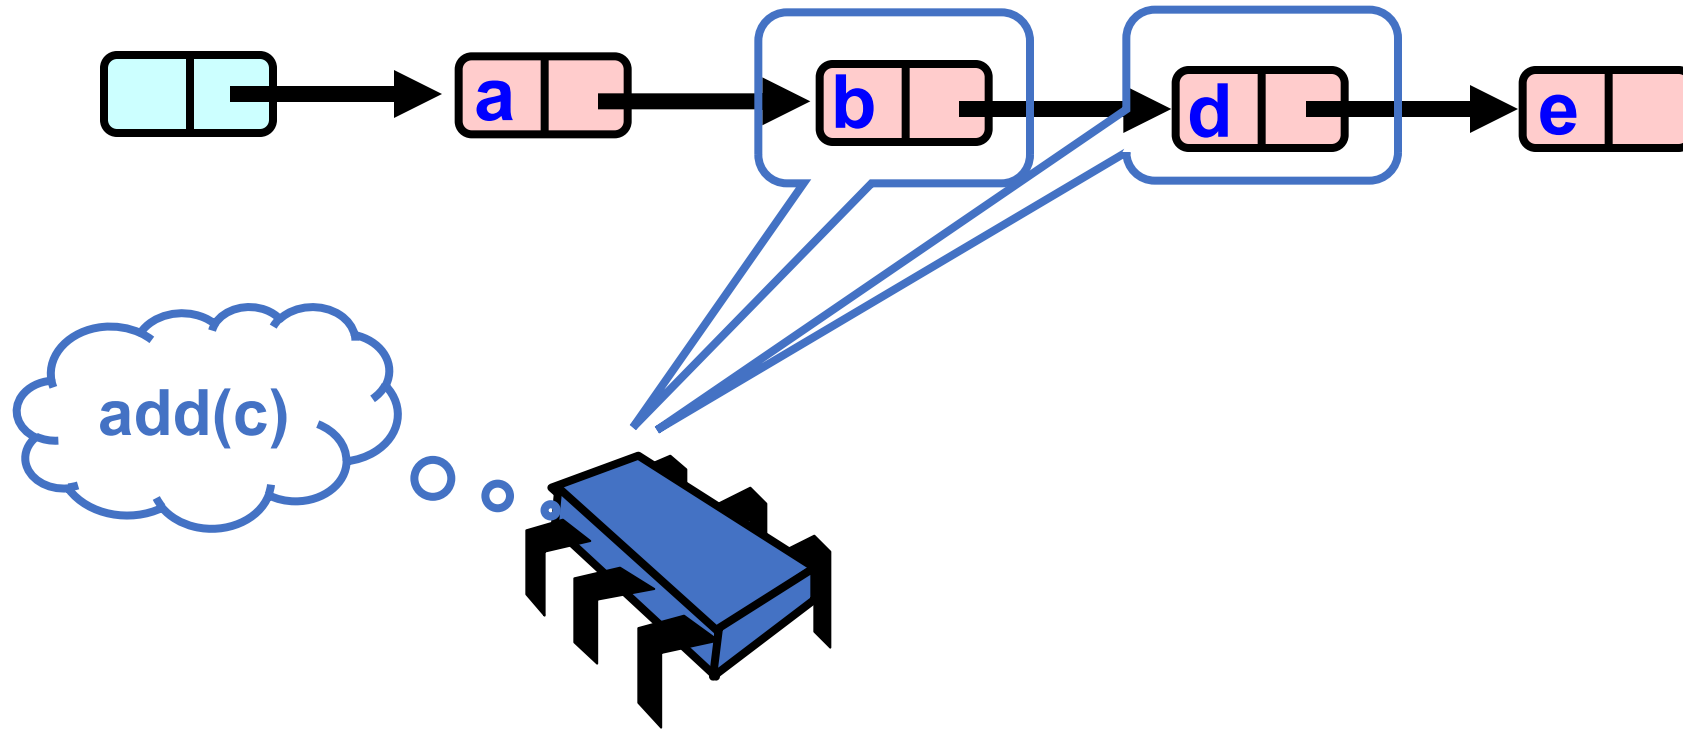
Optimistic: Insert without checking



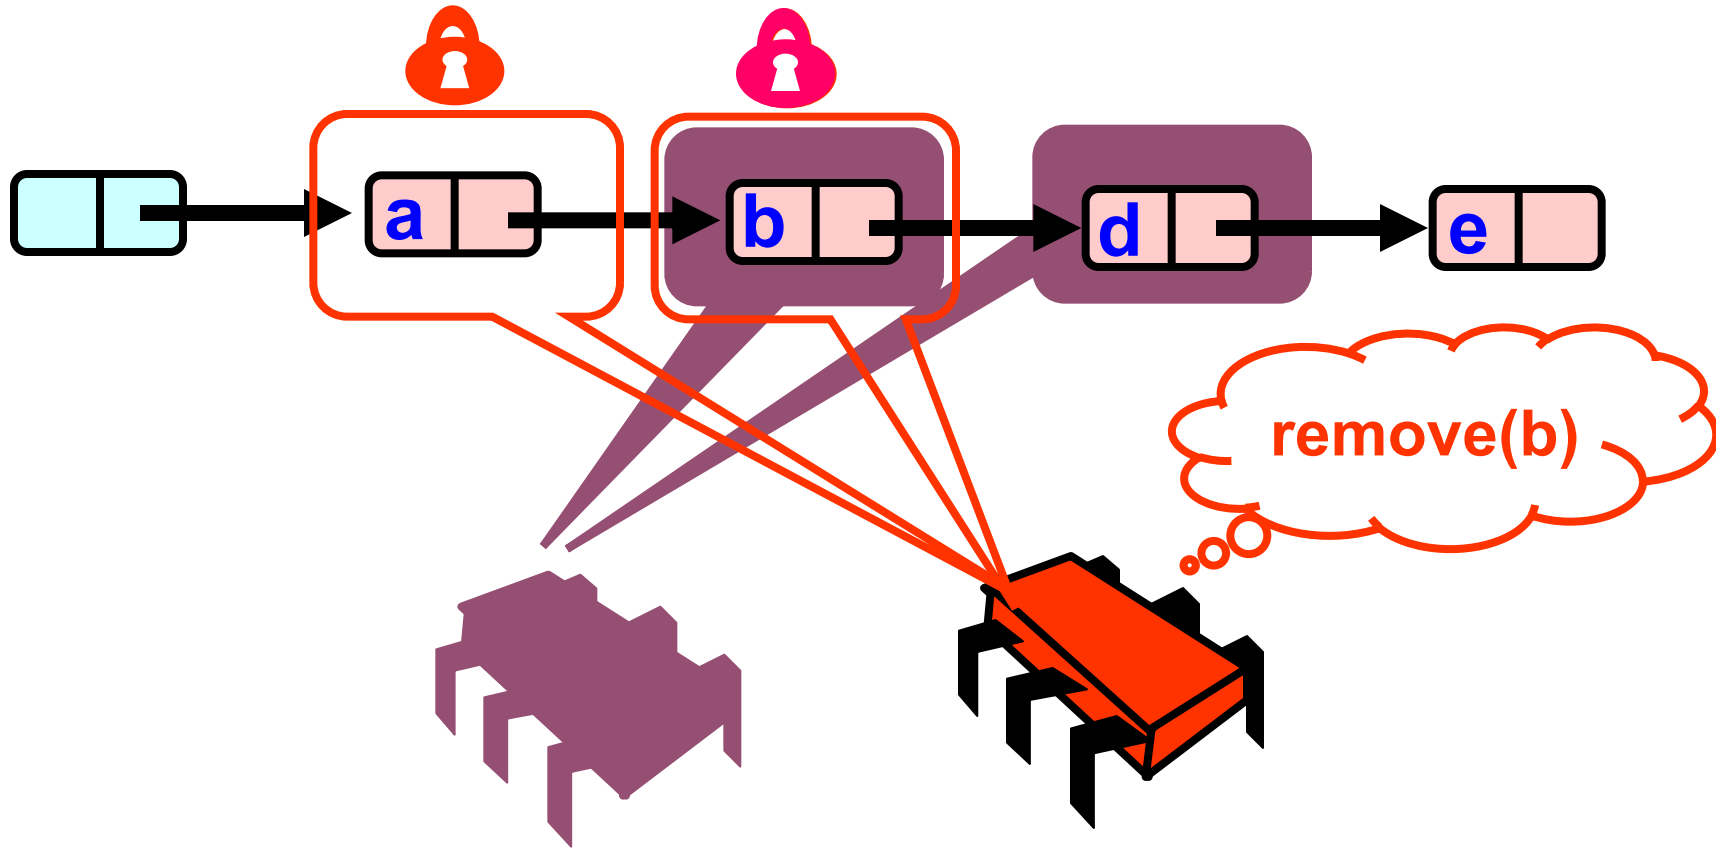
What could go wrong?



What could go wrong?



What could go wrong?



Data conflict (Race)!

- Red thread has the lock on a node (so it can modify the node)
- Blue thread is traversing without locks
- What do we do?

Race-freedom

- We can use atomic variables

Race-freedom

```
class Node {  
    public:  
        Value v;  
        int key;  
        Node *next;  
}
```

Race-freedom

```
class Node {  
    public:  
        Value v;  
        int key;  
        atomic<Node*> next;  
}
```

Create an atomic pointer type using C++ templates

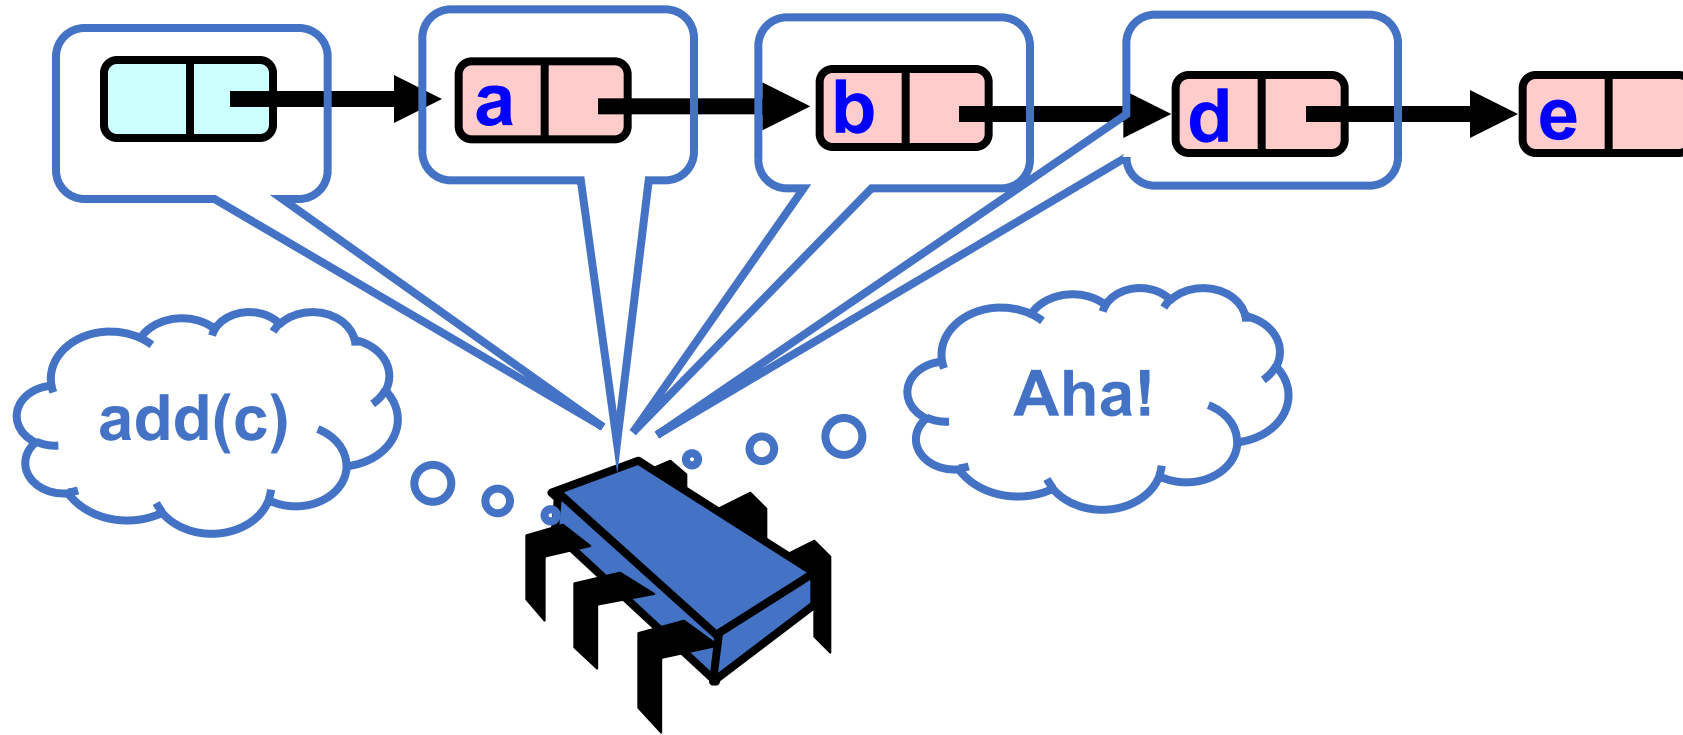
Race-freedom

```
void traverse(node *n) {  
    while (n->next != NULL) {  
        n = n->next;  
    }  
}
```

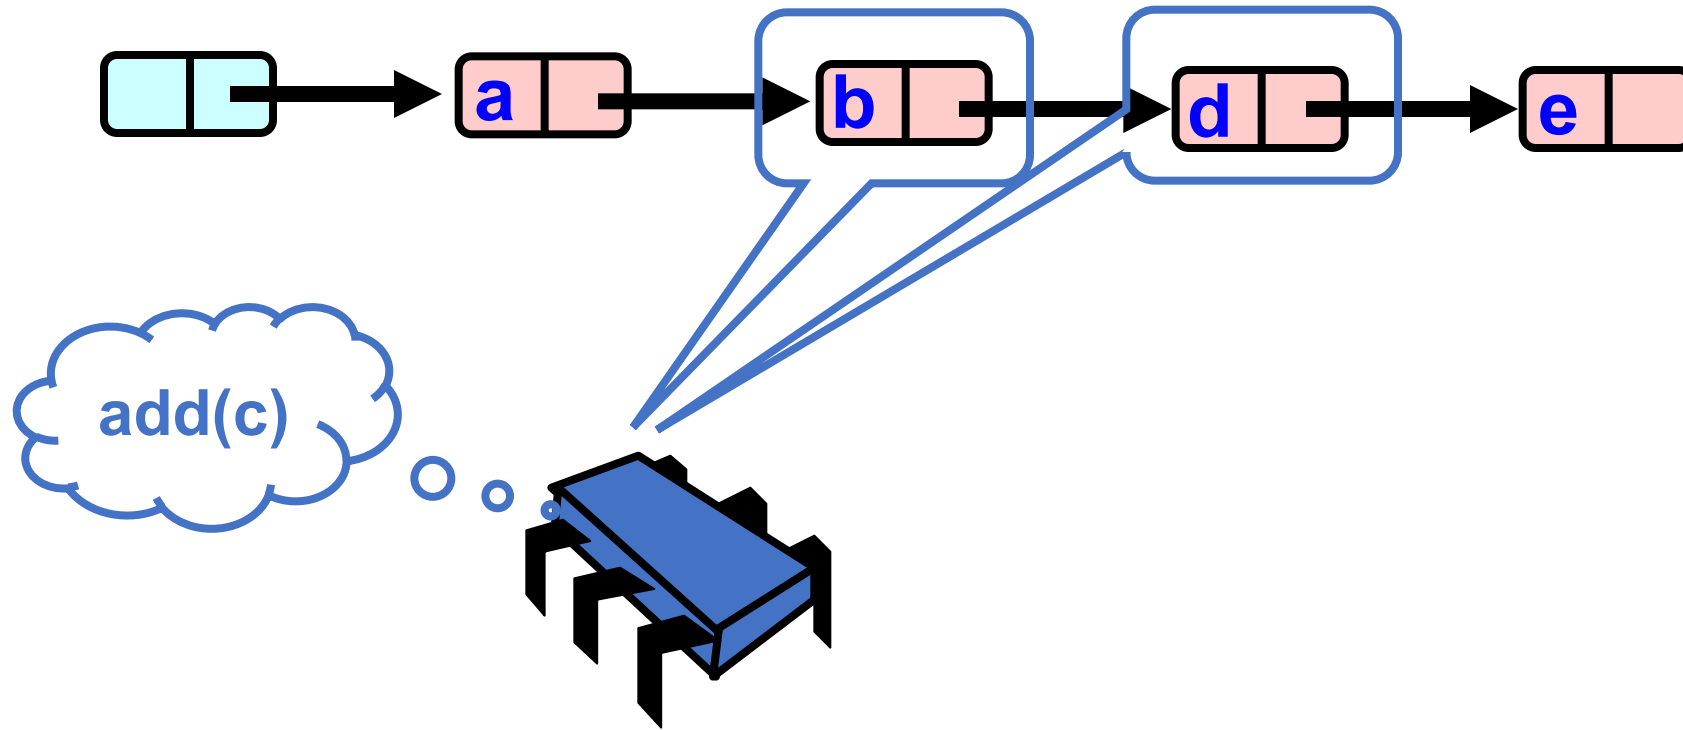

Race-freedom

```
void traverse(node *n) {  
    while (n->next.load() != NULL) {  
        n = n->next.load();  
    }  
}
```

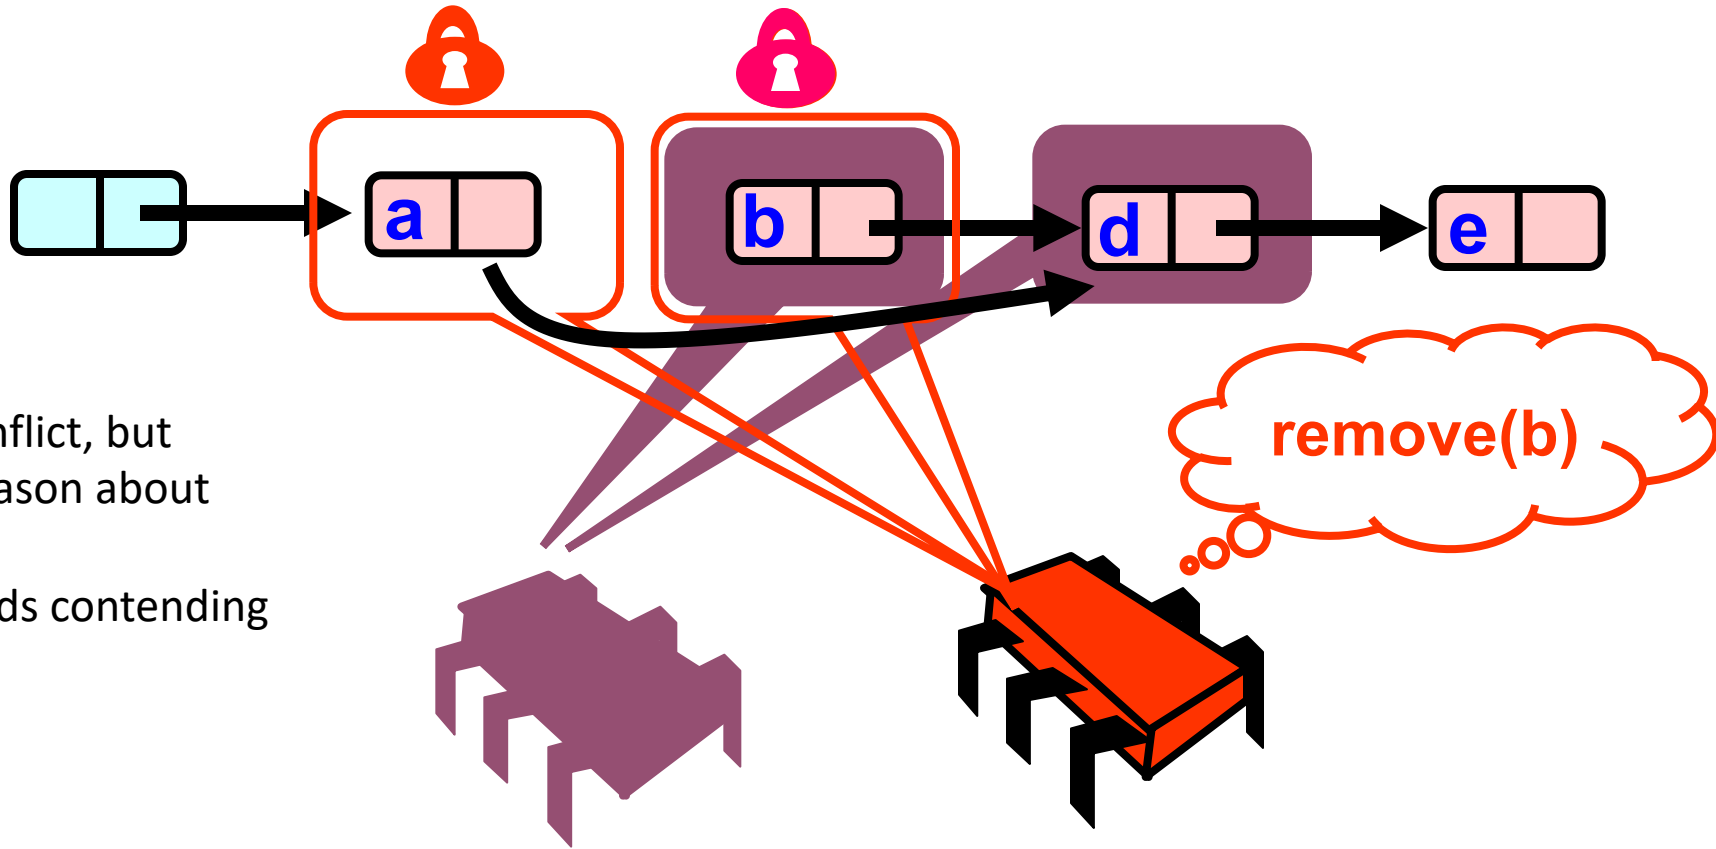
What could go wrong?



What could go wrong?




What could go wrong?

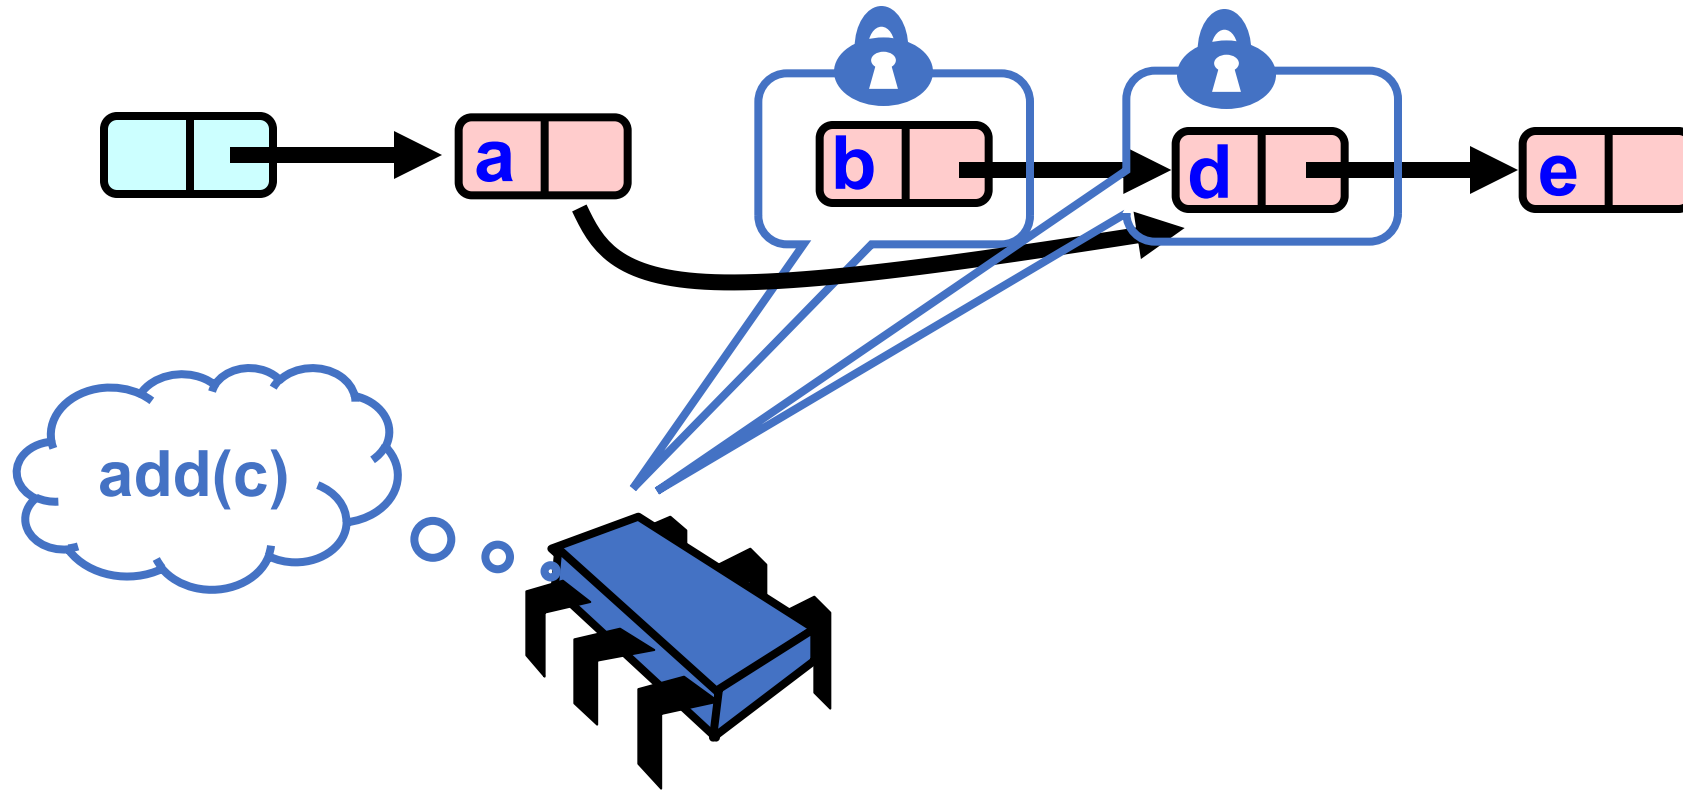


No more data conflict, but
we do need to reason about
interleavings.
Concurrent threads contending
for values.

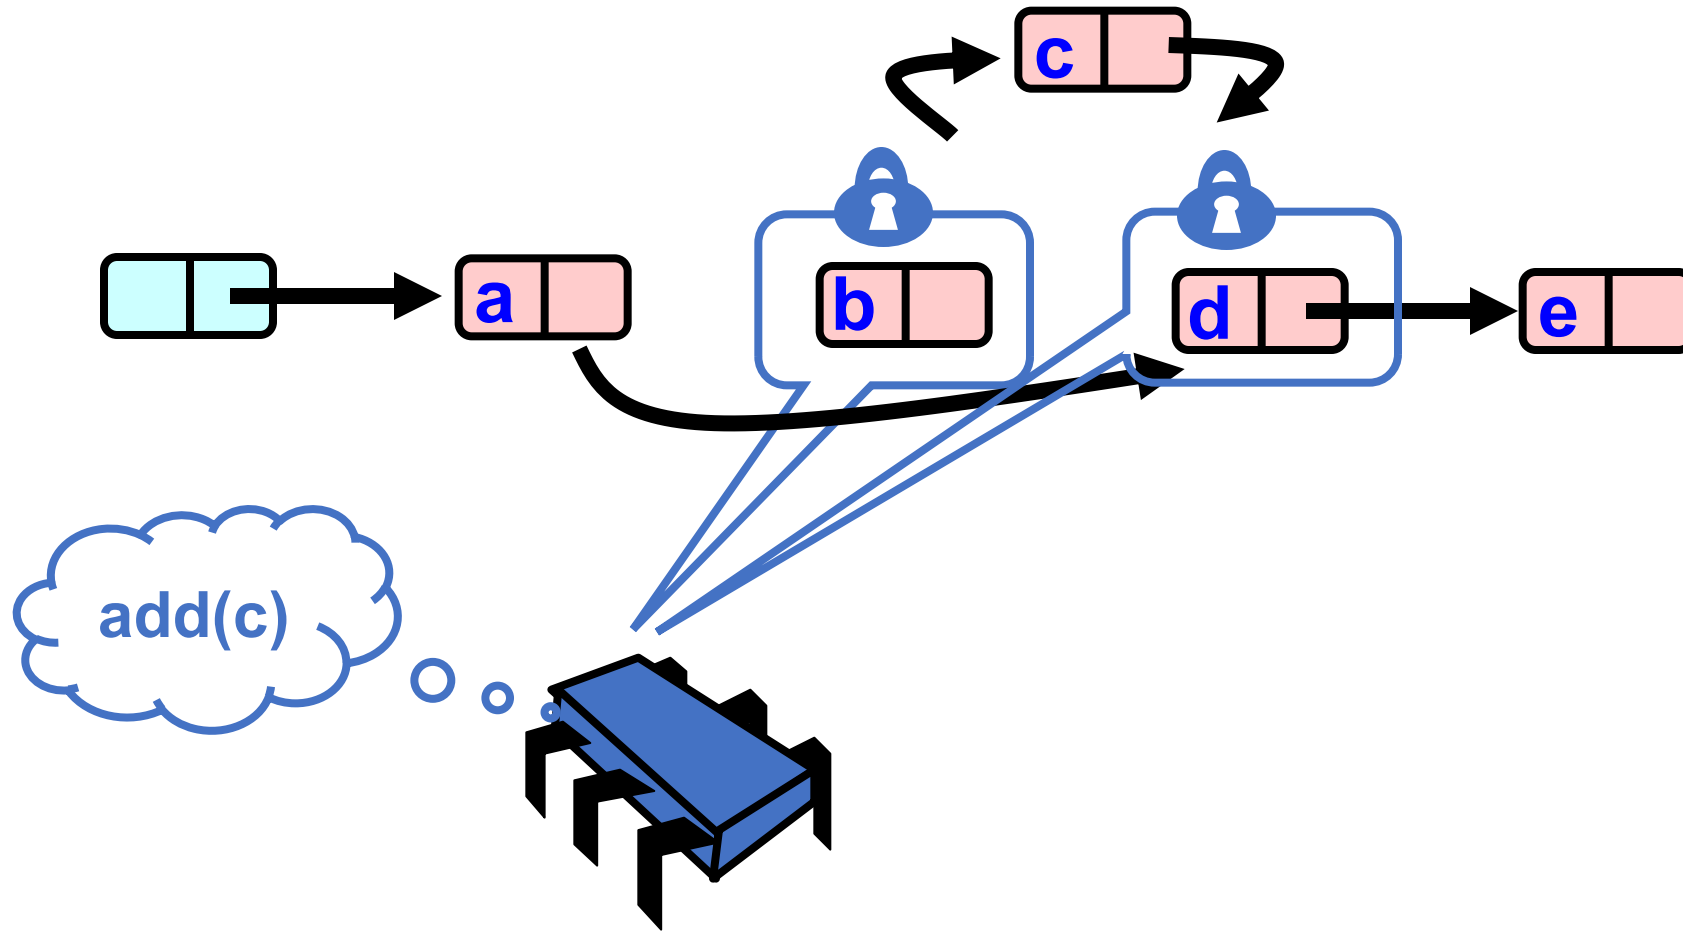
What could go wrong?



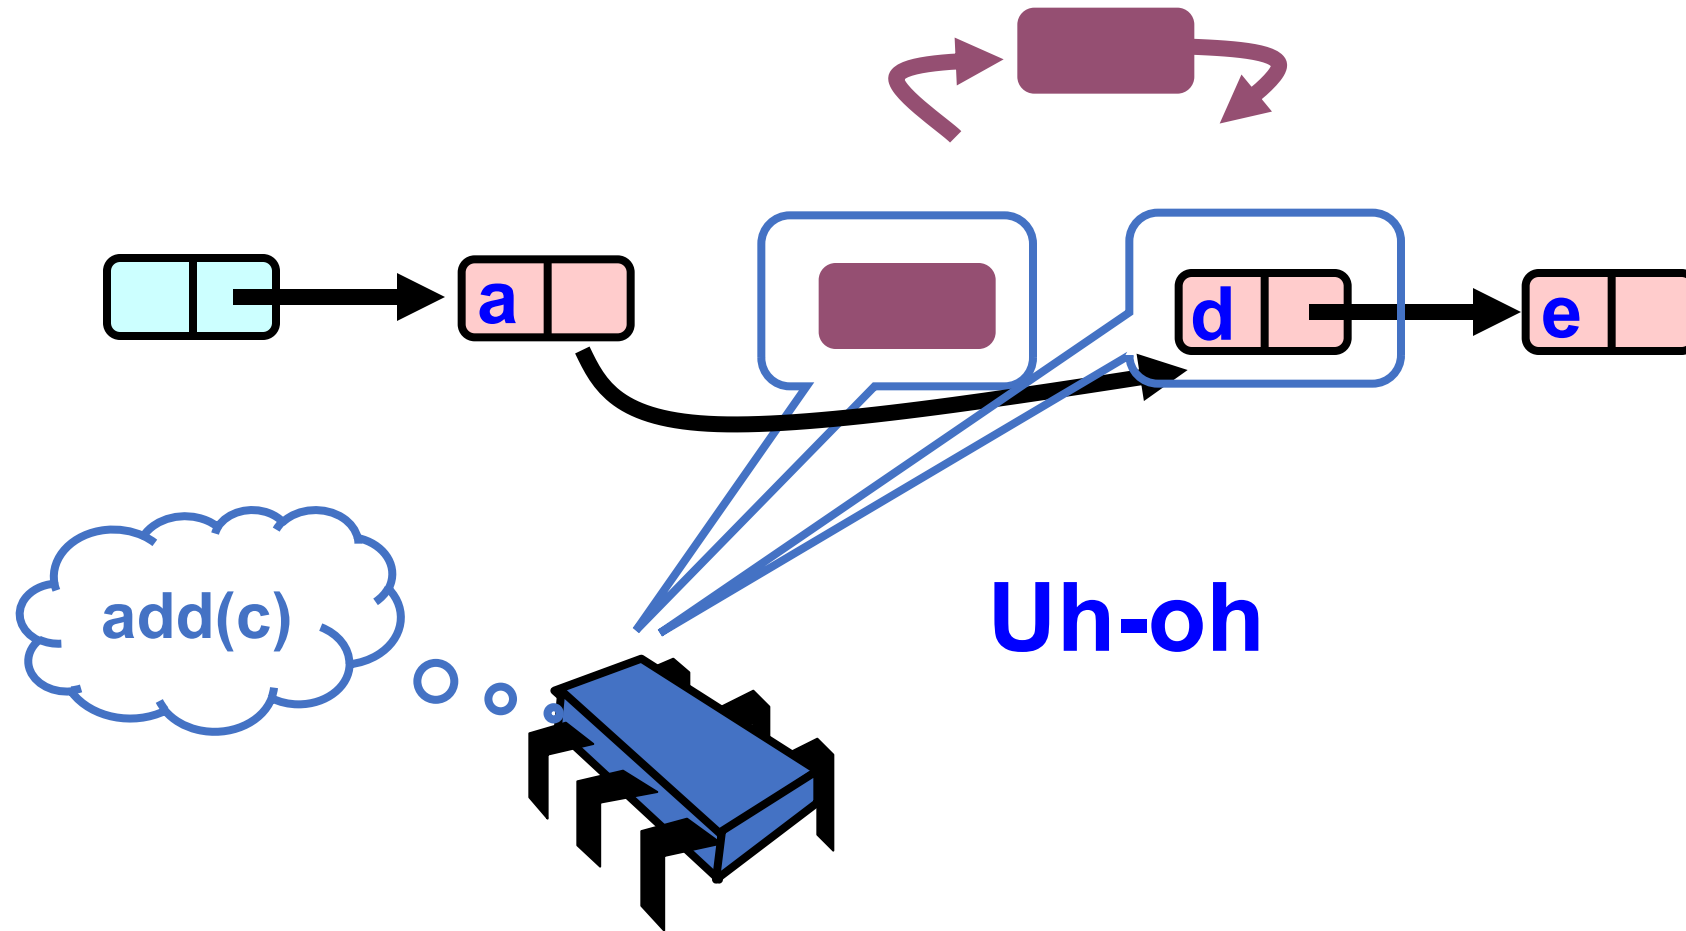
What could go wrong?



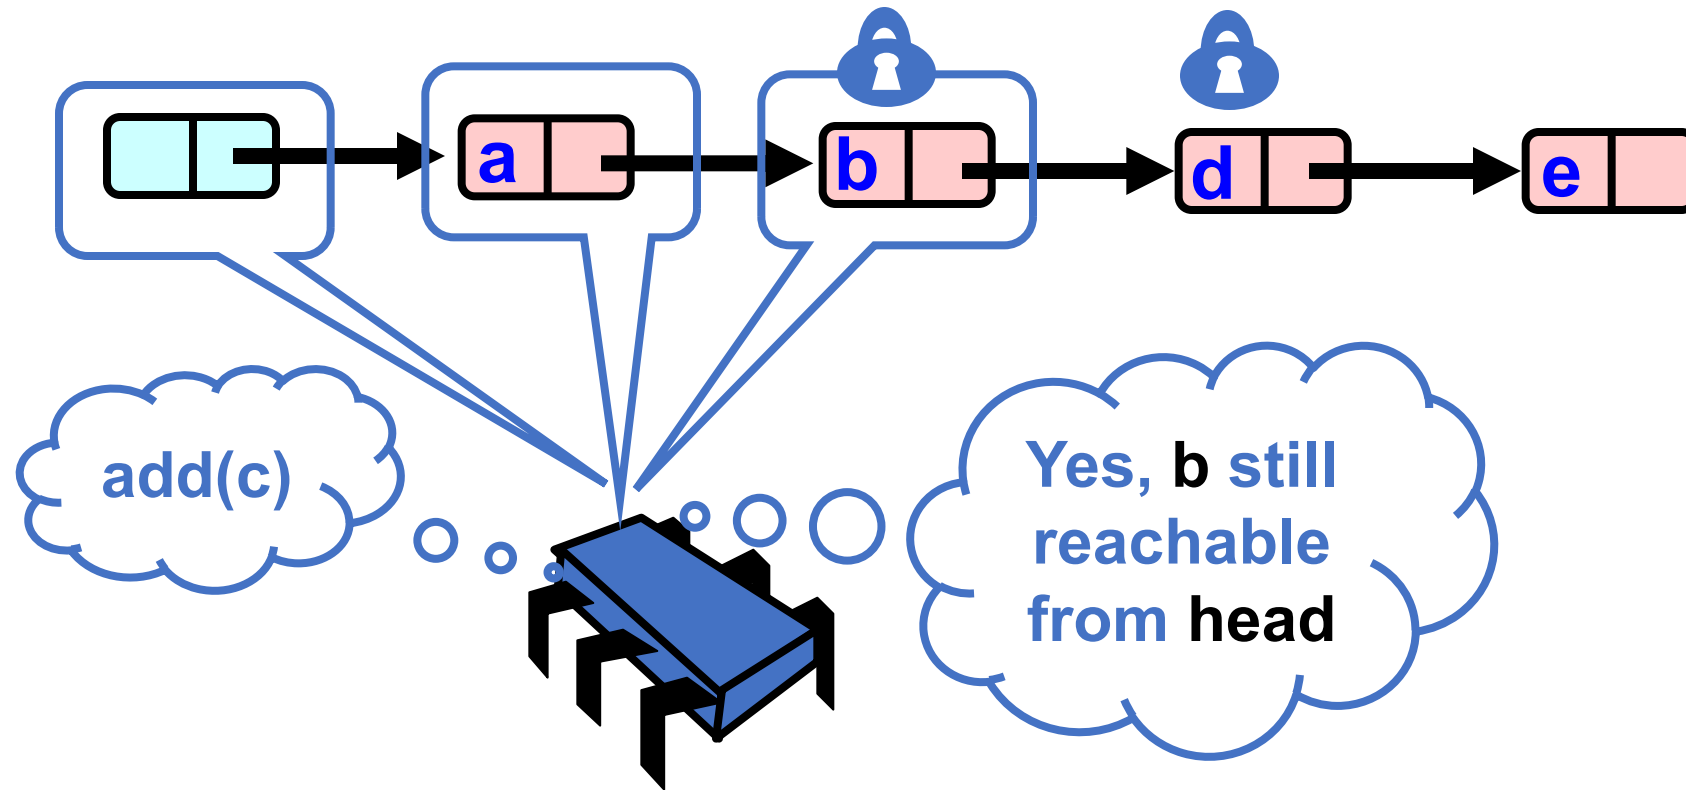
What could go wrong?



What could go wrong?



Validate – Part 1



What happens if failure?

- Ideas?

What happens if failure?

- Could try to recover? Back up a node?
 - Very tricky!
 - Just start over!

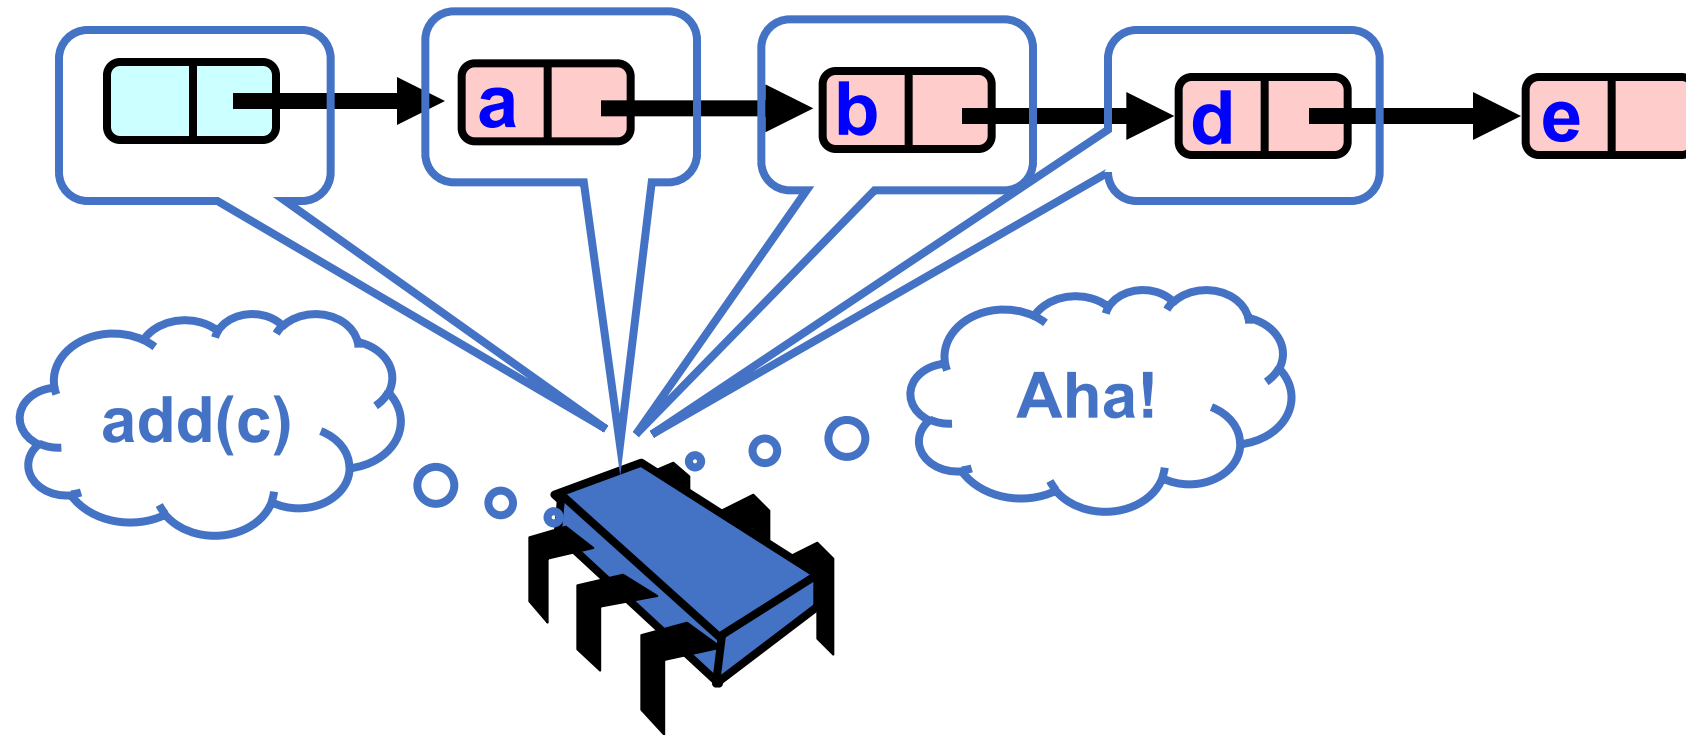
What happens if failure?

- Could try to recover? Back up a node?
 - Very tricky!
 - Just start over!
- Private method:
 - `try_remove`
 - remove loops on `try_remove` until it succeeds


What about concurrent adds?

What if 2 threads try to add a node in the same position?


What Else Could Go Wrong?




What Else Could Go Wrong?




What Else Could Go Wrong?




What Else Could Go Wrong?




What Else Could Go Wrong?




What Else Could Go Wrong?




What Else Could Go Wrong?




Validate Part 2 (while holding locks)




What Else Could Go Wrong?




Can threads that remove a node delete (free) it?




Can threads that remove a node delete it?




Can threads that remove a node delete it?



Can threads that remove a node delete it?



Our own garbage collector



Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:


Our own garbage collector



Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:





Our own garbage collector




Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:
can be thread-local





Our own garbage collector



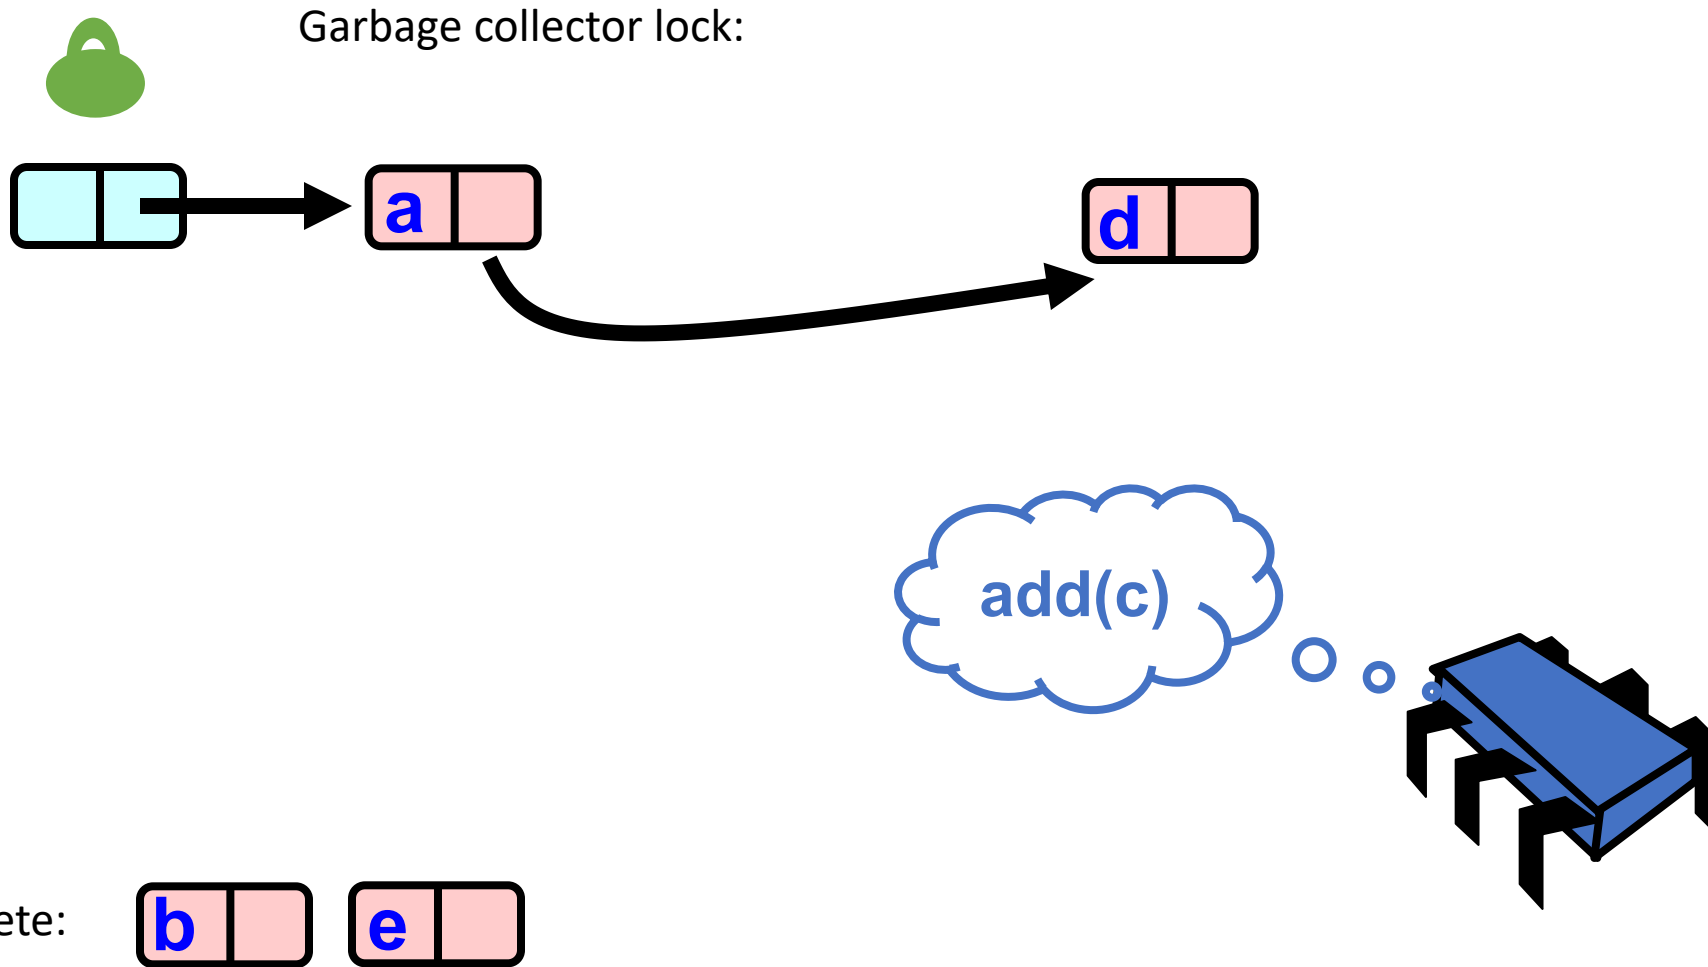
Java's garbage collection will remove b

We are using a better™ language though...


maintain a list to delete:




Our own garbage collector



Our own garbage collector



Similar to a reader/writer lock:
Allows an arbitrary number of threads that operate on the list
Only 1 garbage collector thread erases the list of nodes



maintain a list to delete:

The diagram shows a list of nodes to be deleted. It consists of two nodes, 'b' and 'e'. Each node is a rectangle divided into two sections: the left section contains the letter in blue, and the right section is pink. Node 'b' is on the left, and node 'e' is on the right.

Garbage collector lock

- Many strategies!
 - A big research area ~10 years ago
- **Strat 1:** Threads always try once to take the garbage collector lock:
 - if failed, no worries, the next operation will get a chance
 - can starve garbage collection
- **Strat 2:** Wait until size grows to a threshold:
 - Wait on the lock
 - Can cause performance spikes

Summary

- We traverse without lock
 - Traversal may access nodes that are locked
 - Its okay because we have atomic pointers!
- We avoid traversing the old structure
 - We validate after we obtain locks.
 - Our node is still reachable (it was not deleted)
 - Our insertion point is still valid (no thread has inserted in the meantime)
- To free nodes
 - We put them in a list to be freed later.